

P.PORTO

Microservices Lesson 3

Methods and Techniques for
Software Development

2019/2020

P.PORTO

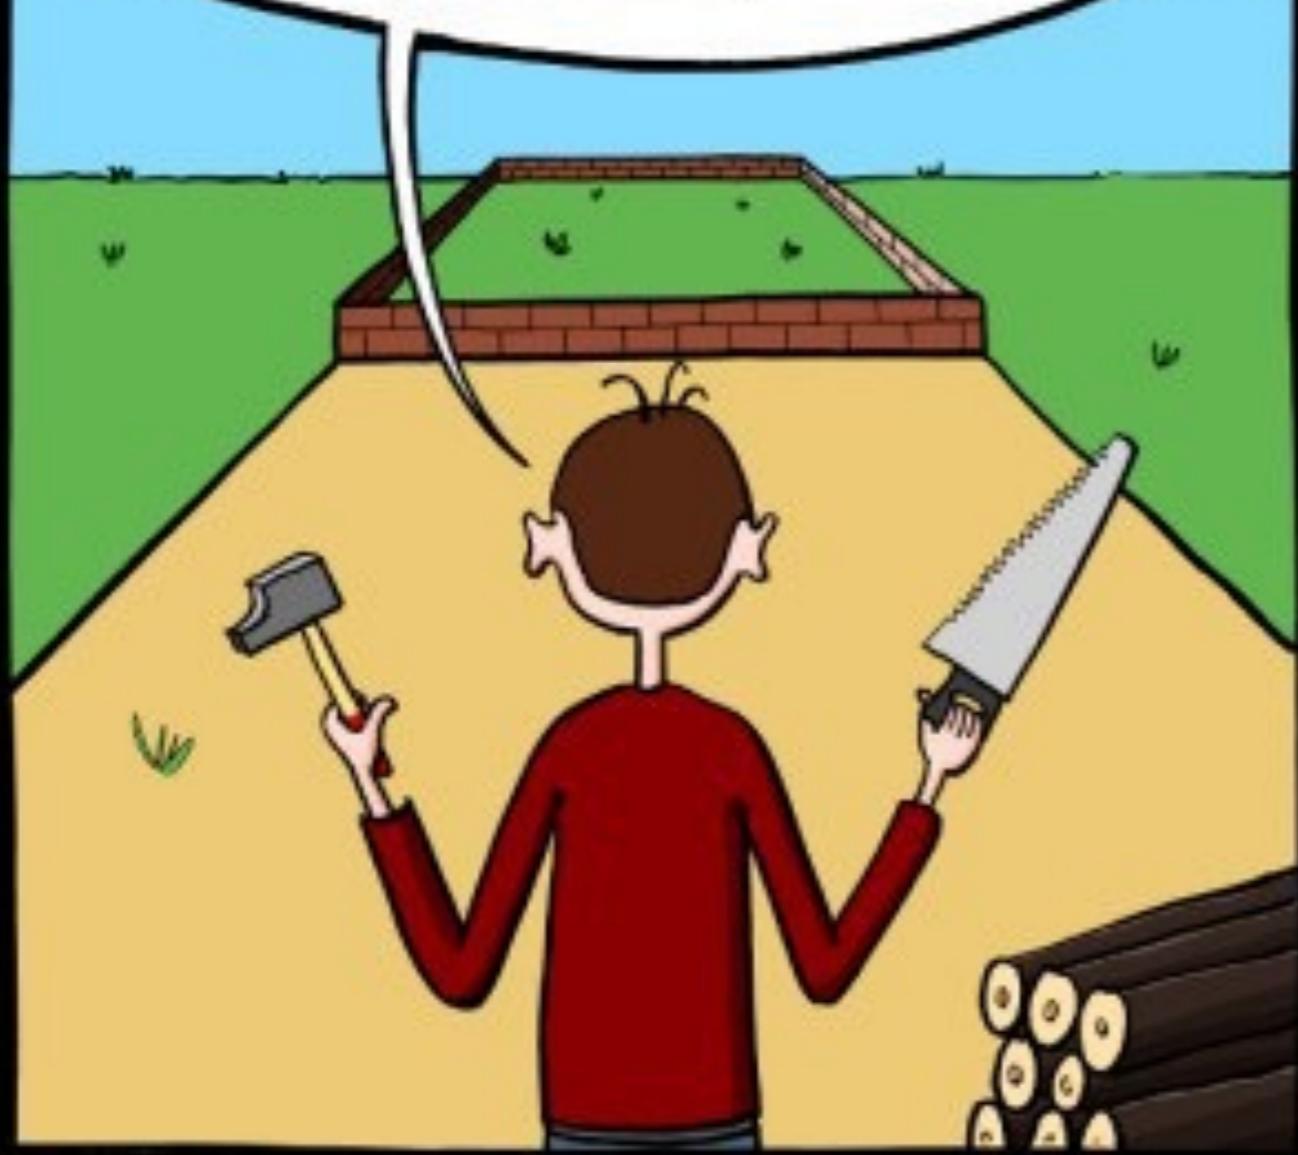
Microserviços Aula 3

Methods and Techniques for
Software Development

2019/2020

THE LIFE OF A SOFTWARE ENGINEER.

CLEAN SLATE. SOLID
FOUNDATIONS. THIS TIME
I WILL BUILD THINGS THE
RIGHT WAY.



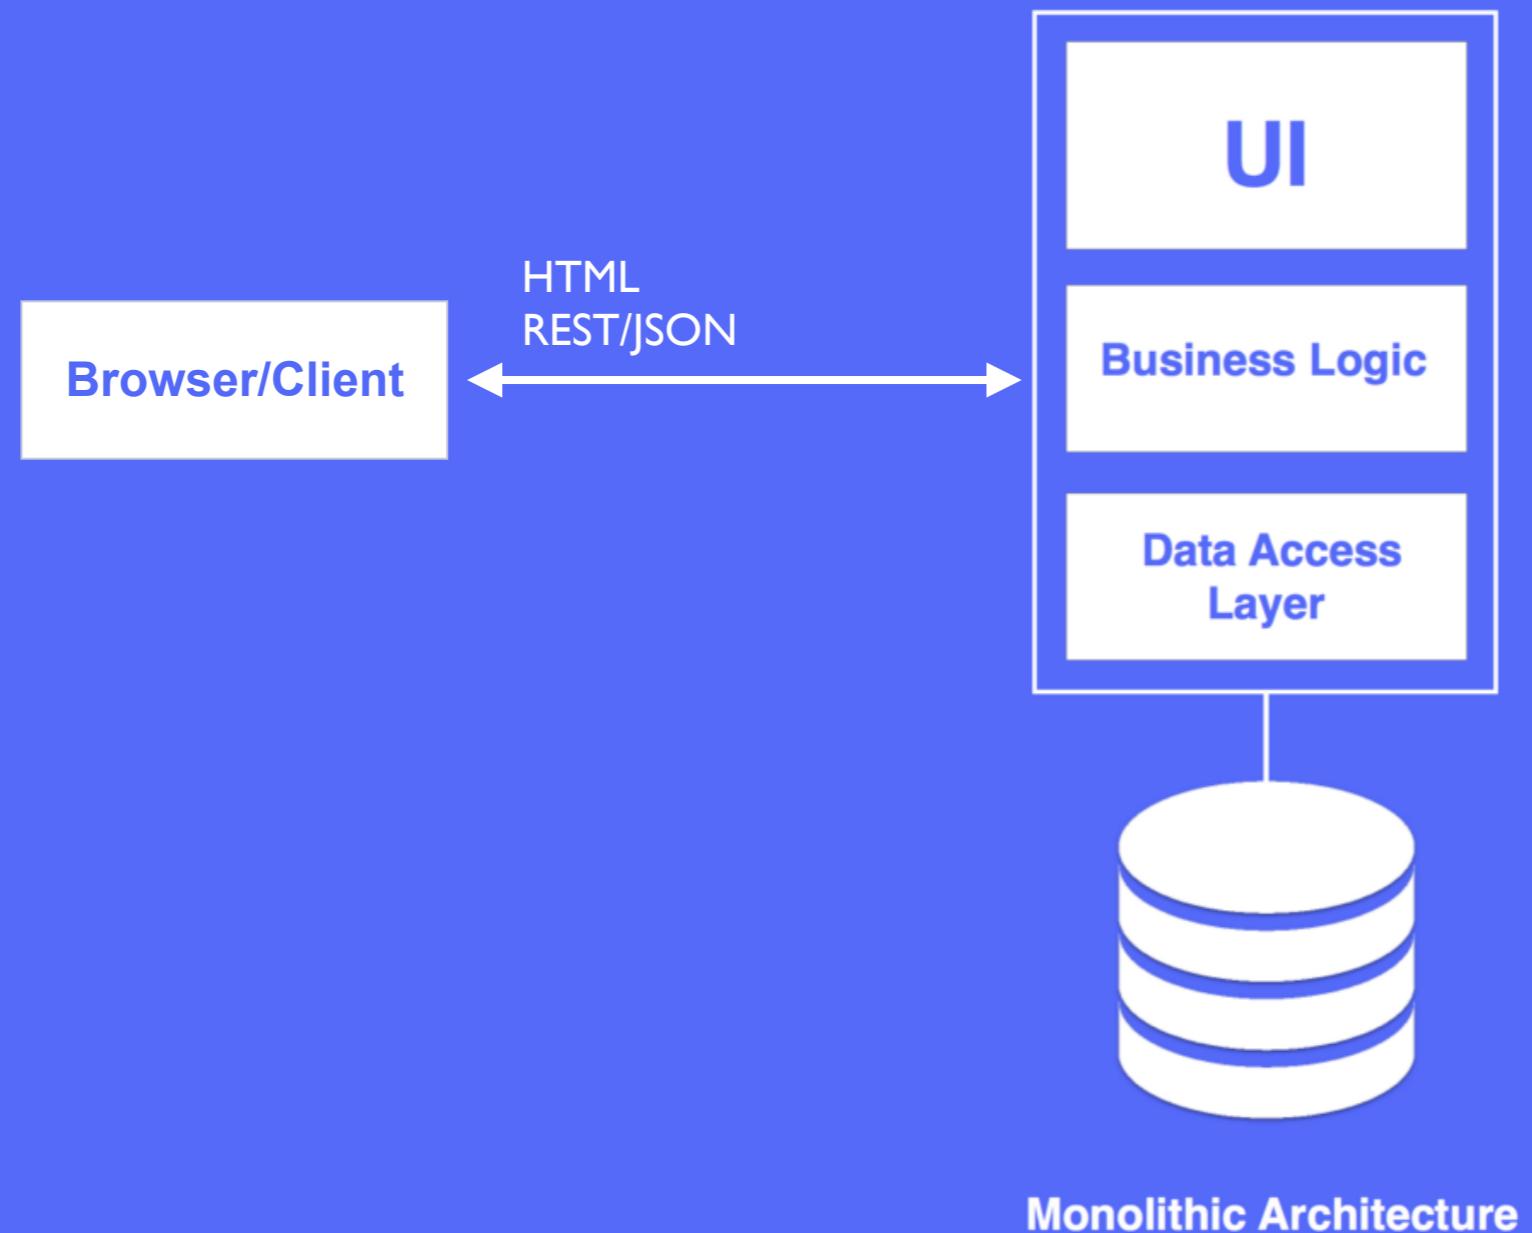
MUCH LATER...

OH MY. I'VE
DONE IT AGAIN,
HAVEN'T I ?

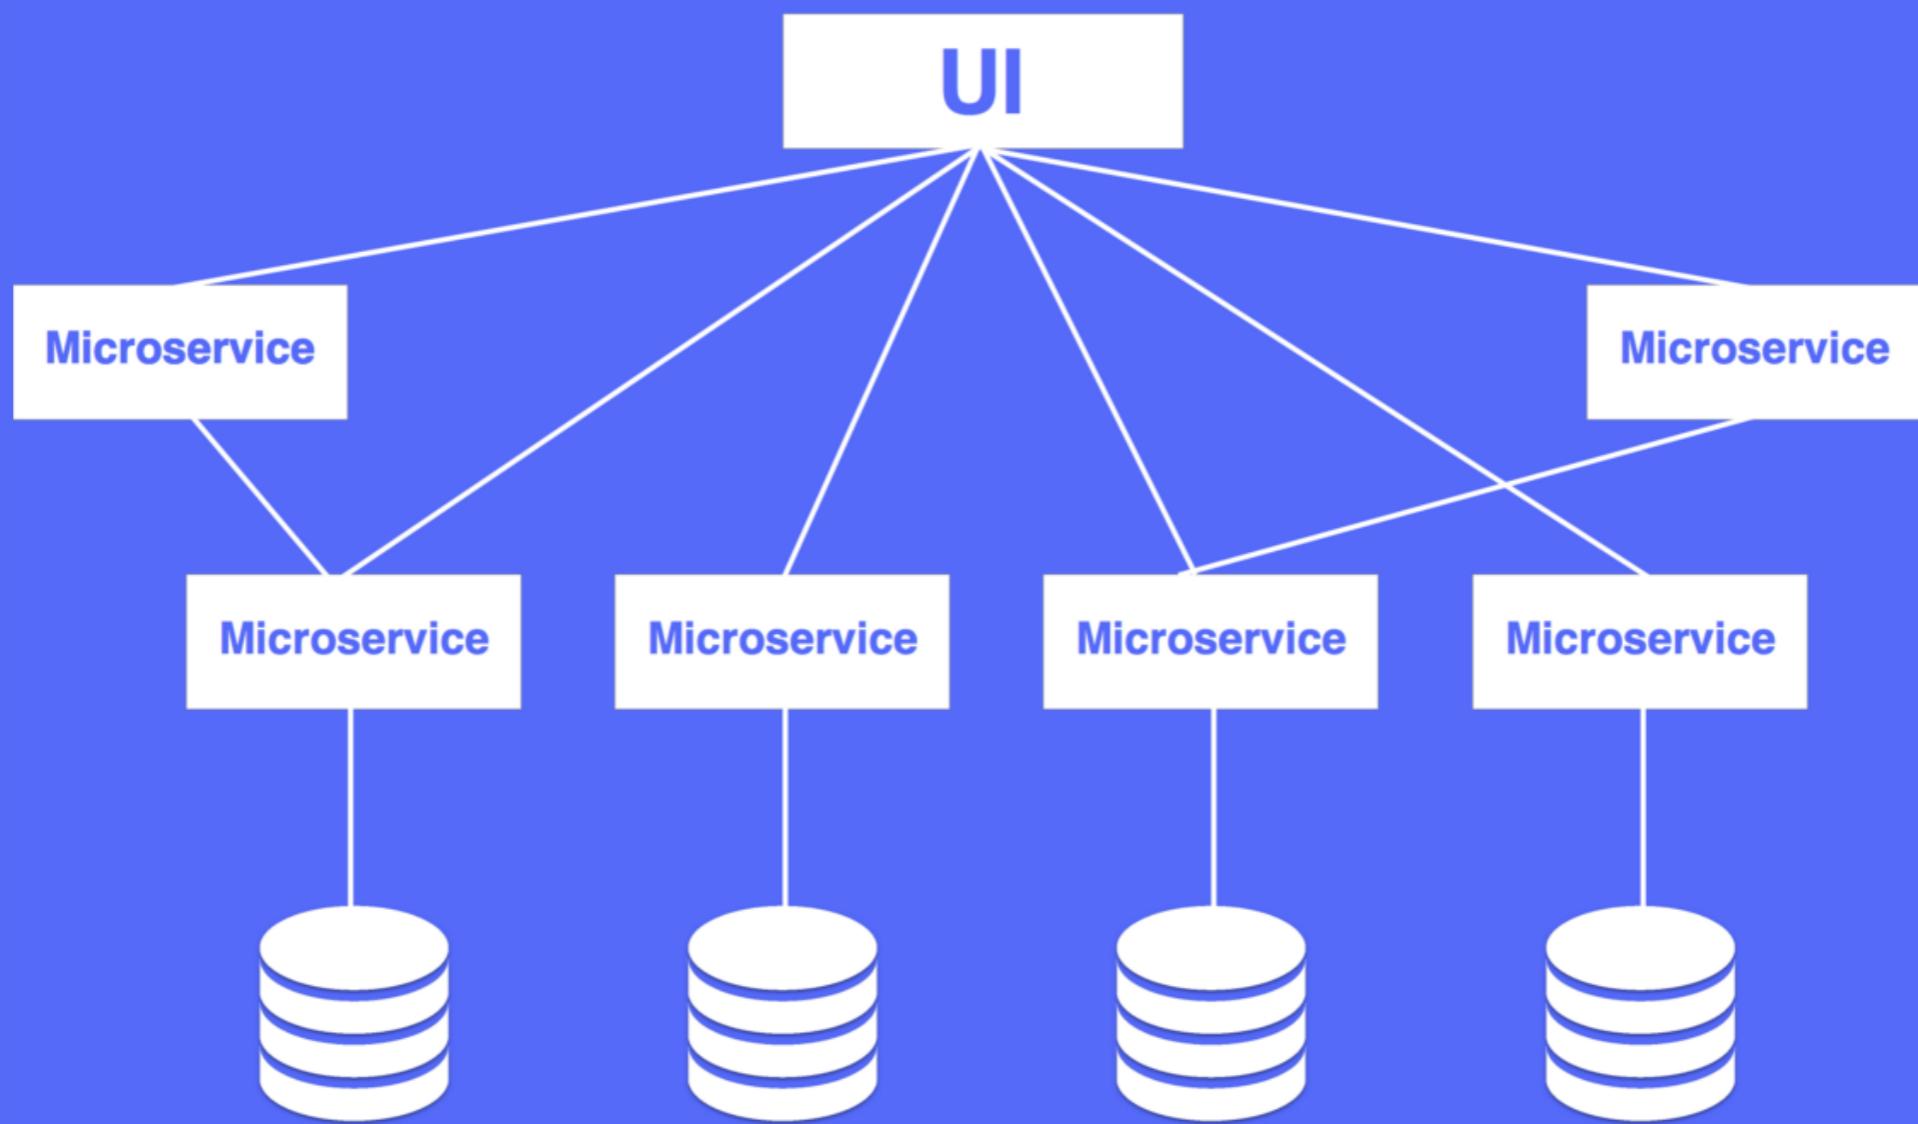




WHAT?

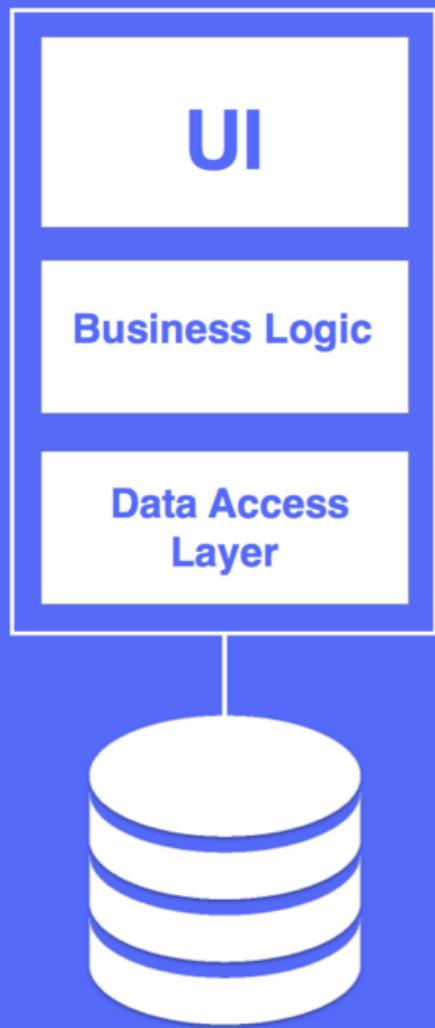




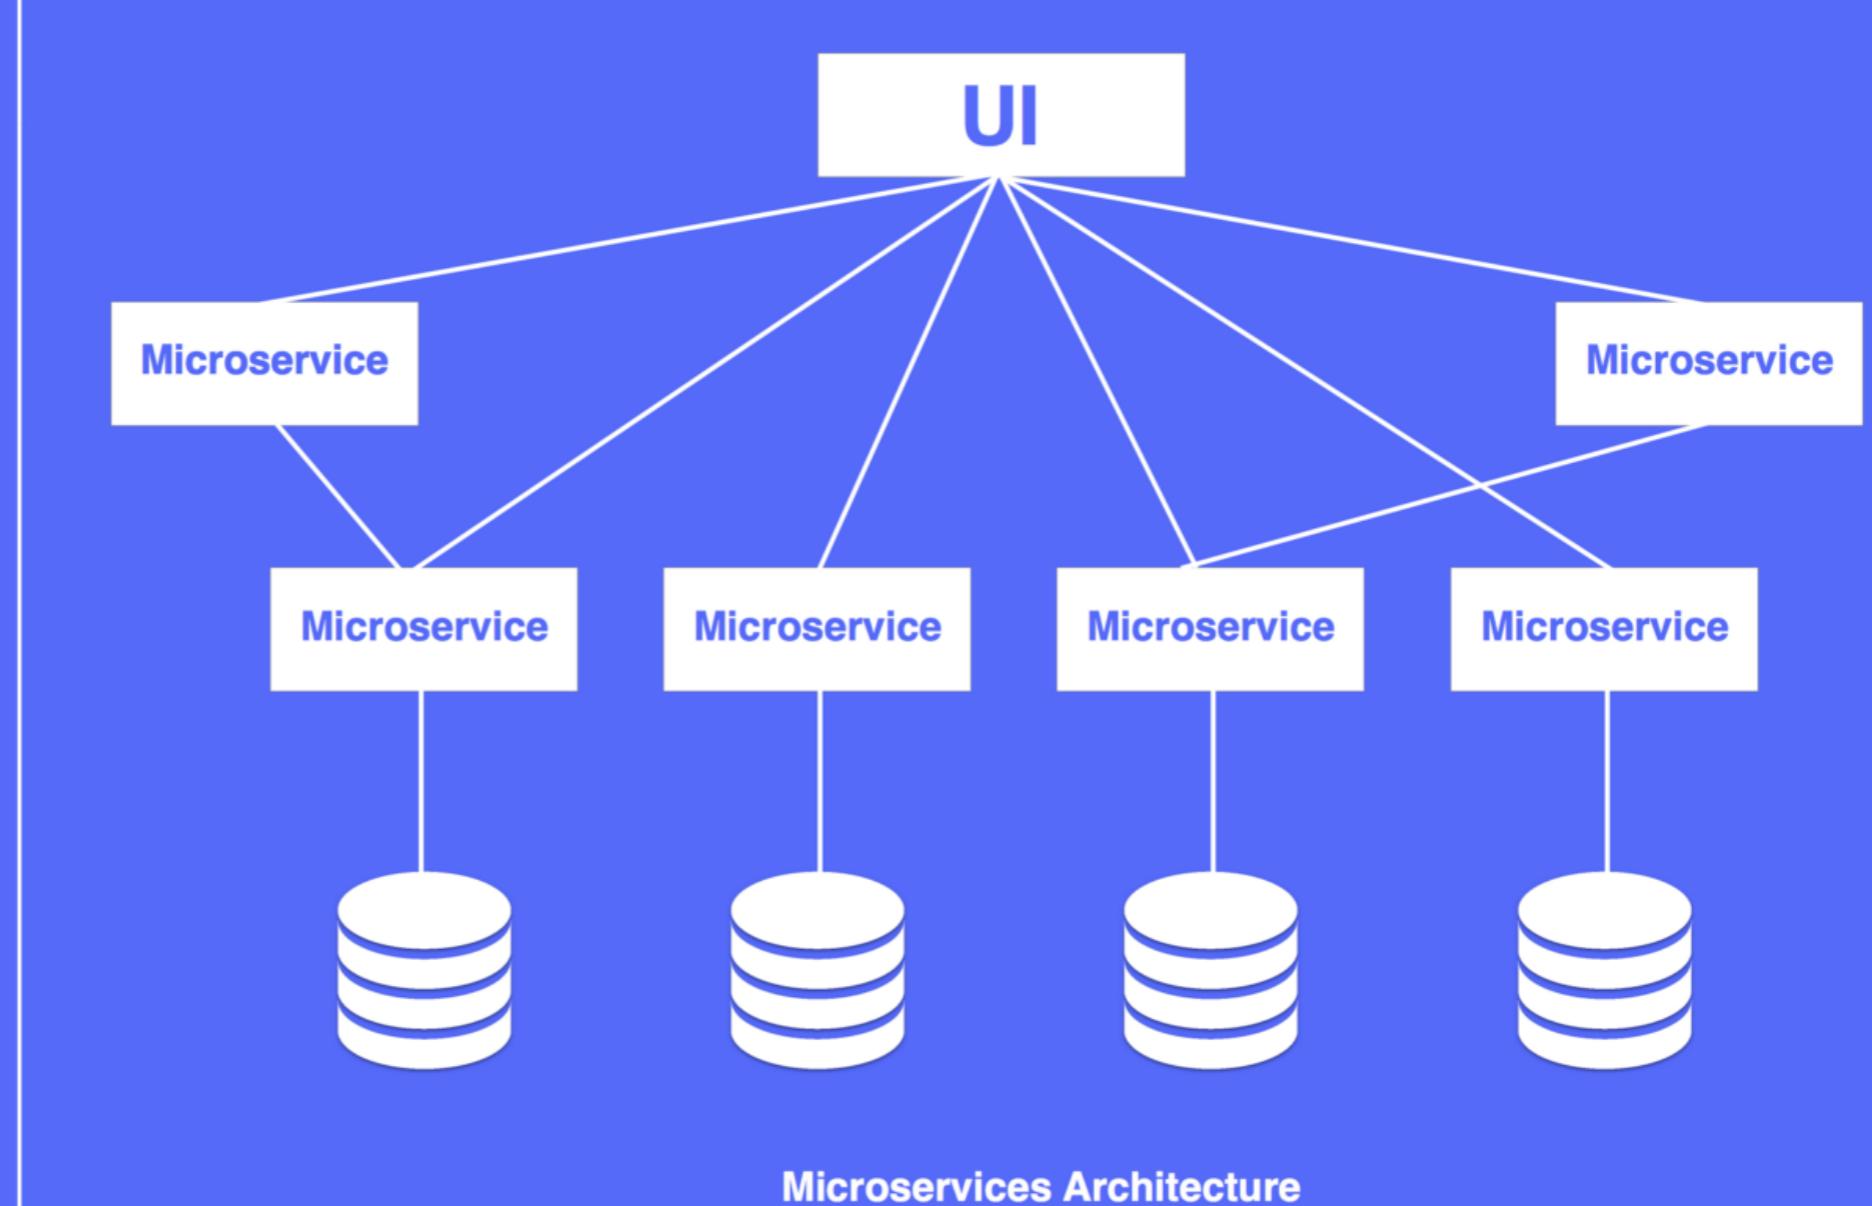


Microservices Architecture





Monolithic Architecture



Microservices Architecture

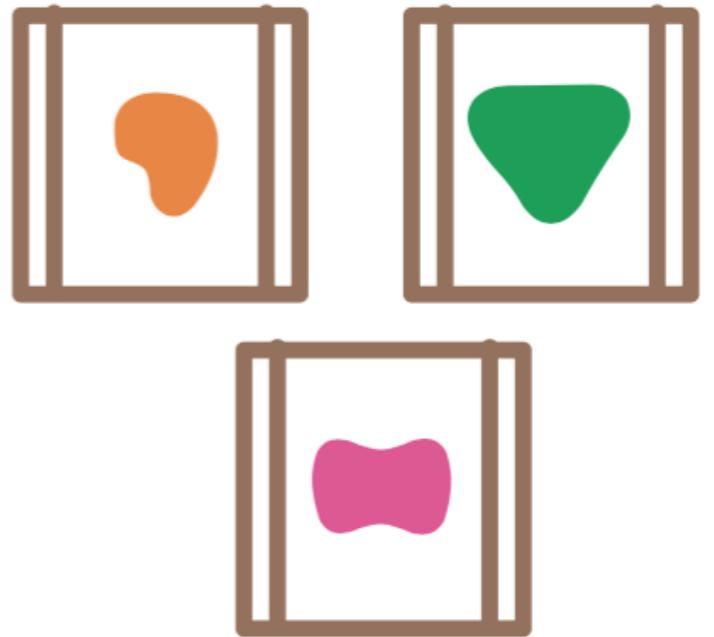
*The microservice architecture
is an architectural style that
structures an application as a
set of loosely coupled, services
organized around business
capabilities*

Simple and Lightweight

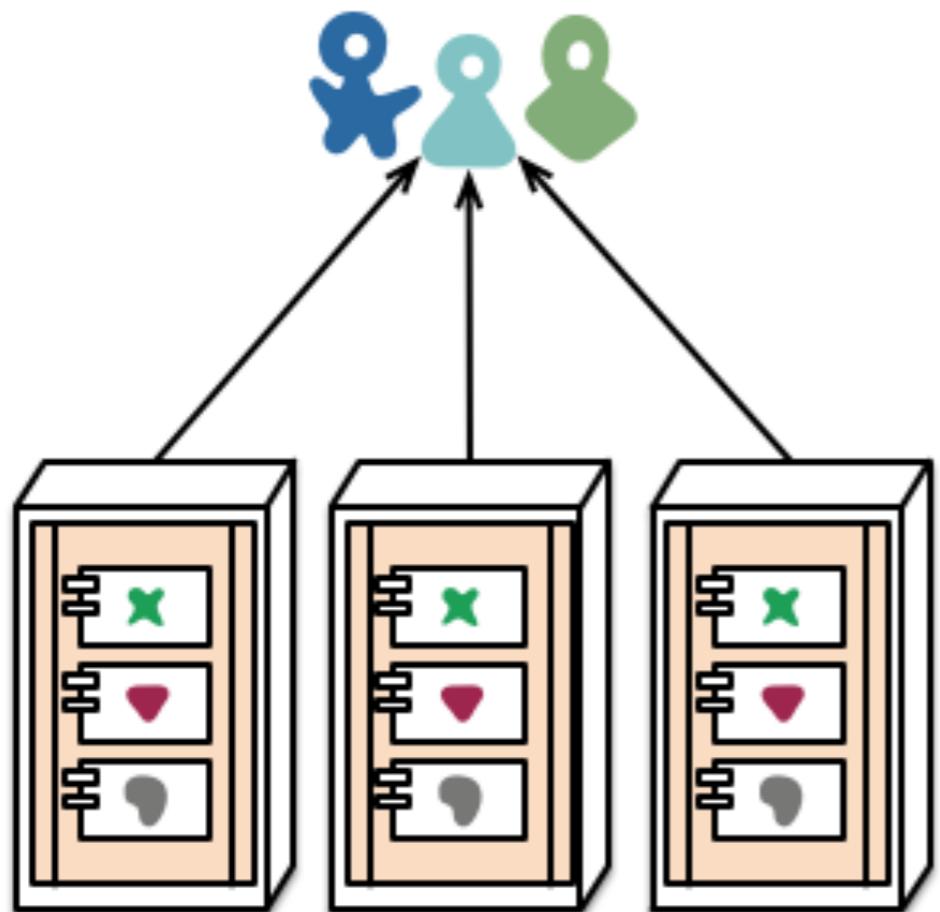
A monolithic application puts all its functionality into a single process...



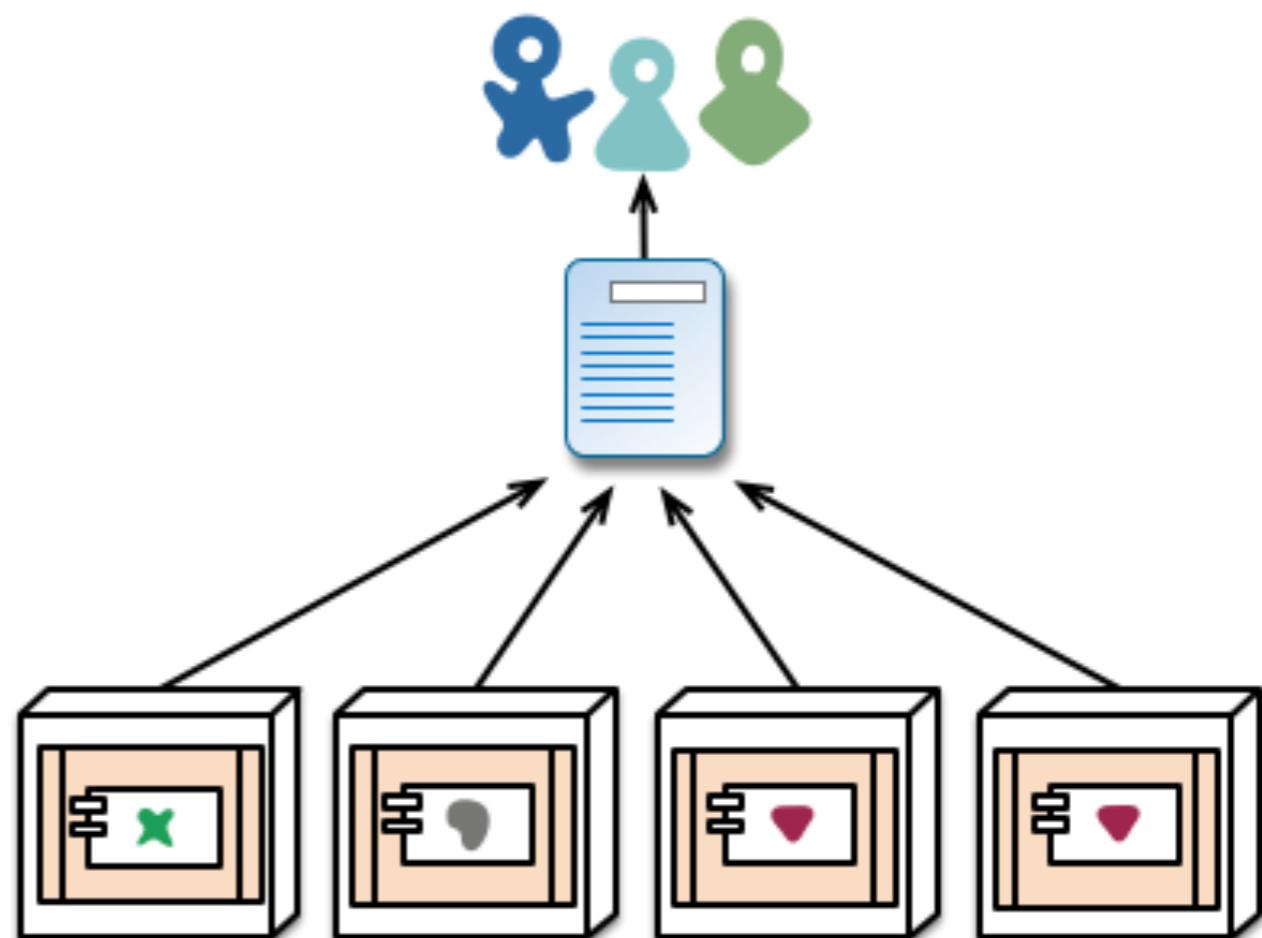
A microservices architecture puts each element of functionality into a separate service...



Independent Processes



monolith - multiple modules in the same process

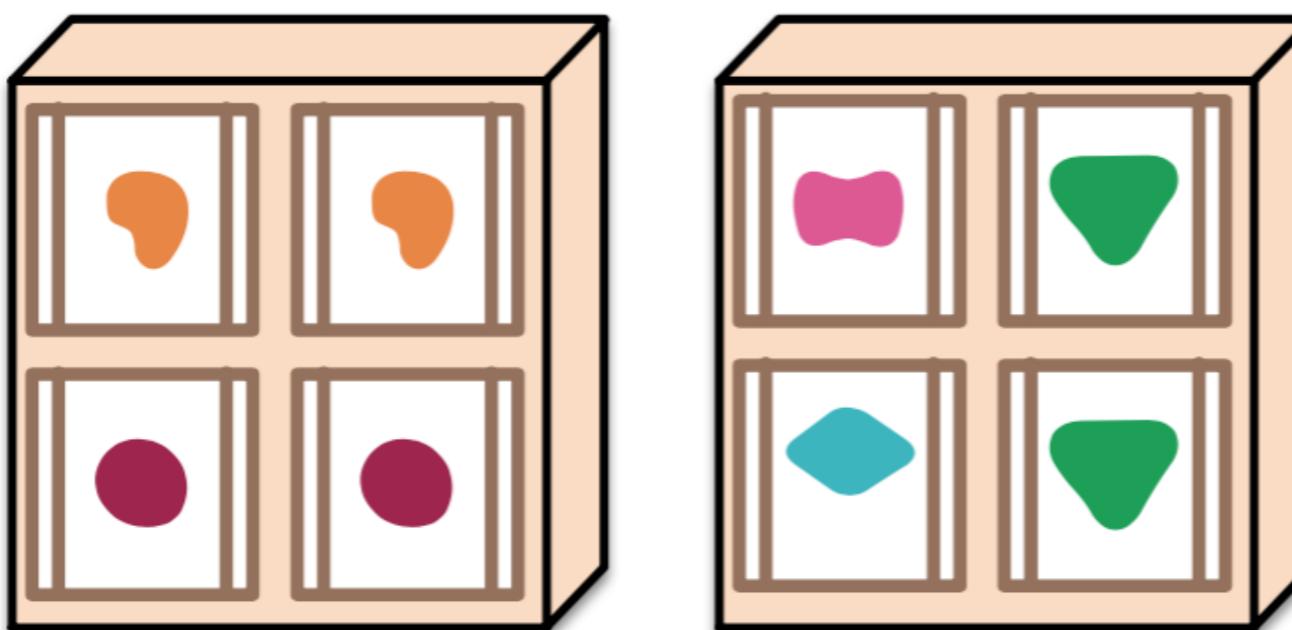
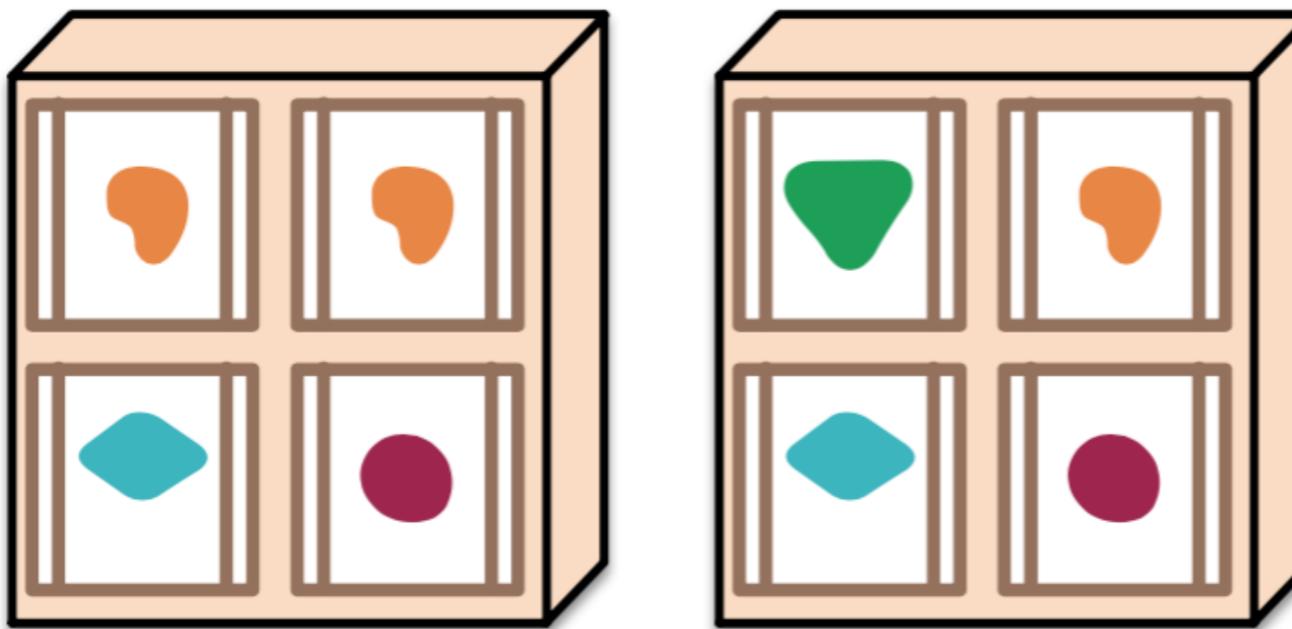


microservices - modules running in different processes

Language Agnostic

“be of the
web”

Decoupled



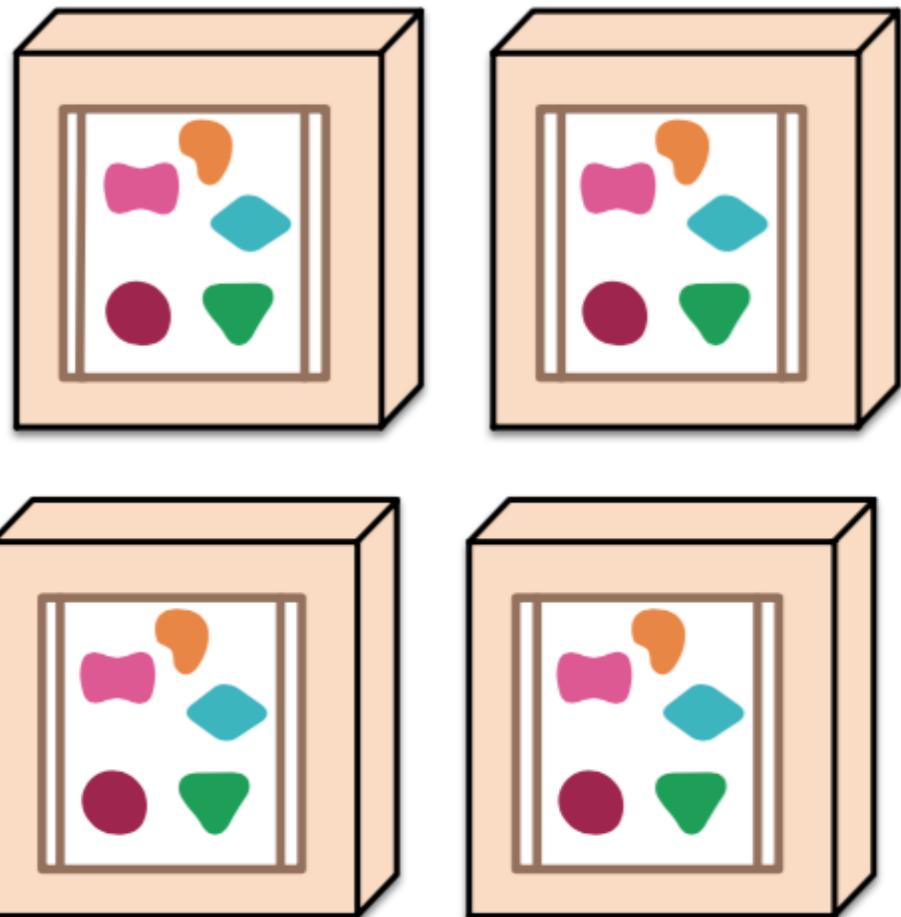
Scaling



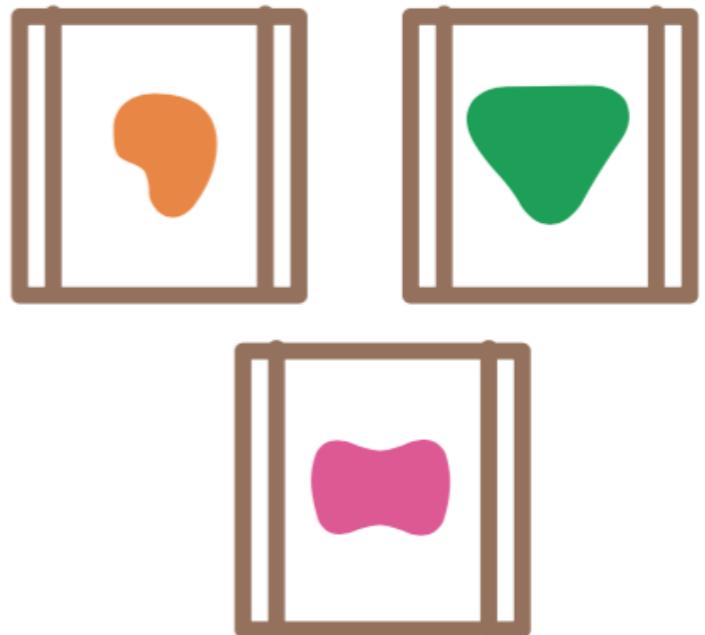
A monolithic application puts all its functionality into a single process...



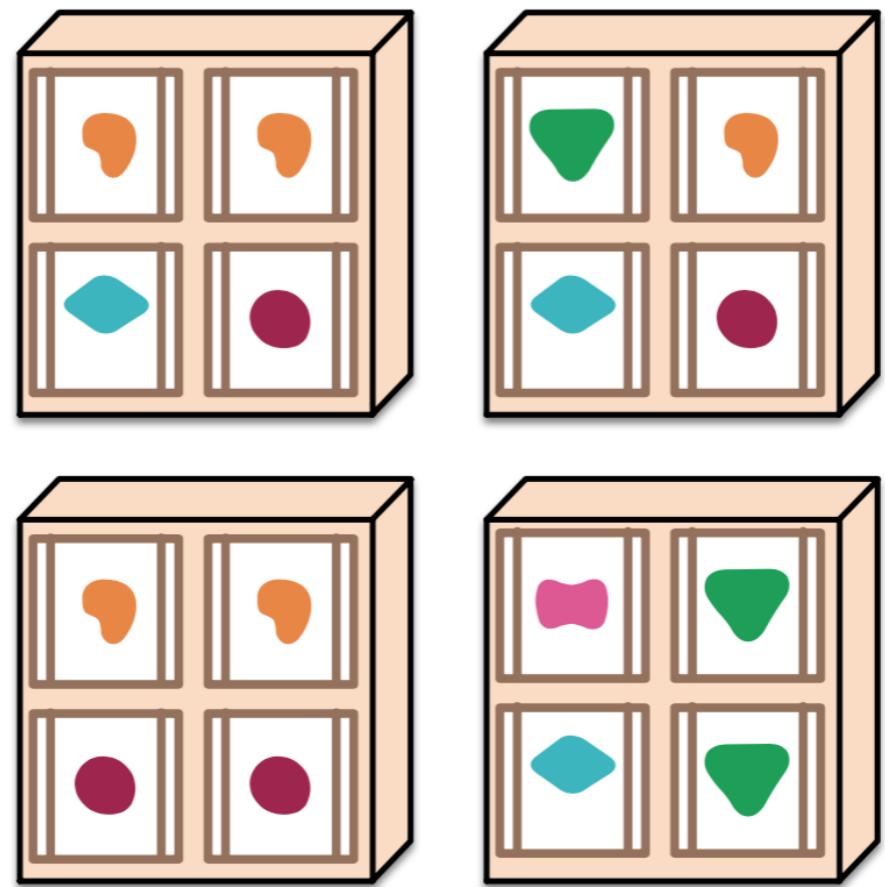
... and scales by replicating the monolith on multiple servers



A microservices architecture puts each element of functionality into a separate service...



... and scales by distributing these services across servers, replicating as needed.



Two levels of architecture

Application-level

Services

Inter-service glue: interfaces and communication mechanisms

Slow changing

Service-level

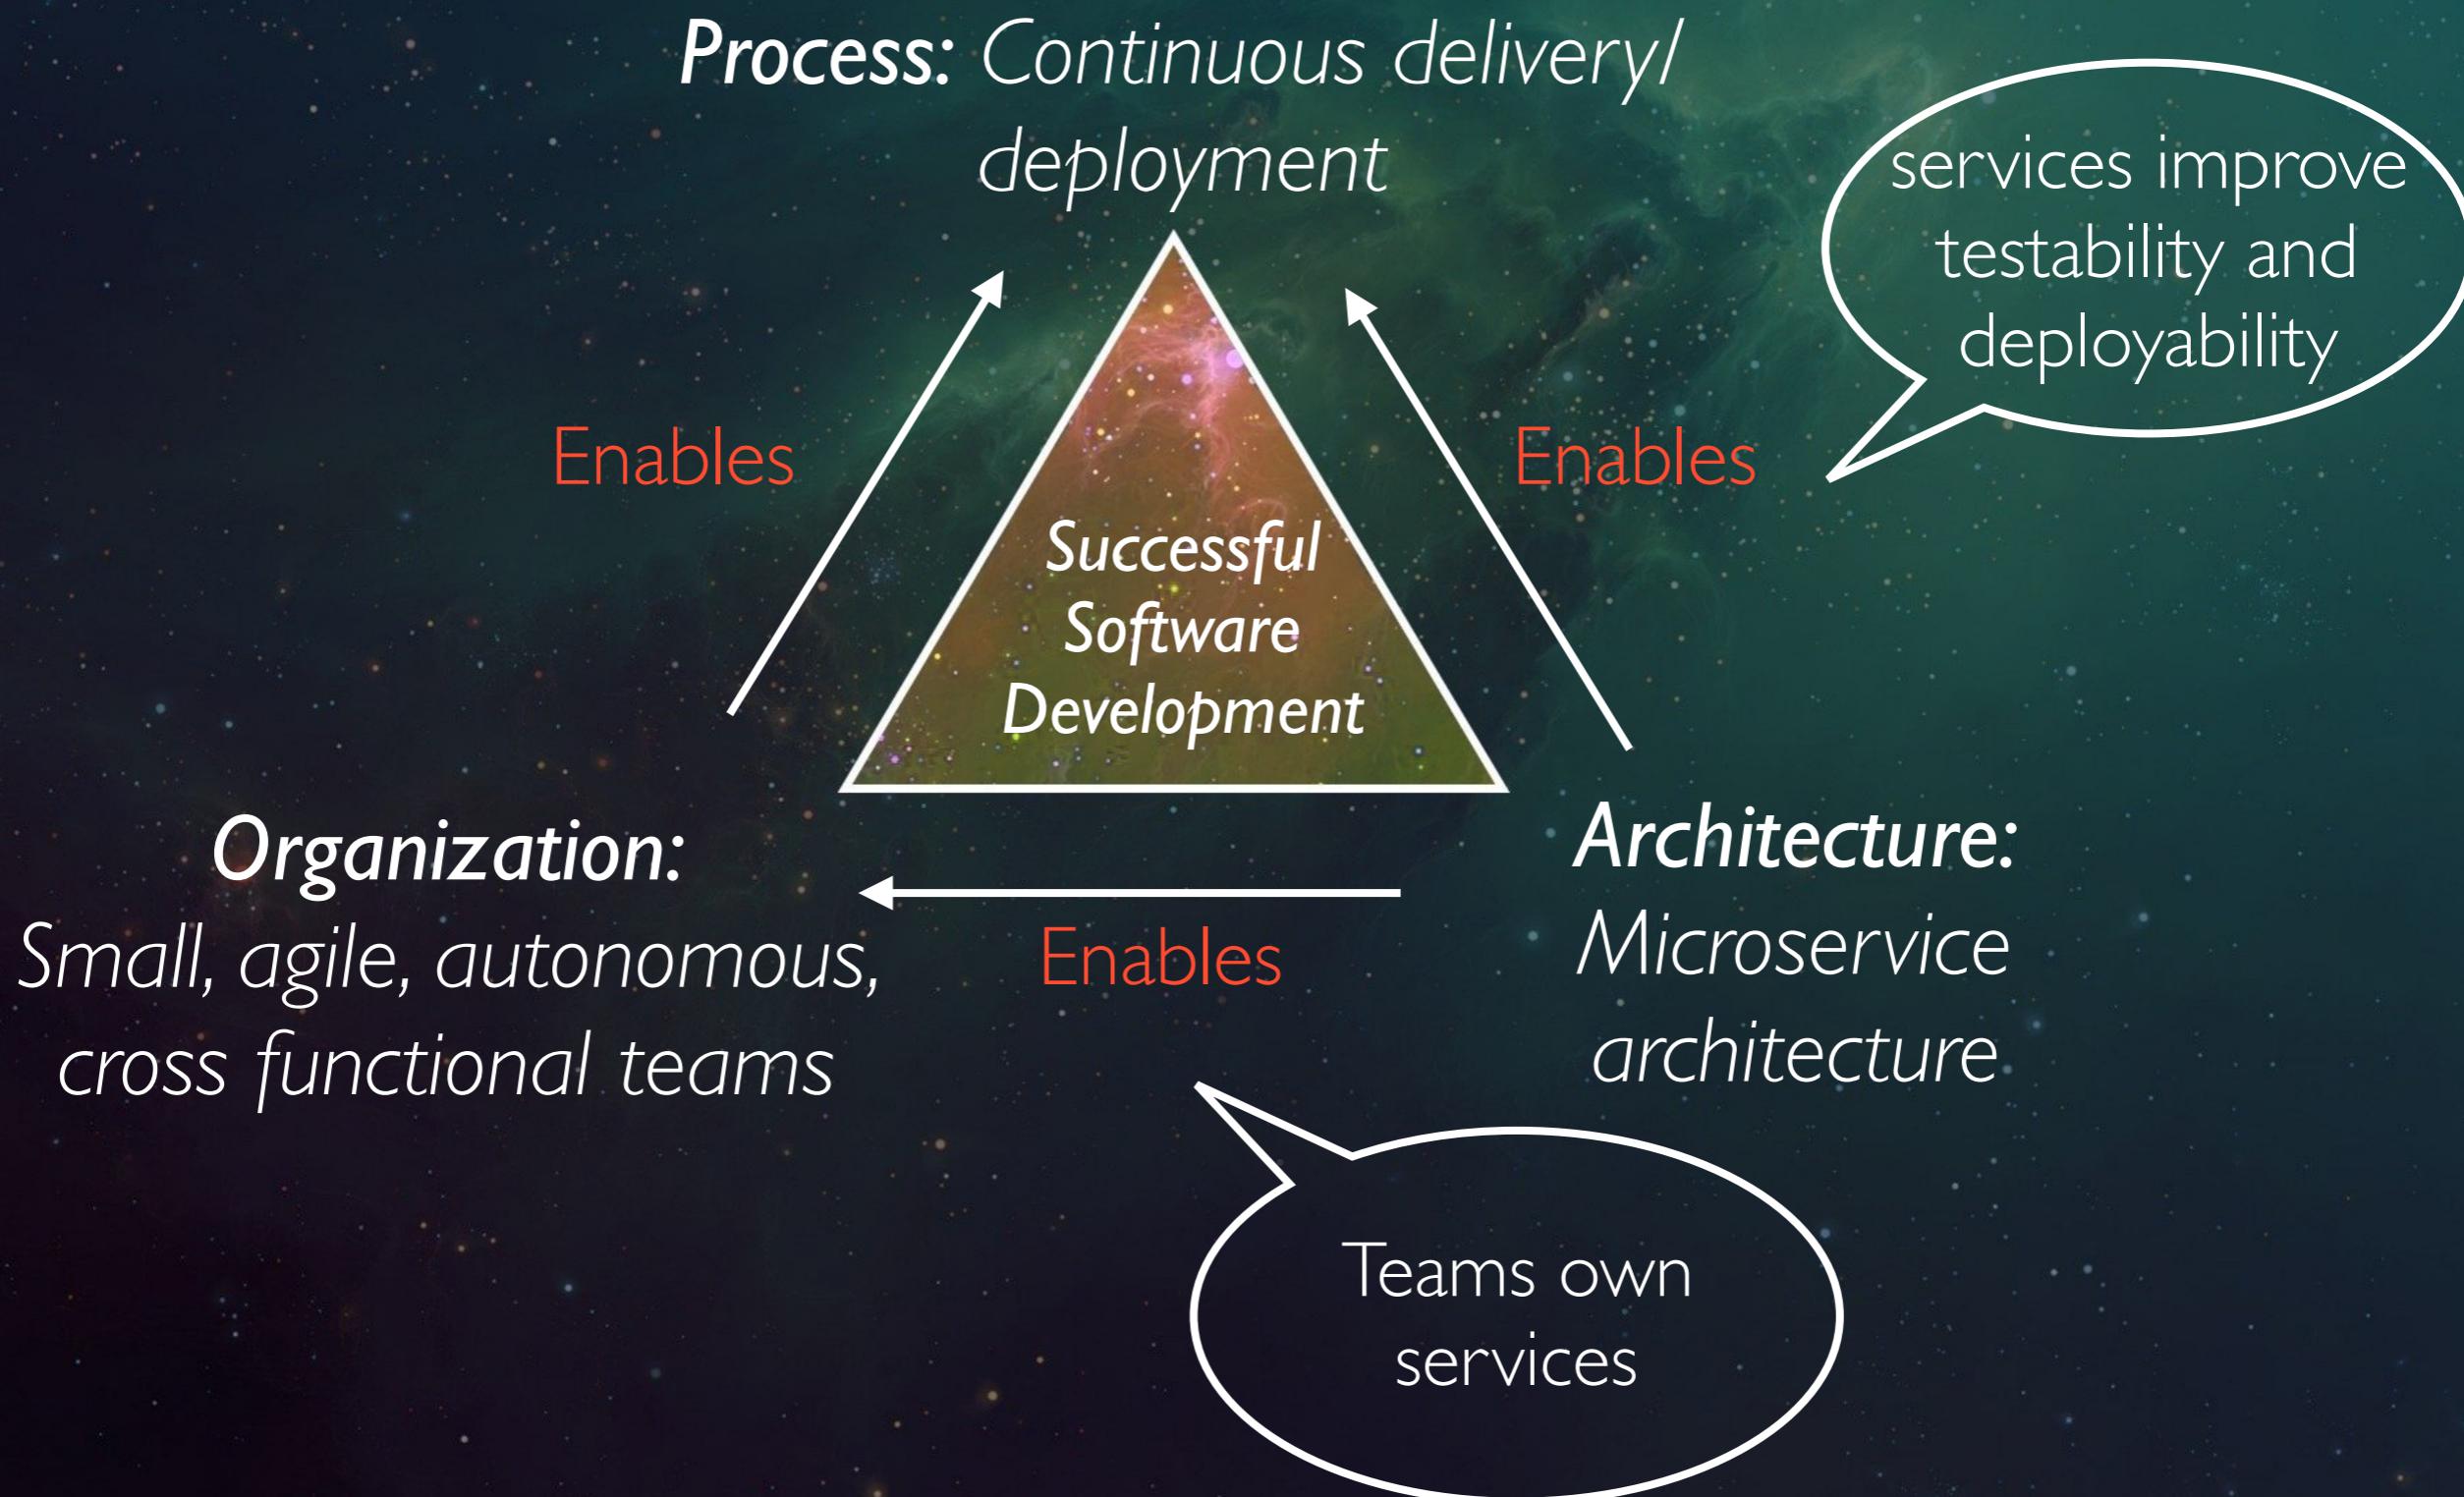
Internal architecture of each service

Each service could use a different technology stack

Pick the best tool for the job

Rapidly evolving

Microservices enable
continuous delivery/deployment



Evolve the technology stack

- ▶ Pick a new technology when
 - Writing a new service
 - Making major changes to an existing service
- ▶ Let's you experiment and fail safely

Drawbacks of micro services

- ▶ Development:
 - Interprocess communication - IPC
 - Partial failure
 - distributed data
- ▶ Testing
- ▶ Deployment
- ▶ ...



- ▶ How to decompose an application into services?
- ▶ How to deploy an application's services?
- ▶ How to handle cross cutting concerns?
- ▶ Which communication mechanisms to use?
- ▶ How do external clients communicate with the services?
- ▶ How does a client discover the network location of a service instance?

- ▶ How to prevent a network or service failure from cascading to other services?
- ▶ How to maintain data consistency and implement queries?
- ▶ How to understand the behavior of an application and troubleshoot problems?
- ▶ How to make testing easier?
- ▶ How to implement a UI screen or page that displays data from multiple services?

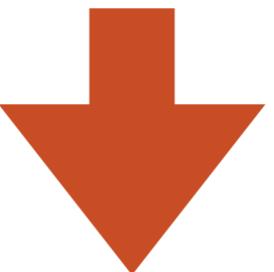


kubernetes

Manage applications, not machines

- ▶ Open source, open API container orchestrator
- ▶ Supports multiple cloud and **bare-metal** environments
- ▶ Inspired and informed by Google's experiences and internal systems

What we need is a system to handle that complexity for us



without locking us into any one vendor or way of doing things

Cattle vs Pets



- ▶ Has a name
- ▶ Is unique or rare
- ▶ Personal attention
- ▶ If it gets ill, you make it better
- ▶ Has a number
- ▶ One is much like any other
- ▶ Run as a group
- ▶ If it gets ill, you make hamburgers



kubernetes

Desired State

One of the core concepts of Kubernetes



Tell kubernetes what you want, not what to do!

Desired States

`$./create_docker_images.sh`

`$./lauch_frontend.sh` **x 3**

`$./launch_services.sh` **x 2**

`$./lauch_backend.sh` **x 1**



► Under and imperative system, you have a series of tasks

- Create images
- Create 3 frontends
- Create 2 services, etc

► If something dies, something has to react

- Under your worst case it's an admin that has to react (ok, maybe you have something automated)

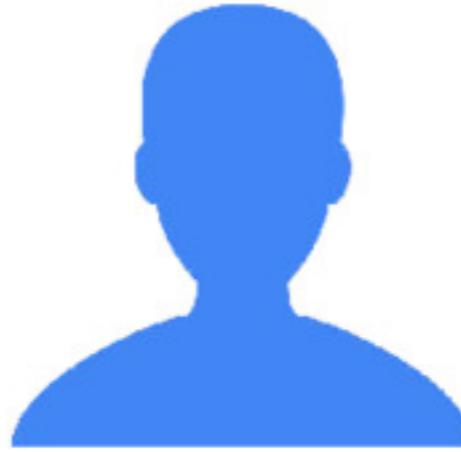


Under a desired state you just say
I want 3 2 |

*Then if something blows up, you're no
longer in the desired state, so
kubernetes will fix it!*

Employees, not Children

Children vs Employees



- ▶ Go upstairs
- ▶ Get undressed
- ▶ Put on pijamas
- ▶ Brush your teeth
- ▶ Pick out 2 stories
- ▶ go get some sleep



kubernetes

External Services

Services

Rolling Update

Pods

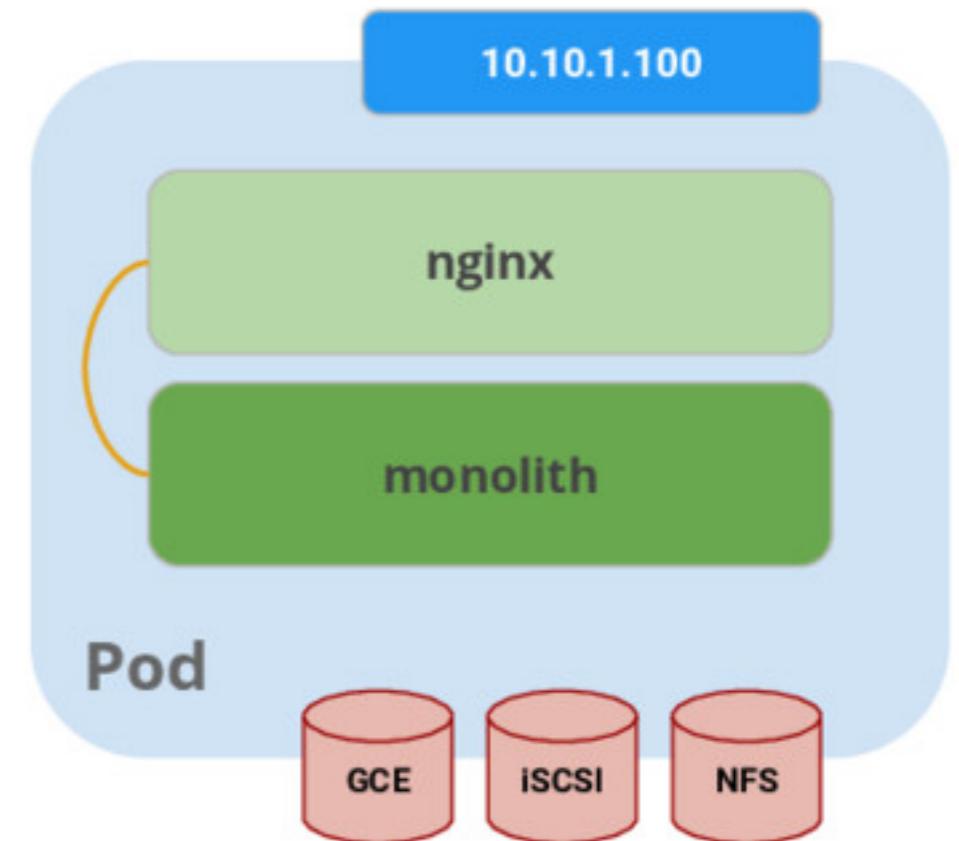
ReplicaSets

Labels & Selectors

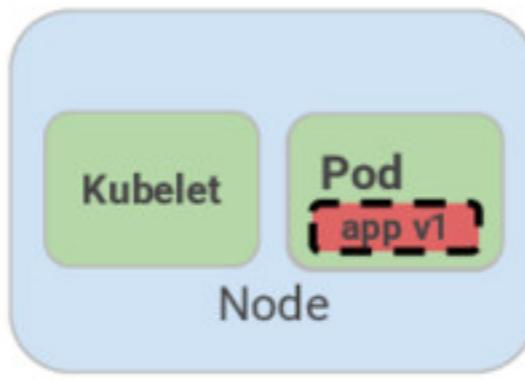
Monitor & Health Checks

Pods

- ▶ Small group of containers & volumes
- ▶ Tightly coupled
- ▶ The atom of scheduling & placement
- ▶ Shared namespace
 - share IP address & localhost
 - share IPC, etc
- ▶ Managed lifecycle
 - bound to a node, restart in place
 - can die, cannot reborn with same ID

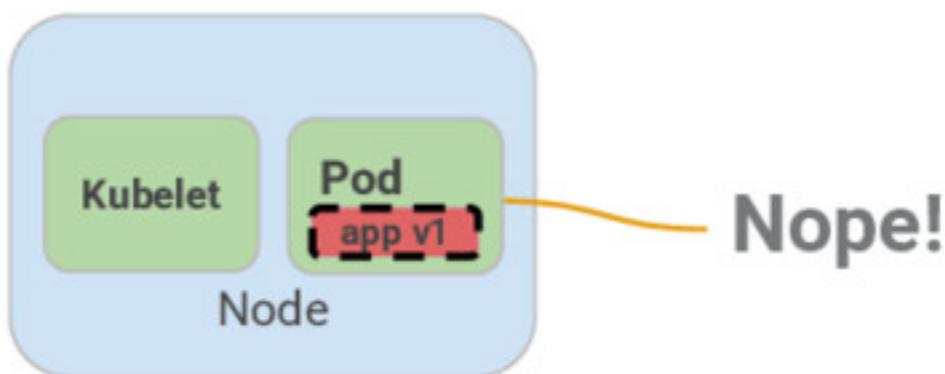
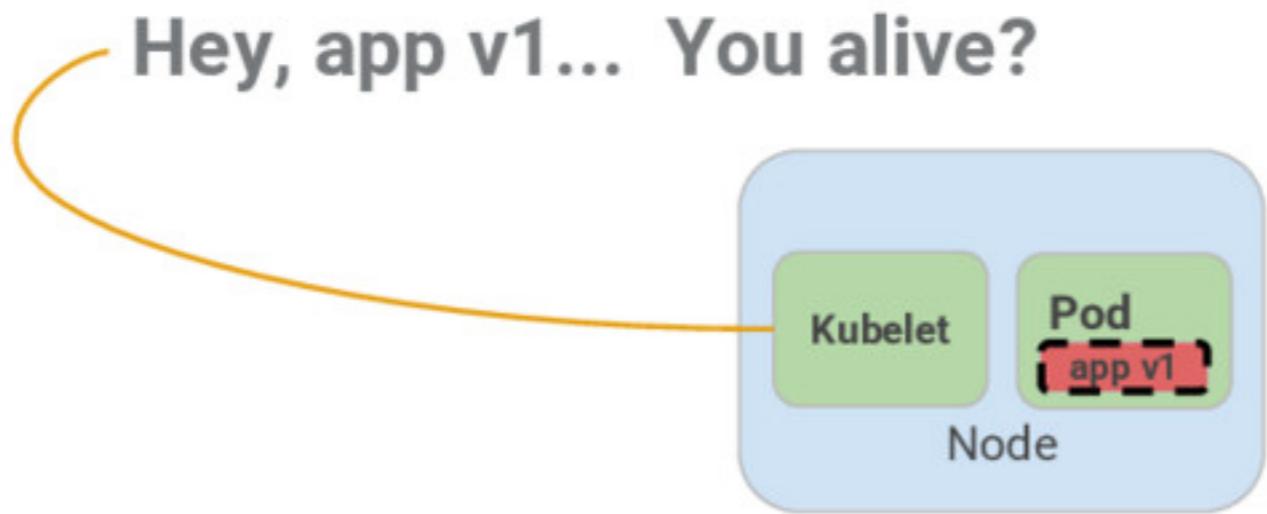


Monitoring and Health Checks



- ▶ On every node is a daemon called a Kubelet
- ▶ One of the Kubelete's jobs is to ensure that pods are healthy

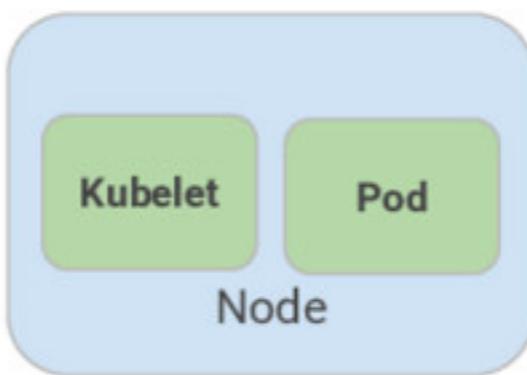
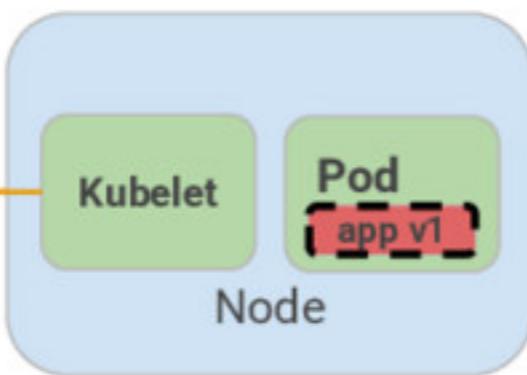
Monitoring and Health Checks



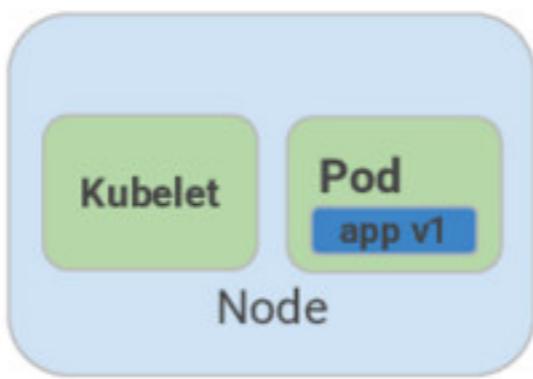
Nope!

Monitoring and Health Checks

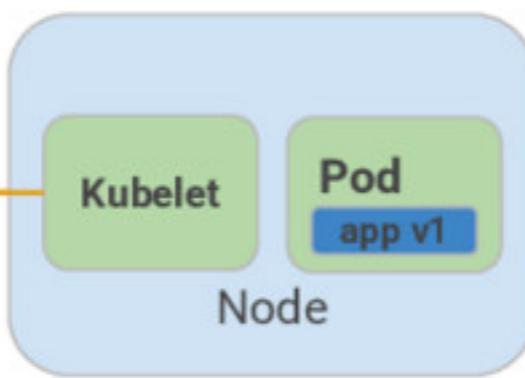
OK, then I'm going to restart you...



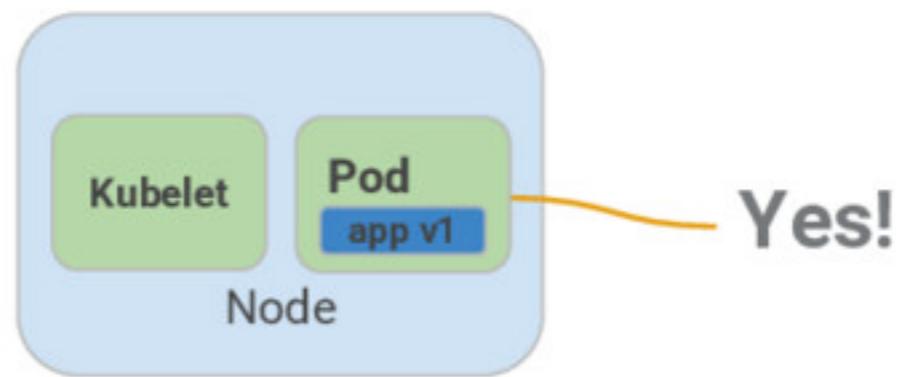
Monitoring and Health Checks



Hey, app v1... You alive?



Monitoring and Health Checks



Labels

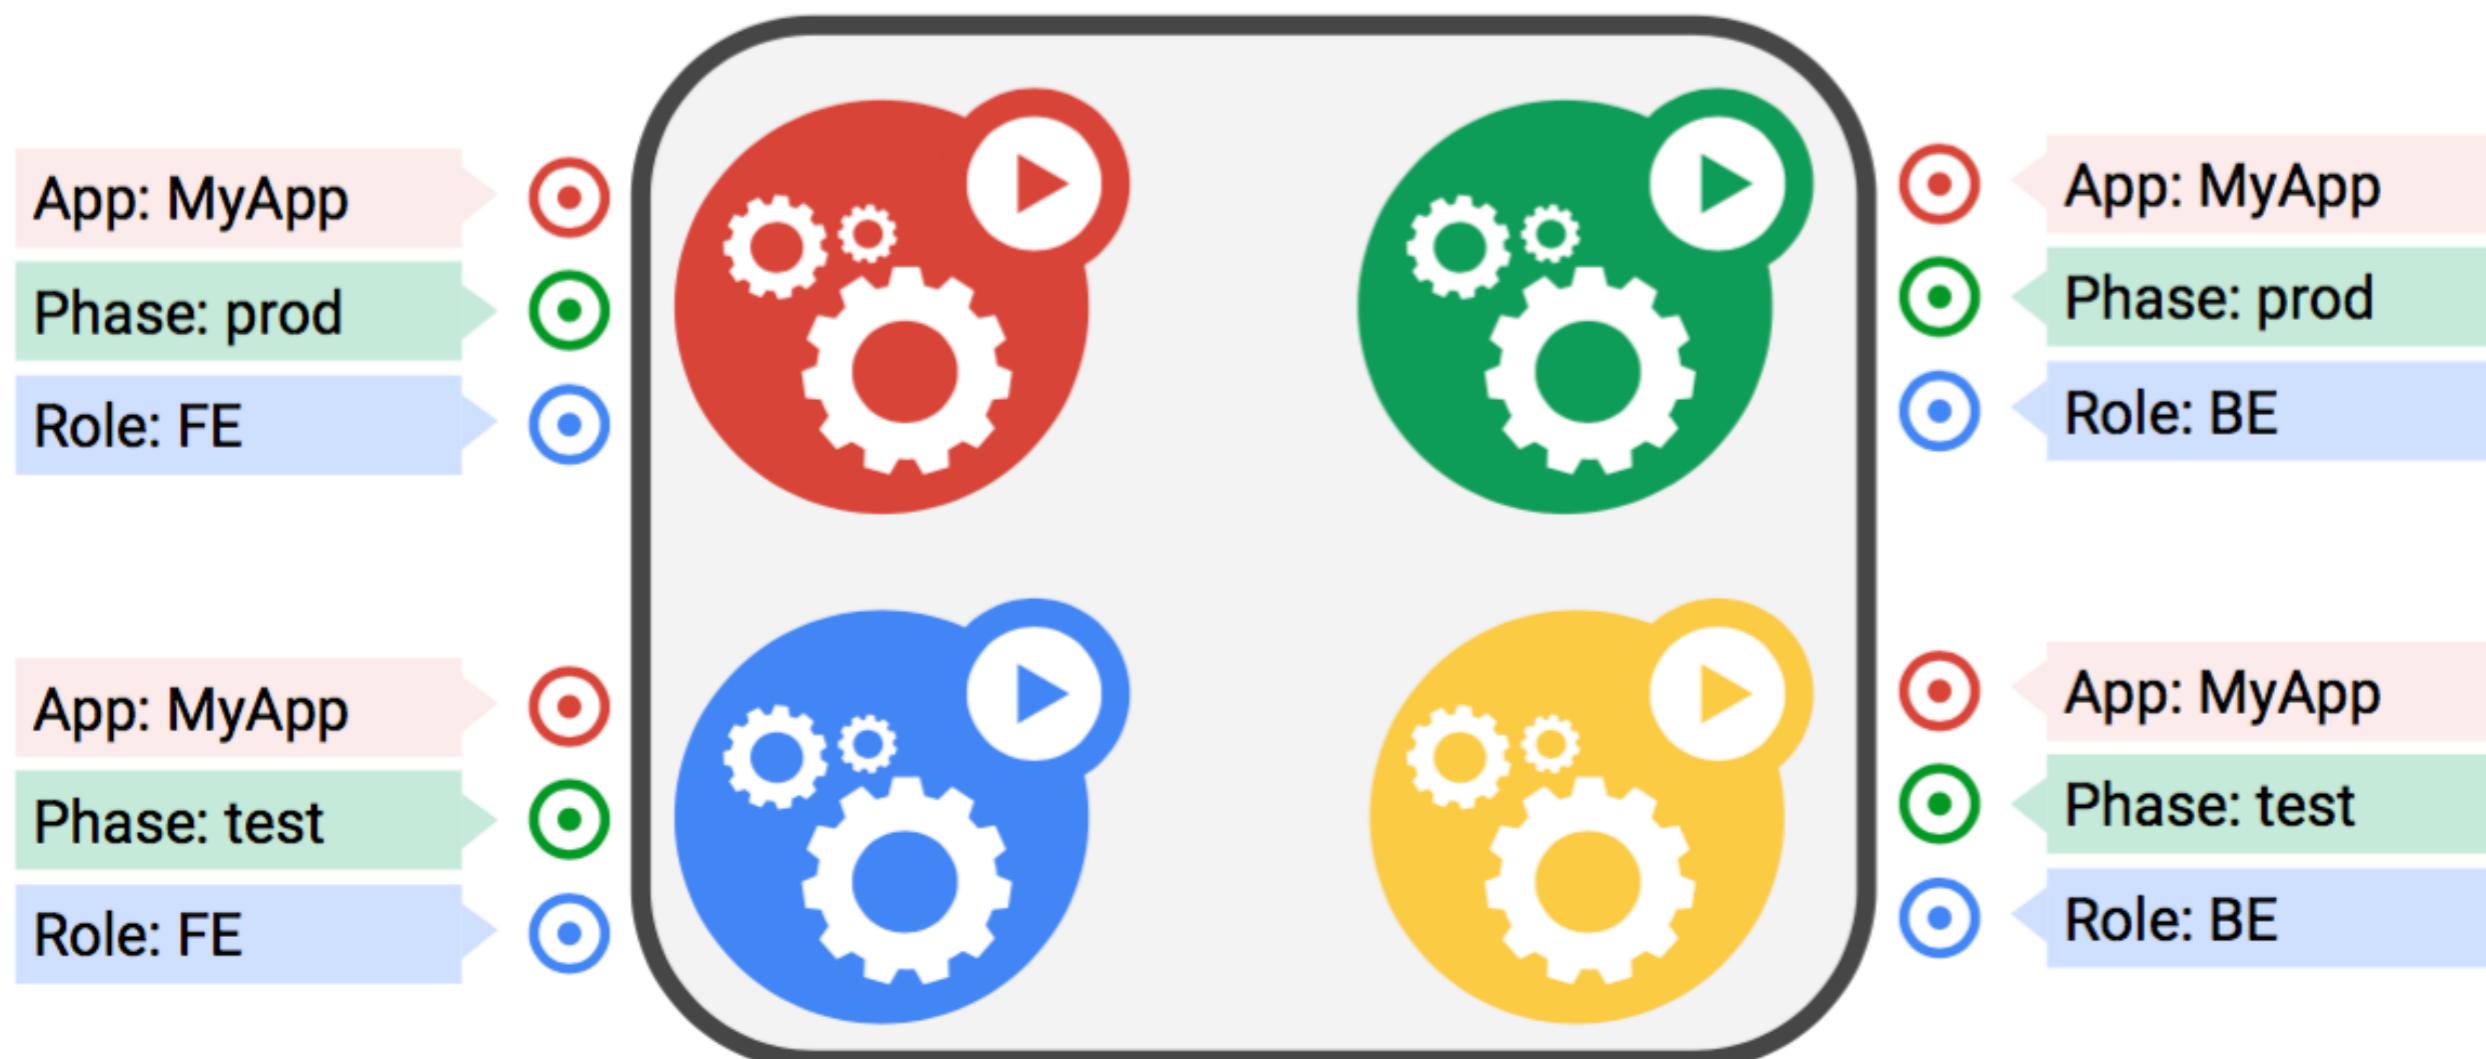
- ▶ Arbitrary metadata
- ▶ Attached to any API object
- ▶ Generally represent **identity**
- ▶ Queryable by **selectors**
 - think SQL 'select... where...'
- ▶ The **only** grouping mechanism
 - pods under a ReplicaSet
 - pods in a Service
 - capabilities of a node (constraints)



Selectors



Selectors



App = MyApp

Selectors



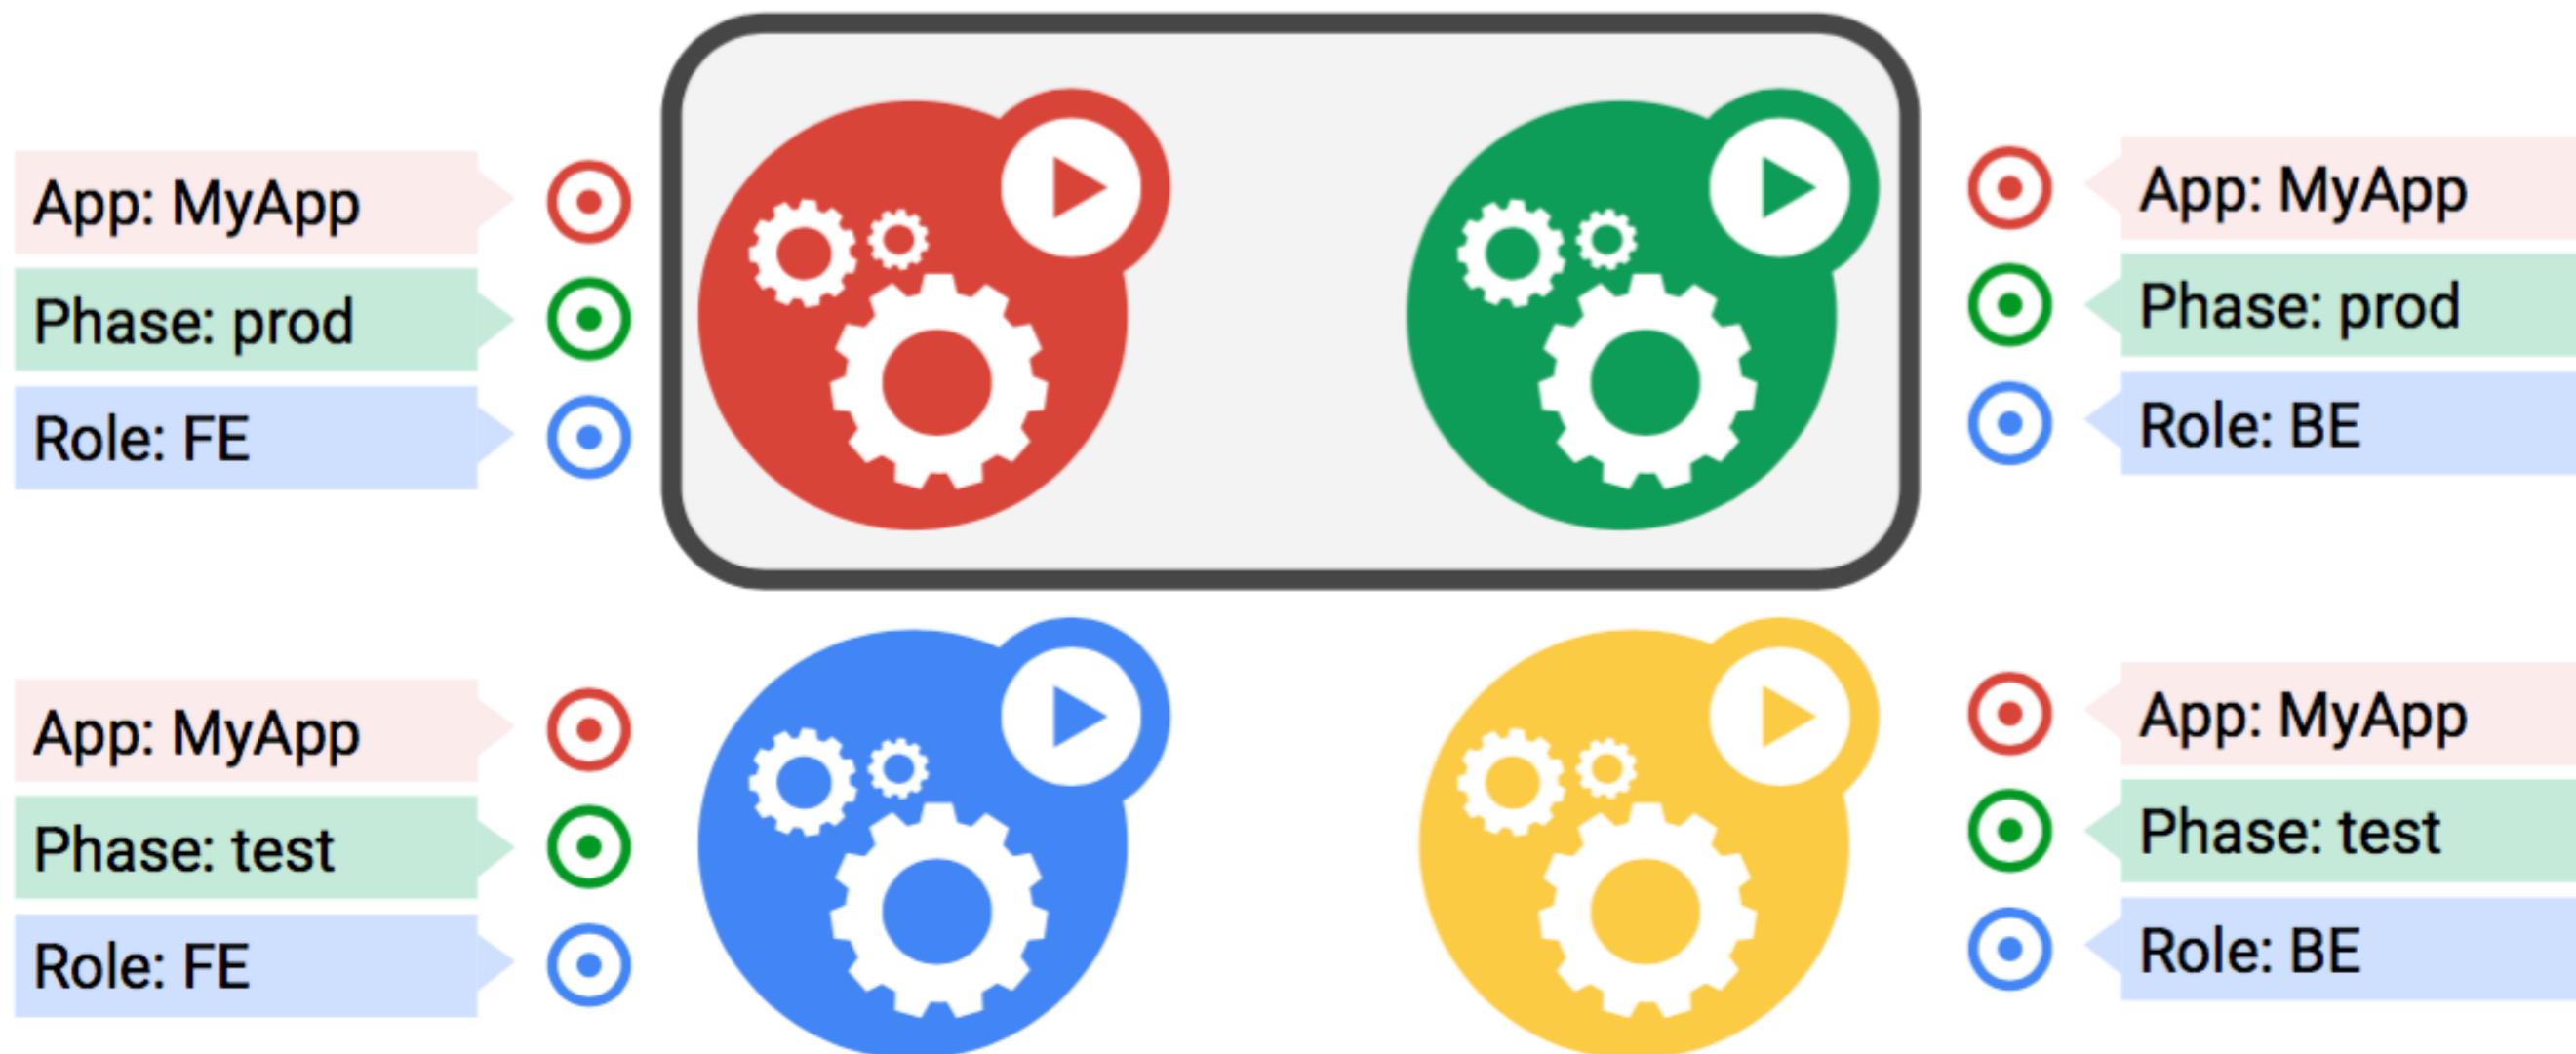
App = MyApp, Role = FE

Selectors



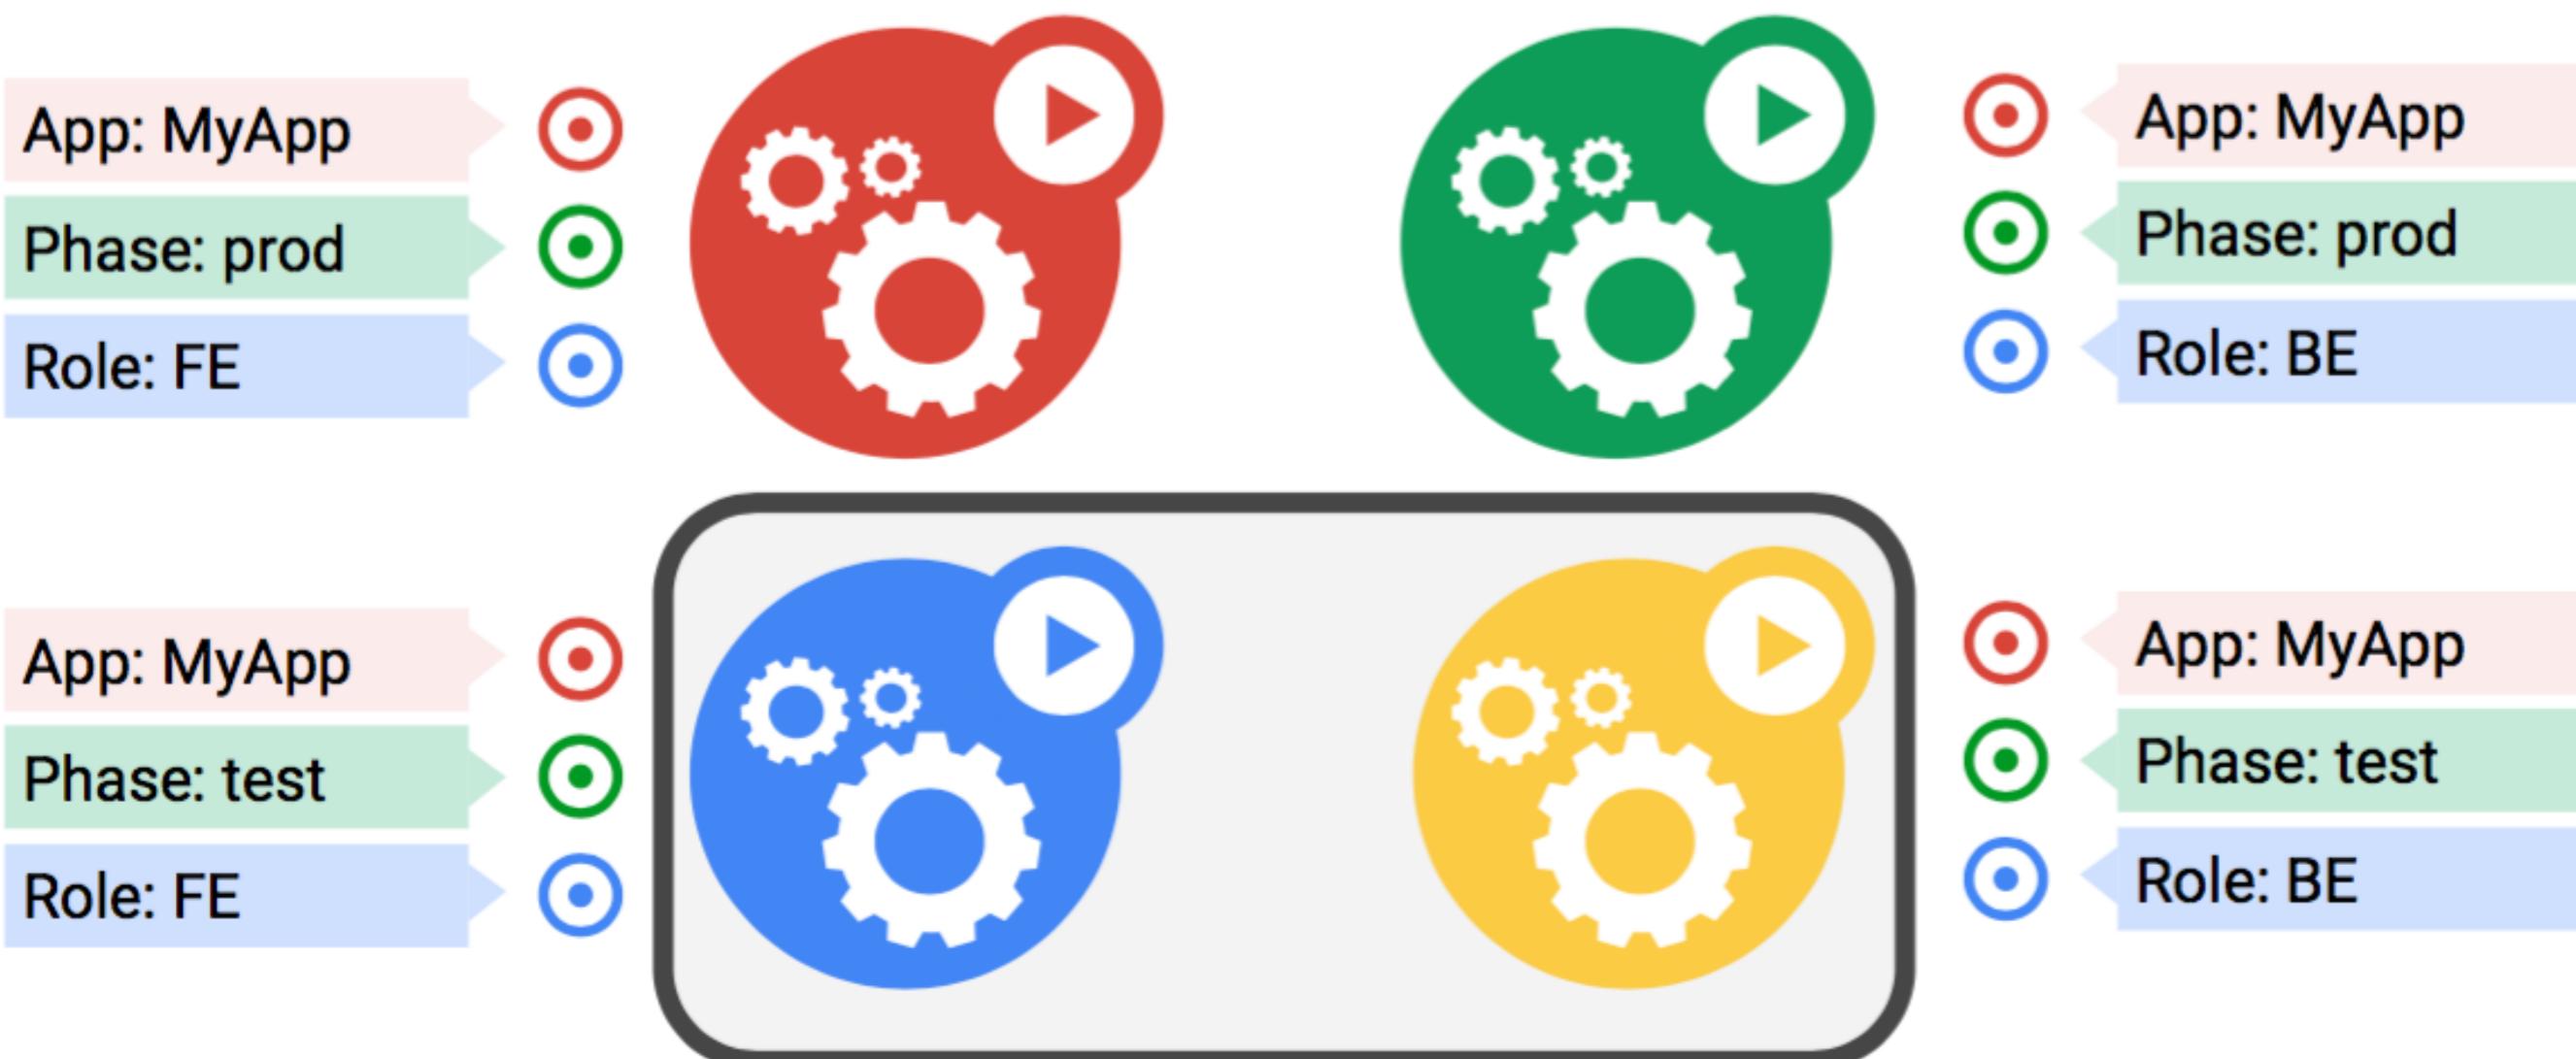
App = MyApp, Role = BE

Selectors



App = MyApp, Phase = prod

Selectors



App = MyApp, Phase = test

ReplicaSets

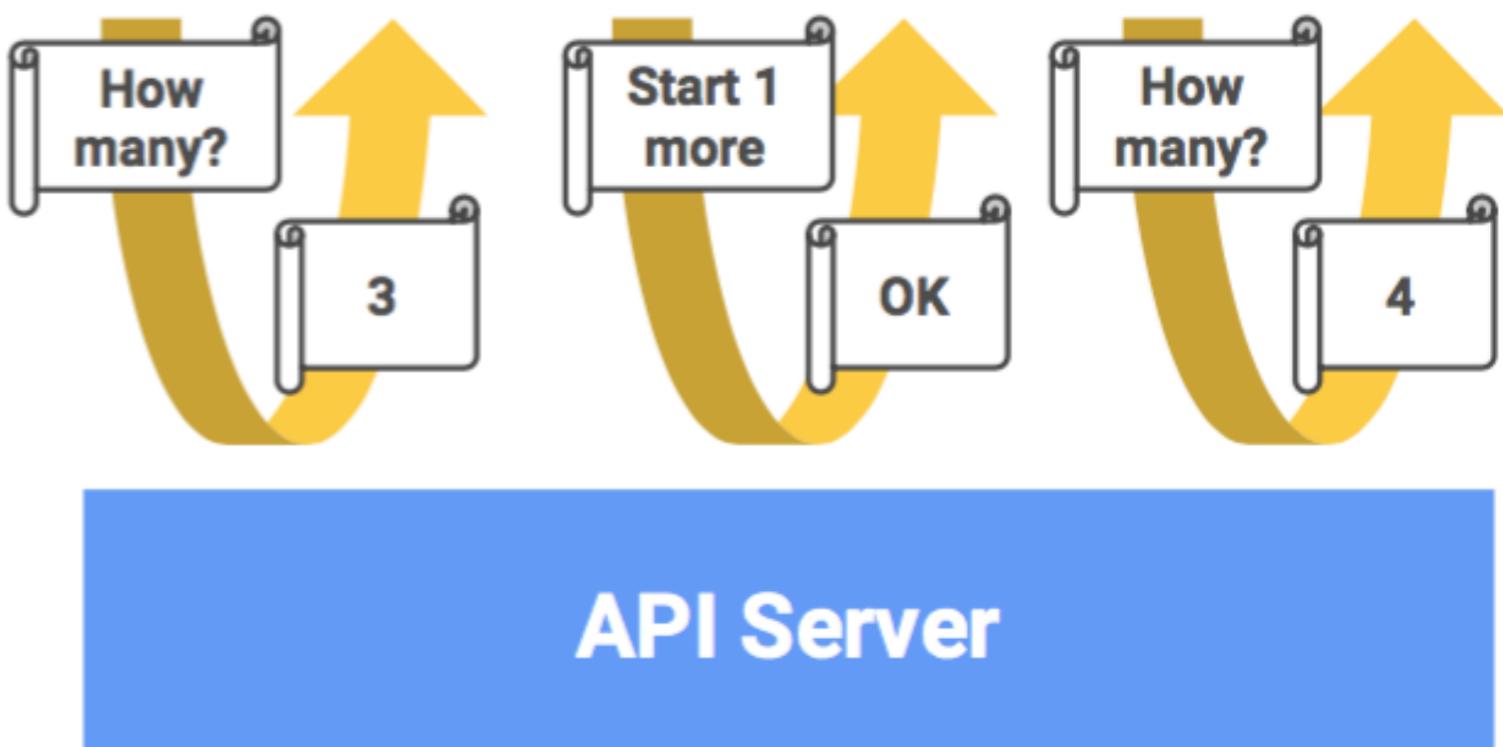
► A simple control loop

► One job: ensure N copies
of a pod

- grouped by a selector
- too few? start some
- too many? kill some

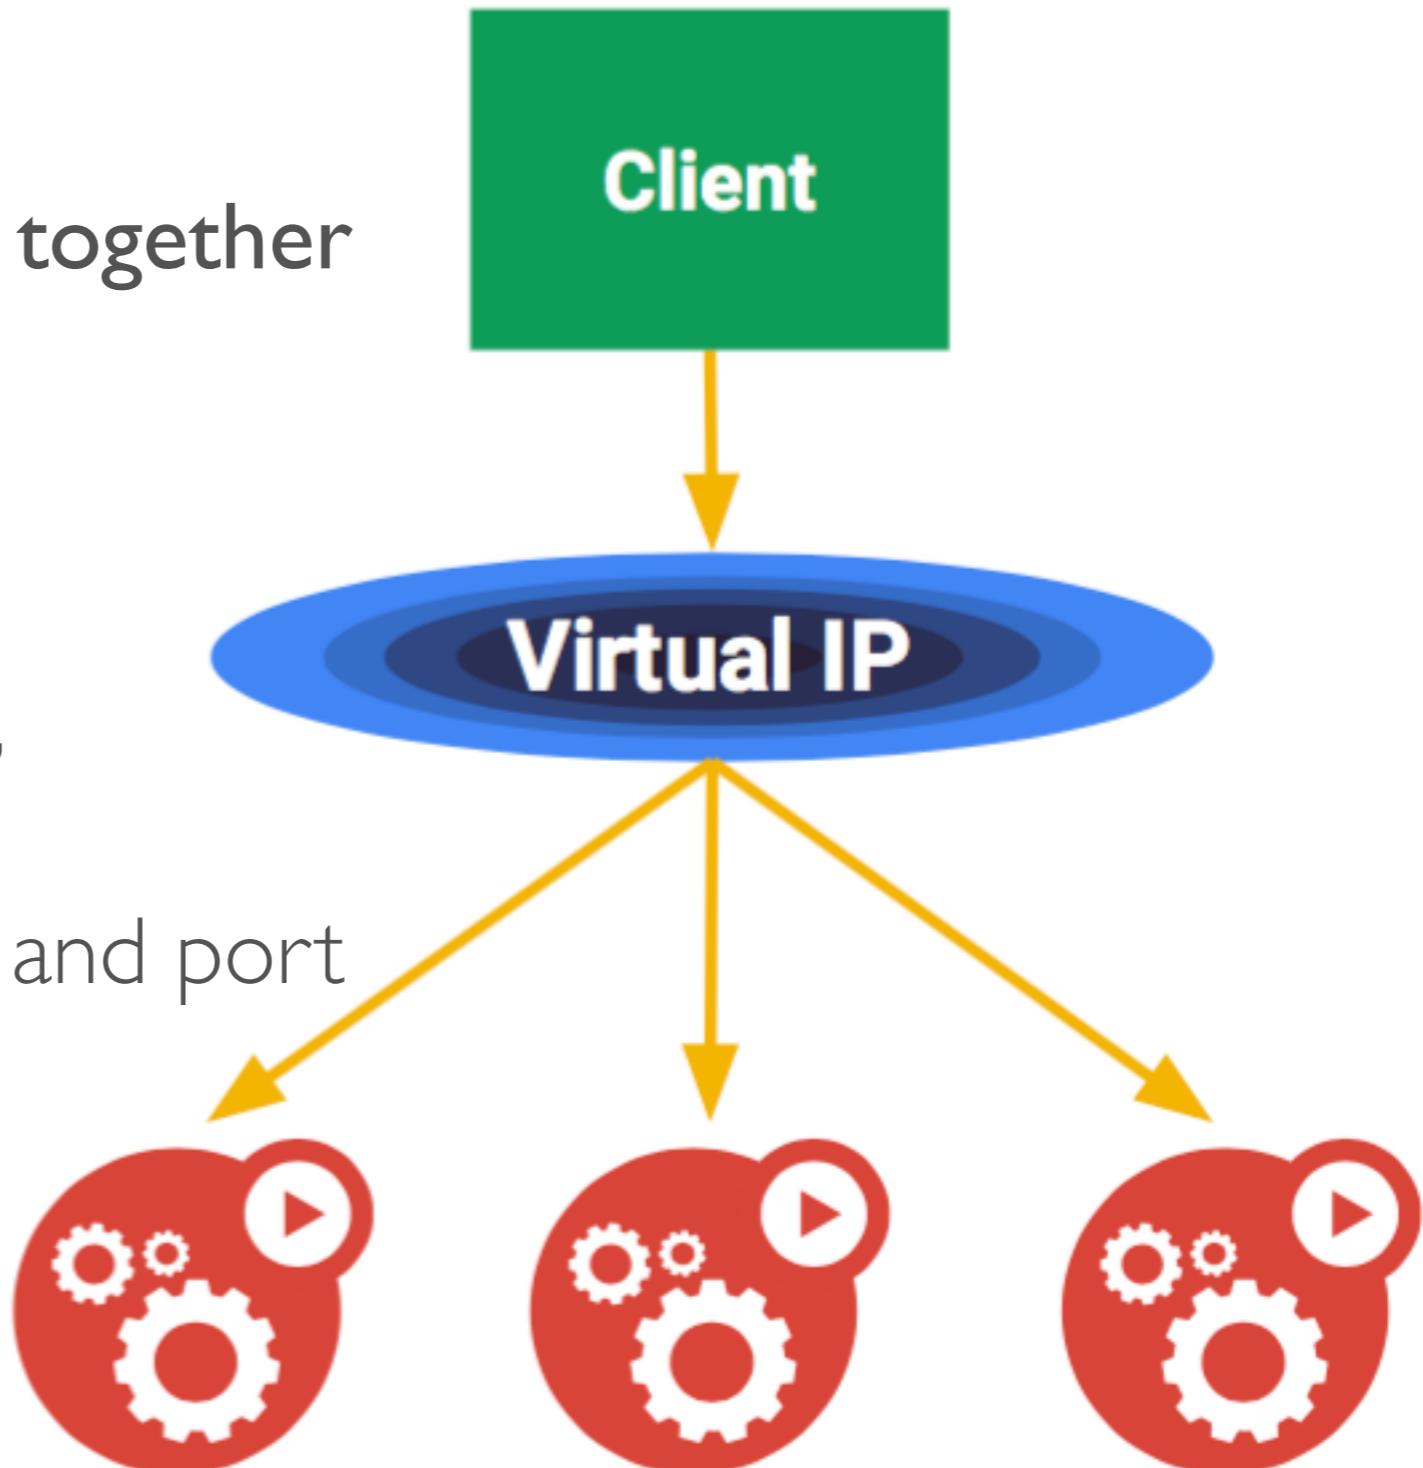
ReplicaSet

- **name = "my-rc"**
- **selector = {"App": "MyApp"}**
- **template = { ... }**
- **replicas = 4**



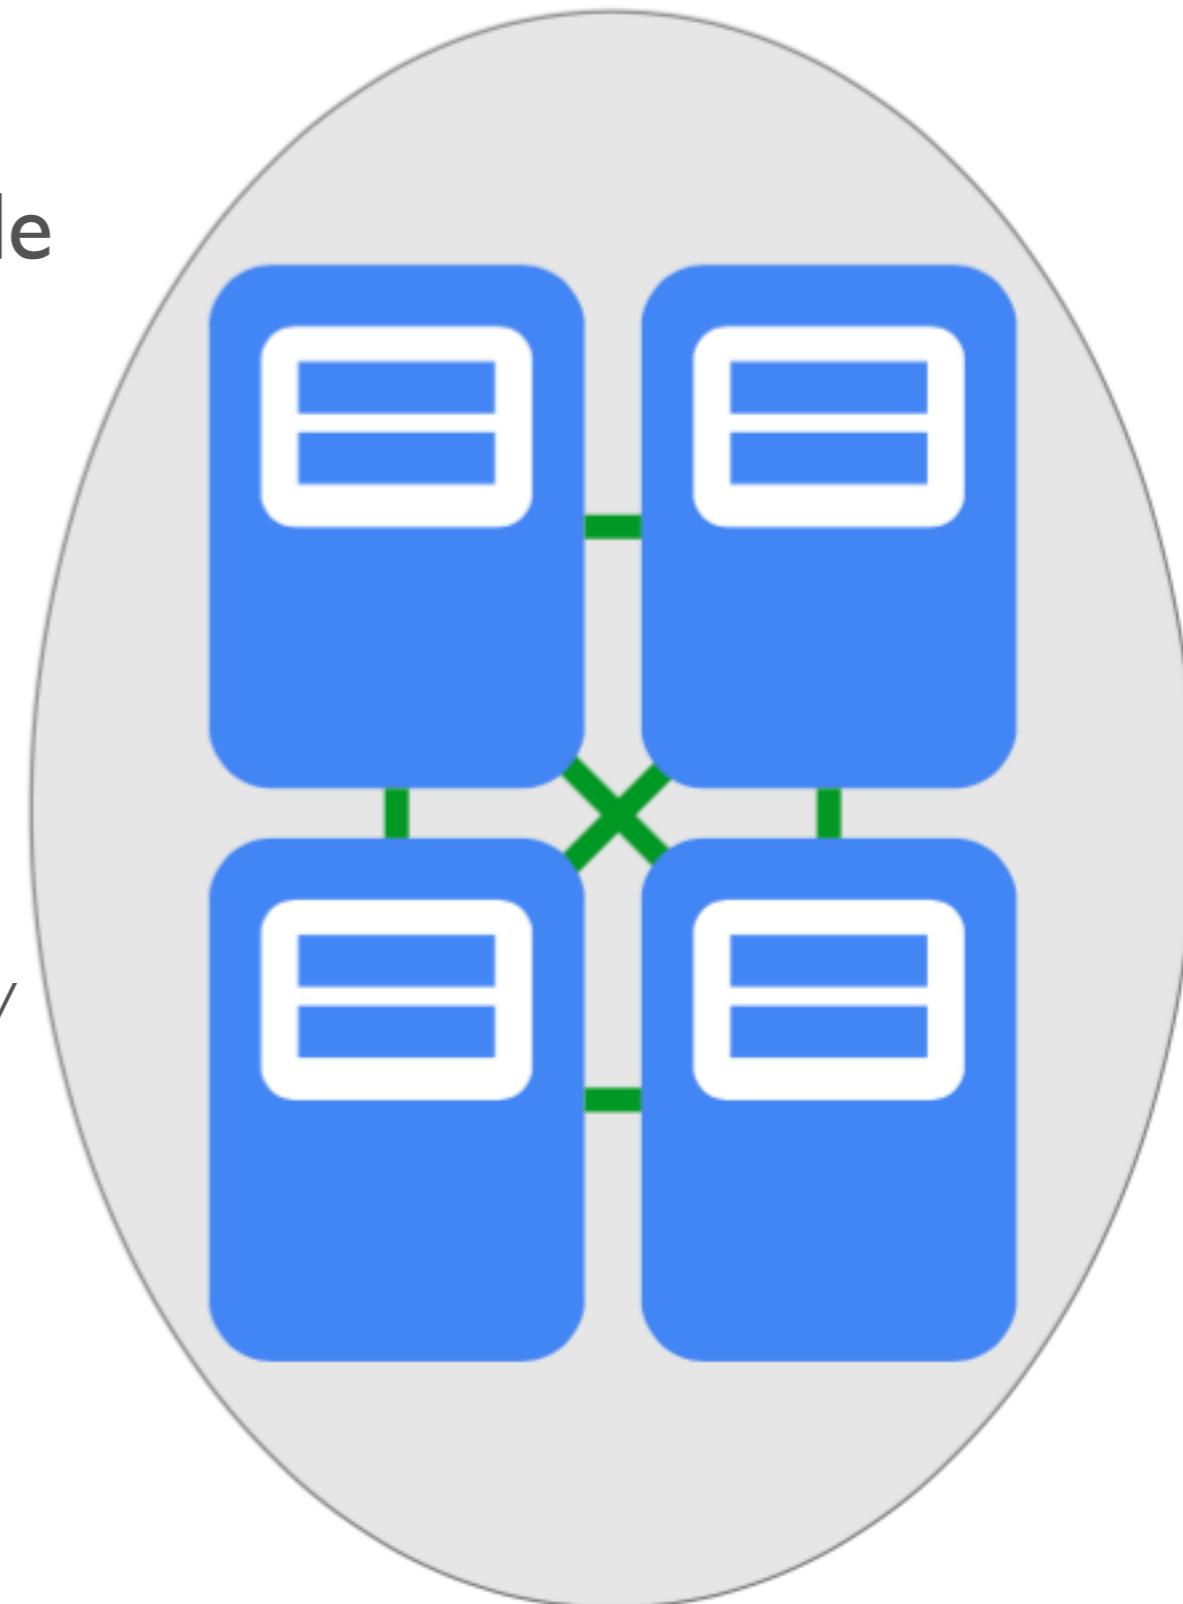
Services

- ▶ A group of pods that work together
 - grouped by a selector
- ▶ Defines access policy
 - “load balanced” or “headless”
- ▶ Can have a stable **virtual IP** and port
 - also a DNS name
- ▶ Hides complexity

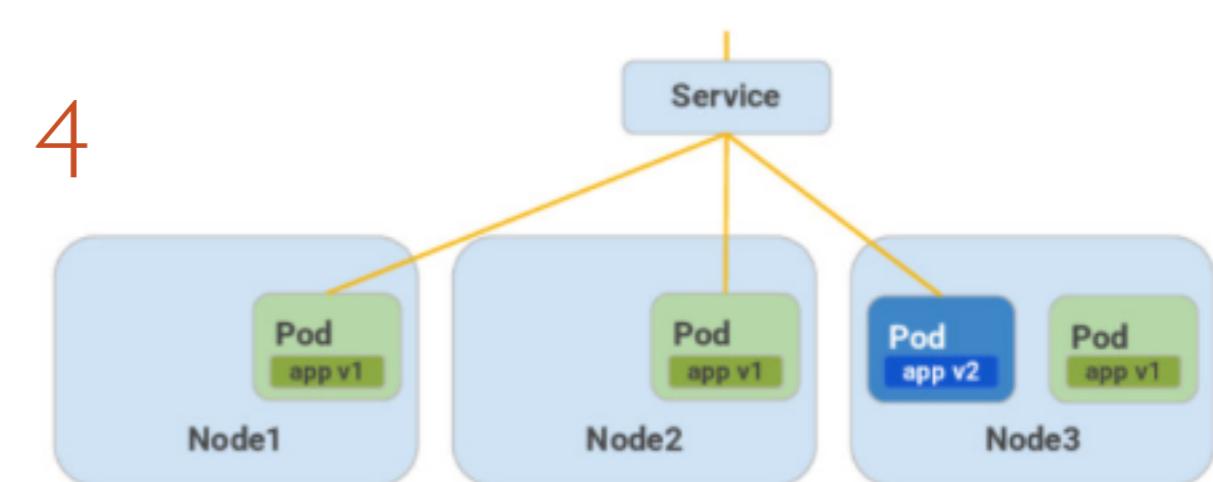
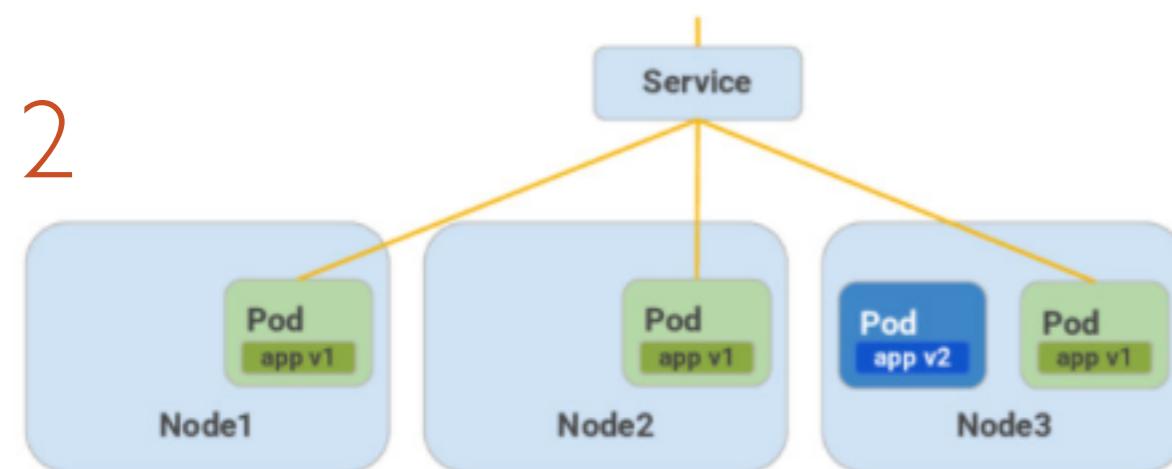
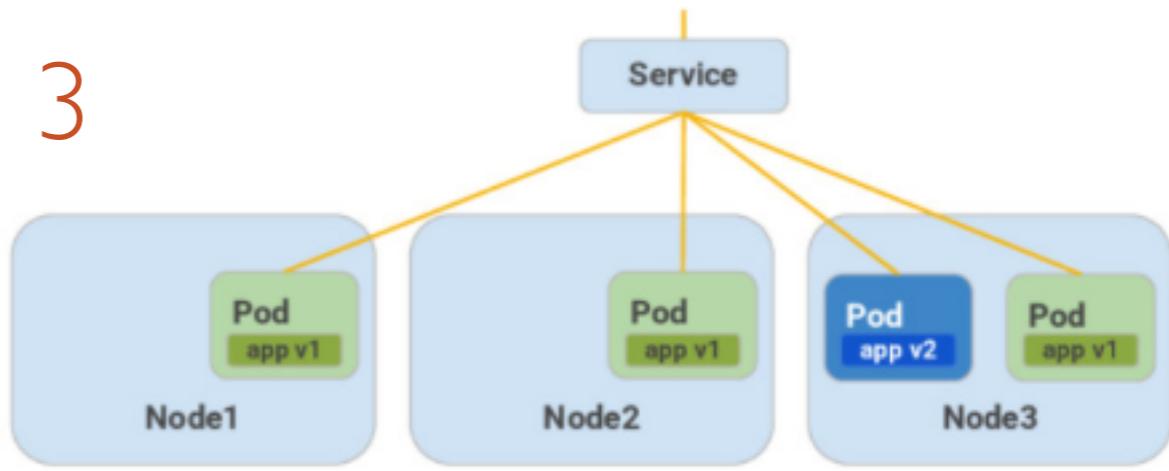
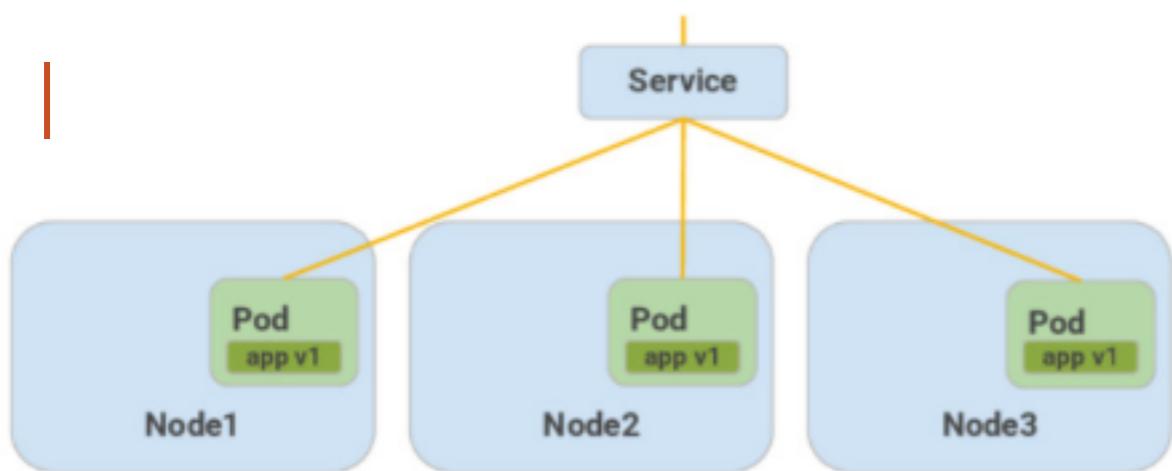


External Services

- ▶ Services VIPs are only available **inside** the cluster
- ▶ Need to receive traffic from “the outside world”
- ▶ Service “type”
 - NodePort: expose on a port on every node
 - LoadBalancer: provision a cloud load-balancer

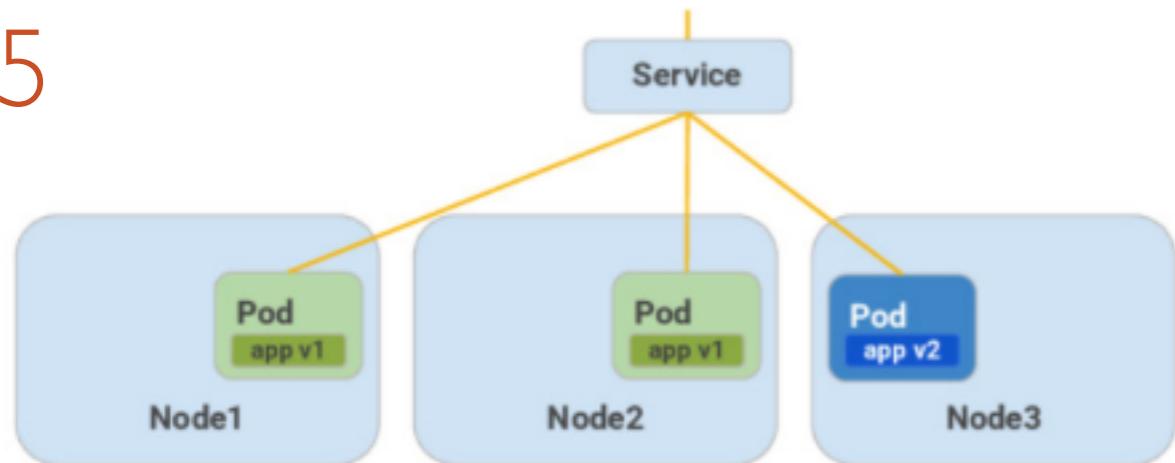


Rolling Update

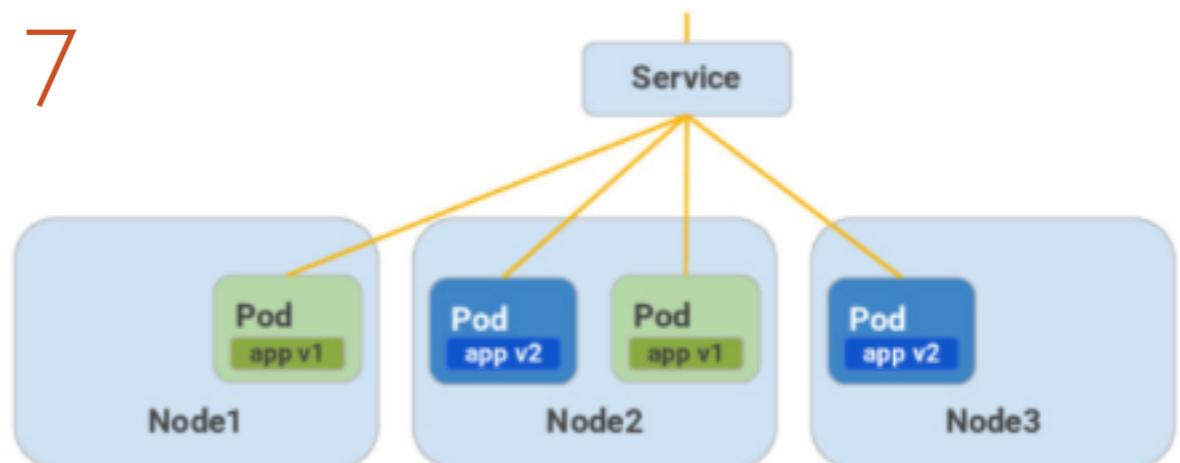


Rolling Update

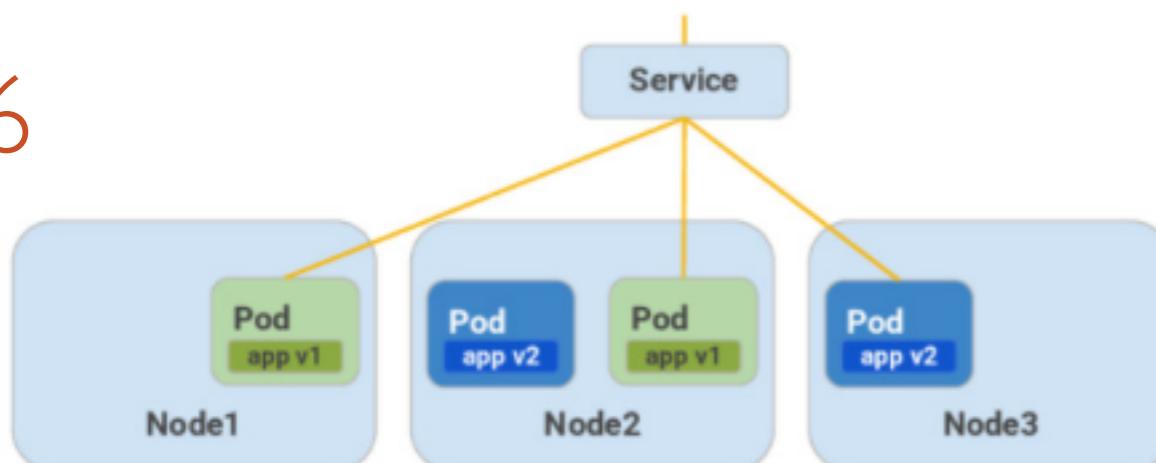
5



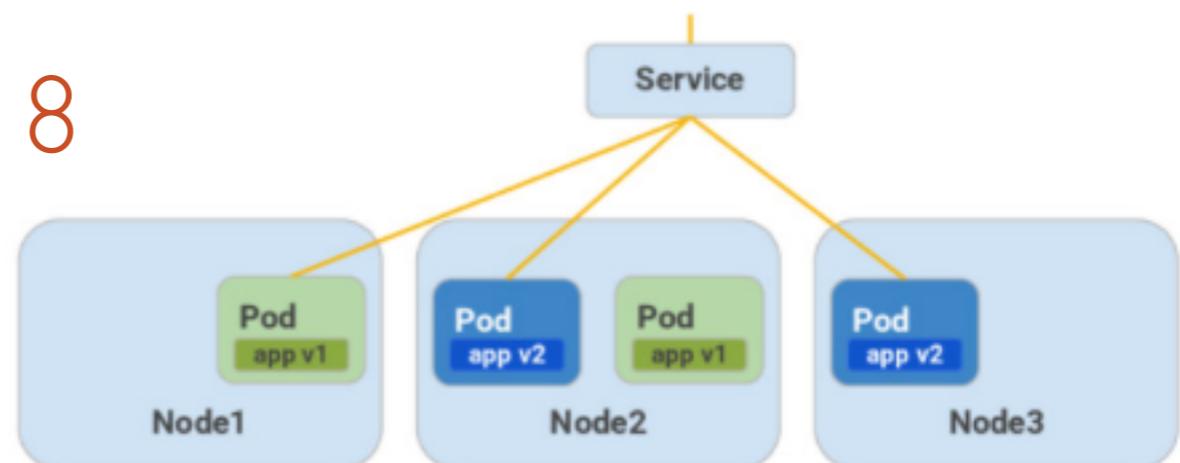
7



6

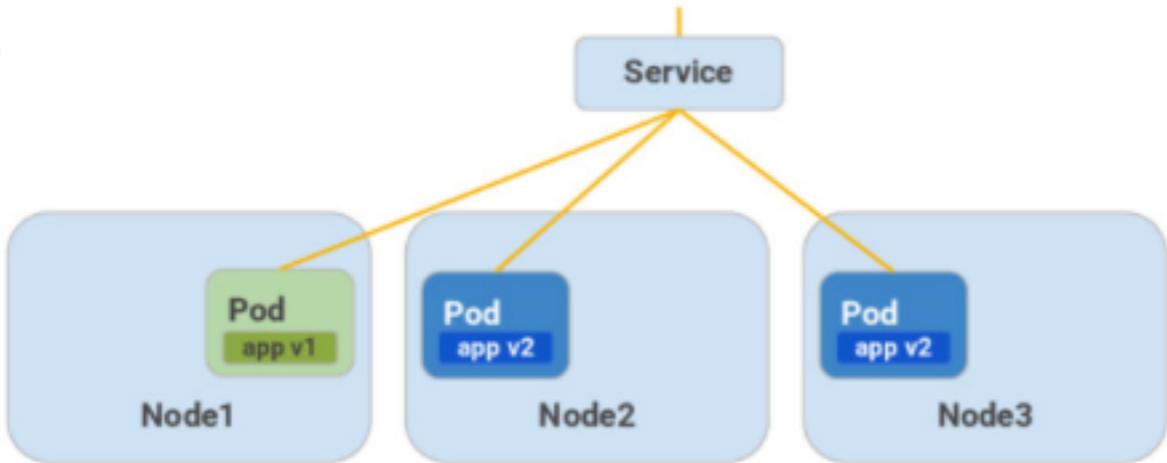


8

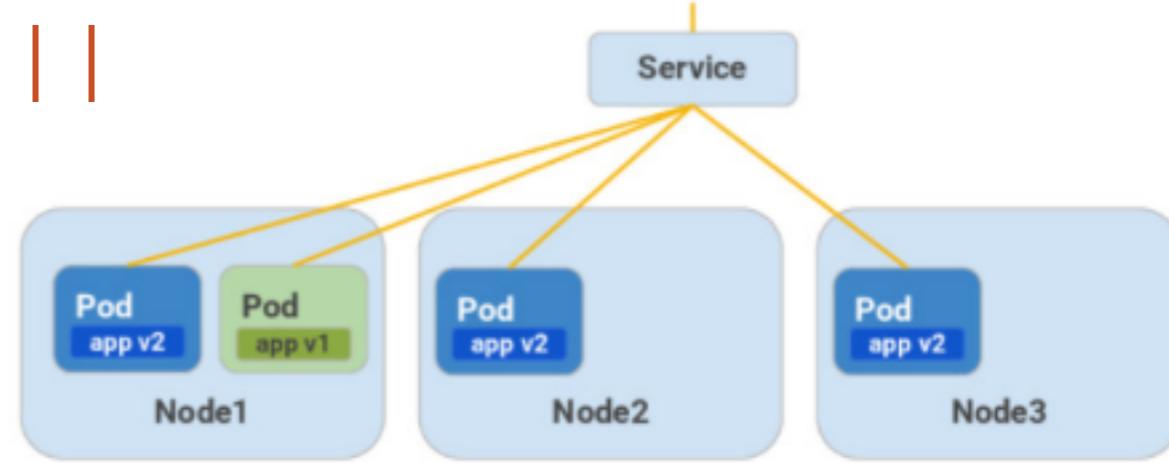


Rolling Update

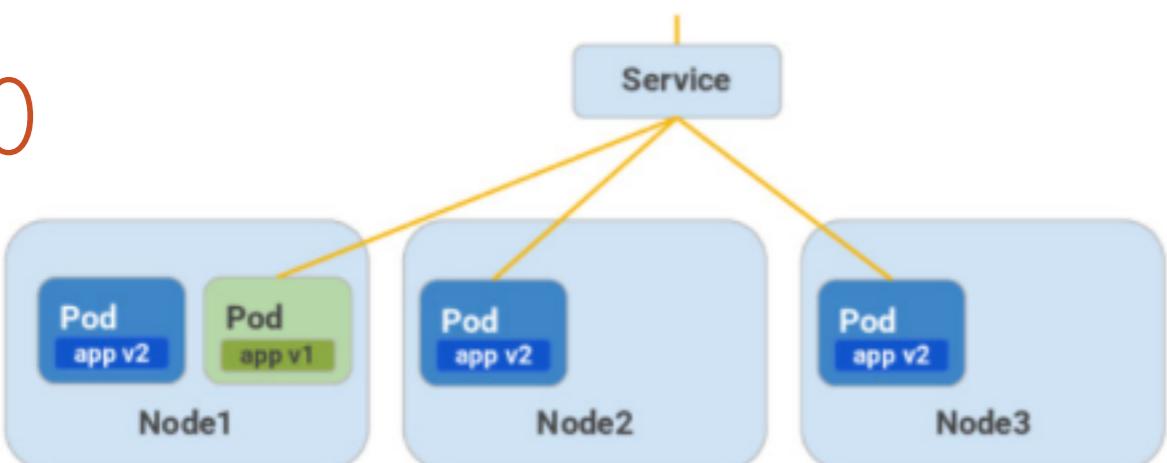
9



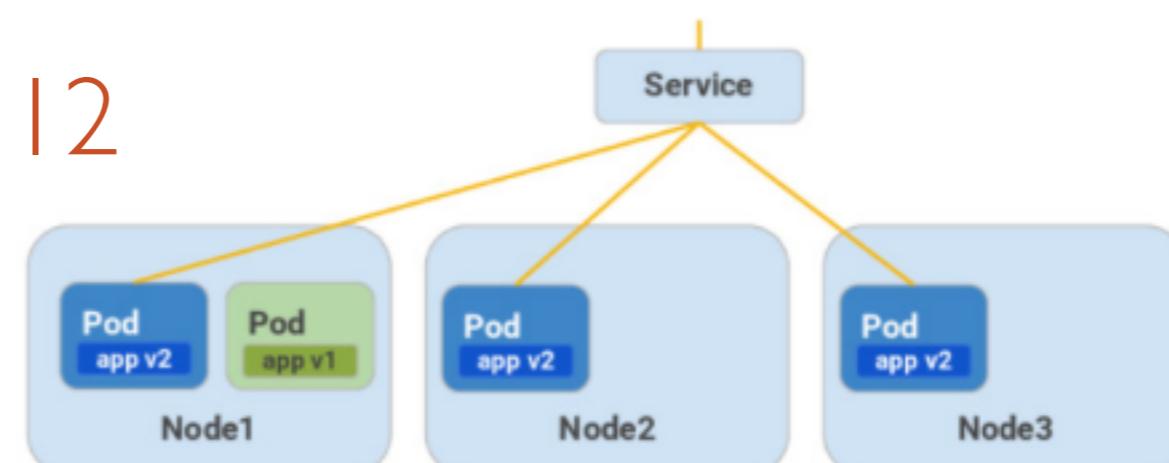
10



10

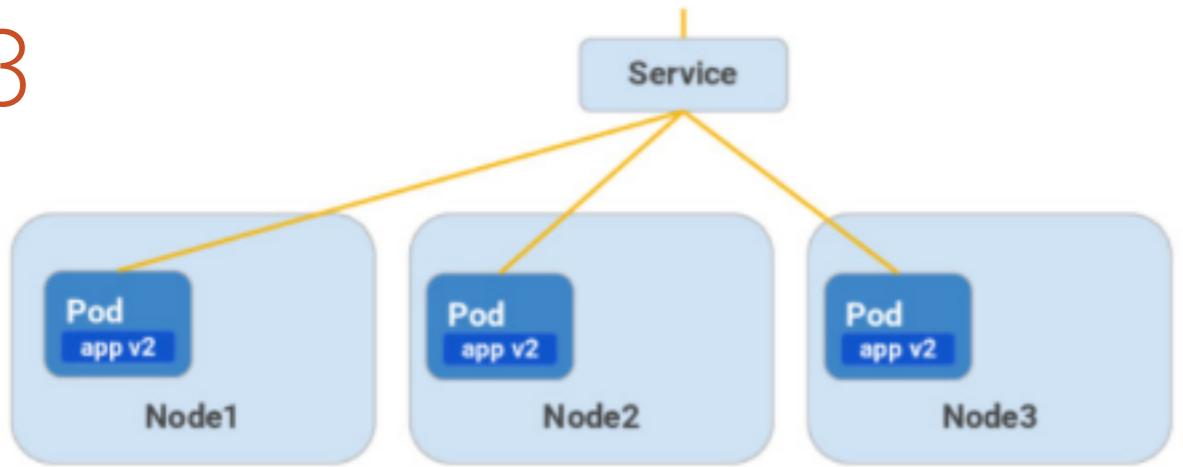


12



Rolling Update

| 3





kubernetes

Hands on lab

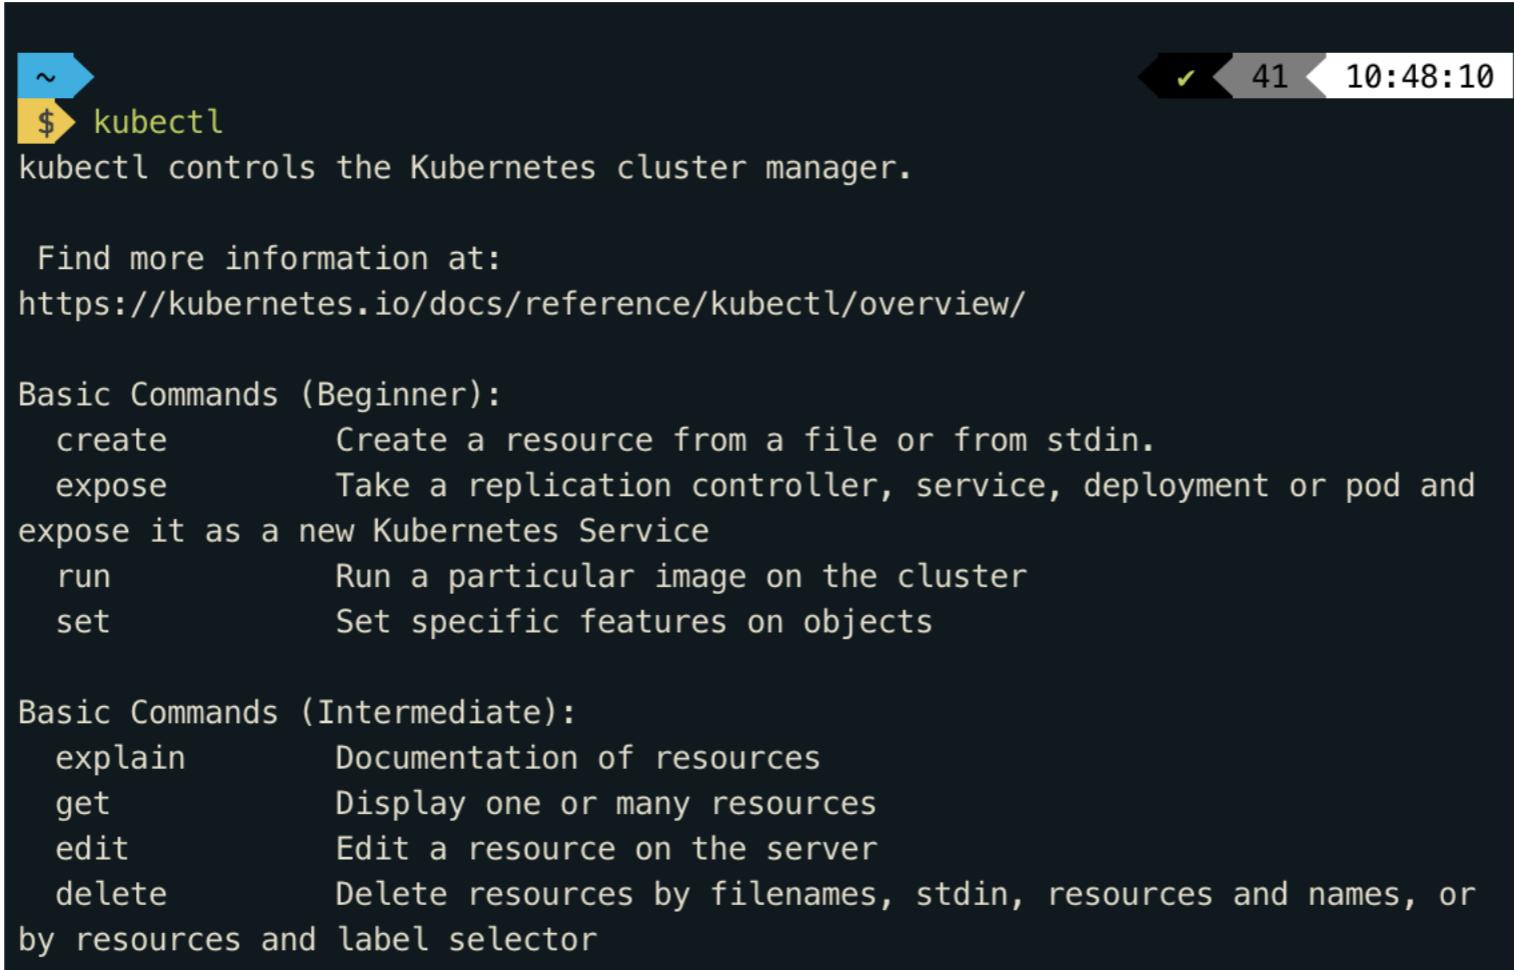
Welcome to Kubernetes

- Minikube - Intended for local kubernetes development
- Docker

<https://kubernetes.io/docs/tasks/tools/install-minikube/>



1. Install **kubectl** binary with **curl** (Linux and MacOS)
2. On Windows download and add the folder to the path
3. Try the the tool on the terminal (on windows, install a different command line) - **kubectl**



A screenshot of a terminal window showing the output of the `kubectl` command. The terminal has a dark theme with a blue status bar at the top showing the time as 10:48:10. The output includes the command usage, basic commands for beginners, and intermediate commands.

```
~ $ kubectl
kubectl controls the Kubernetes cluster manager.

Find more information at:
https://kubernetes.io/docs/reference/kubectl/overview/

Basic Commands (Beginner):
  create      Create a resource from a file or from stdin.
  expose      Take a replication controller, service, deployment or pod and
  expose it as a new Kubernetes Service
  run         Run a particular image on the cluster
  set         Set specific features on objects

Basic Commands (Intermediate):
  explain     Documentation of resources
  get         Display one or many resources
  edit        Edit a resource on the server
  delete     Delete resources by filenames, stdin, resources and names, or
  by resources and label selector
```



1. Install **minikube** binary with **curl** (Linux and MacOS)
2. On Windows download and add the folder where you downloaded **kubectl**
3. Try the the tool on the terminal (on windows, install a different command line) - **minikube**

A screenshot of a terminal window showing the output of the `minikube` command. The terminal has a dark theme with a light-colored status bar at the top right showing a checkmark, the number 44, and the time 10:50:39. The main text area shows the following:

```
~ $ minikube
Minikube is a CLI tool that provisions and manages single-node Kubernetes clusters optimized for development workflows.

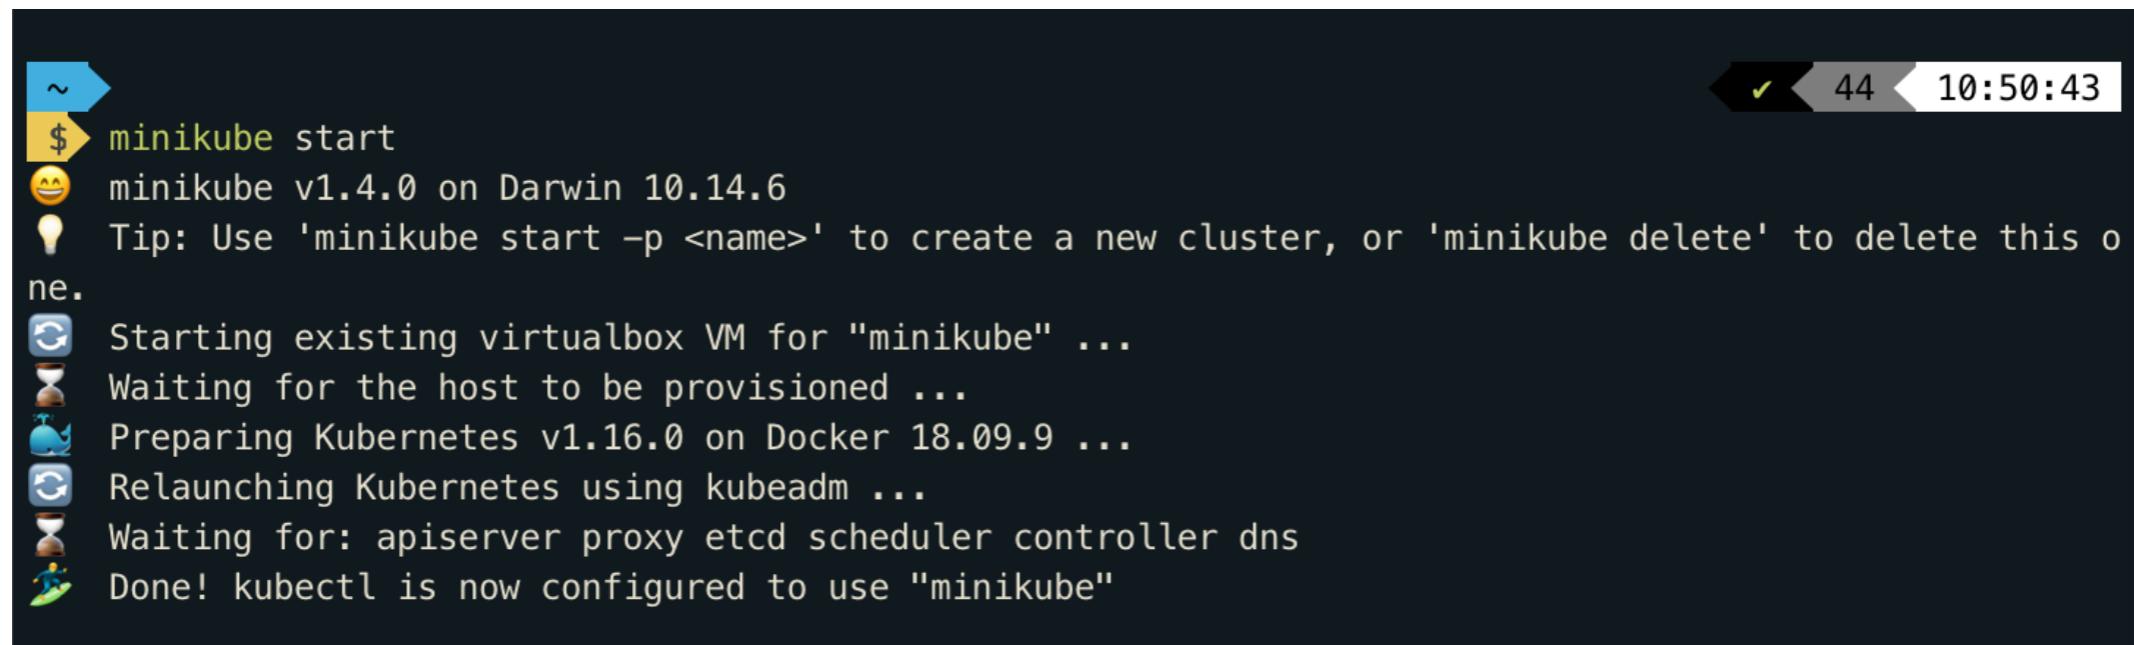
Basic Commands:
  start      Starts a local kubernetes cluster
  status     Gets the status of a local kubernetes cluster
  stop       Stops a running local kubernetes cluster
  delete     Deletes a local kubernetes cluster
  dashboard  Access the kubernetes dashboard running within the minikube cluster

Images Commands:
  docker-env Sets up docker env variables; similar to '$(docker-machine env)'
  cache      Add or delete an image from the local cache.
```



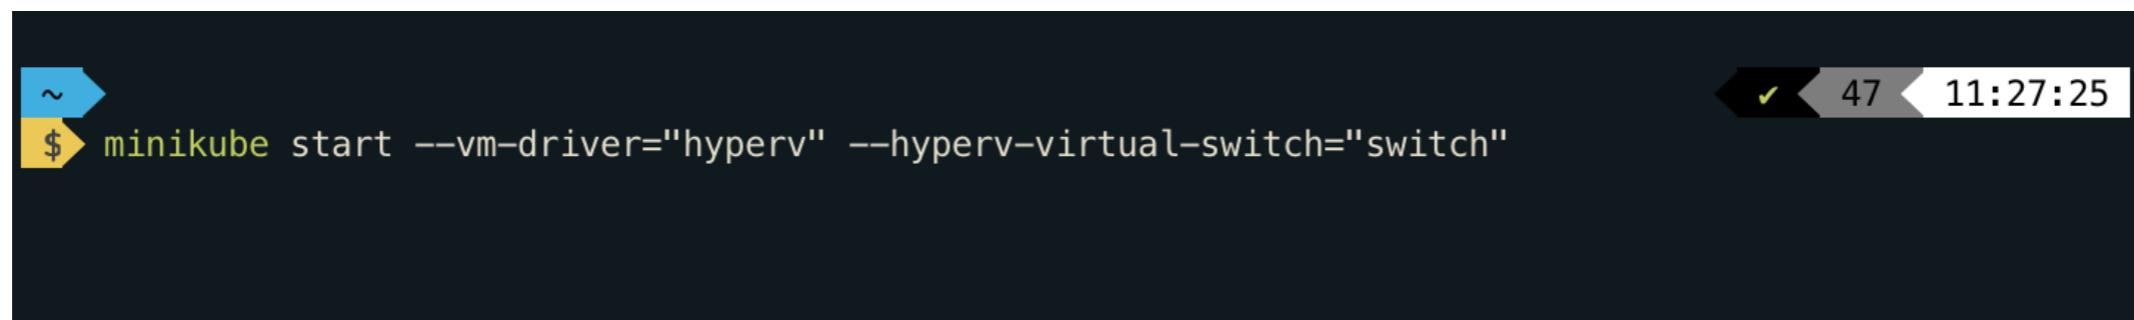
Try minikube

I. On linux and MacOS start minikube - `minikube start`



```
~ $ minikube start
😄 minikube v1.4.0 on Darwin 10.14.6
💡 Tip: Use 'minikube start -p <name>' to create a new cluster, or 'minikube delete' to delete this one.
🔄 Starting existing virtualbox VM for "minikube" ...
⏳ Waiting for the host to be provisioned ...
🐳 Preparing Kubernetes v1.16.0 on Docker 18.09.9 ...
🔄 Relaunching Kubernetes using kubeadm ...
⏳ Waiting for: apiserver proxy etcd scheduler controller dns
🏄 Done! kubectl is now configured to use "minikube"
```

2. On Windows (with HyperV)



```
~ $ minikube start --vm-driver="hyperv" --hyperv-virtual-switch="switch"
```



3. Kubectl version

```
~ $ kubectl version
Client Version: version.Info{Major:"1", Minor:"16", GitVersion:"v1.16.0", GitCommit:"2bd9643cee5b3b3a5
ecbd3af49d09018f0773c77", GitTreeState:"clean", BuildDate:"2019-09-18T14:36:53Z", GoVersion:"go1.12.9"
, Compiler:"gc", Platform:"darwin/amd64"}
Server Version: version.Info{Major:"1", Minor:"16", GitVersion:"v1.16.0", GitCommit:"2bd9643cee5b3b3a5
ecbd3af49d09018f0773c77", GitTreeState:"clean", BuildDate:"2019-09-18T14:27:17Z", GoVersion:"go1.12.9"
, Compiler:"gc", Platform:"linux/amd64"}
```

There should be a Client and Server version

Docker: Important things...

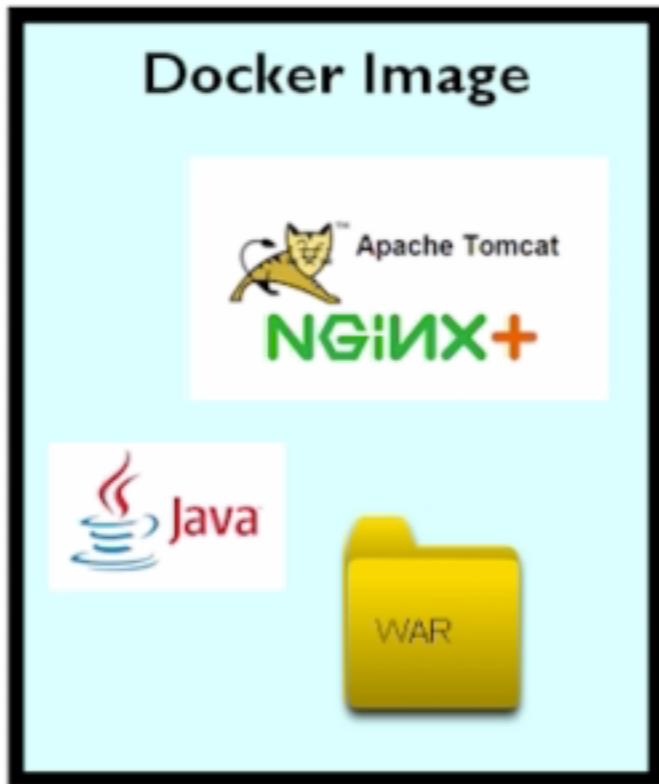
- ▶ Most important concept in docker

containers vs images

- ▶ An image e a definition of a container
- ▶ binary file with all the software a things like ambiente variables, general definitions, etc

Let's imagine i'm a developer given the task to do a website





***In the end I would
deliver the war file to
someone responsible to
deploy it***

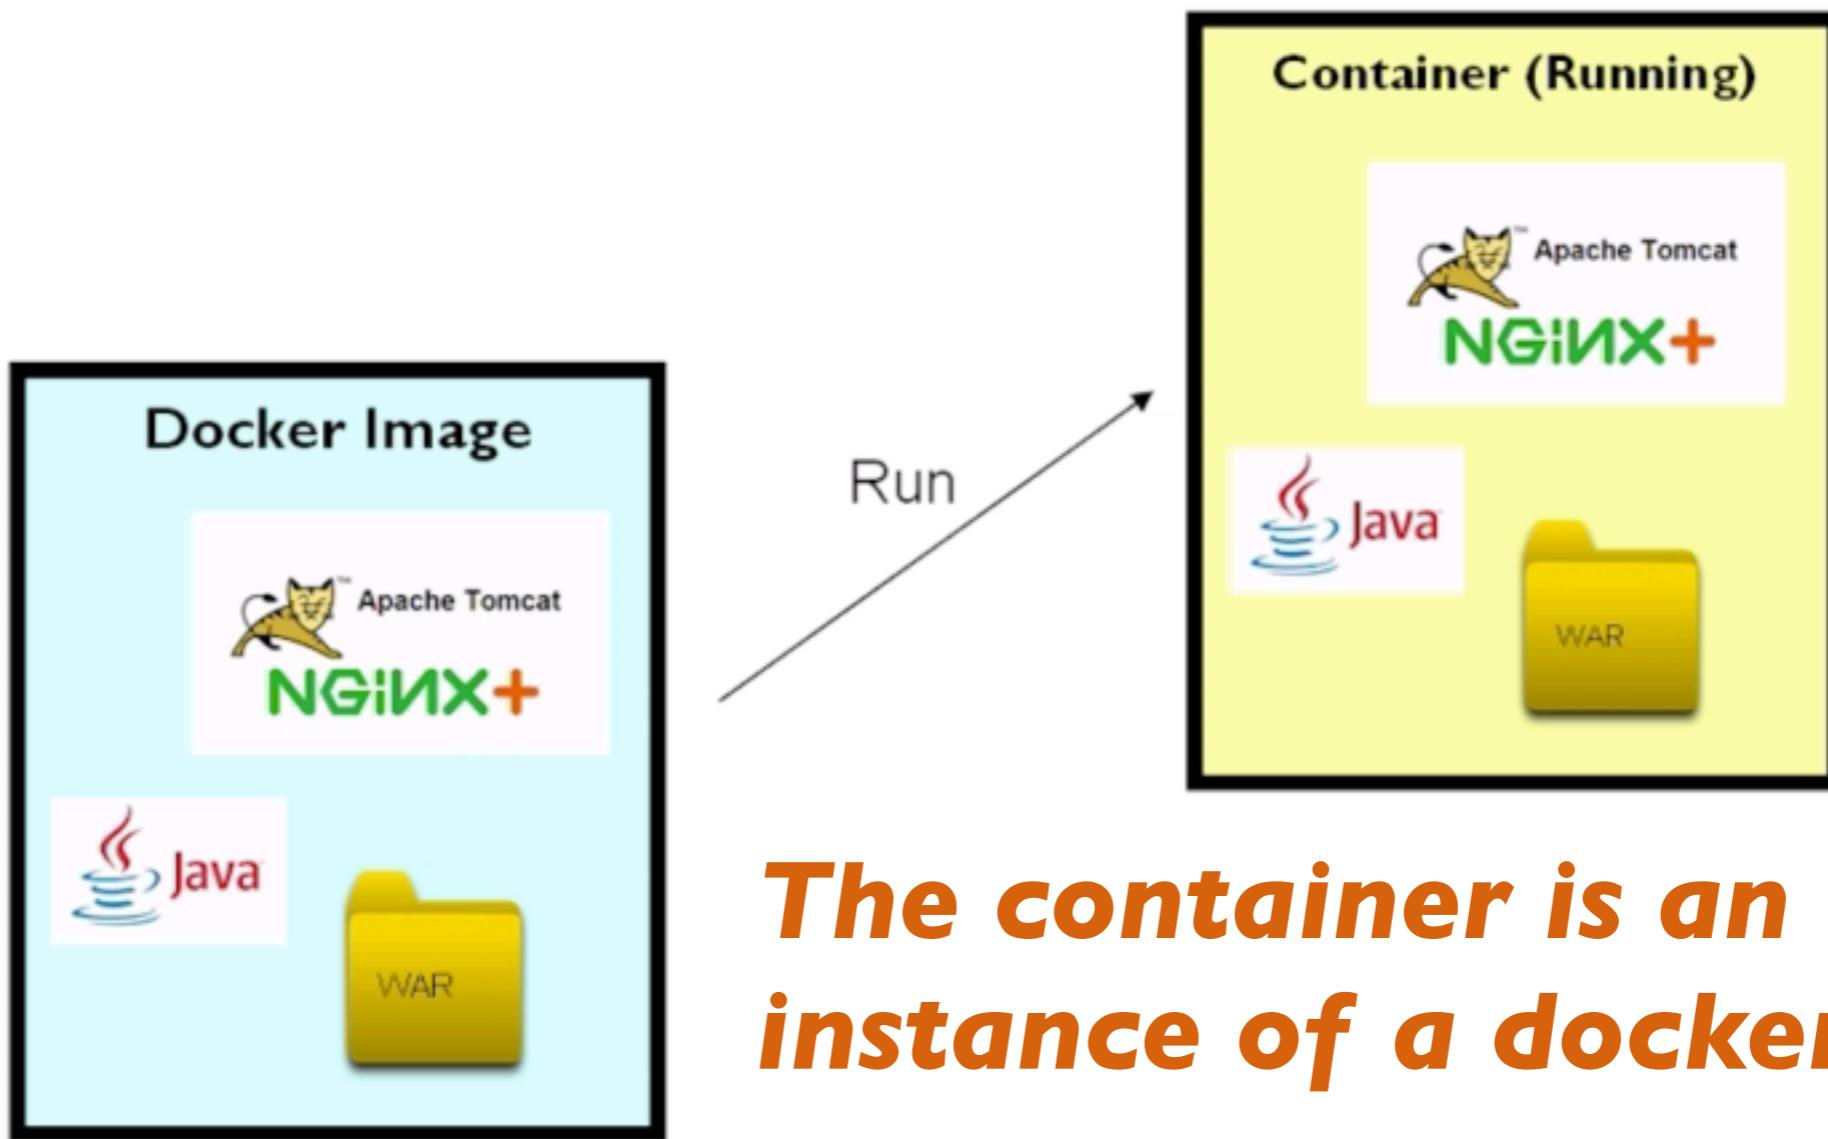
will it work?



I would expect them to configure java, tomcat and nginx

With an image, I as a developer can package everything that is needed every single dependency required to make that application run

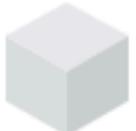
- ▶ The idea then is to pass the image to the deployer - then he has to run that image
- ▶ The running image is the **container**



- ▶ We can publish our images on the docker ecosystem for other developers

hub.docker.com

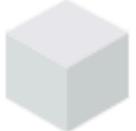
- We will use for the next lessons the following repository



[richardchesterwood/k8s-fleetman-position-simulator](#) 8.6K Downloads 4 Stars

By [richardchesterwood](#) • Updated 5 months ago

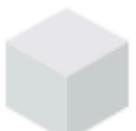
Container



[richardchesterwood/k8s-fleetman-webapp-angular](#) 10K+ Downloads 26 Stars

By [richardchesterwood](#) • Updated 7 months ago

Container

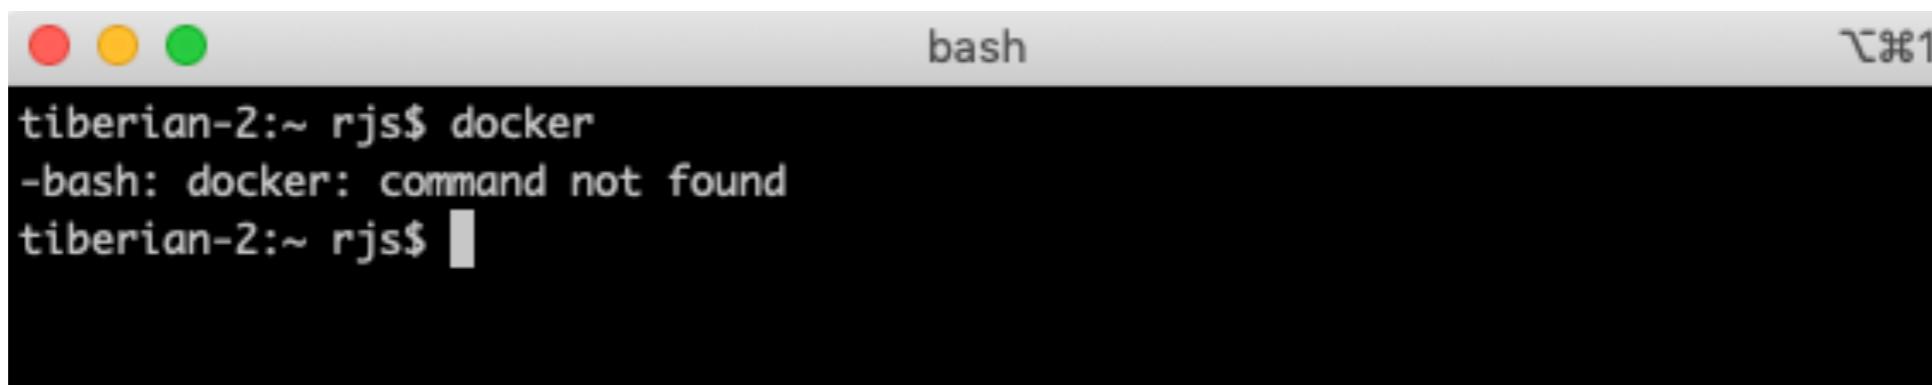


[richardchesterwood/k8s-fleetman-queue](#) 10K+ Downloads 8 Stars

By [richardchesterwood](#) • Updated 8 months ago

Container

- ▶ You will need docker installed

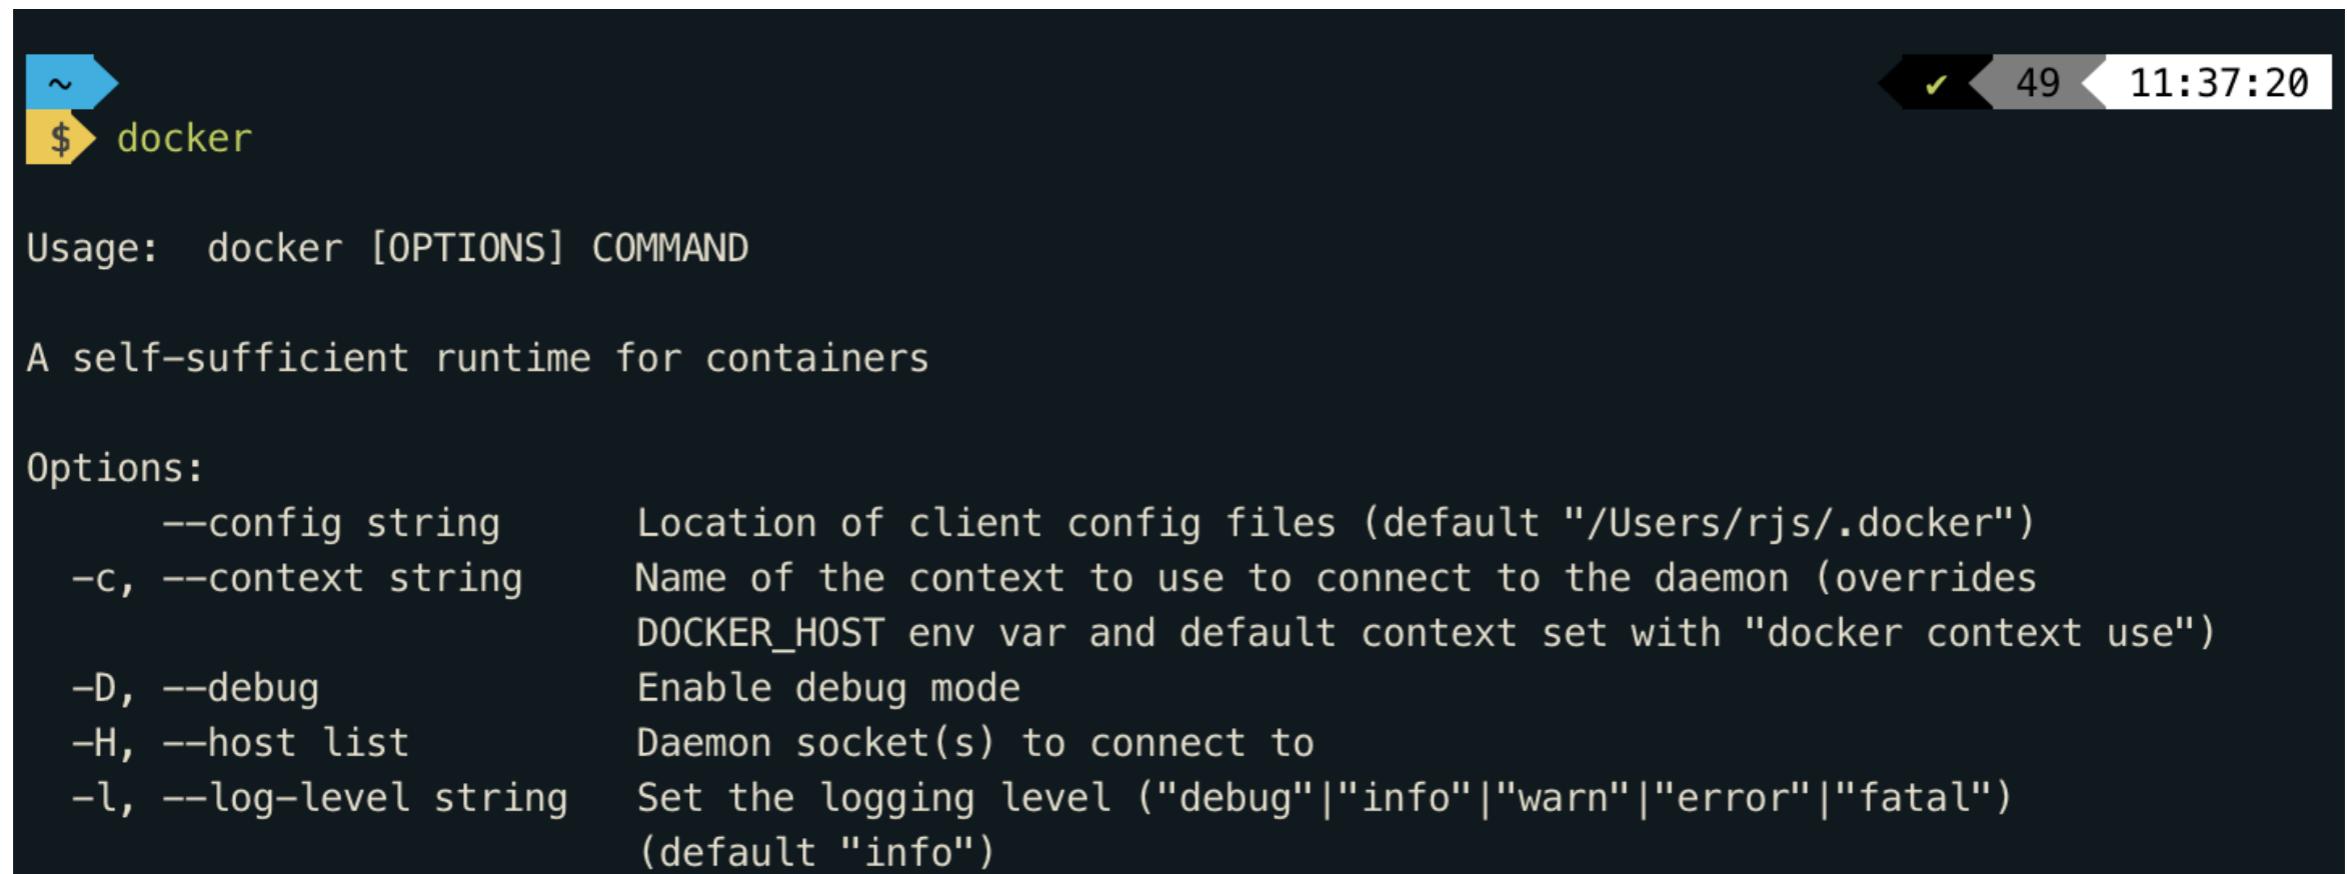


A screenshot of a Mac OS X terminal window titled "bash". The window has three colored status icons (red, yellow, green) at the top left. The title bar says "bash" and there is a small icon at the top right. The terminal window contains the following text:

```
tiberian-2:~ rjs$ docker
-bash: docker: command not found
tiberian-2:~ rjs$ █
```

- ▶ If you are in a Mac - get **Docker for Mac**
- ▶ Otherwise, get **Docker Toolbox**

- ▶ If you are in a Mac - get Docker for Mac
- ▶ Otherwise, get Docker Toolbox



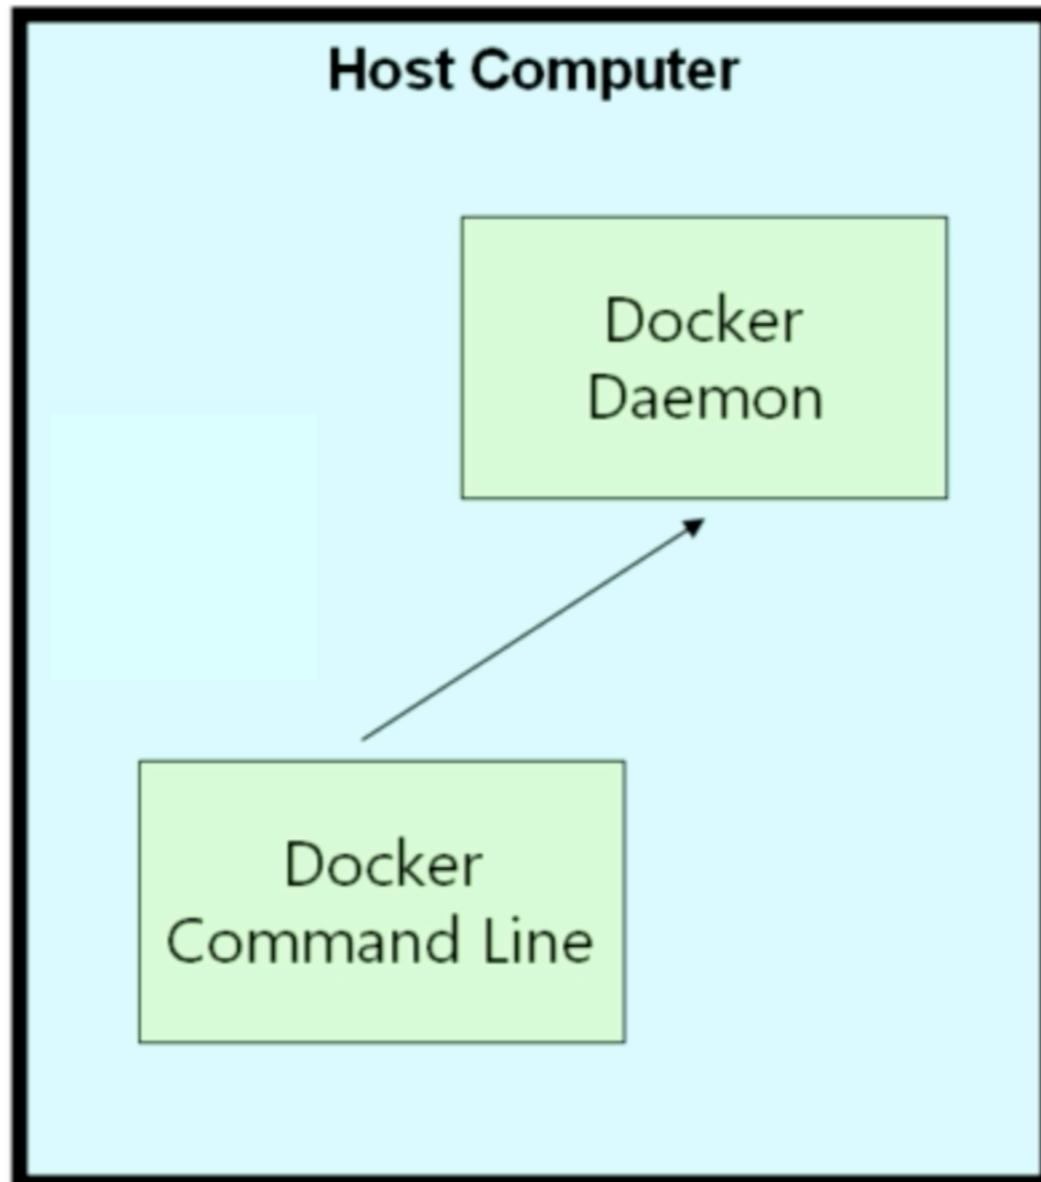
A screenshot of a terminal window on a Mac. The window title bar shows a checkmark icon, the number 49, and the time 11:37:20. The terminal prompt is ~ \$ docker. The output shows the usage of the docker command, its purpose as a self-sufficient runtime for containers, and a detailed list of options:

```
Usage: docker [OPTIONS] COMMAND
A self-sufficient runtime for containers
Options:
  --config string      Location of client config files (default "/Users/rjs/.docker")
  -c, --context string Name of the context to use to connect to the daemon (overrides
                        DOCKER_HOST env var and default context set with "docker context use")
  -D, --debug          Enable debug mode
  -H, --host list       Daemon socket(s) to connect to
  -l, --log-level string Set the logging level ("debug"|"info"|"warn"|"error"|"fatal")
                        (default "info")
```

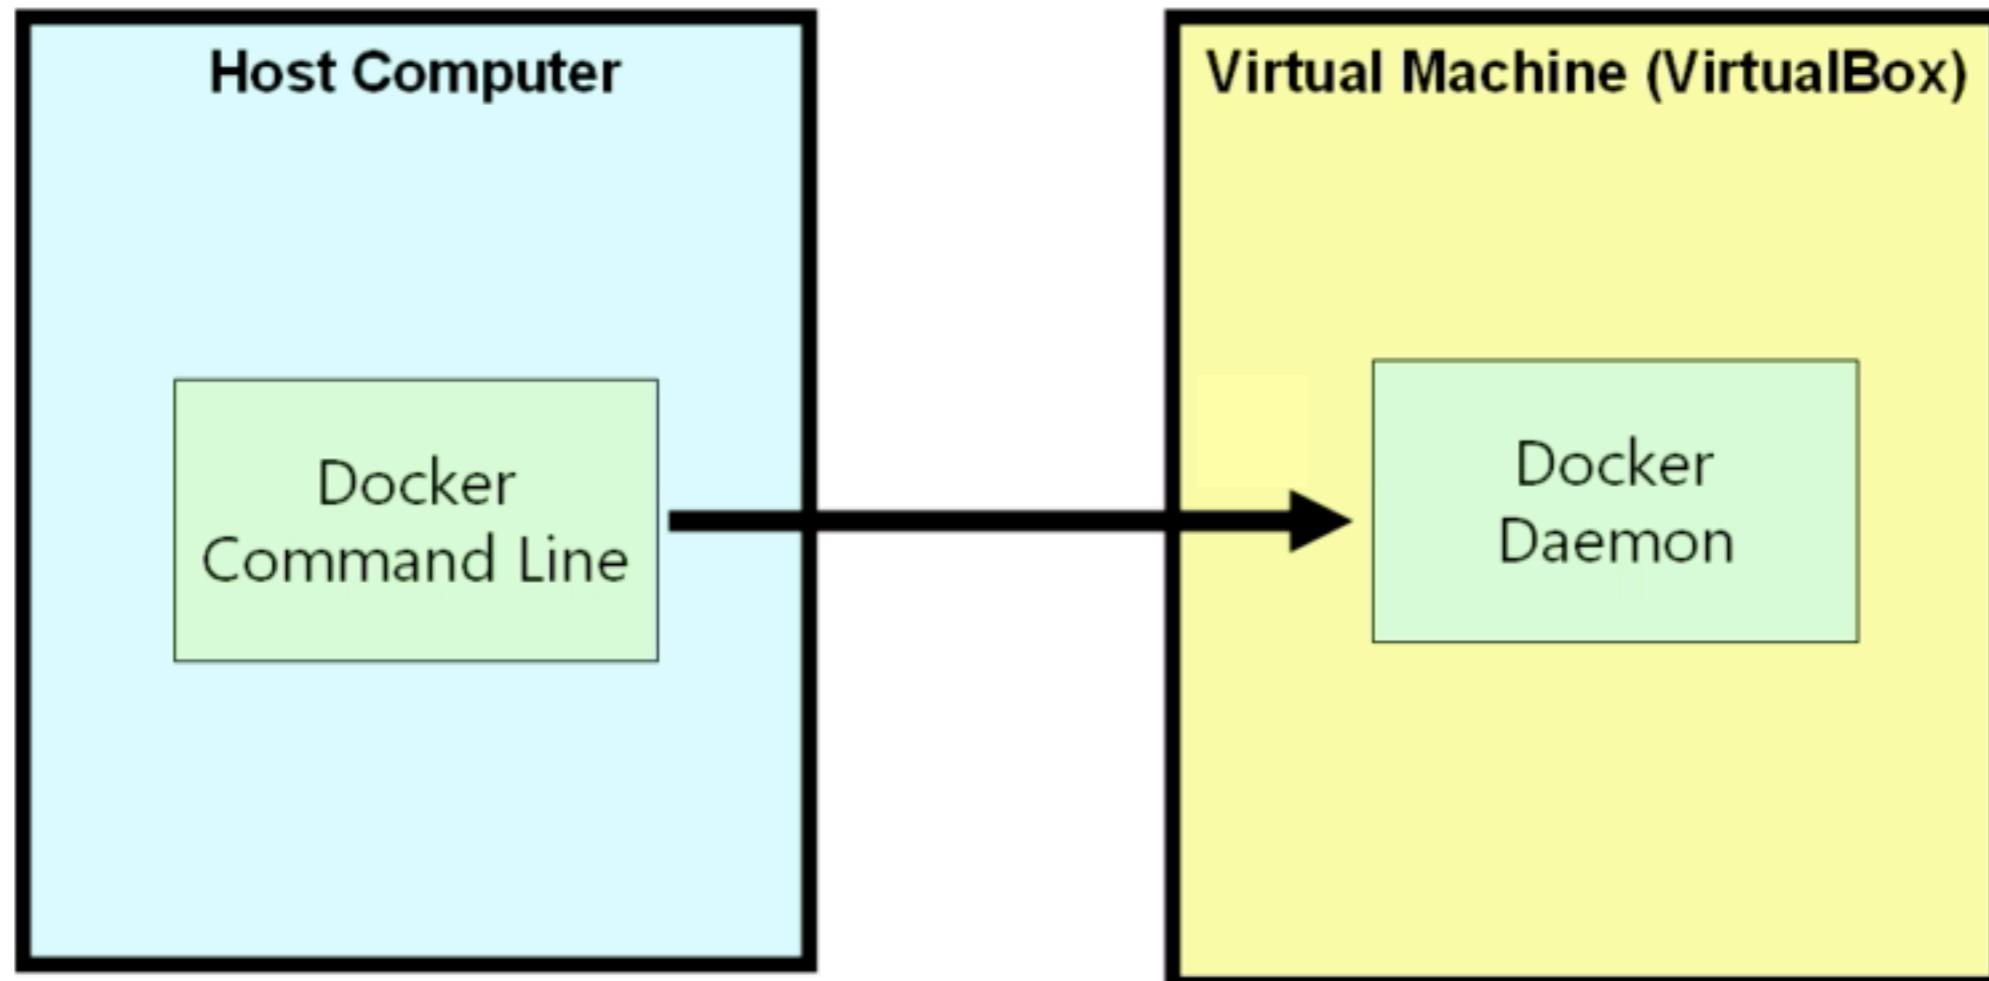
- ▶ It took a lot of time right?

A screenshot of a terminal window with a black background. At the top, there's a blue arrow icon followed by a yellow arrow icon, and at the bottom, there's a blue arrow icon followed by a yellow arrow icon. To the right of the icons are three small navigation arrows (left, right, up) and the number '51' above '11:39:36'. In the center, the command `docker image ls` is typed, followed by an error message: "Cannot connect to the Docker daemon at unix:///var/run/docker.sock. Is the docker daemon running?". Below this, another blue arrow icon and a yellow arrow icon are visible, along with three small navigation arrows (left, right, up) and the number '52' above '11:39:39'.

Why?

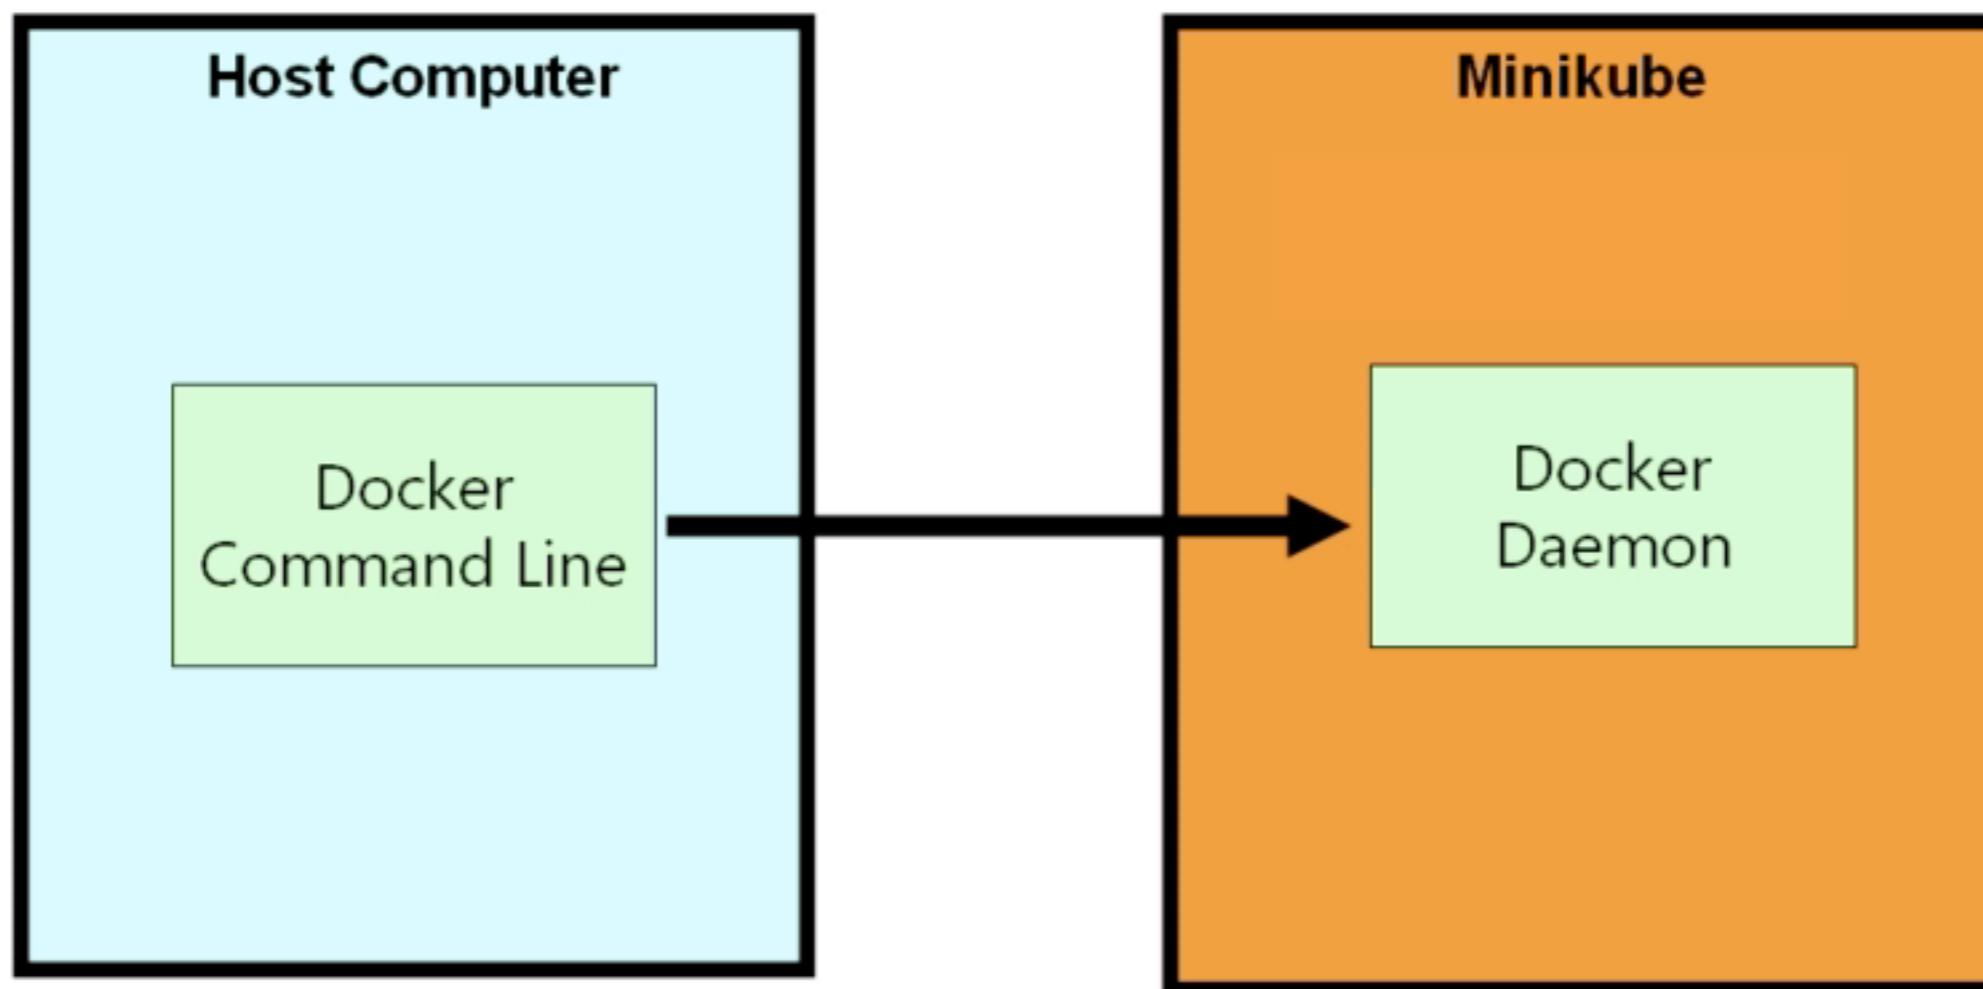


- ▶ Docker daemon need to run on linux - because it uses various features of the linux operation system to support containers

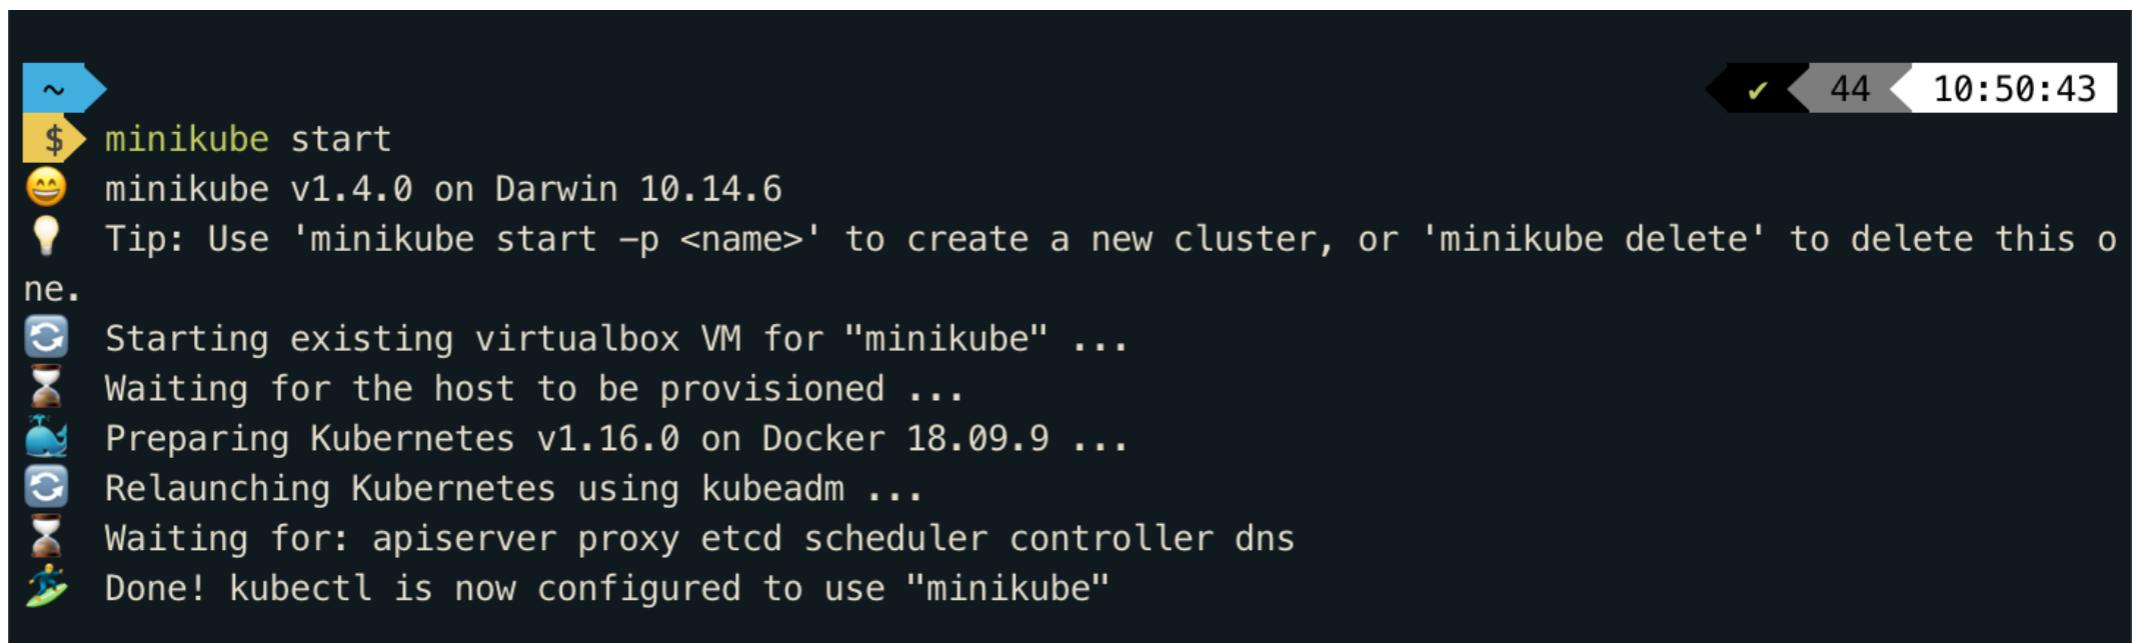


- ▶ Docker tries to look for a virtual machine - looks for a docker daemon on the VM

- ▶ You have to start that VM
- ▶ Because we intend to work on Kubernetes, and we installed Minikube (cut down version of kubernetes)
- ▶ Inside Minikube we have an installation of docker - docker daemon

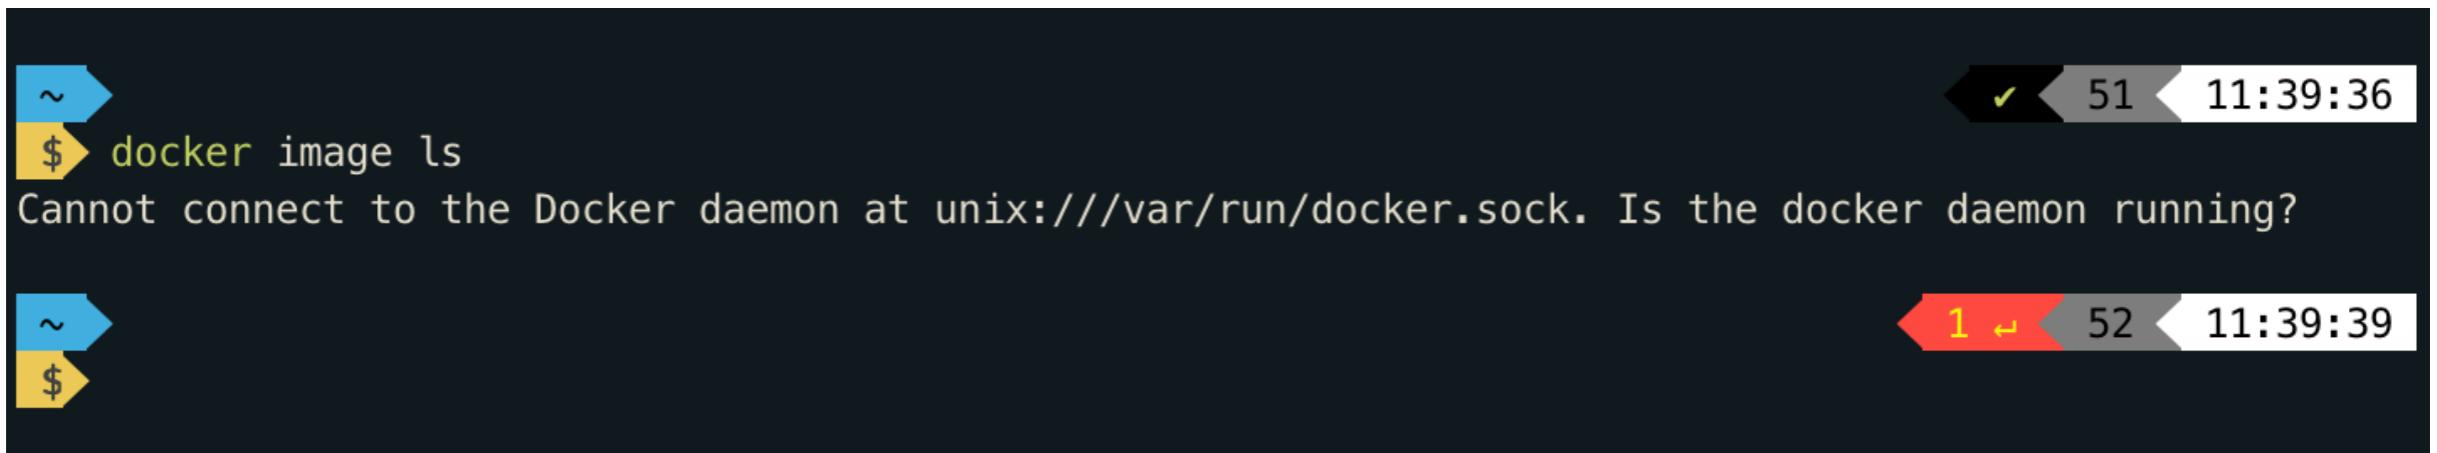


- ▶ Let's make use of that docker installation -
`minikube start`



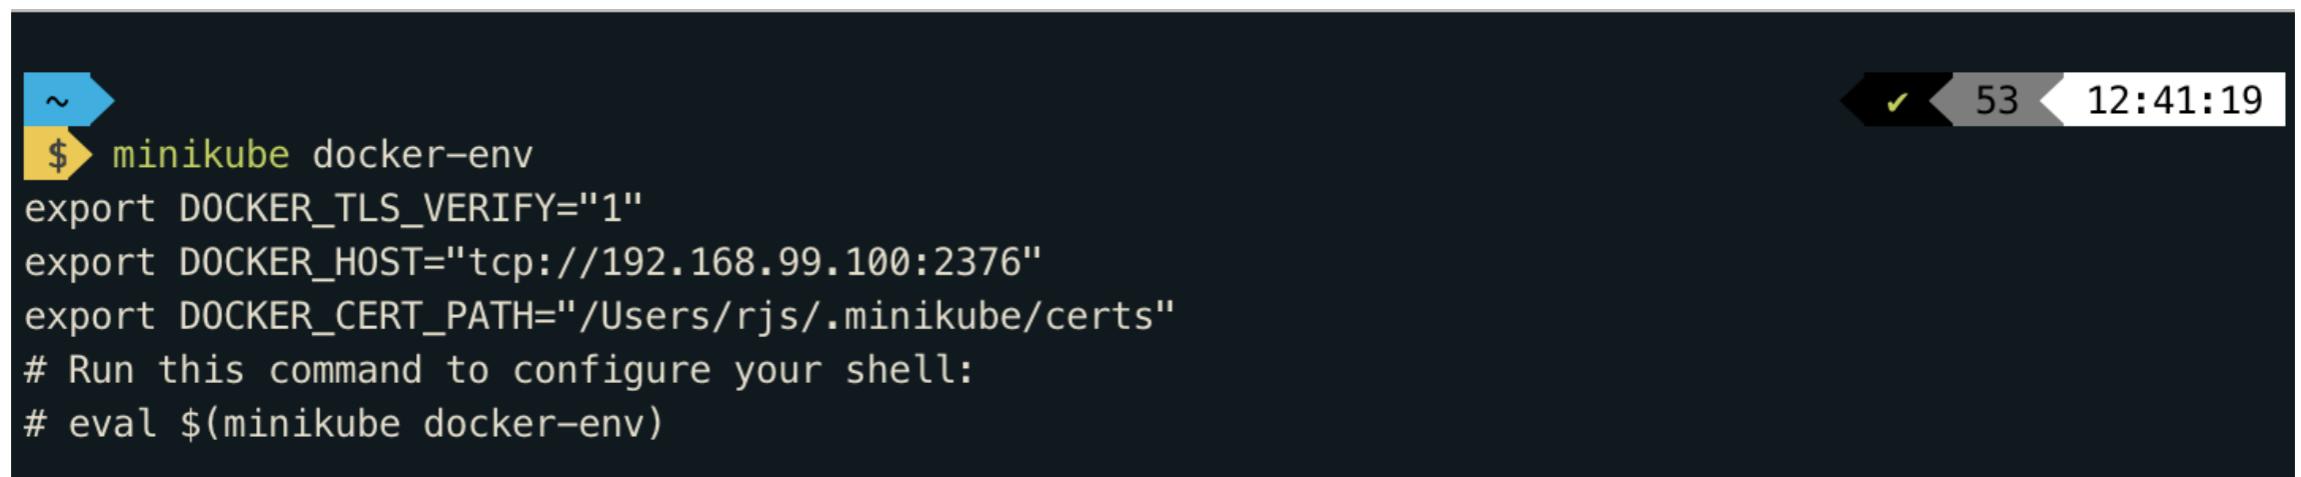
```
~ $ minikube start
😄 minikube v1.4.0 on Darwin 10.14.6
💡 Tip: Use 'minikube start -p <name>' to create a new cluster, or 'minikube delete' to delete this one.
🔄 Starting existing virtualbox VM for "minikube" ...
⏳ Waiting for the host to be provisioned ...
🐳 Preparing Kubernetes v1.16.0 on Docker 18.09.9 ...
🔄 Relaunching Kubernetes using kubeadm ...
⏳ Waiting for: apiserver proxy etcd scheduler controller dns
🏃 Done! kubectl is now configured to use "minikube"
```

- ▶ `docker image ls`



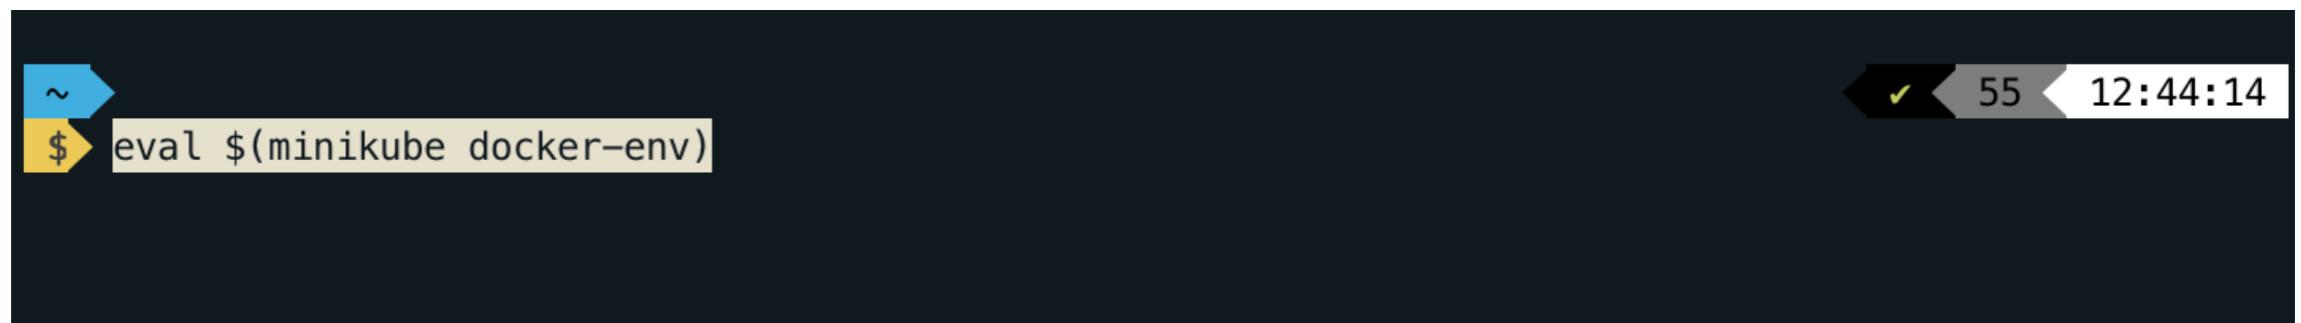
```
~ $ docker image ls
Cannot connect to the Docker daemon at unix:///var/run/docker.sock. Is the docker daemon running?
```

- ▶ We need to configure docker to tell it that we are talking to the Minikube environment rather than the tradition docker running on its virtual machine - **minikube docker-env**



```
~ $ minikube docker-env
export DOCKER_TLS_VERIFY="1"
export DOCKER_HOST="tcp://192.168.99.100:2376"
export DOCKER_CERT_PATH="/Users/rjs/.minikube/certs"
# Run this command to configure your shell:
# eval $(minikube docker-env)
```

Environment variable needed to setup in order to use the docker inside Minikube



```
~ $ eval $(minikube docker-env)
```

```
~$ echo $DOCKER_HOST  
tcp://192.168.99.100:2376
```

✓ 59 15:08:09

```
~$ minikube docker-env  
export DOCKER_TLS_VERIFY="1"  
export DOCKER_HOST="tcp://192.168.99.100:2376"  
export DOCKER_CERT_PATH="/Users/rjs/.minikube/certs"  
# Run this command to configure your shell:  
# eval $(minikube docker-env)
```

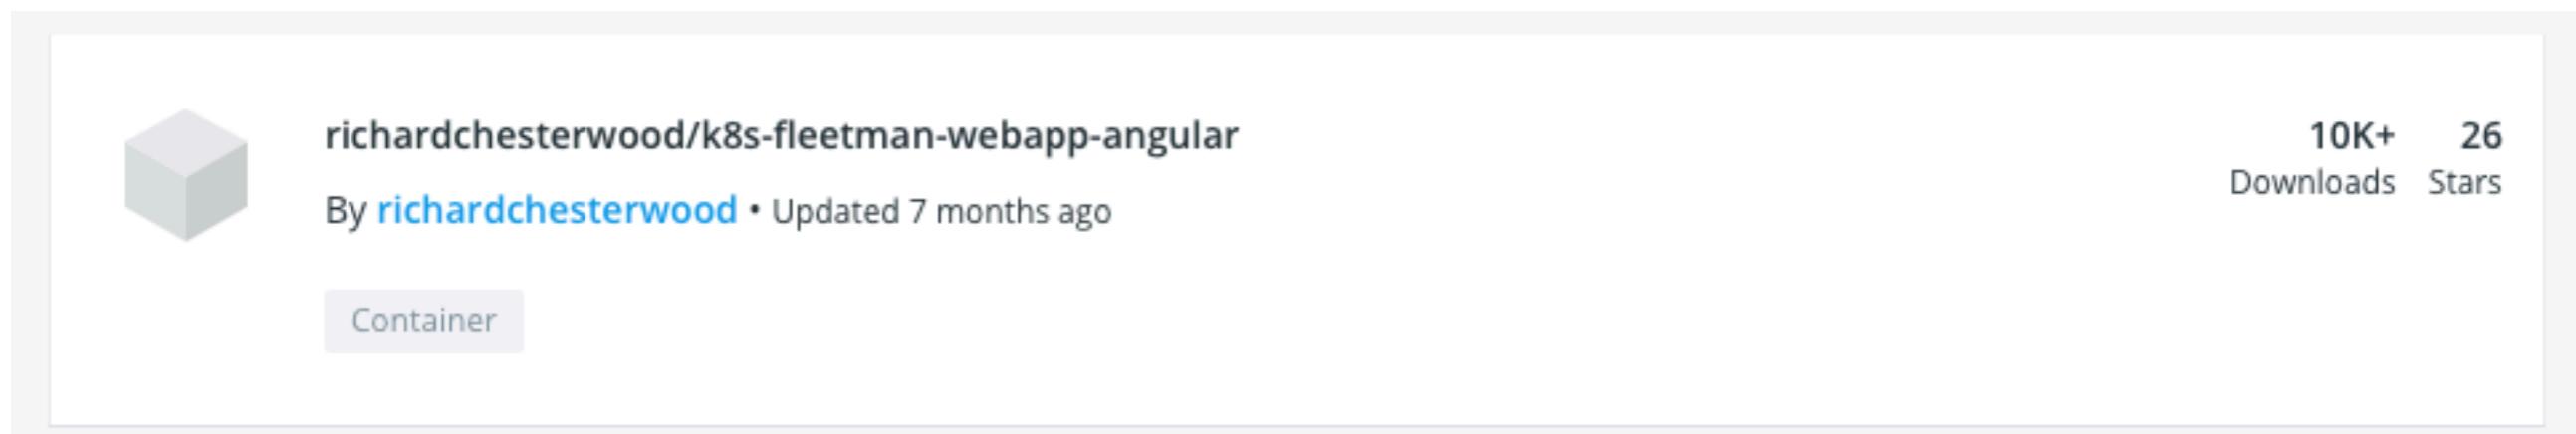
✓ 53 12:41:19

```
~ $ docker image ls
REPOSITORY           TAG      IMAGE ID      CREATED
SIZE
k8s.gcr.io/kube-apiserver   v1.16.0   b305571ca60a   4 weeks ago
  217MB
k8s.gcr.io/kube-controller-manager   v1.16.0   06a629a7e51c   4 weeks ago
  163MB
k8s.gcr.io/kube-proxy     v1.16.0   c21b0c7400f9   4 weeks ago
  86.1MB
k8s.gcr.io/kube-scheduler   v1.16.0   301ddc62b80b   4 weeks ago
  87.3MB
```

If you don't use linux this configuration is much better in term of performance

- ▶ This ways you don't have two VMs running (docker toolbox + minikube)

- ▶ Go to [hub.docker.com](https://hub.docker.com/r/richardchesterwood/k8s-fleetman-webapp-angular) and find the following image



- ▶ How can we deploy this?
- ▶ First we need to see the releases

release0 9.76 MBLast updated **a year ago** by [richardchesterwood](#)

DIGEST

[9b98fec20772](#)

ARCHITECTURE

amd64

OS

SIZE

9.76 MB

release0-5 9.76 MBLast updated **a year ago** by [richardchesterwood](#)

DIGEST

[636f57ab20c3](#)

ARCHITECTURE

amd64

OS

SIZE

9.76 MB

release2 25.07 MBLast updated **a year ago** by [richardchesterwood](#)

DIGEST

[ed7d720878ac](#)

ARCHITECTURE

amd64

OS

SIZE

25.07 MB

release1 25.07 MBLast updated **a year ago** by [richardchesterwood](#)

DIGEST

[46bef951fc3c](#)

ARCHITECTURE

amd64

OS

SIZE

25.07 MB

- ▶ First download the image for our computer (optional step)

```
~ $ docker image pull richardchesterwood/k8s-fleetman-webapp-angular:release0-5
```

```
~ $ docker image pull richardchesterwood/k8s-fleetman-webapp-angular:release0-5
release0-5: Pulling from richardchesterwood/k8s-fleetman-webapp-angular
Digest: sha256:636f57ab20c367bcab78d1994d4990b72bc8f69686fc716bc06386c2894fb8dd
Status: Image is up to date for richardchesterwood/k8s-fleetman-webapp-angular:release0-5
docker.io/richardchesterwood/k8s-fleetman-webapp-angular:release0-5
```

- ▶ You should see it in your images list -
`docker image ls`

► Let's run it!



~ 77 17:12:20

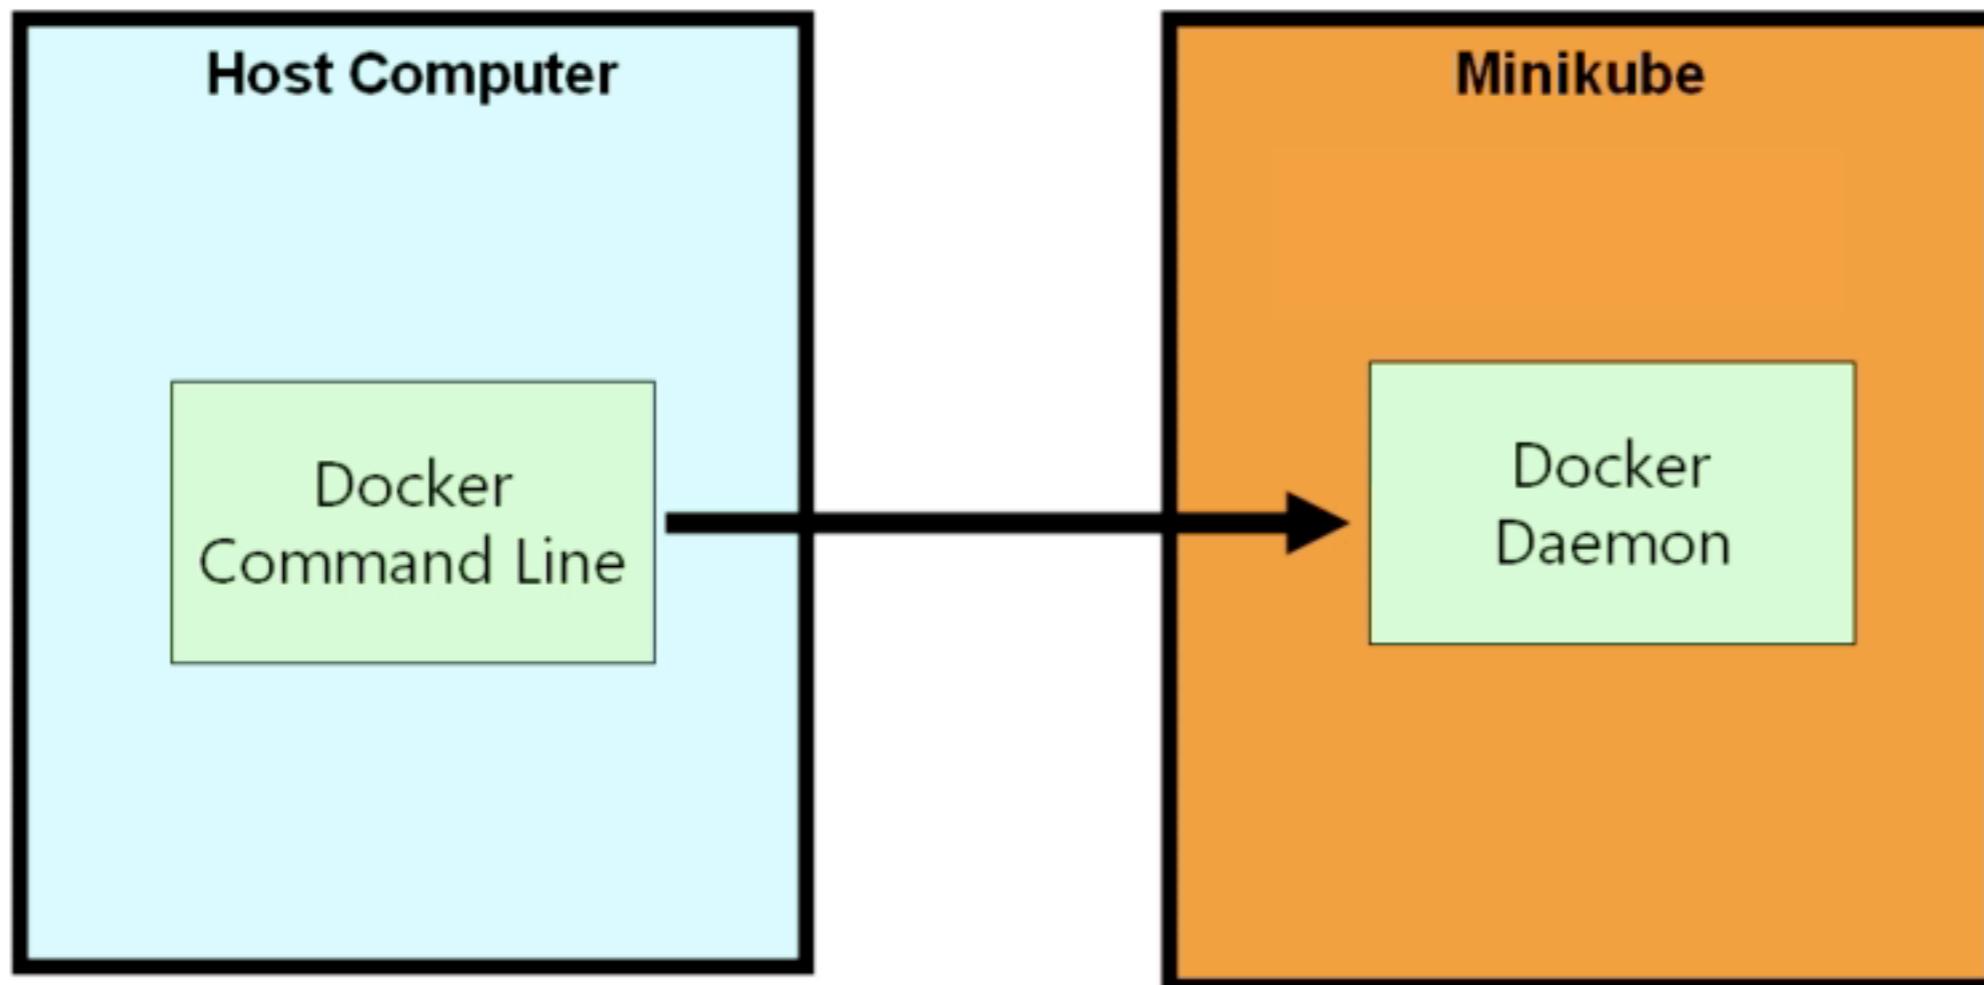
```
$ docker container run -p 8080:80 -d richardchesterwood/k8s-fleetman-webapp-angular:release0-5
cc4be1f5ab9d82be5f8be5c16a2ff0e73155c073852b413c2f9d740ecfc196c3
```

- 8080 - the port we want to expose to the outside world
- Internally the traffic will be routed to port 80

- ▶ Let's see if the container is running

```
~ $ docker container ls
CONTAINER ID        IMAGE                               COMMAND      CREATED          STATUS              PORTS
-----  -----
cc4be1f5ab9d        richardchesterwood/k8s-fleetman-webapp-angular:release0-5    "nginx -g 'daemon off...'"   33 seconds ago
o          Up 32 seconds           0.0.0.0:8080->80/tcp   dazzling_goldwasser
d1a55236610d        25be029ad6d0                  "nginx -g 'daemon off...'"   6 hours ago
          Up 6 hours
1dac098d046_3
4f2eddd97539        25be029ad6d0                  "nginx -g 'daemon off...'"   6 hours ago
          Up 6 hours
883f263bccb_3
4e416a36e6d0        383867b75fd2                "docker-entrypoint.s..."   6 hours ago
```

- ▶ Let's try <http://localhost:8080>



- ▶ The Minikube docker daemon can't be accessible from our computer

- ▶ Let's figure out the Minikube ip

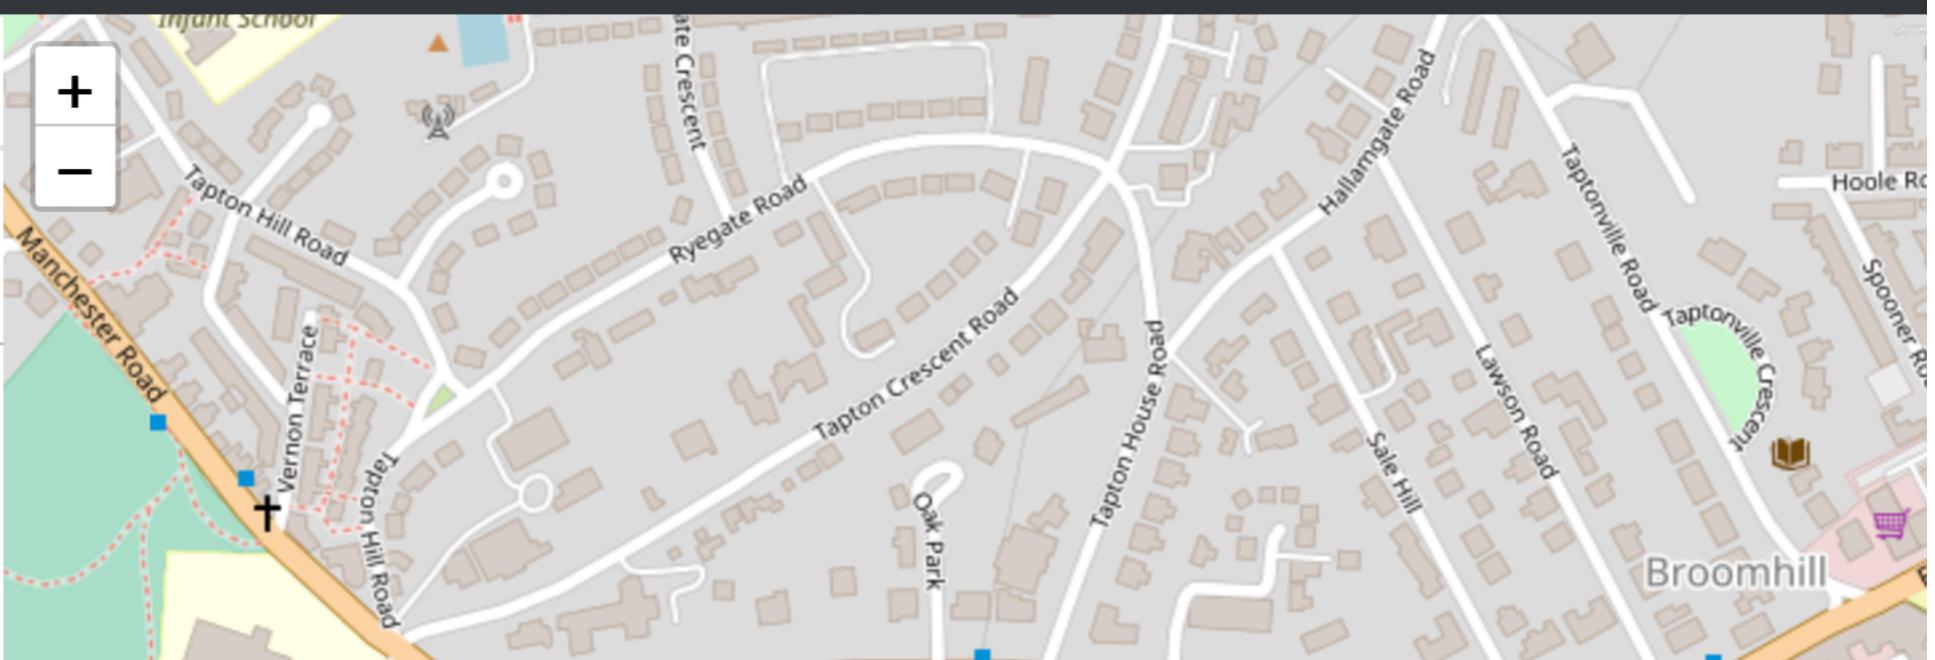
```
~ $ minikube ip
192.168.99.100
```

- ▶ Let's try <http://192.168.99.100:8080>

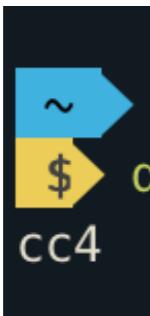
Fleet Management System PROTOTYPE. Release 0.5

Name	Last seen	Speed mph
------	-----------	-----------

Live vehicle updates will appear here.
Once we've implemented it!



- ▶ How to stop it?



```
$ docker container stop cc4
cc4
```

- ▶ And remove it



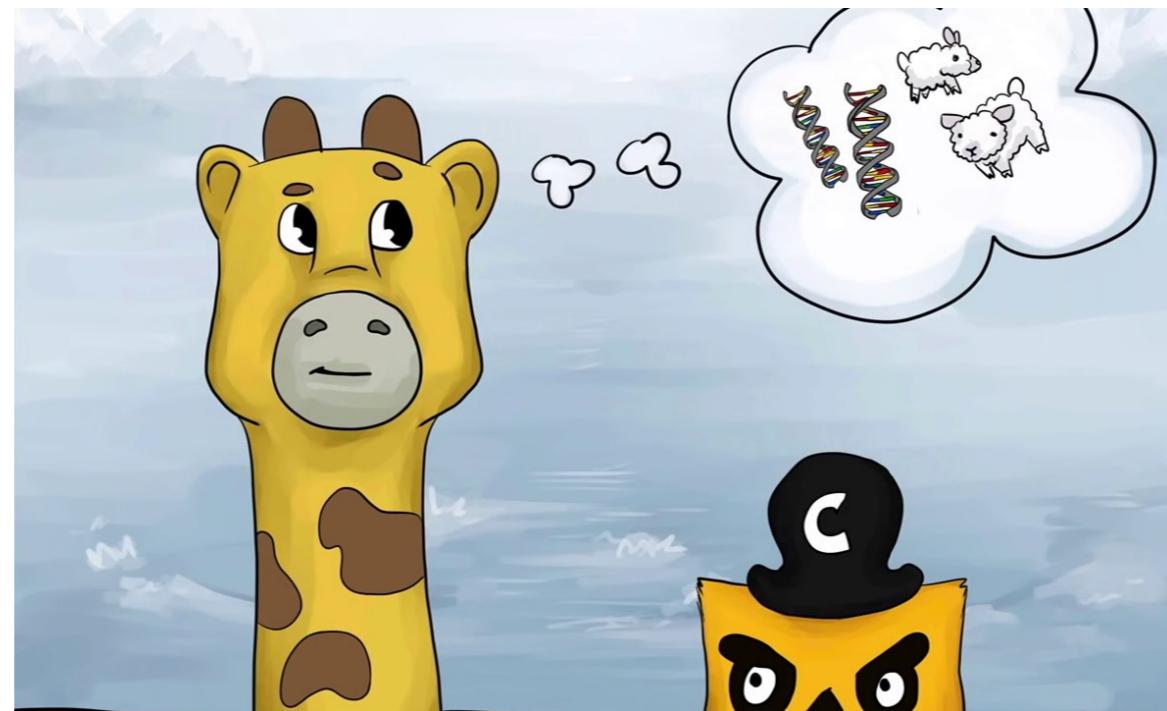
```
$ docker container rm cc4
cc4
```

- ▶ Let's try <http://192.168.99.100:8080>

Safari Can't Connect to the Server

Safari can't open the page "192.168.99.100:8080" because Safari can't connect to the server "192.168.99.100".

- <https://martinfowler.com/articles/microservices.html>
- <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
- <http://microservices.io>



- <https://www.youtube.com/watch?v=Q4W8Z-D-gcQ>

P.PORTO

Polytechnic of
Porto

ESTG - School of
Management and Technology