

P.PORTO

Methods and Techniques for
Software Development

Ricardo Santos | 2019/2020
rjs@estg.ipp.pt

P.PORTO

From Monoliths to Microservices

Lesson 4

Methods and Techniques for
Software Development

2019/2020

P.PORTO

Evoluir Aplicações Monolíticas para Microserviços

Aula 4

Methods and Techniques for
Software Development

2019/2020

It is very difficult to manage complex enterprise systems

- ▶ Understanding how any one change will affect the entire system
- ▶ New developers spend months learning the system's codebase before they can even begin to work on it
- ▶ Even the most knowledgeable of development teams is fearful of making changes or adding new code that would disrupt operation in some unforeseen way

Microservices can turn this scenario, enabling a new, more agile world in which developers and operations teams work hand in hand to deliver small, loosely coupled bundles of software quickly and safely

*The key challenge is to design and develop
the integration between the existing system
and the new microservices*





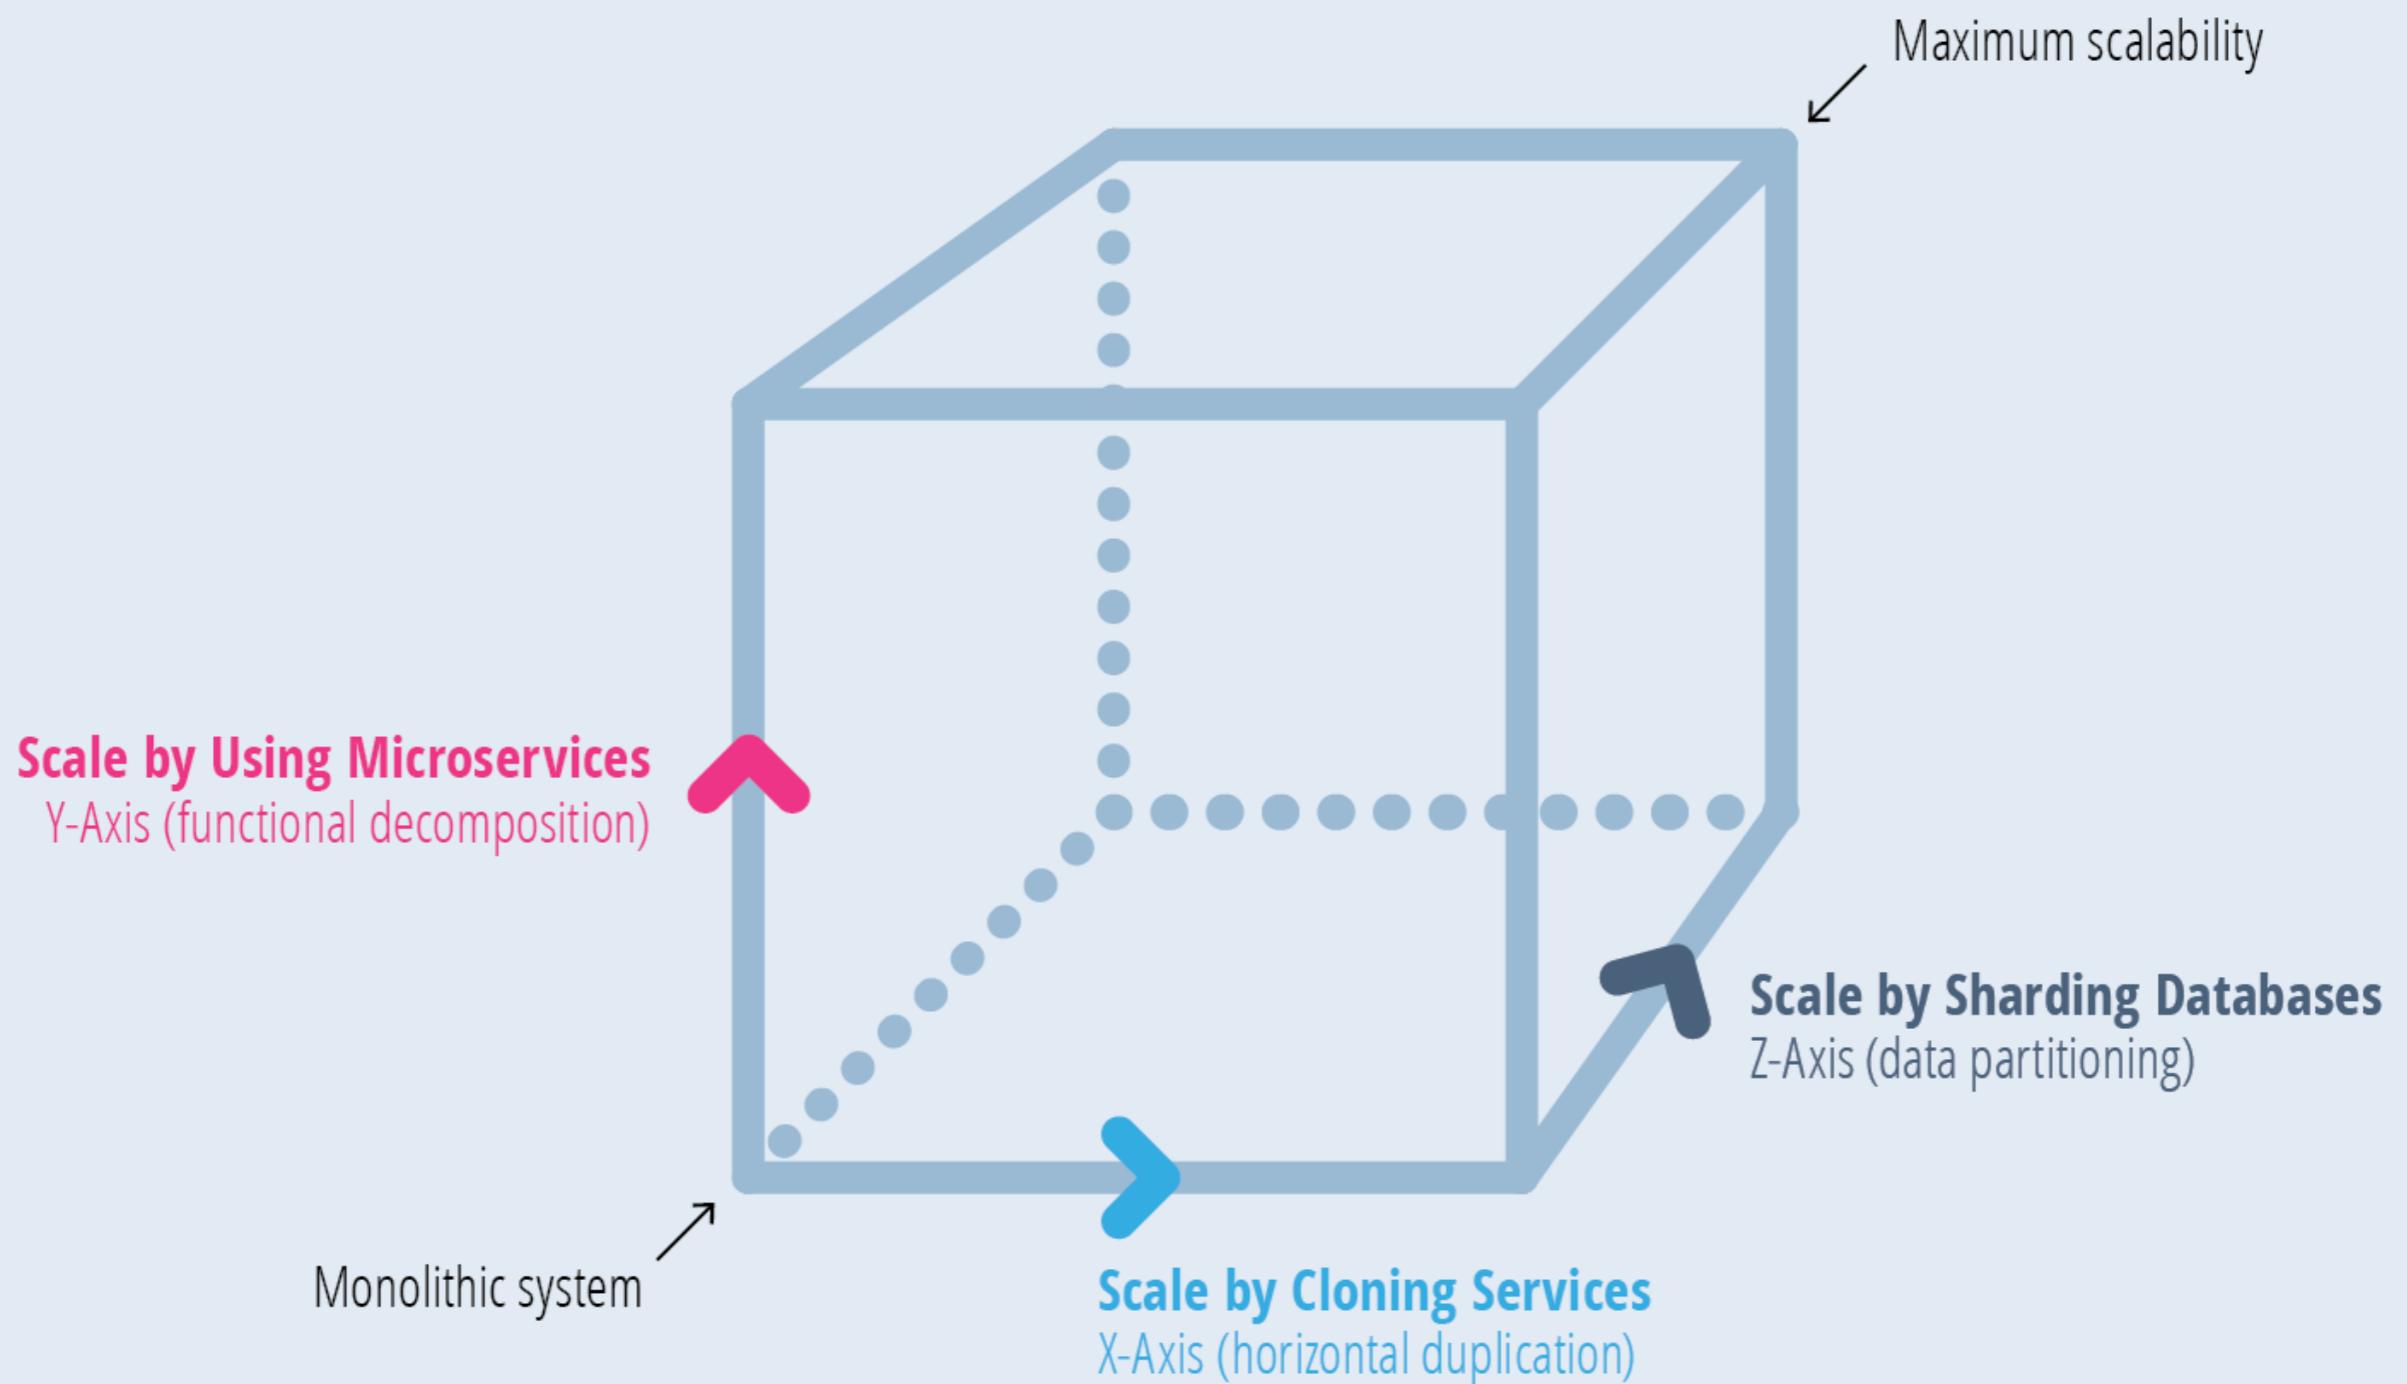
If you can't feed a
team with two
pizzas, it's too large.

Startup Quote!

JEFF BEZOS

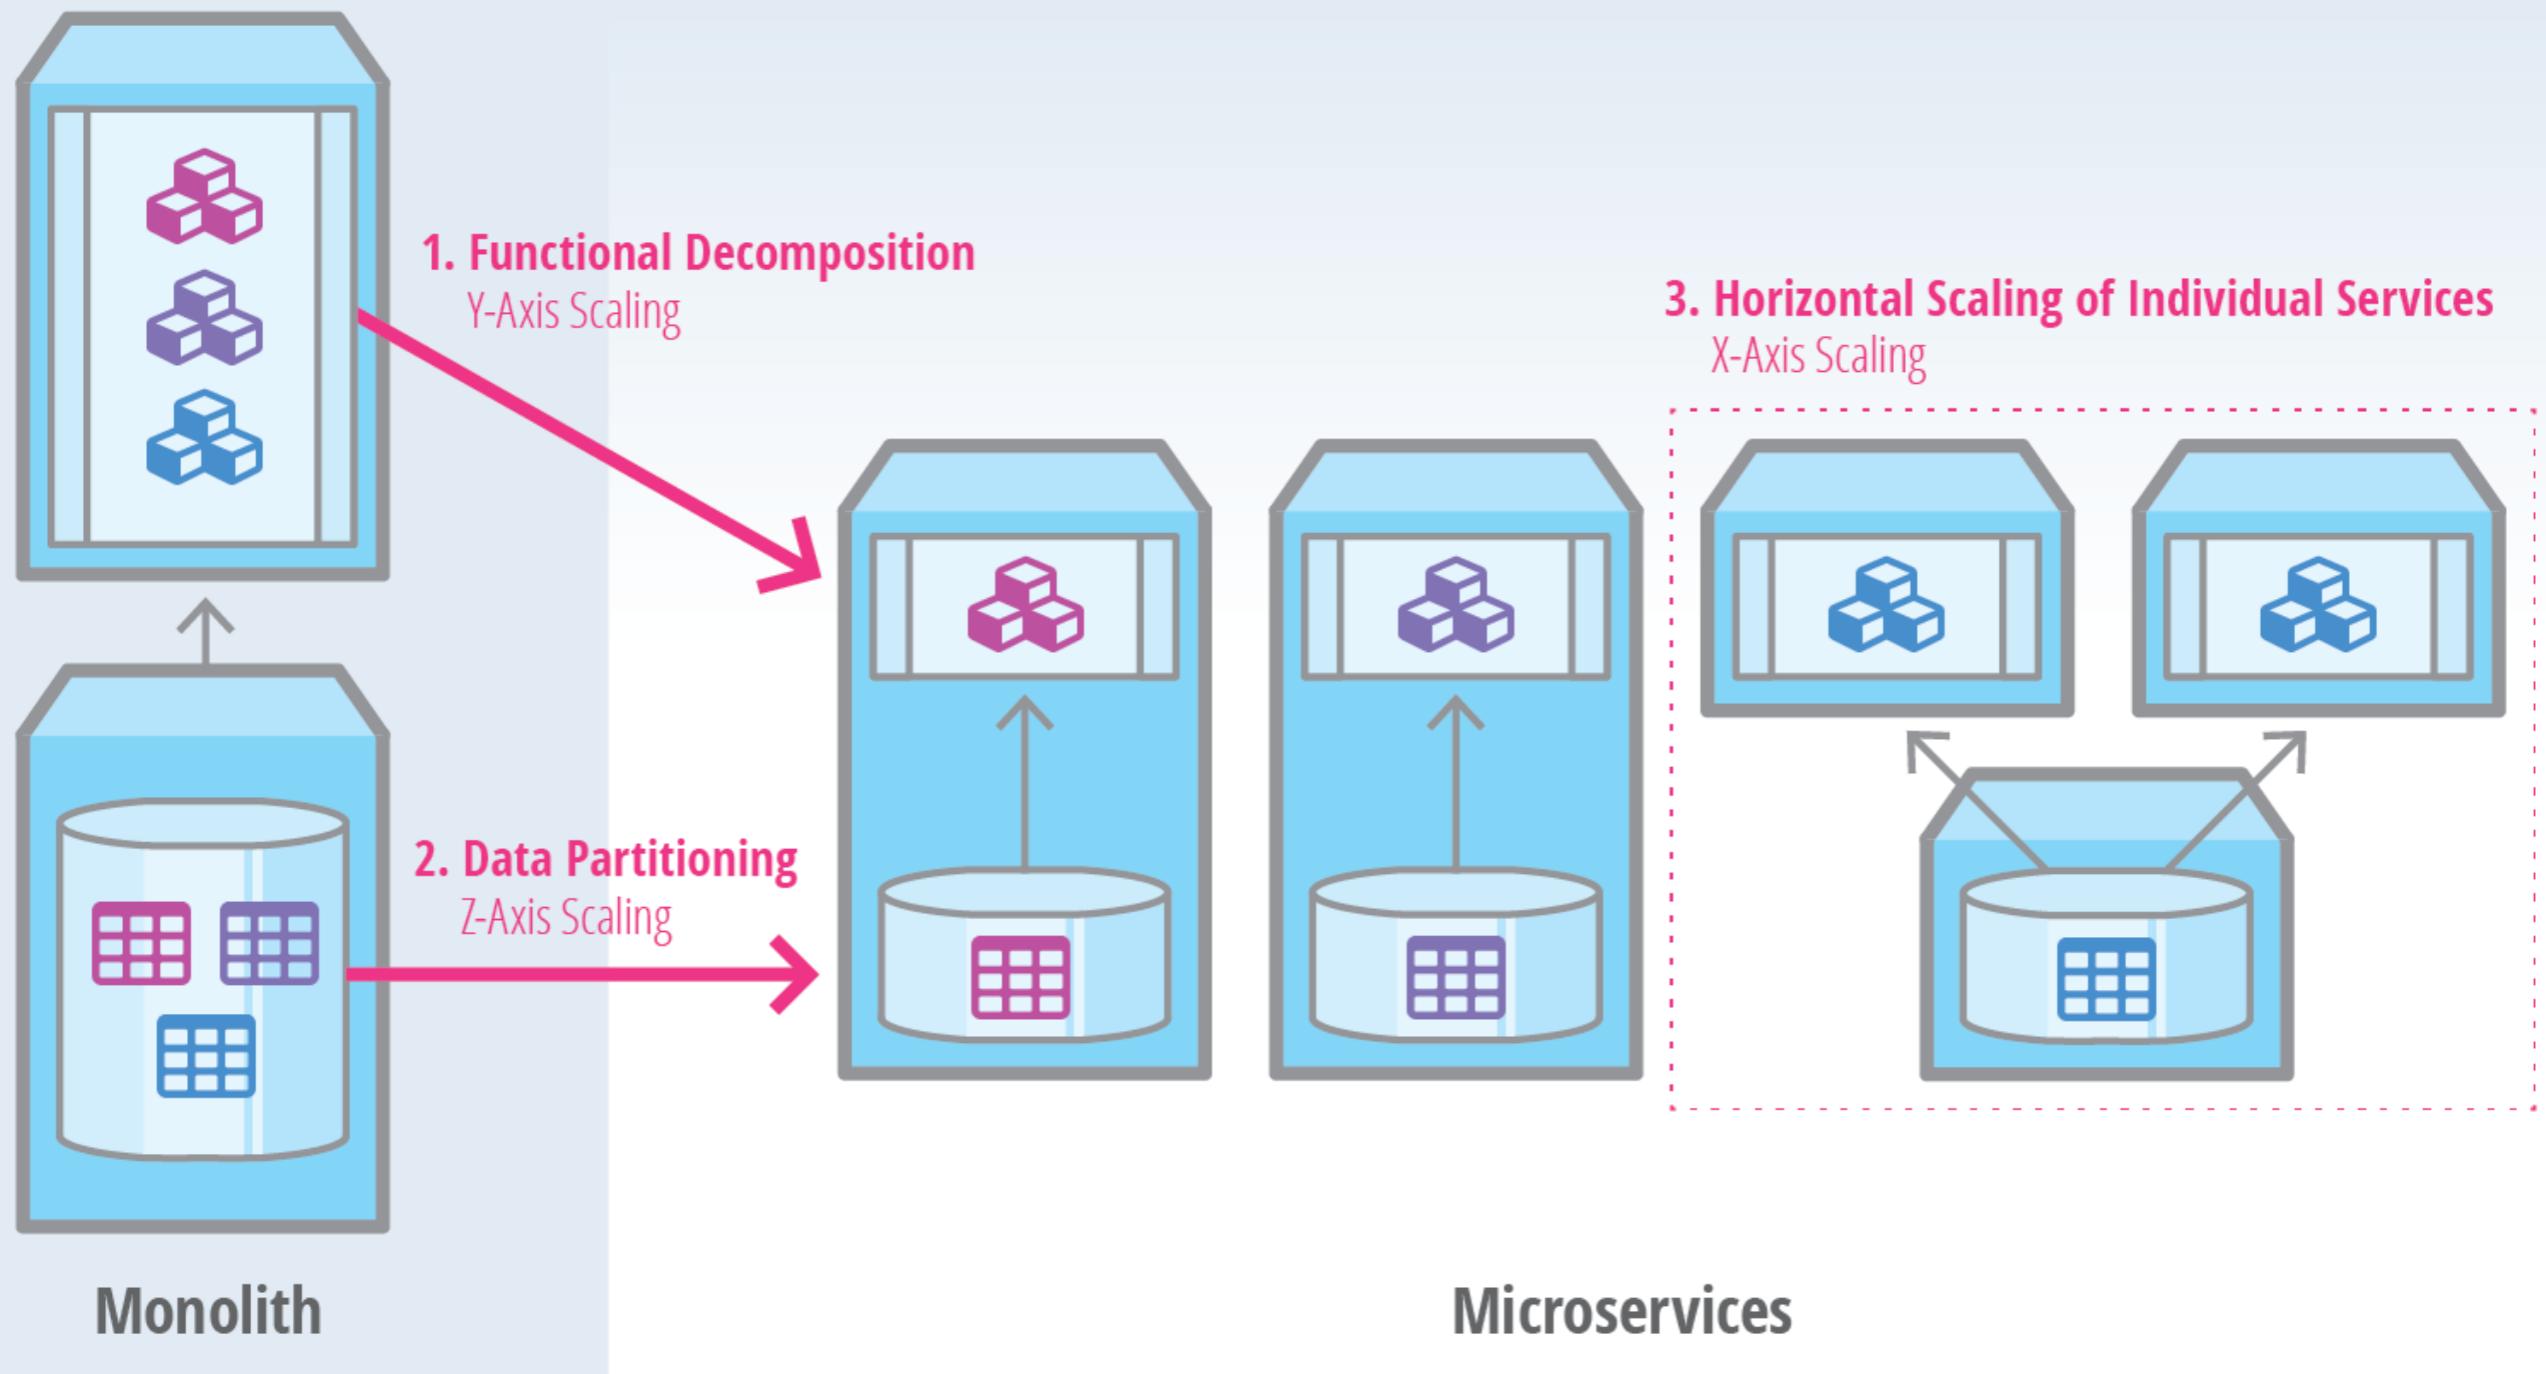
FOUNDER, AMAZON

The Scale Cube and Microservices: 3 Dimensions to Scaling



Scaling With Microservices

Microservice architectures have 3 dimensions of scalability



Source: The New Stack. Based on a diagram appearing in "Microservices," by James Lewis & Martin Fowler.

THE NEW STACK

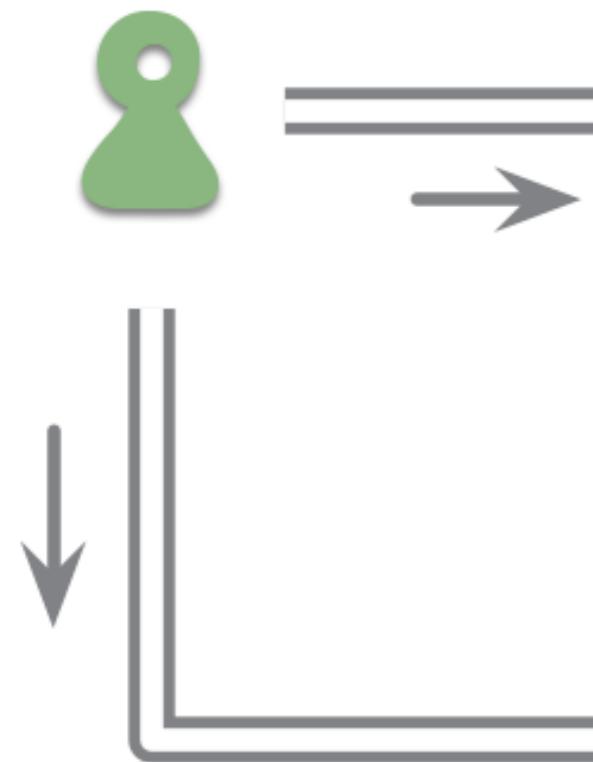
Monolith First

you shouldn't start a new project with microservices, even if you're sure your application will be big enough to make it worthwhile

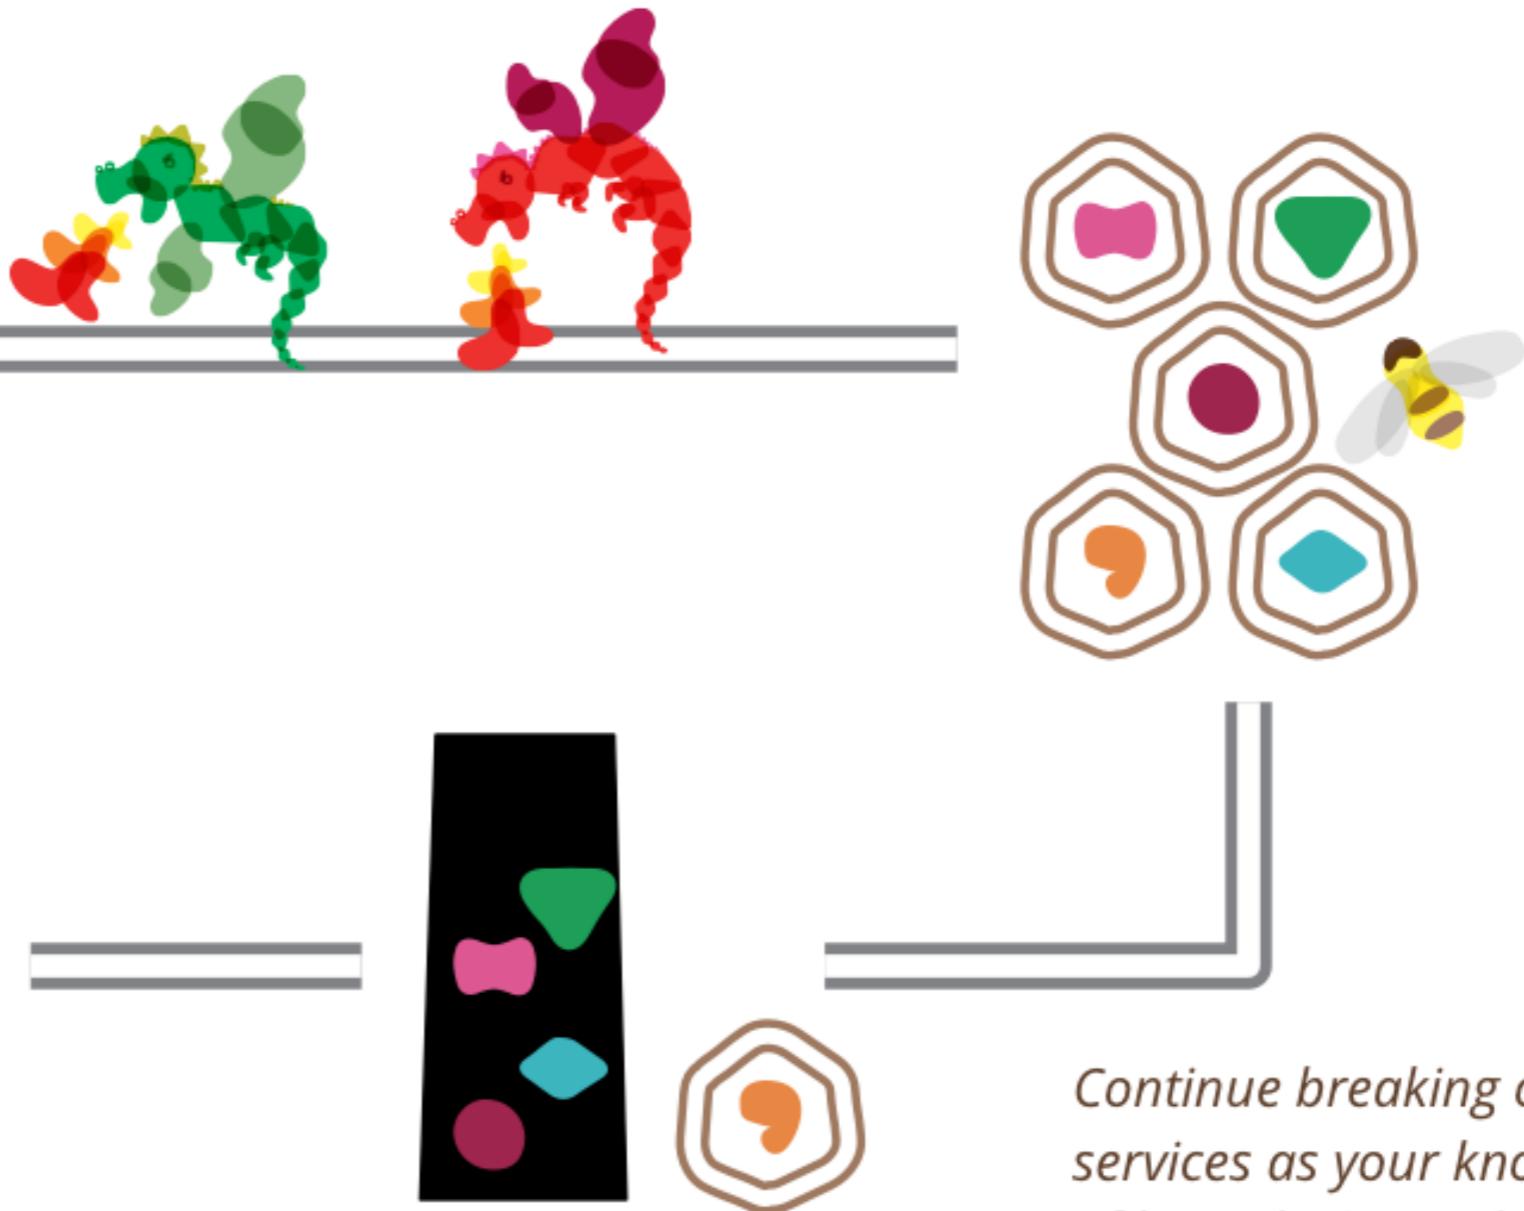
Almost all the successful microservice stories have started with a monolith that got too big and was broken up

Almost all the cases where I've heard of a system that was built as a microservice system from scratch, it has ended up in serious trouble

Going directly to a microservices architecture is risky



A monolith allows you to explore both the complexity of a system and its component boundaries



As complexity rises start breaking out some microservices

Continue breaking out services as your knowledge of boundaries and service management increases

Refactoring to Microservices

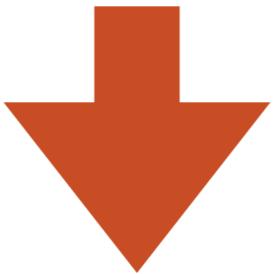
- ▶ The process of transforming a monolithic application into microservices is a **form of application modernization**
- ▶ That is something that developers have been doing for decades
- ▶ There are some ideas that we can reuse when refactoring an application into microservices

Big Bang Strategy

- ▶ you focus all of your development efforts on building a new microservices-based application from scratch
- ▶ it is extremely risky and will likely end in failure

the only thing a Big Bang rewrite guarantees is a Big Bang!"

you should incrementally refactor your monolithic application



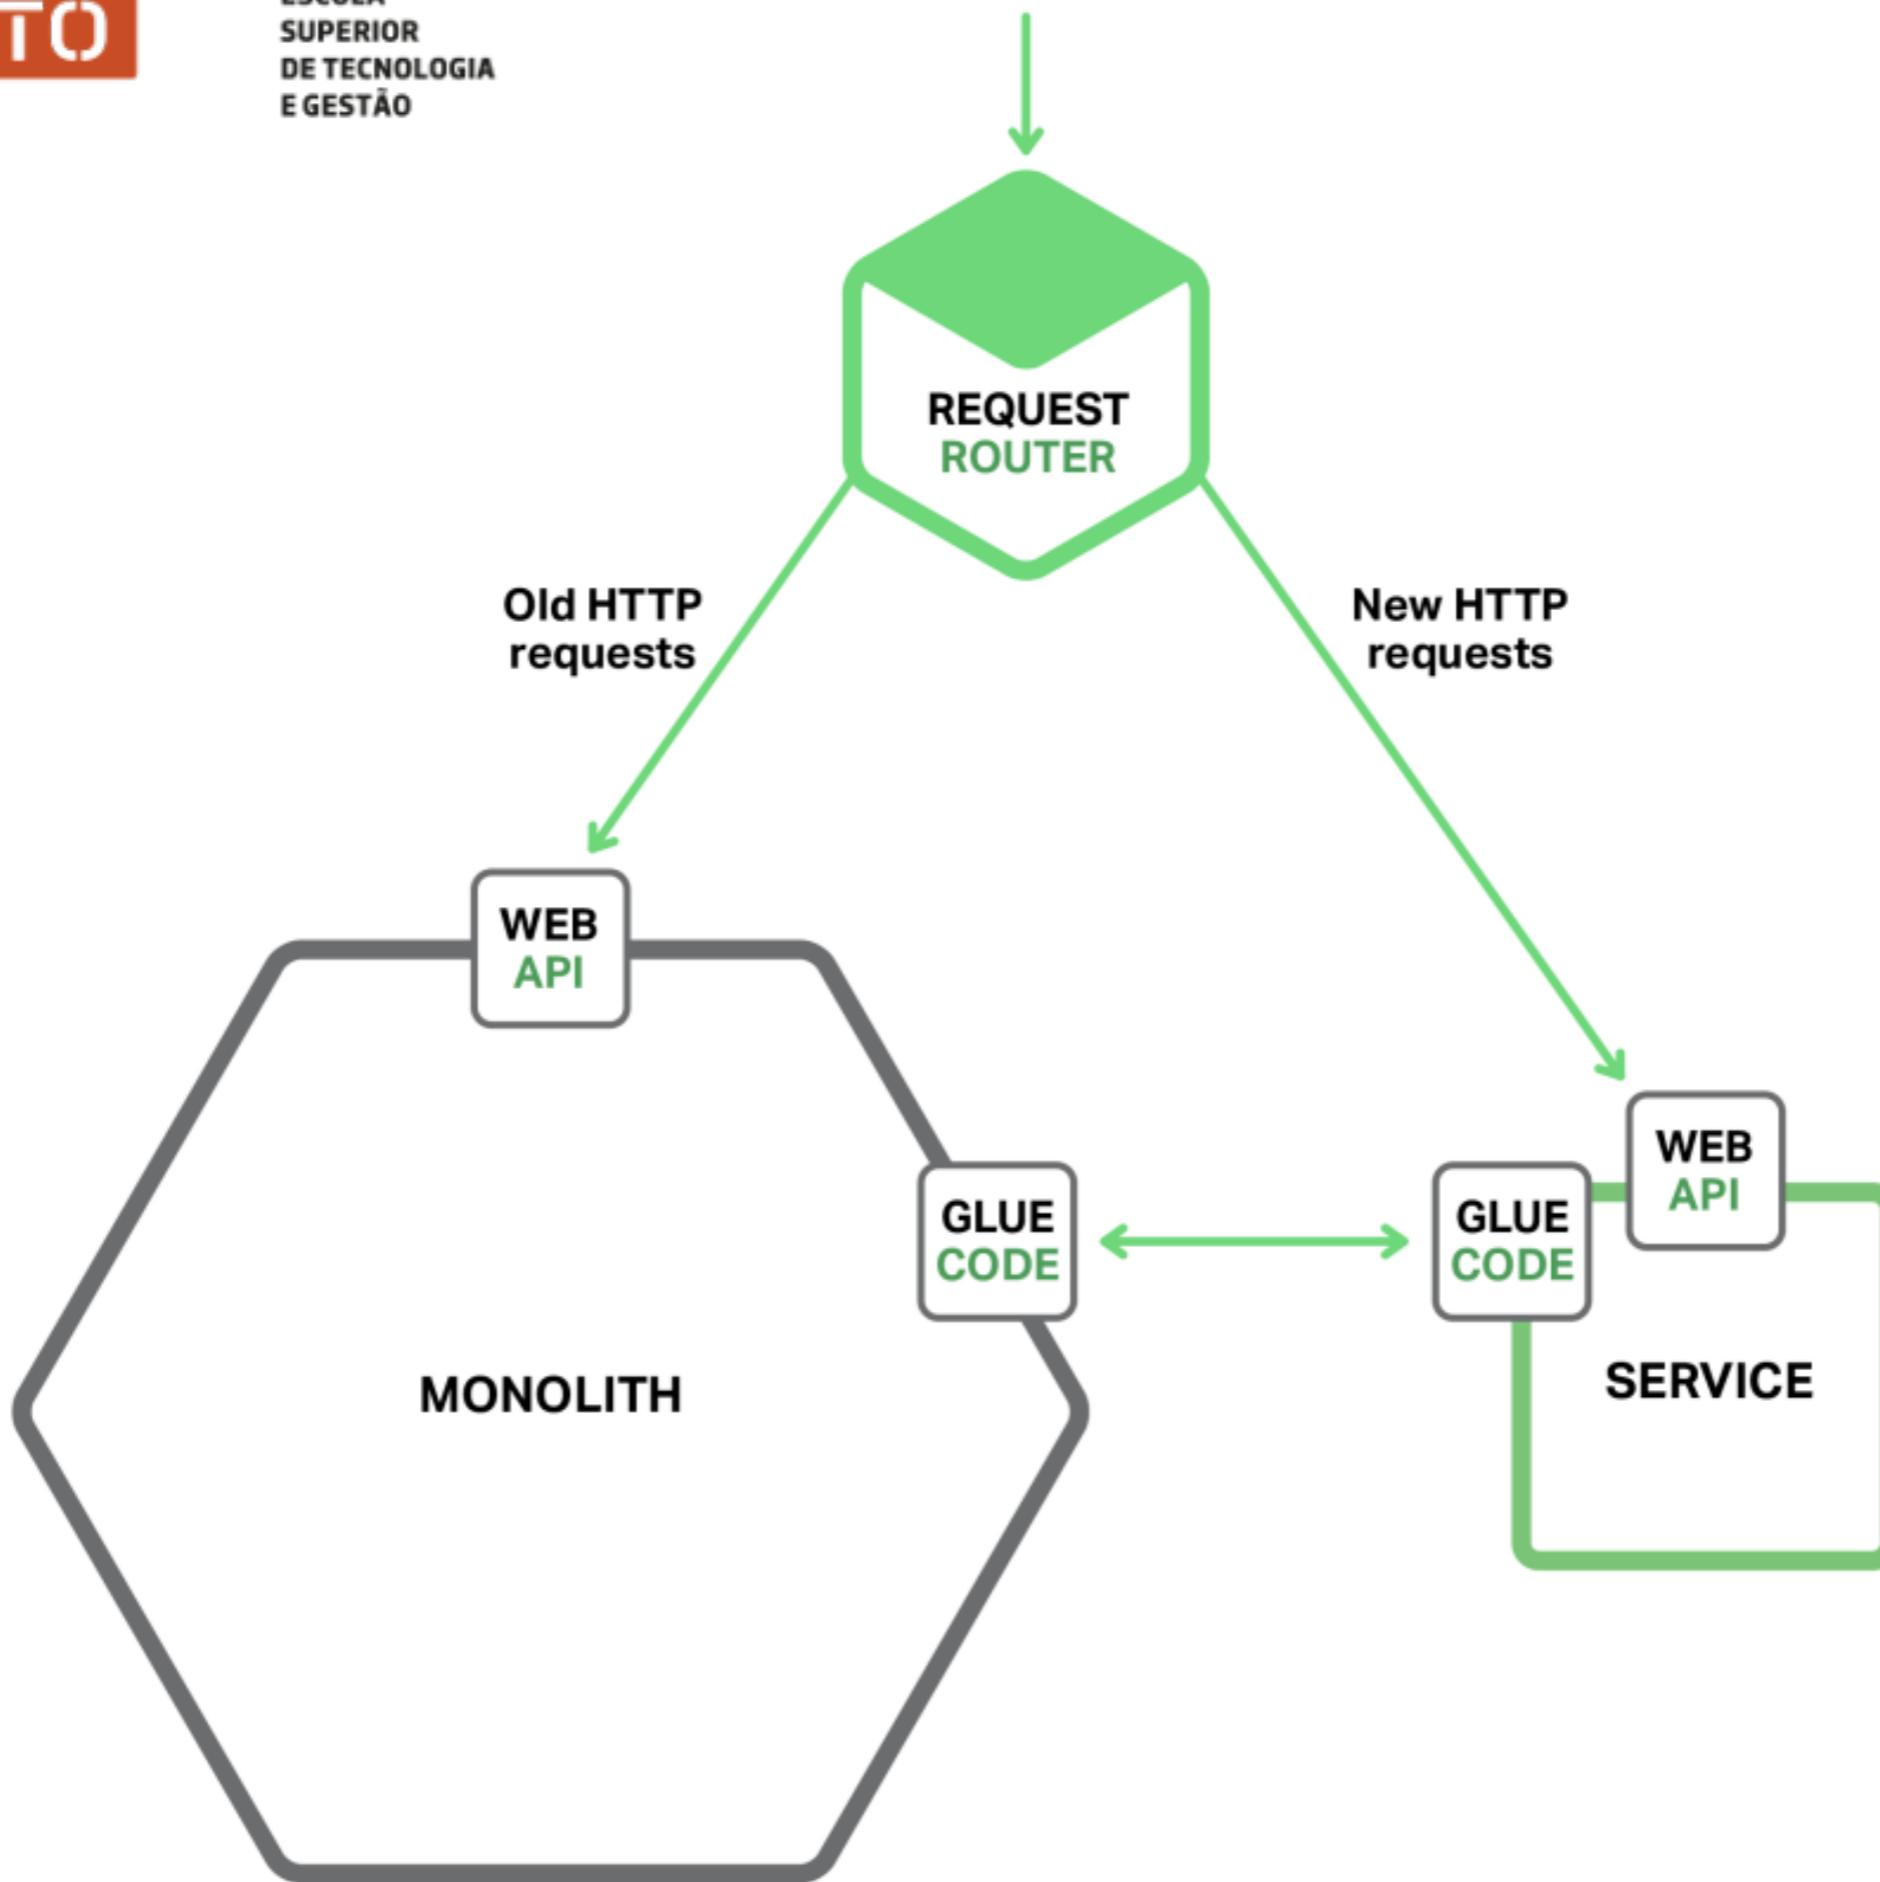
Strangler Application Pattern

Martin Fowler, 2004

Strategy I - Stop Digging

- ▶ you should **stop making the monolith bigger**
- ▶ when you are implementing new functionality you should **not add more code to the monolith**
- ▶ put that **new code in a standalone microservice**

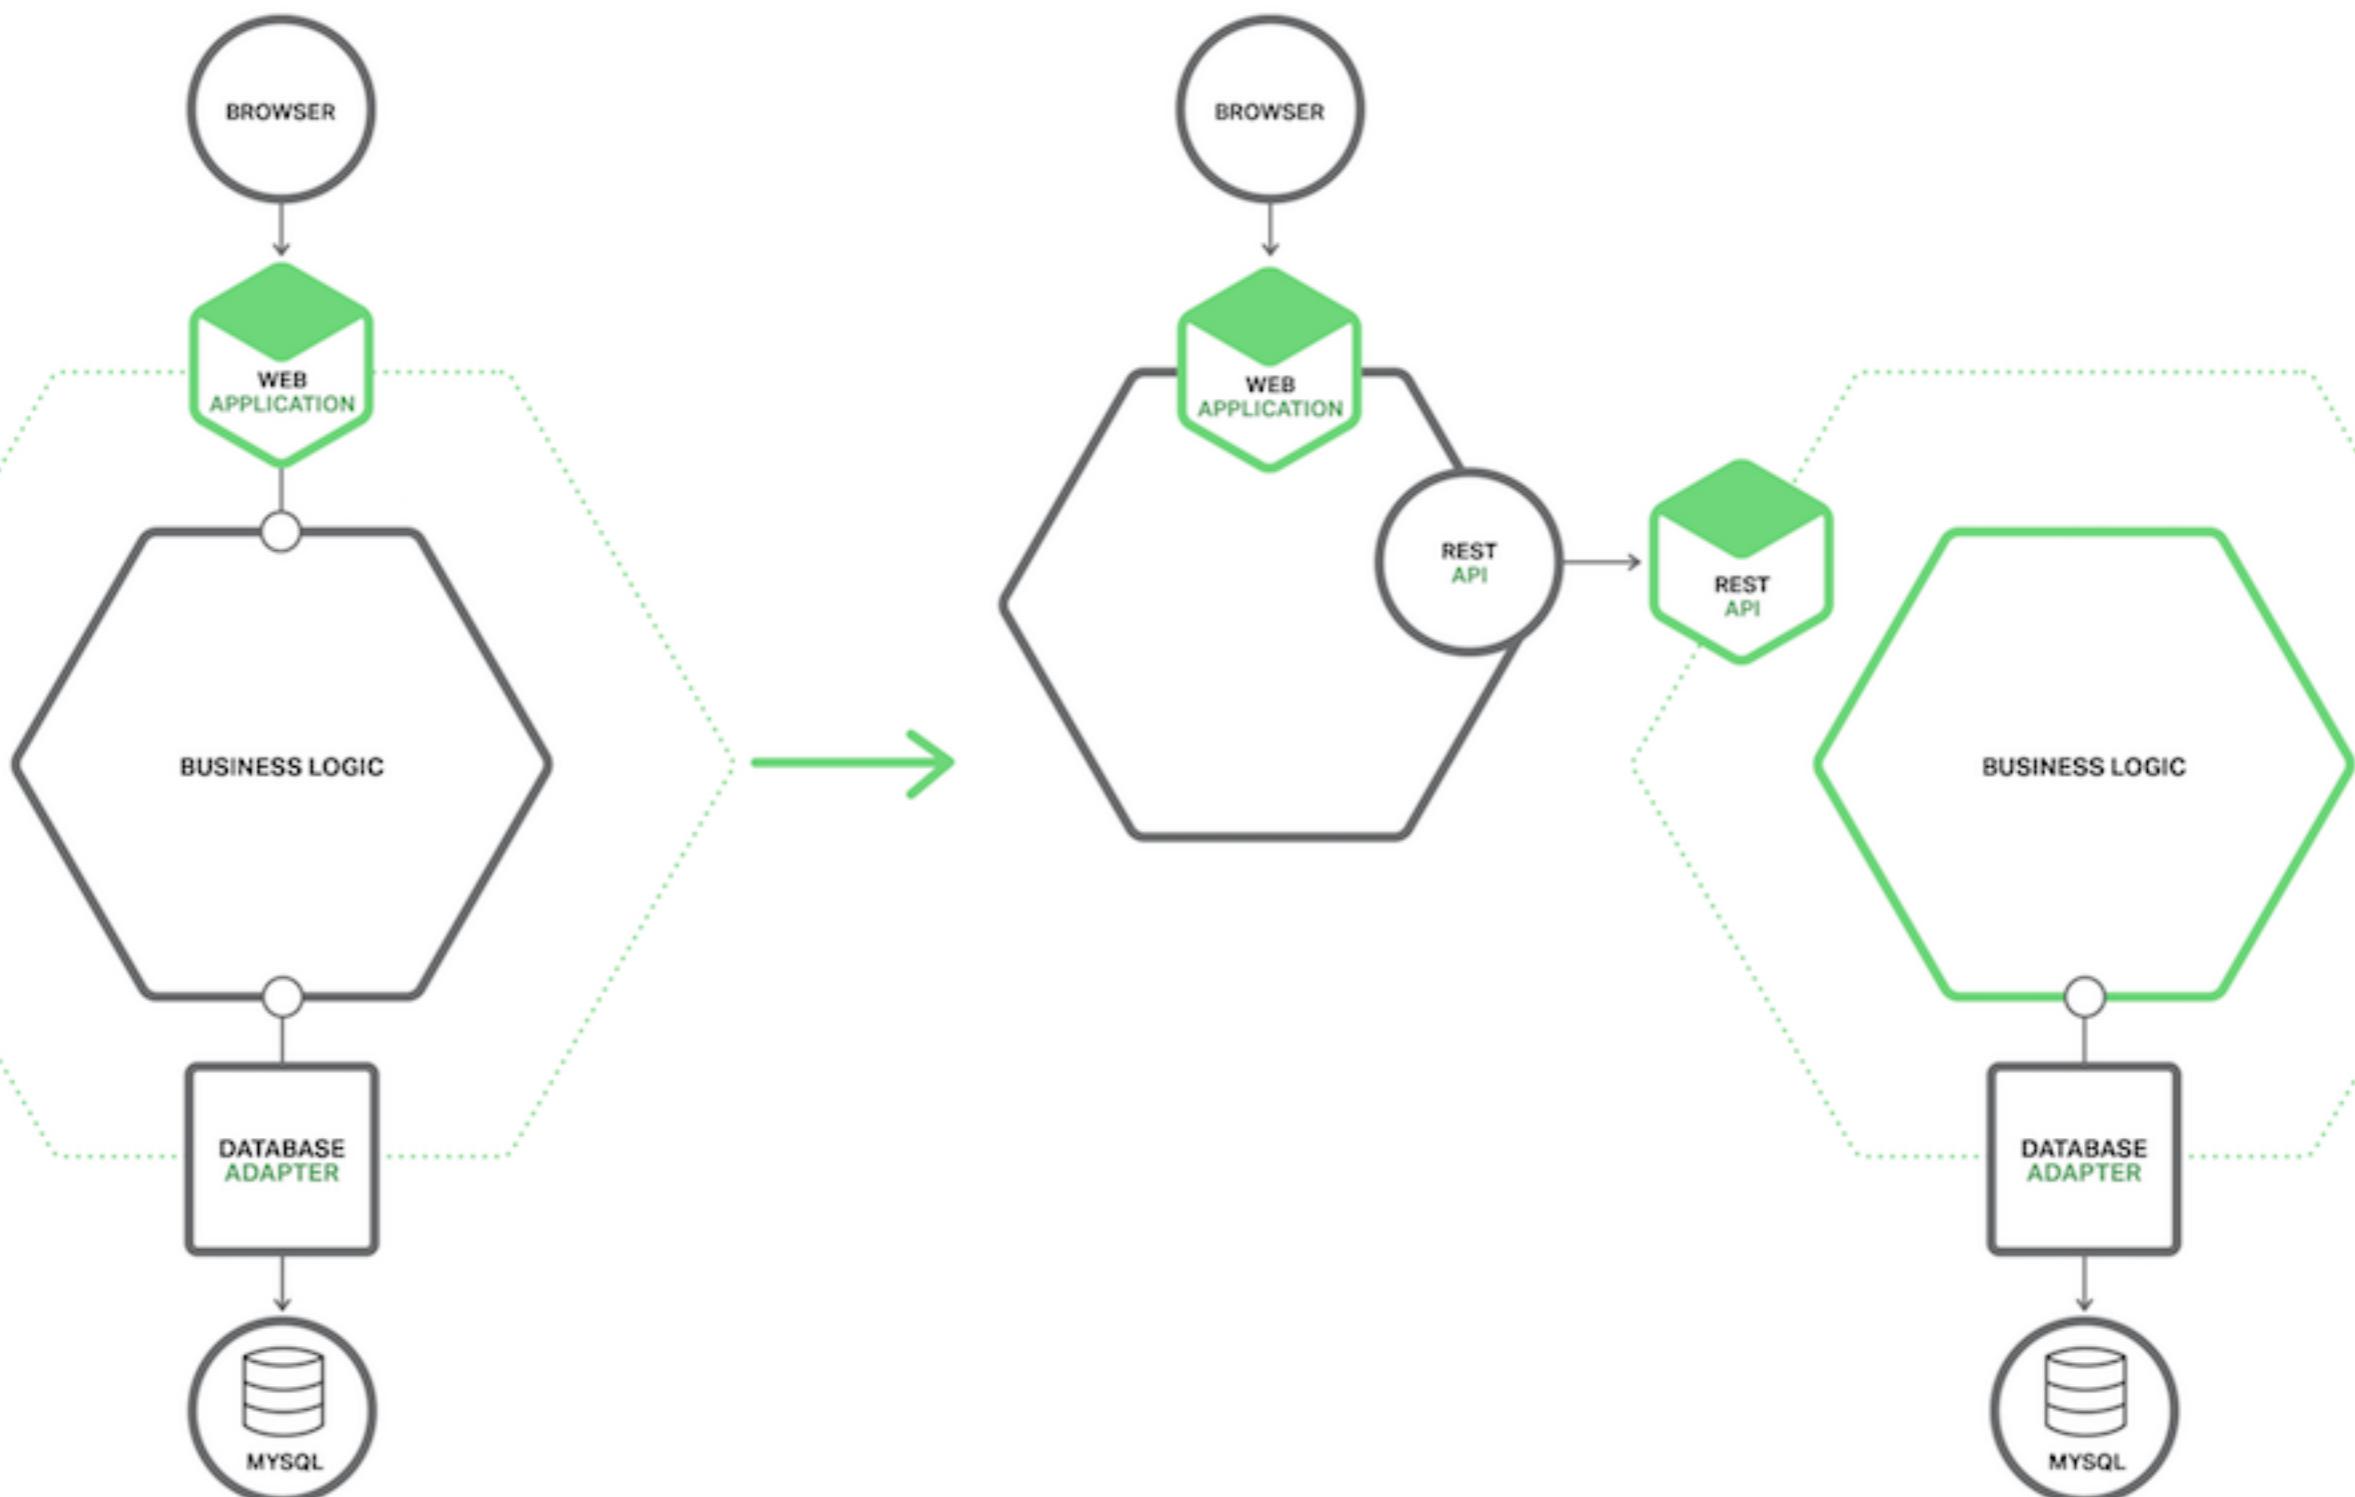
HTTP REQUEST



- ▶ Monolith and new microservice only occasionally live next to each other without a need of accessing their data
- ▶ To ensure that they can talk one to each other, glue code is needed and it can be provided in three ways:
 - Allowing direct database access
 - Maintaining its own copy of the data
 - Creating a remote API

Strategy 2 - Split Frontend and Backend

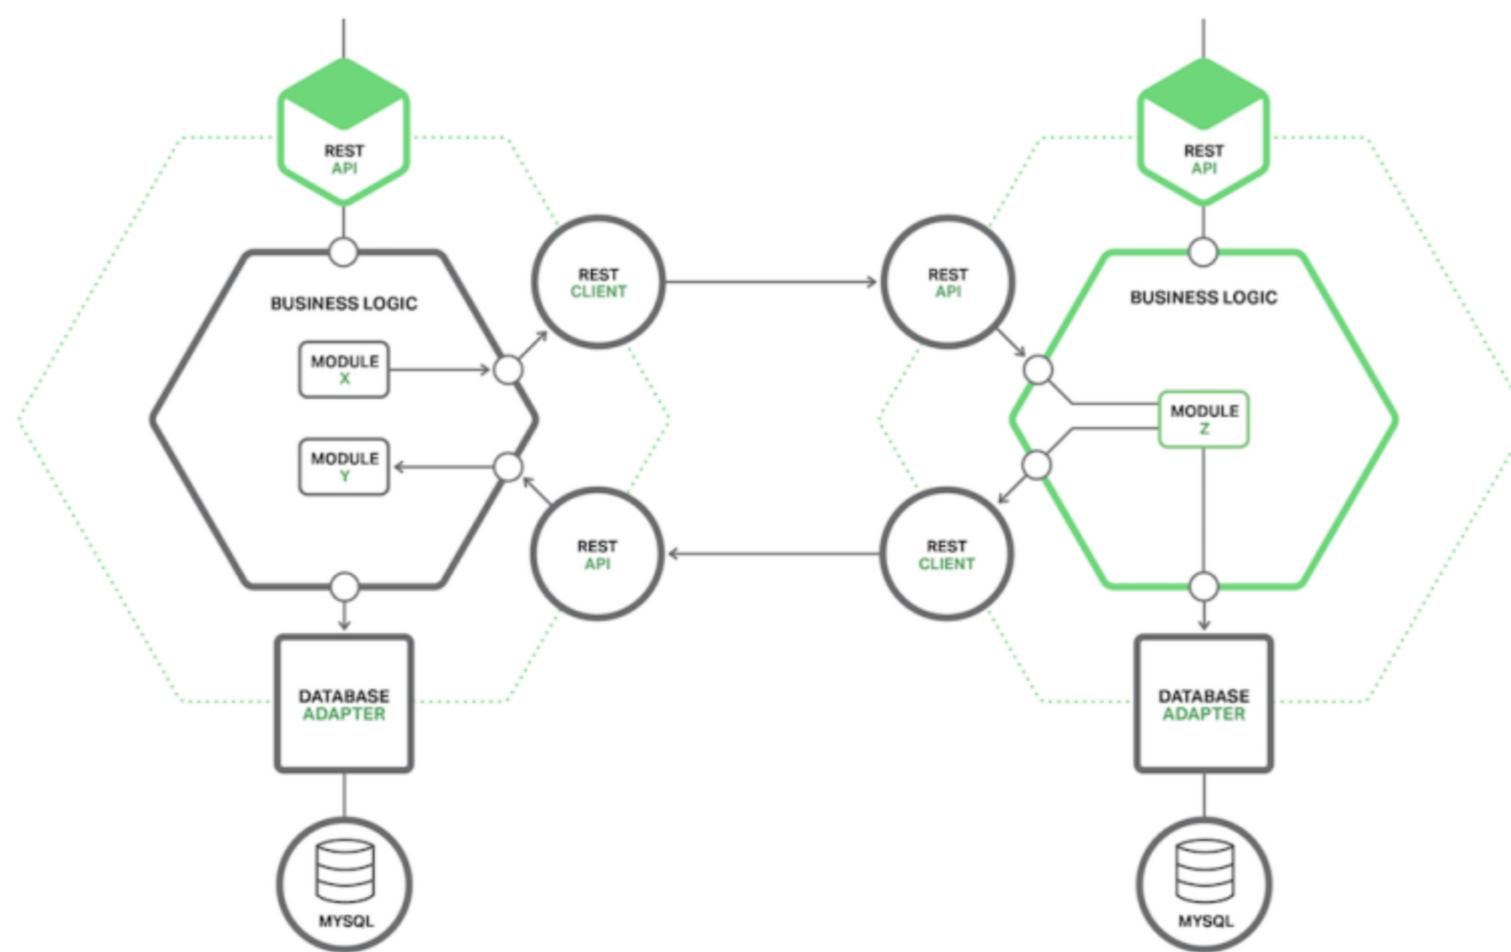
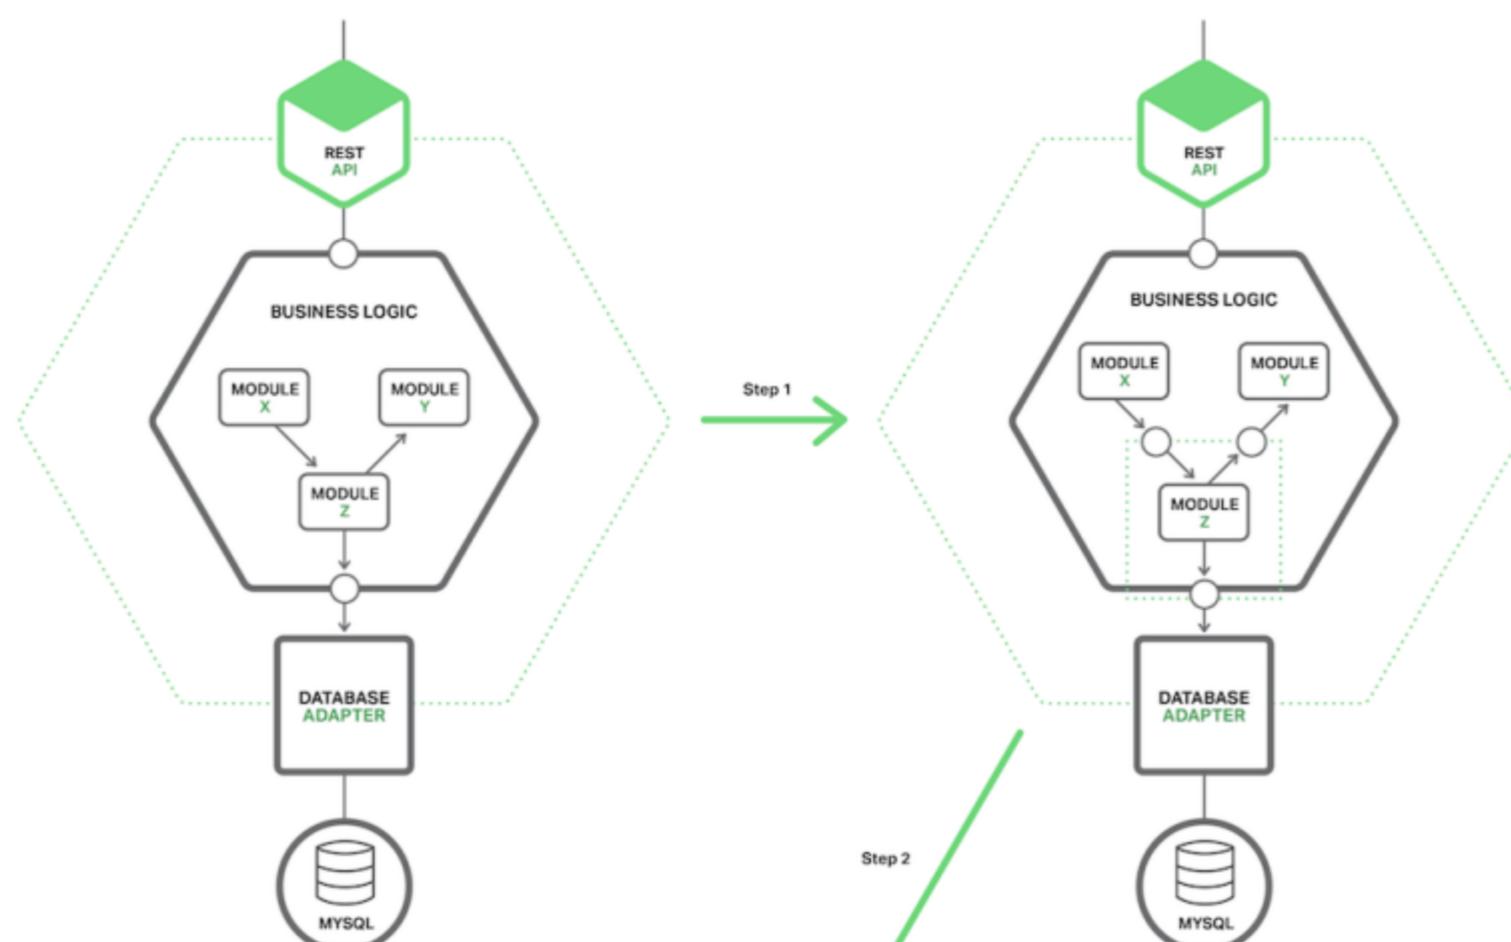
- ▶ split the presentation layer from the business logic and data access layers
- ▶ splitting a monolith in this way has two main benefits:
 - It enables you to develop, deploy, and scale the two applications independently of one another
 - another benefit of this approach is that it exposes a remote API that can be called by the microservices that you develop



Strategy 3 - Extract Services

- ▶ Turn existing modules within the monolith into standalone microservices
- ▶ Each time you extract a module and turn it into a service, the monolith shrinks
- ▶ Once you have converted enough modules, the monolith will cease to be a problem
- ▶ Either it disappears entirely or it becomes small enough that it is just another service.

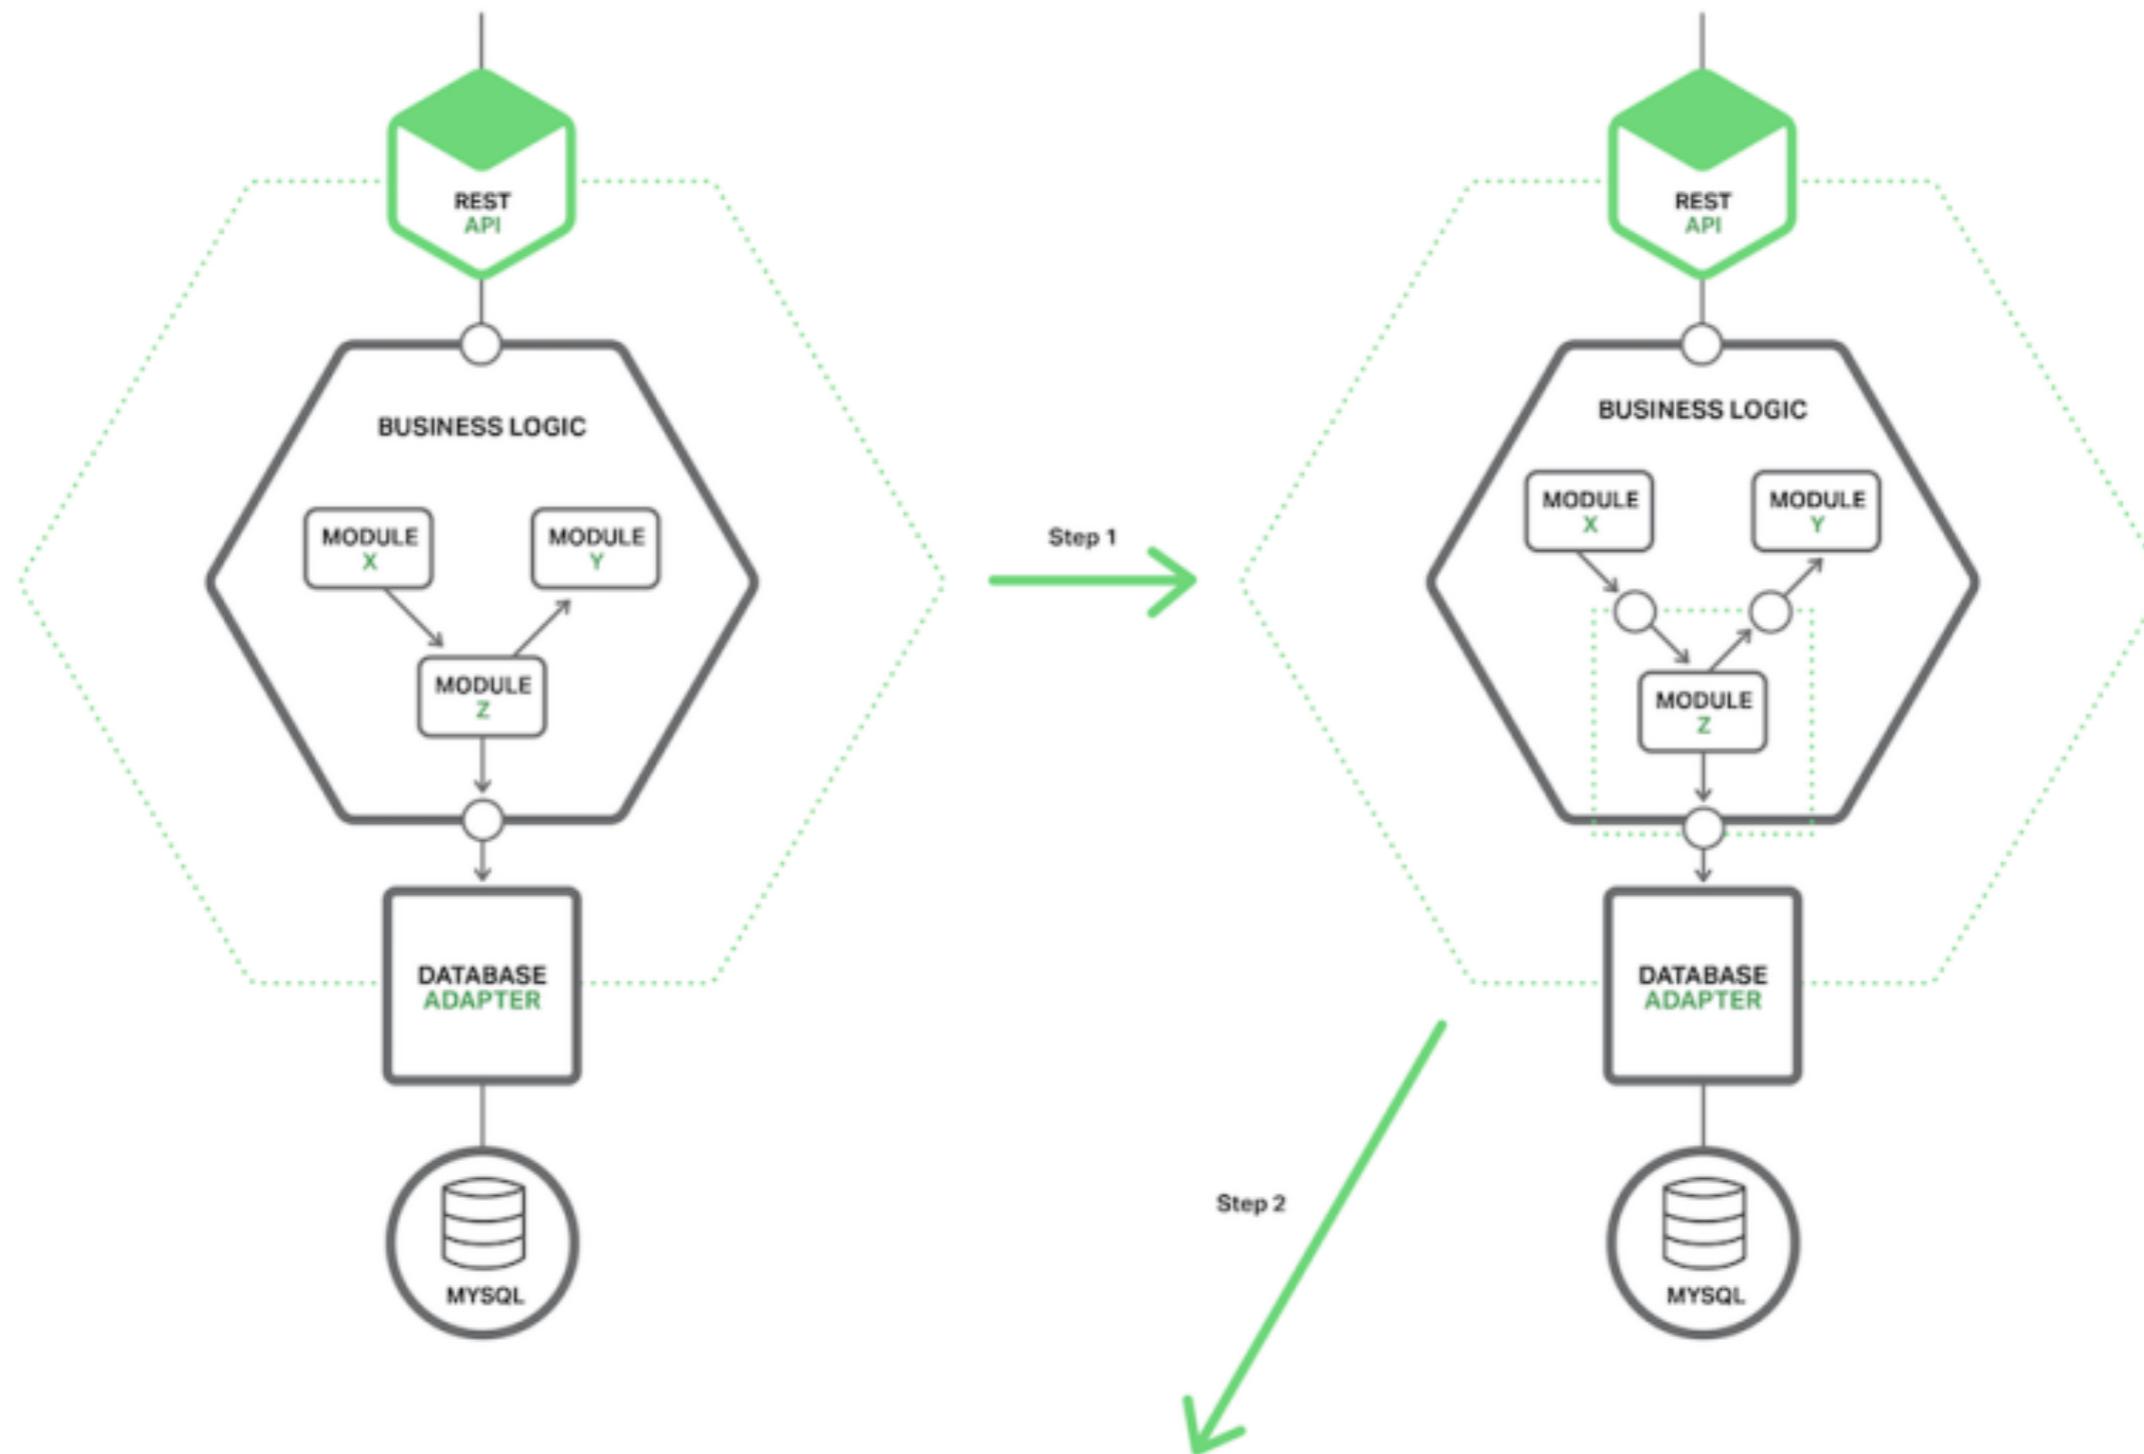
P.PORTO



The final phase of breaking up a monolith to microservices consists of two steps

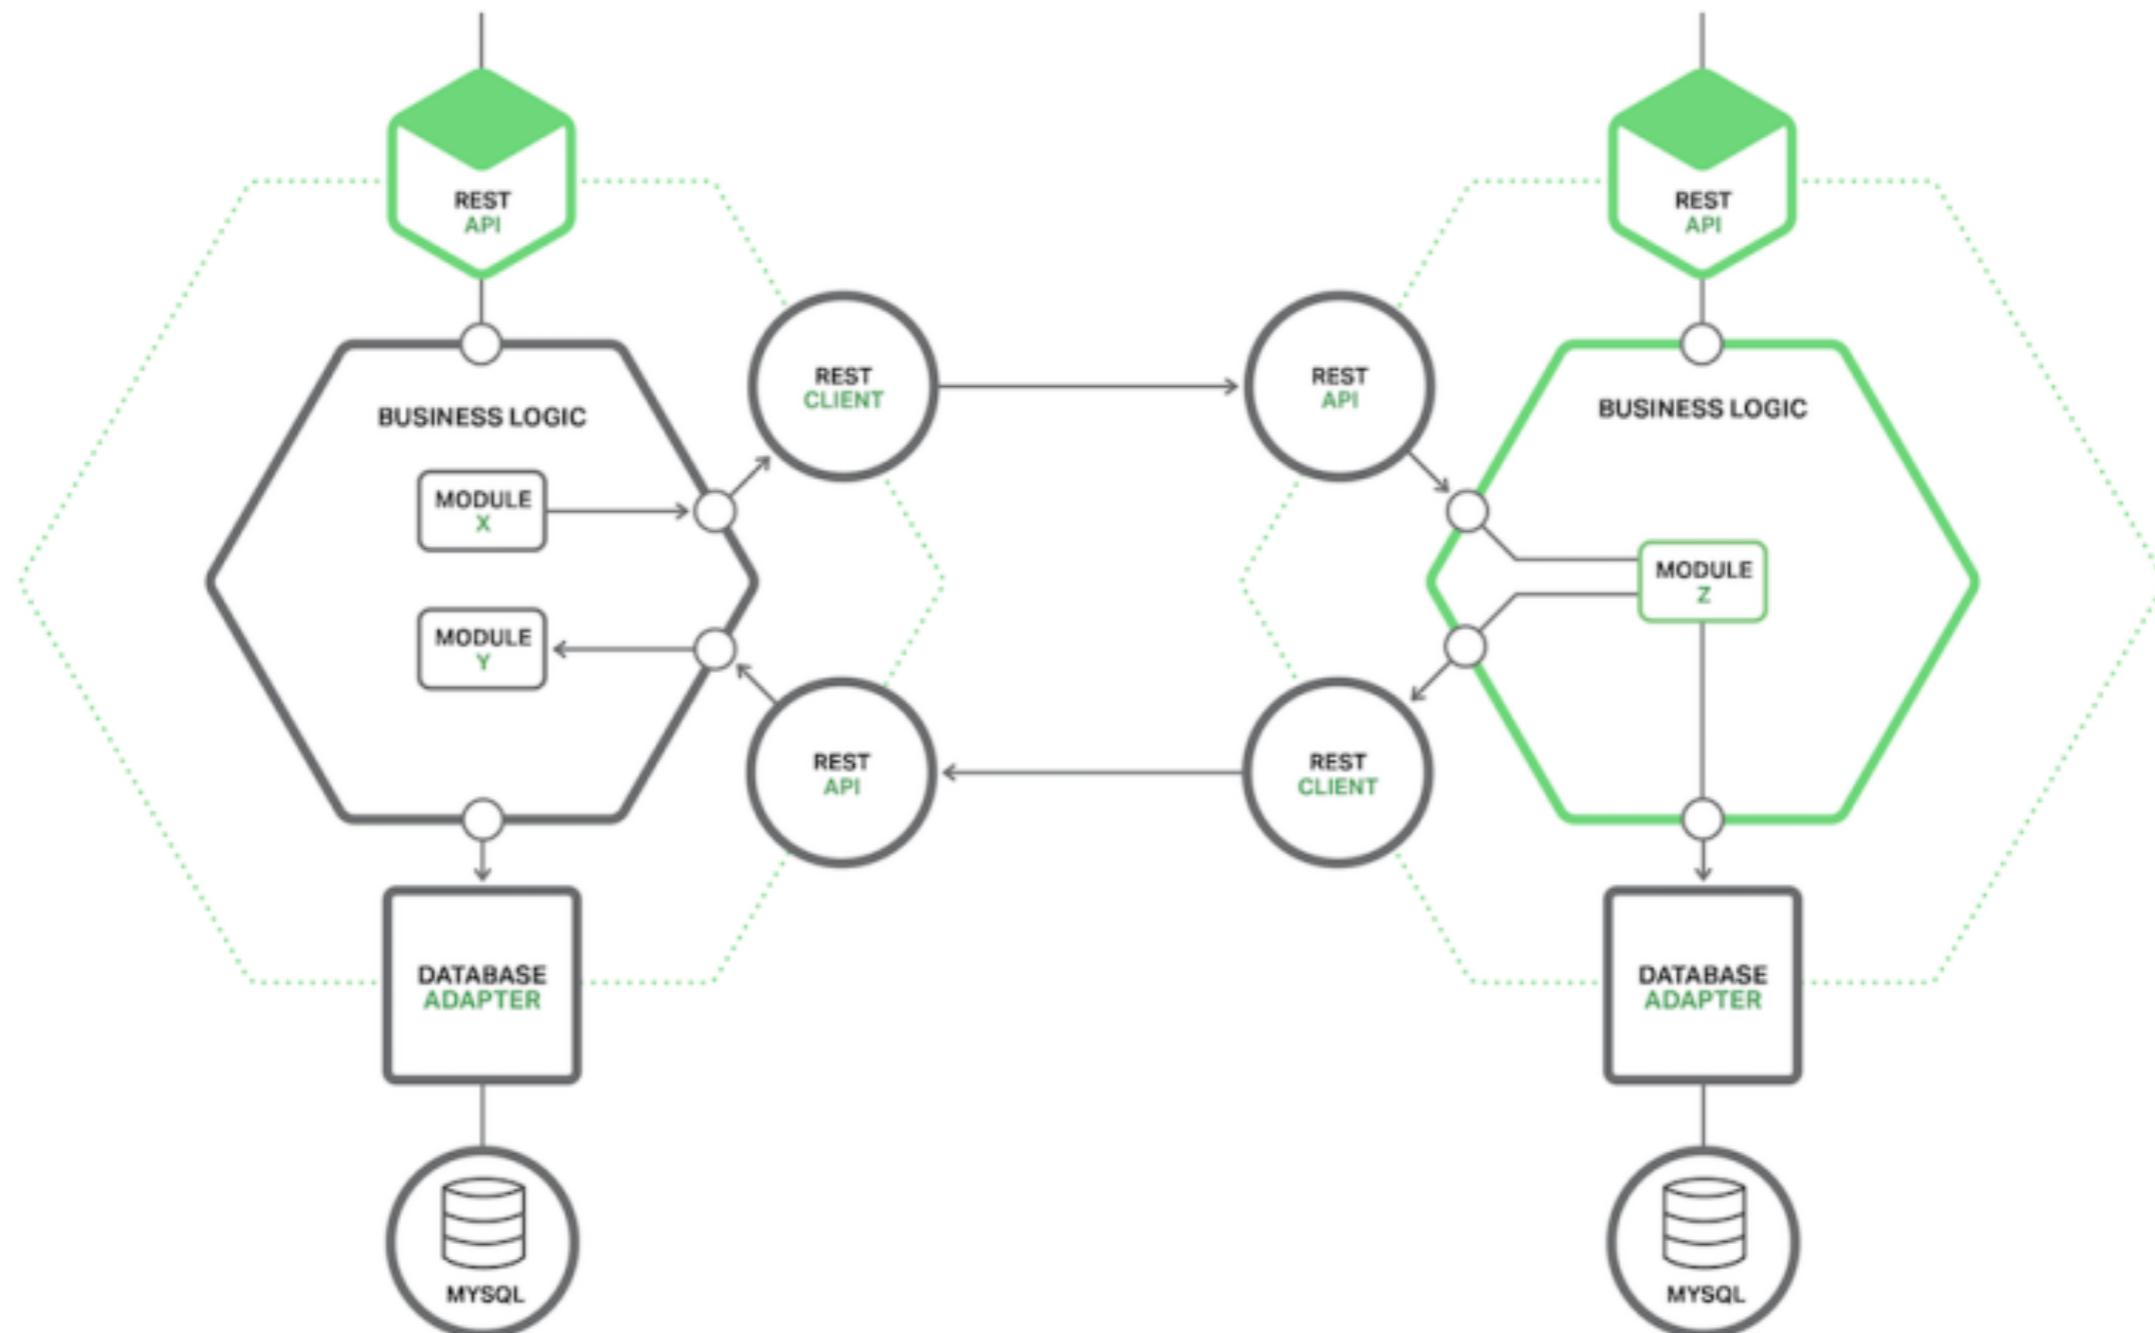
► Step I

- All dependencies between the monolith and the new service are untangled by defining the API between them
- Once the API is defined, communication with the monolith can be restricted to API only, which prepares the service to be completely detached



► Step 2

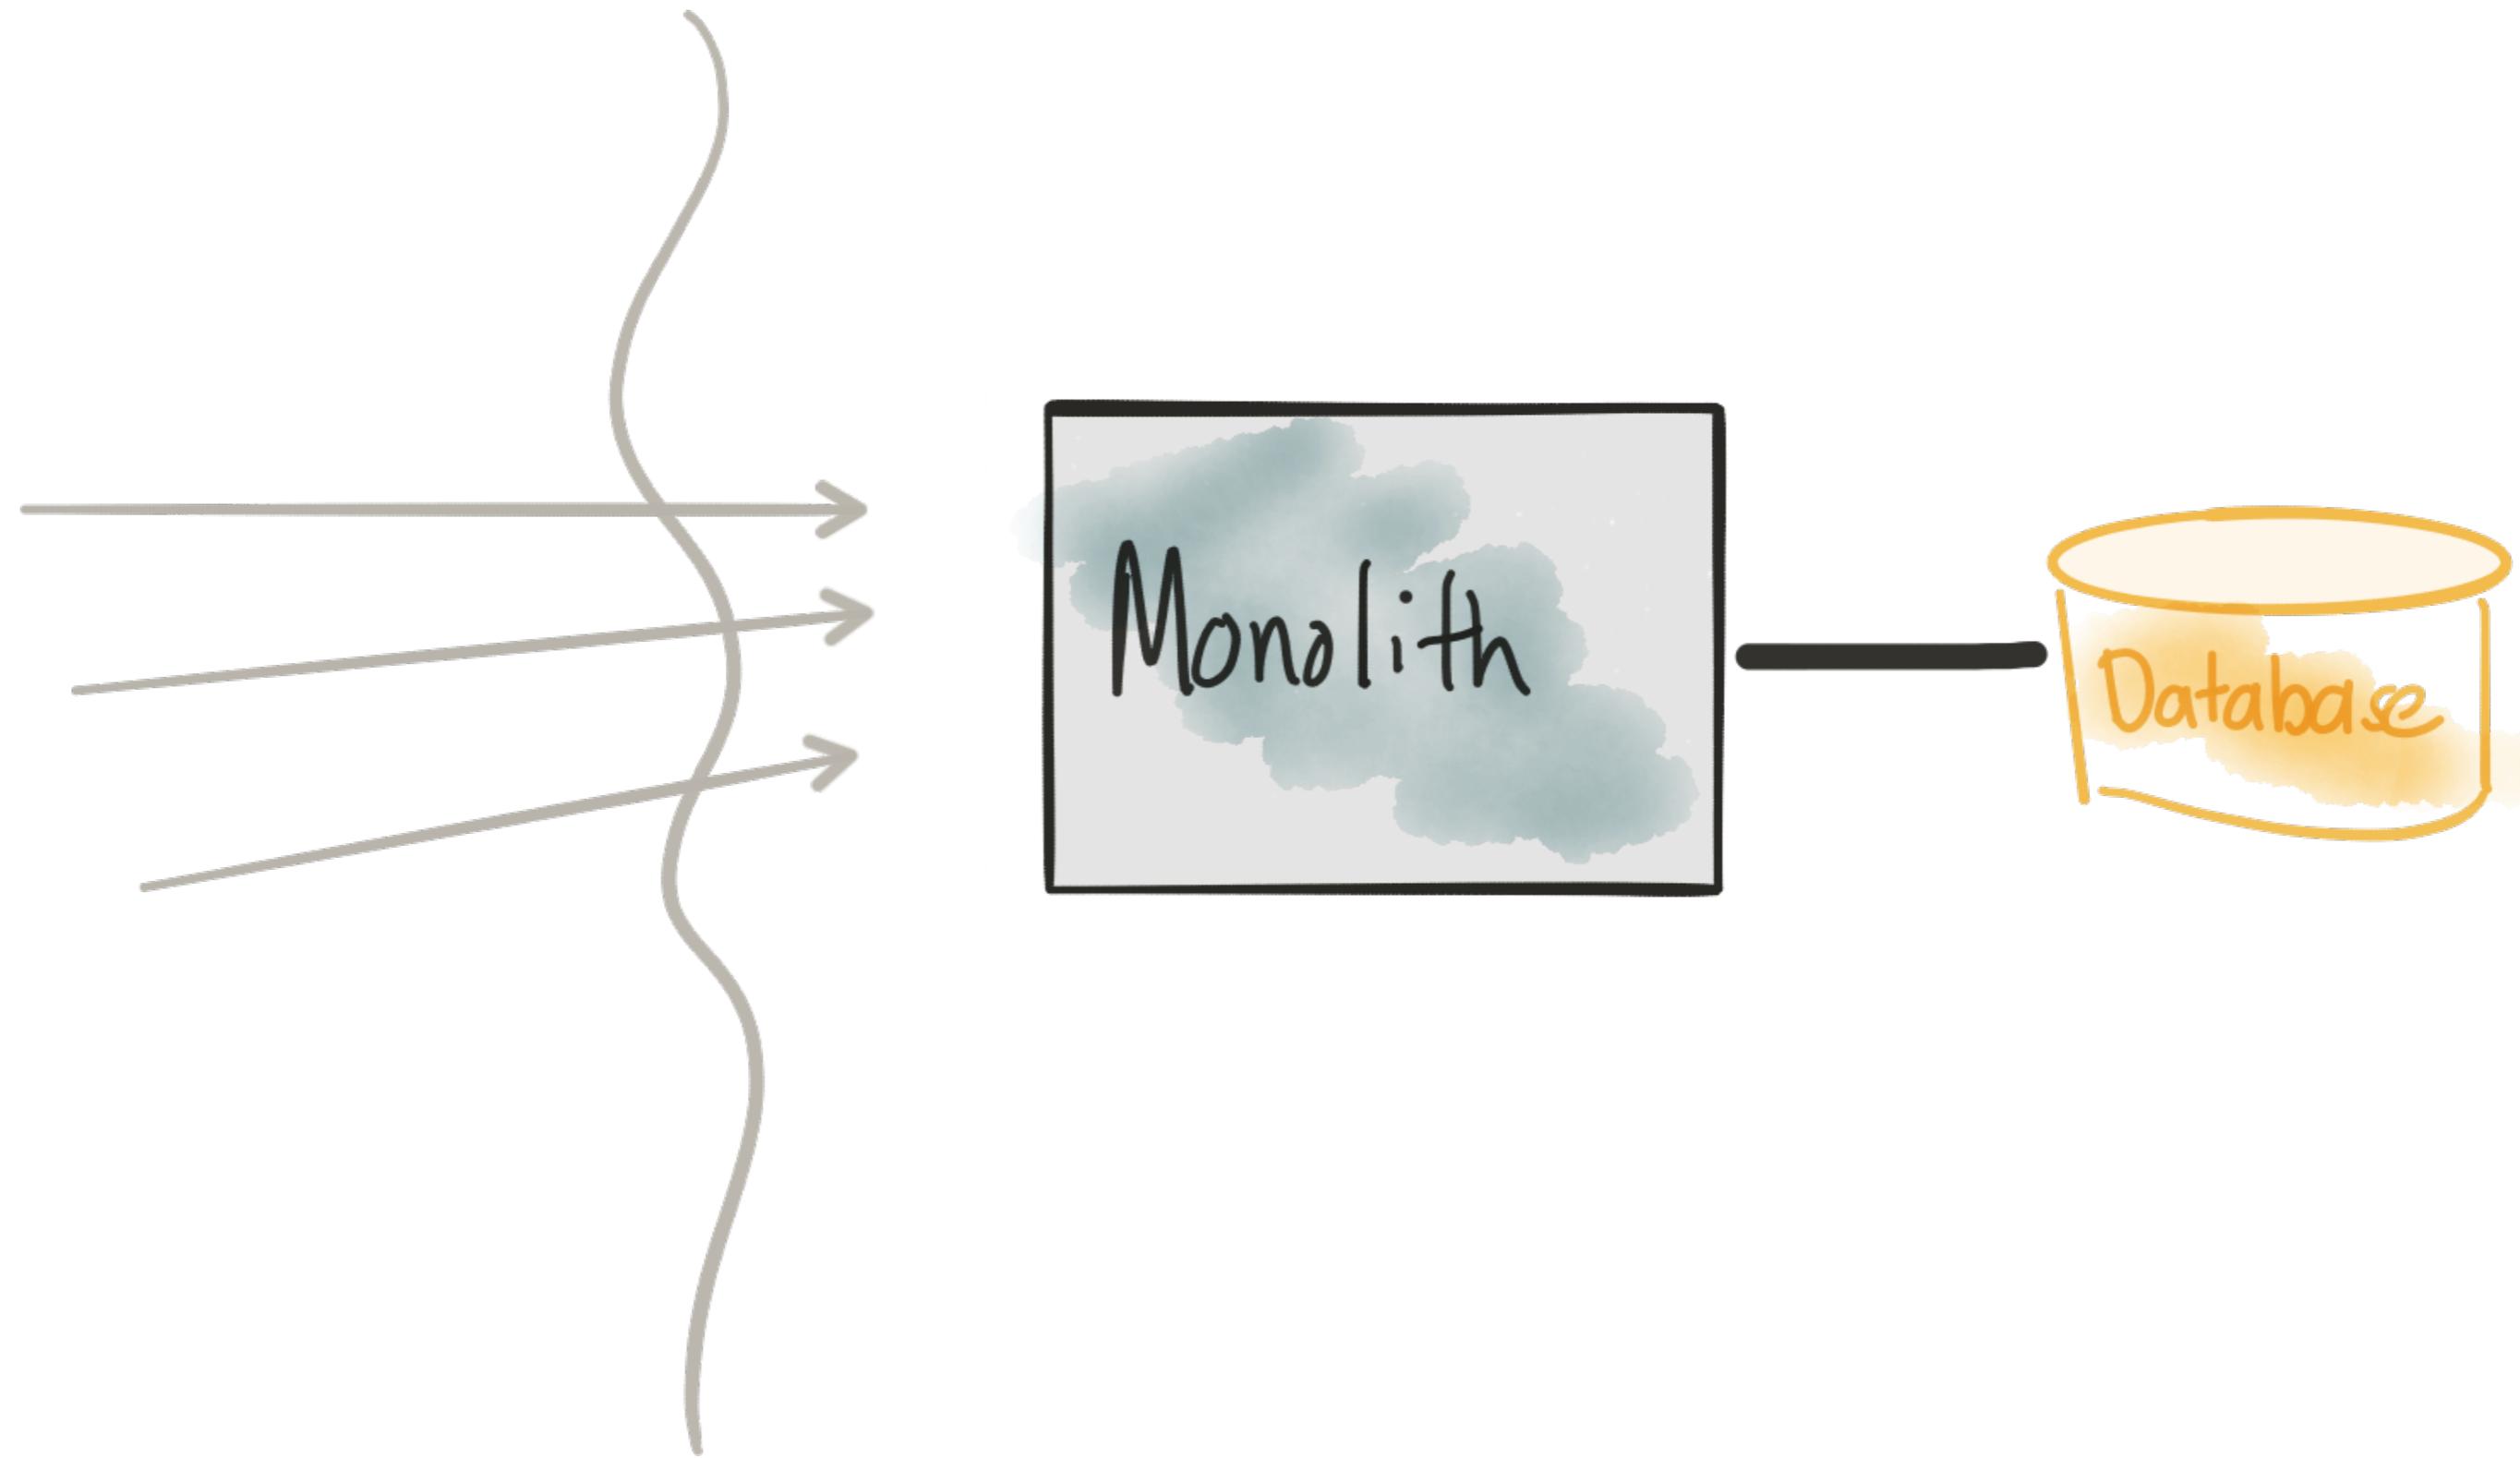
- the service is moved to standalone microservice running in a separate process
- After this step, a new microservice can be developed, scaled and deployed independently to monolith



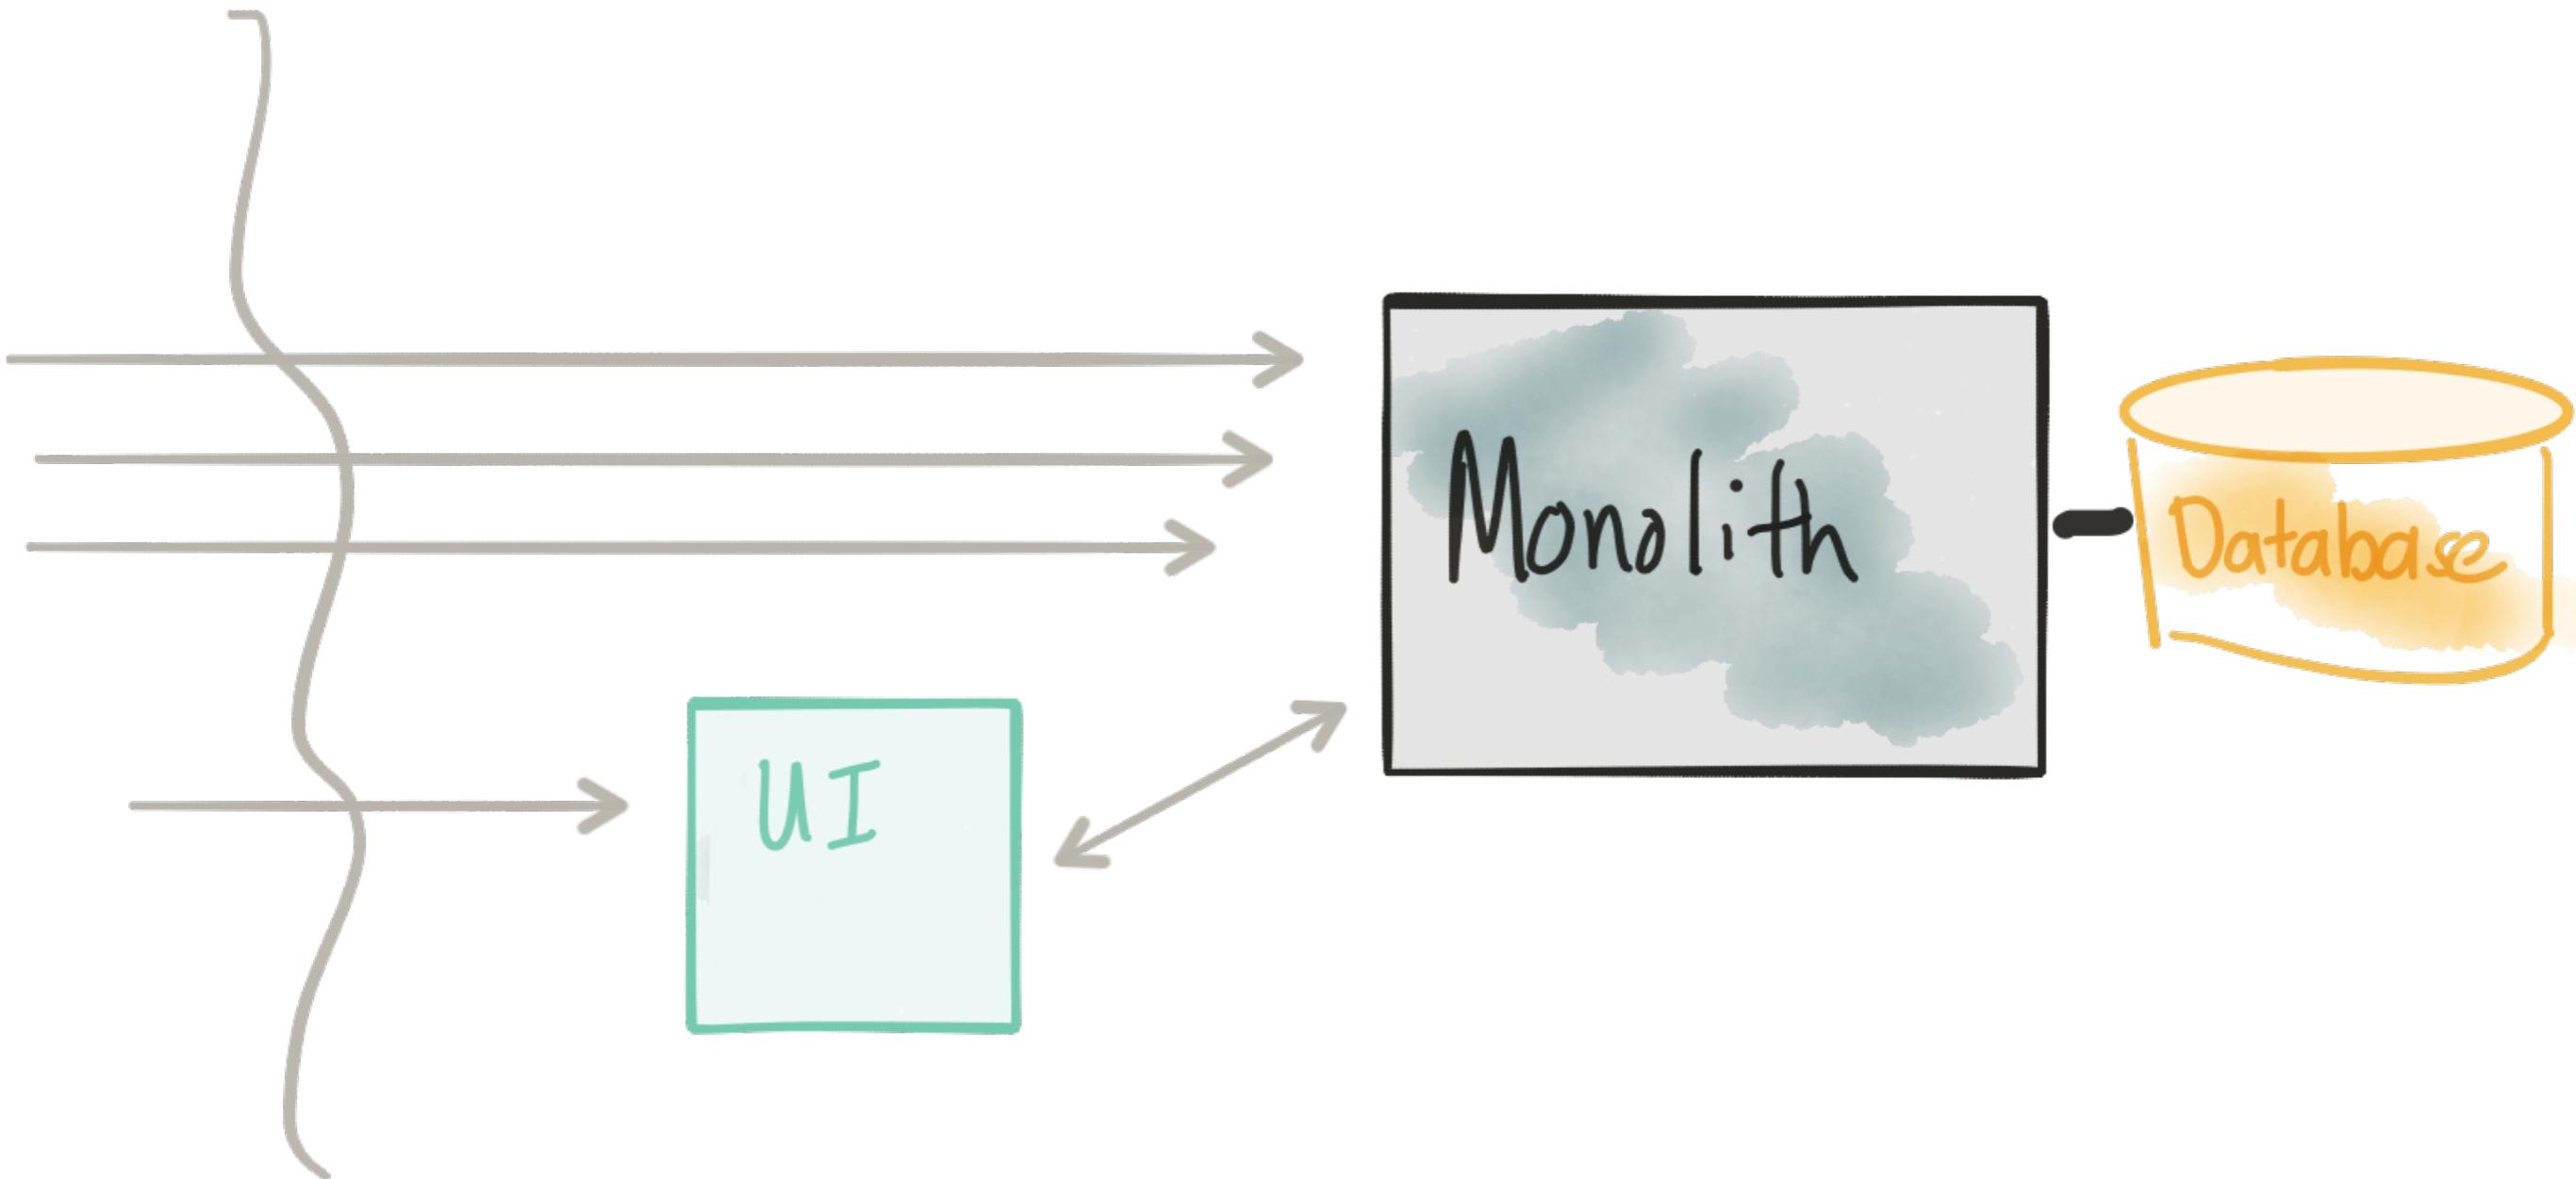
Monolith to Microservice Evolution

Concrete example

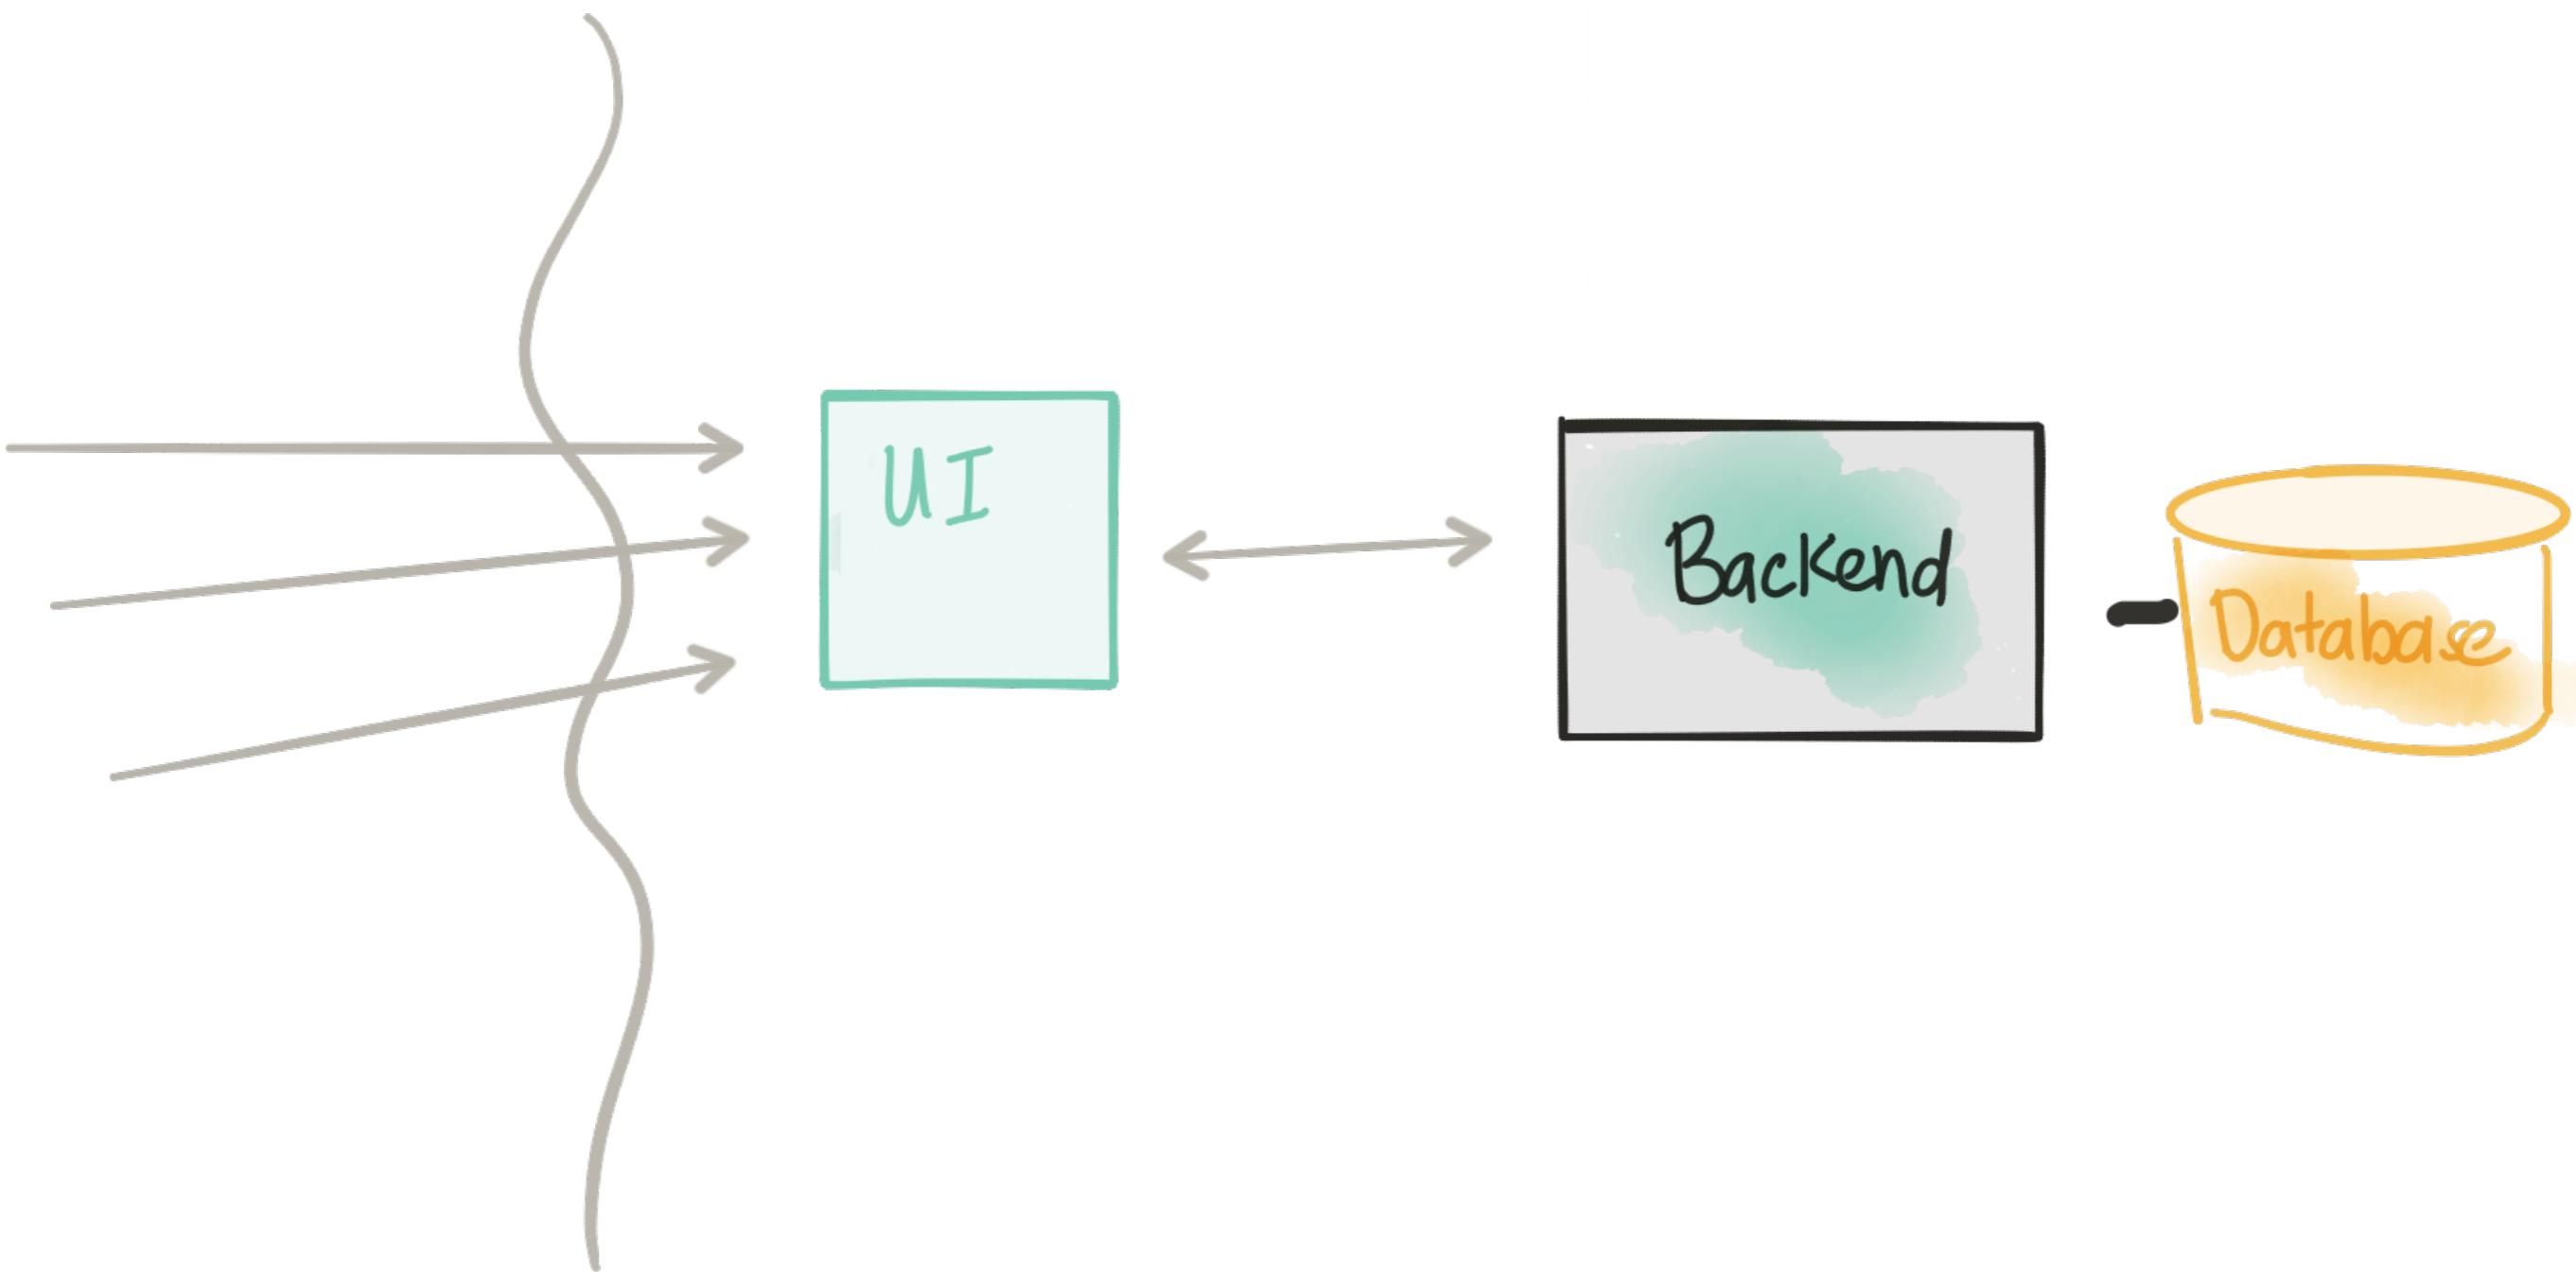
Meet the monolith



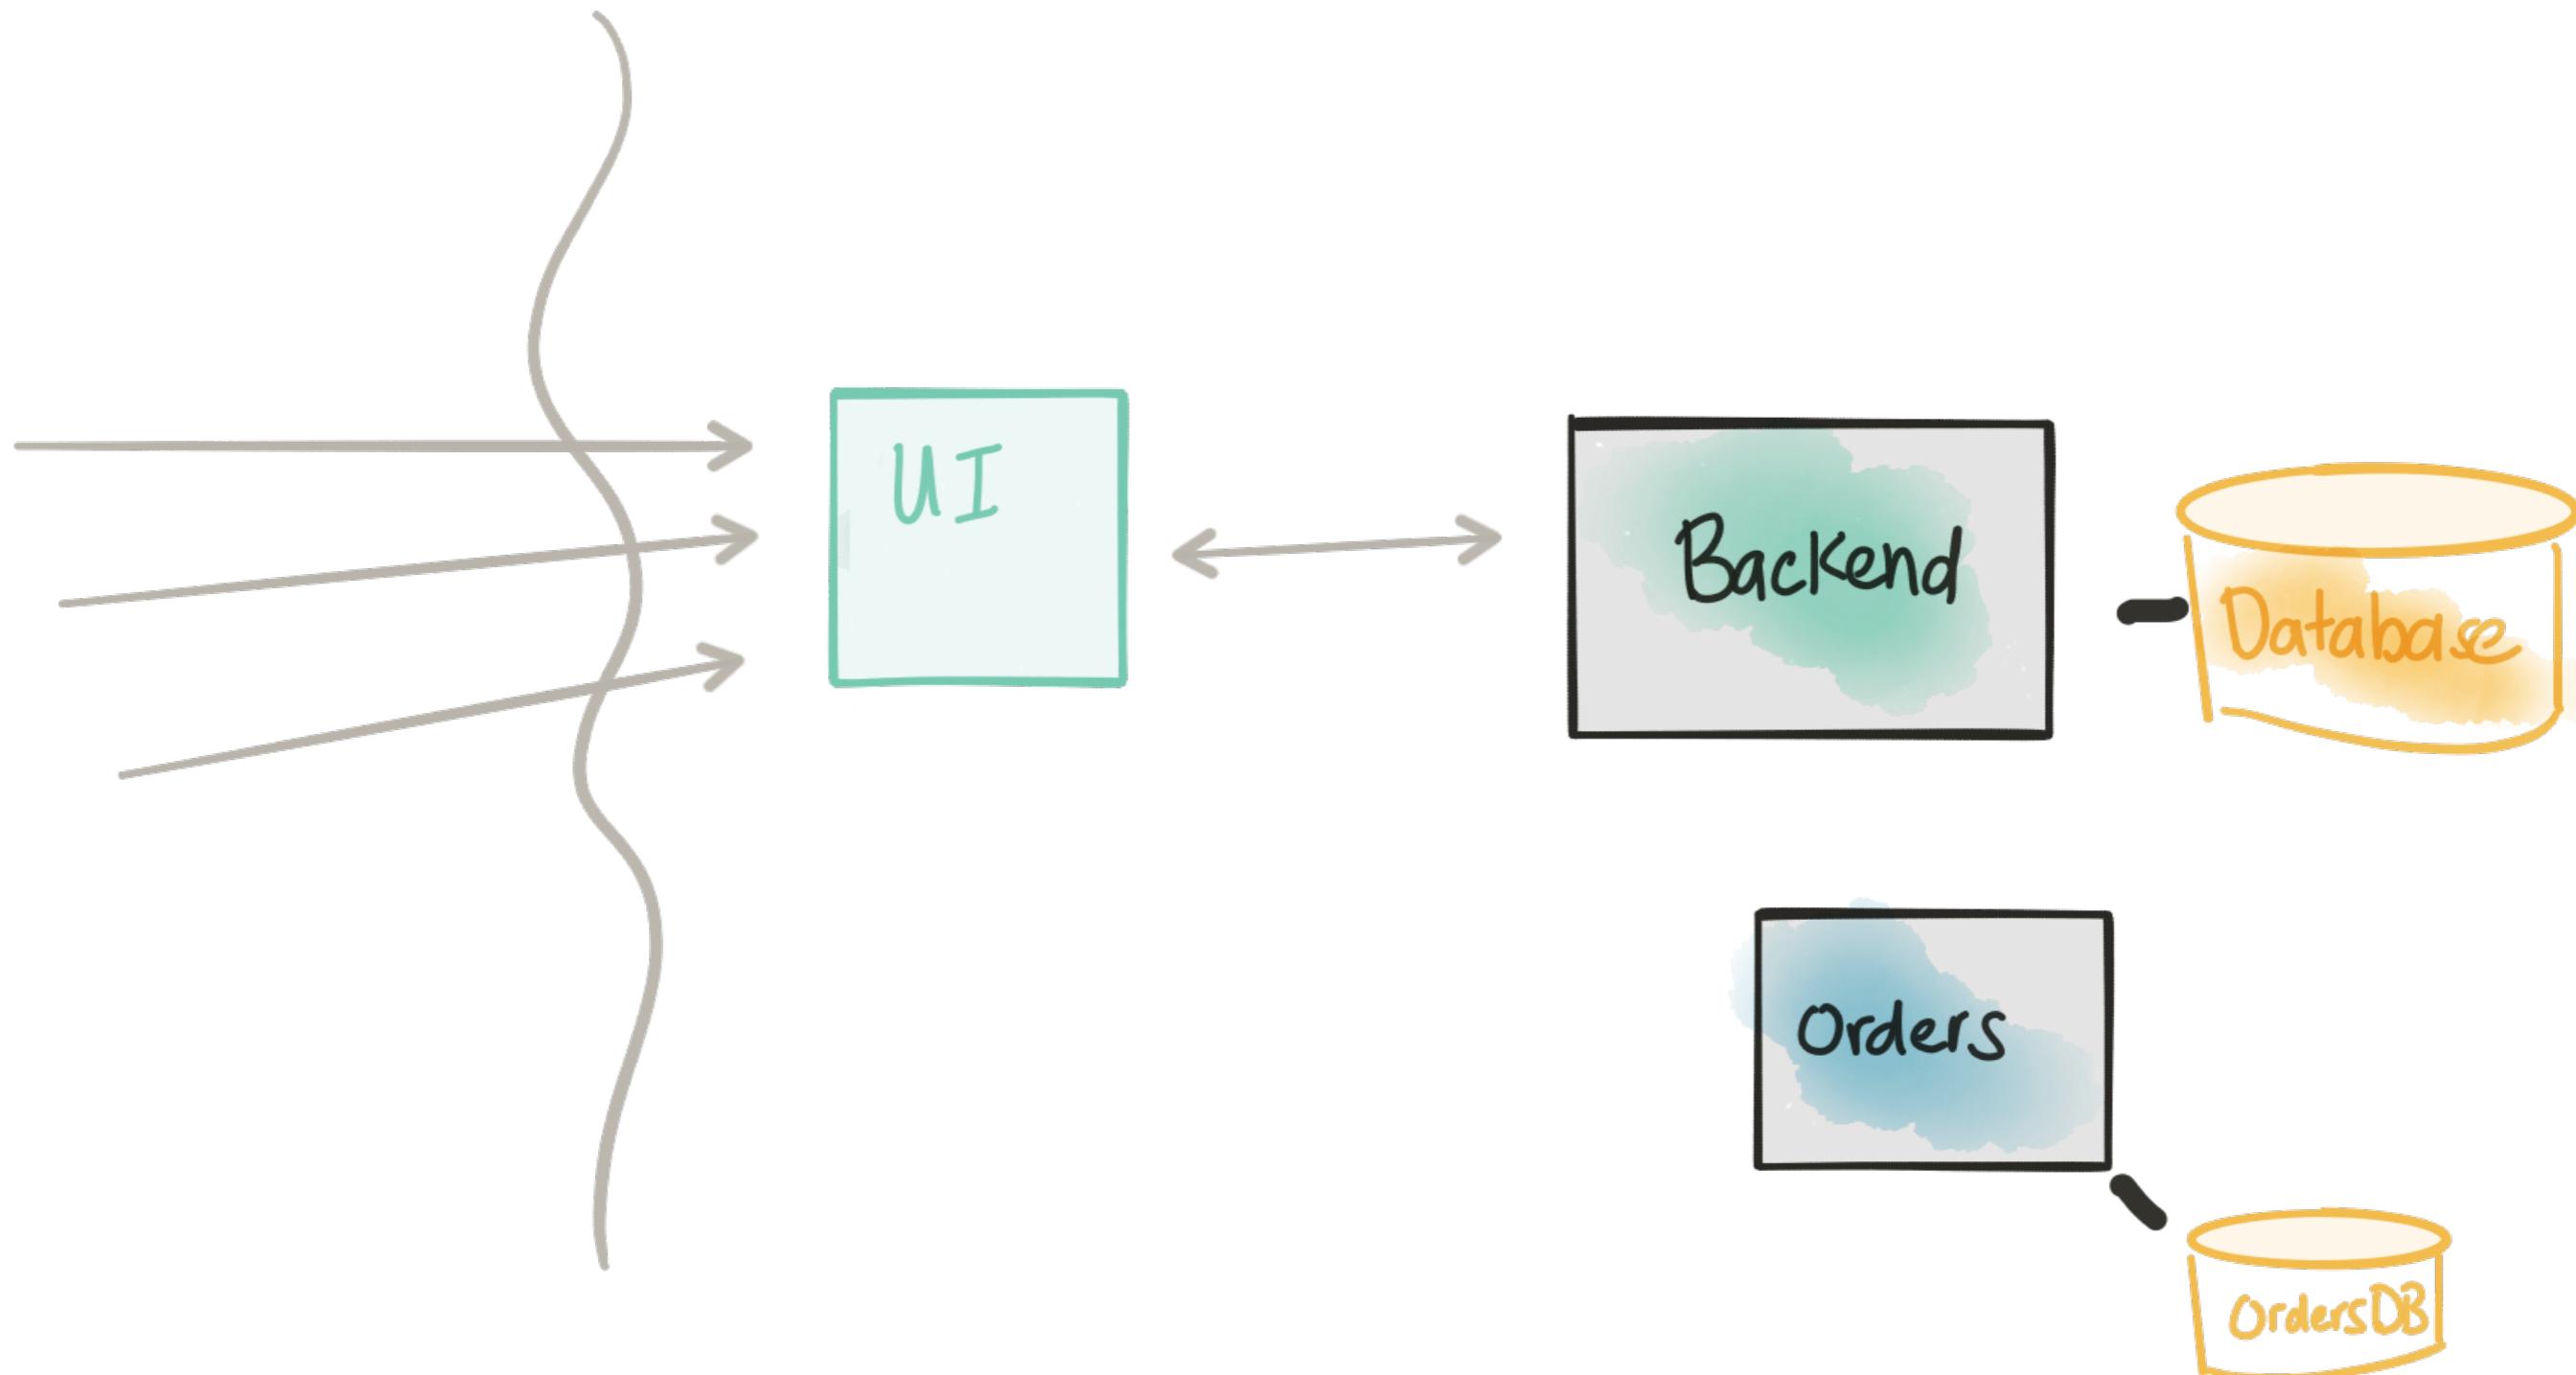
Extract the UI



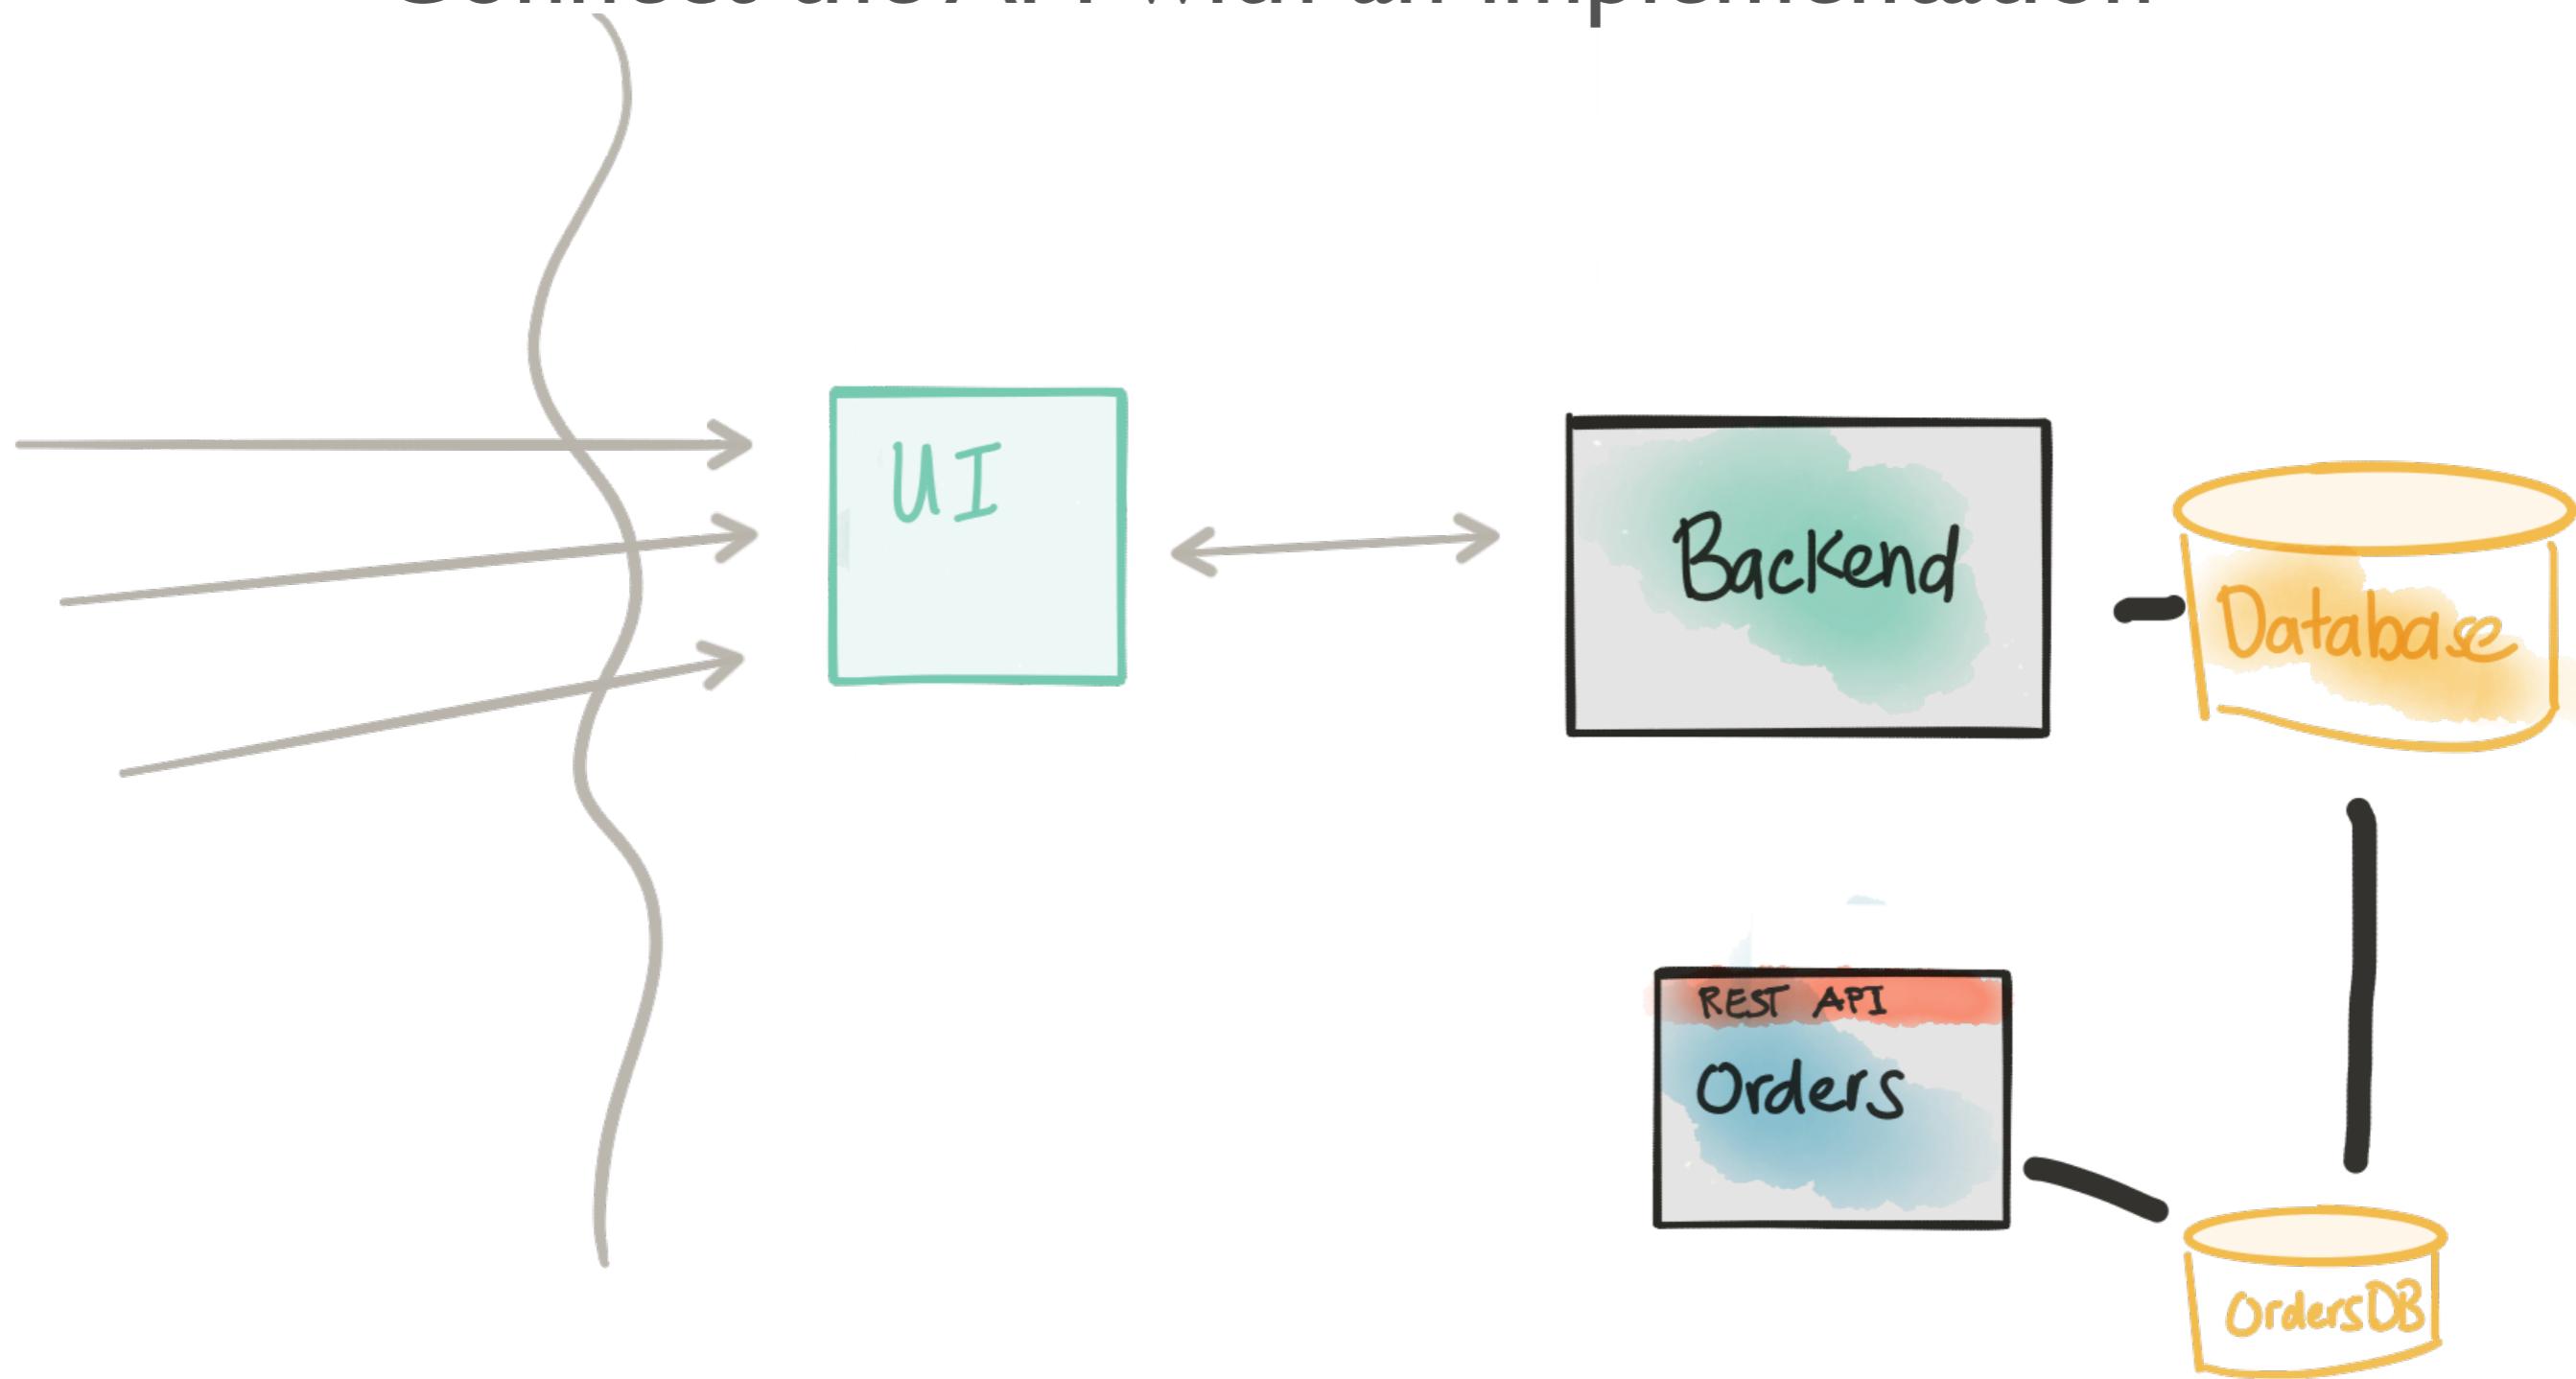
Drop the UI from the monolith



Introduce a new service

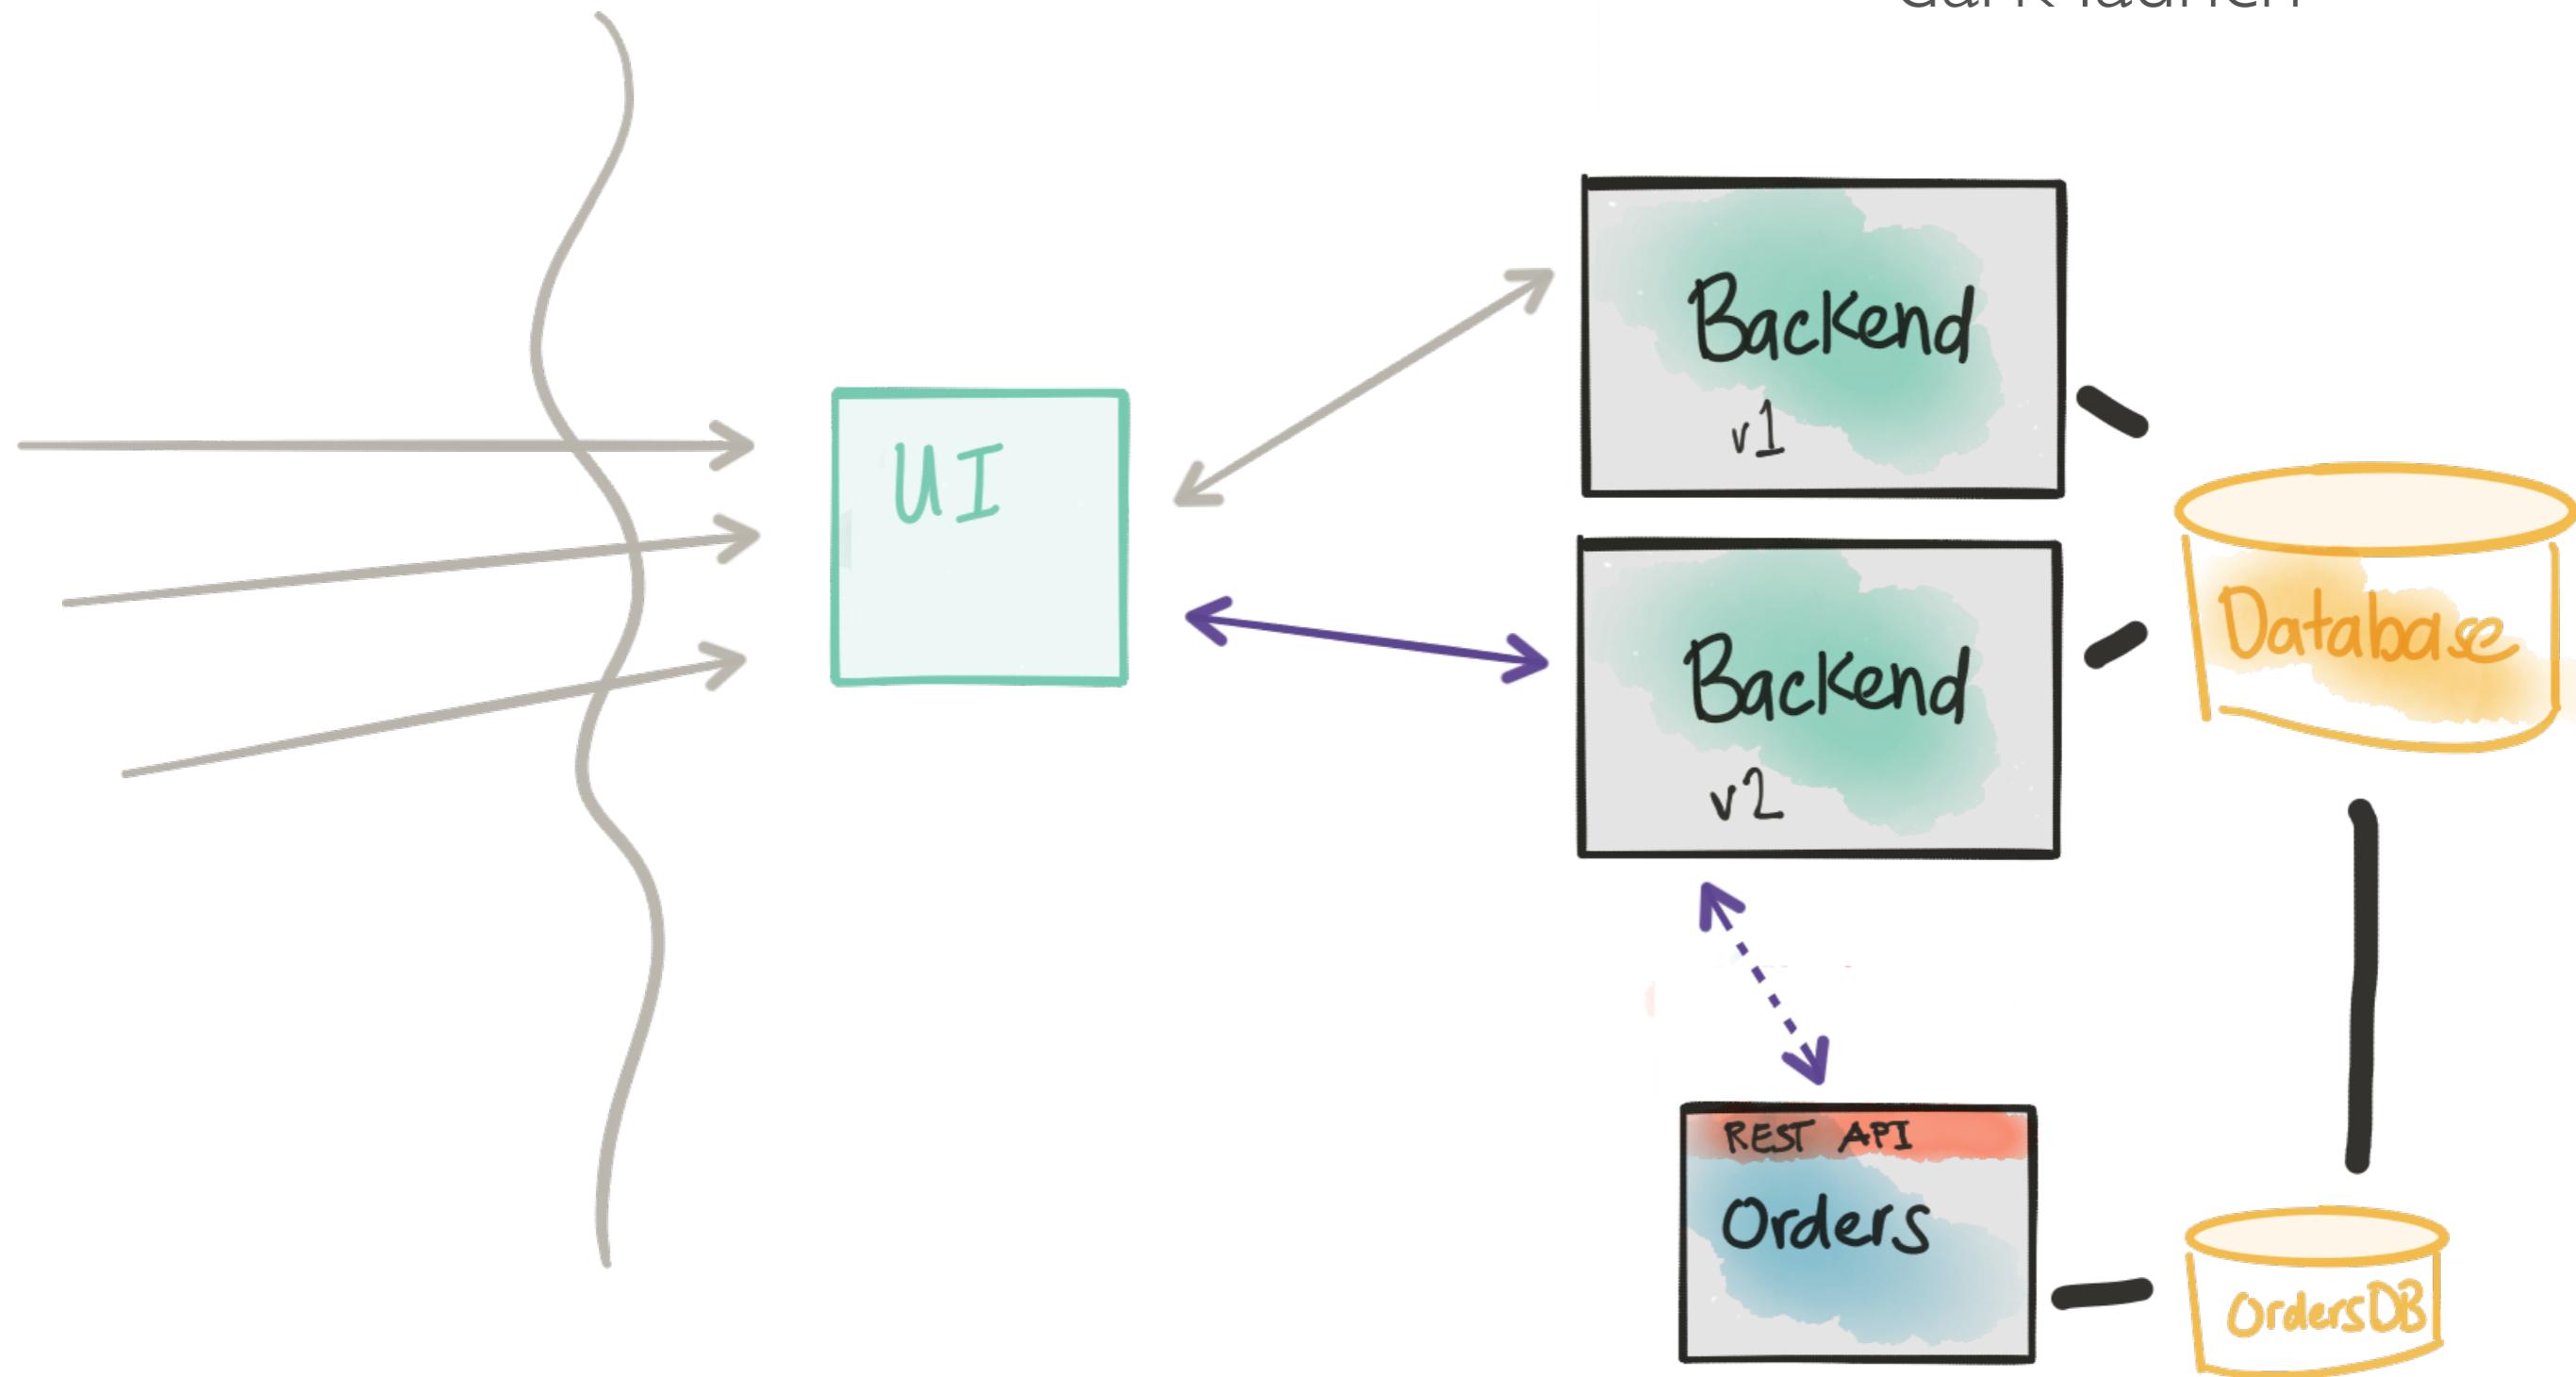


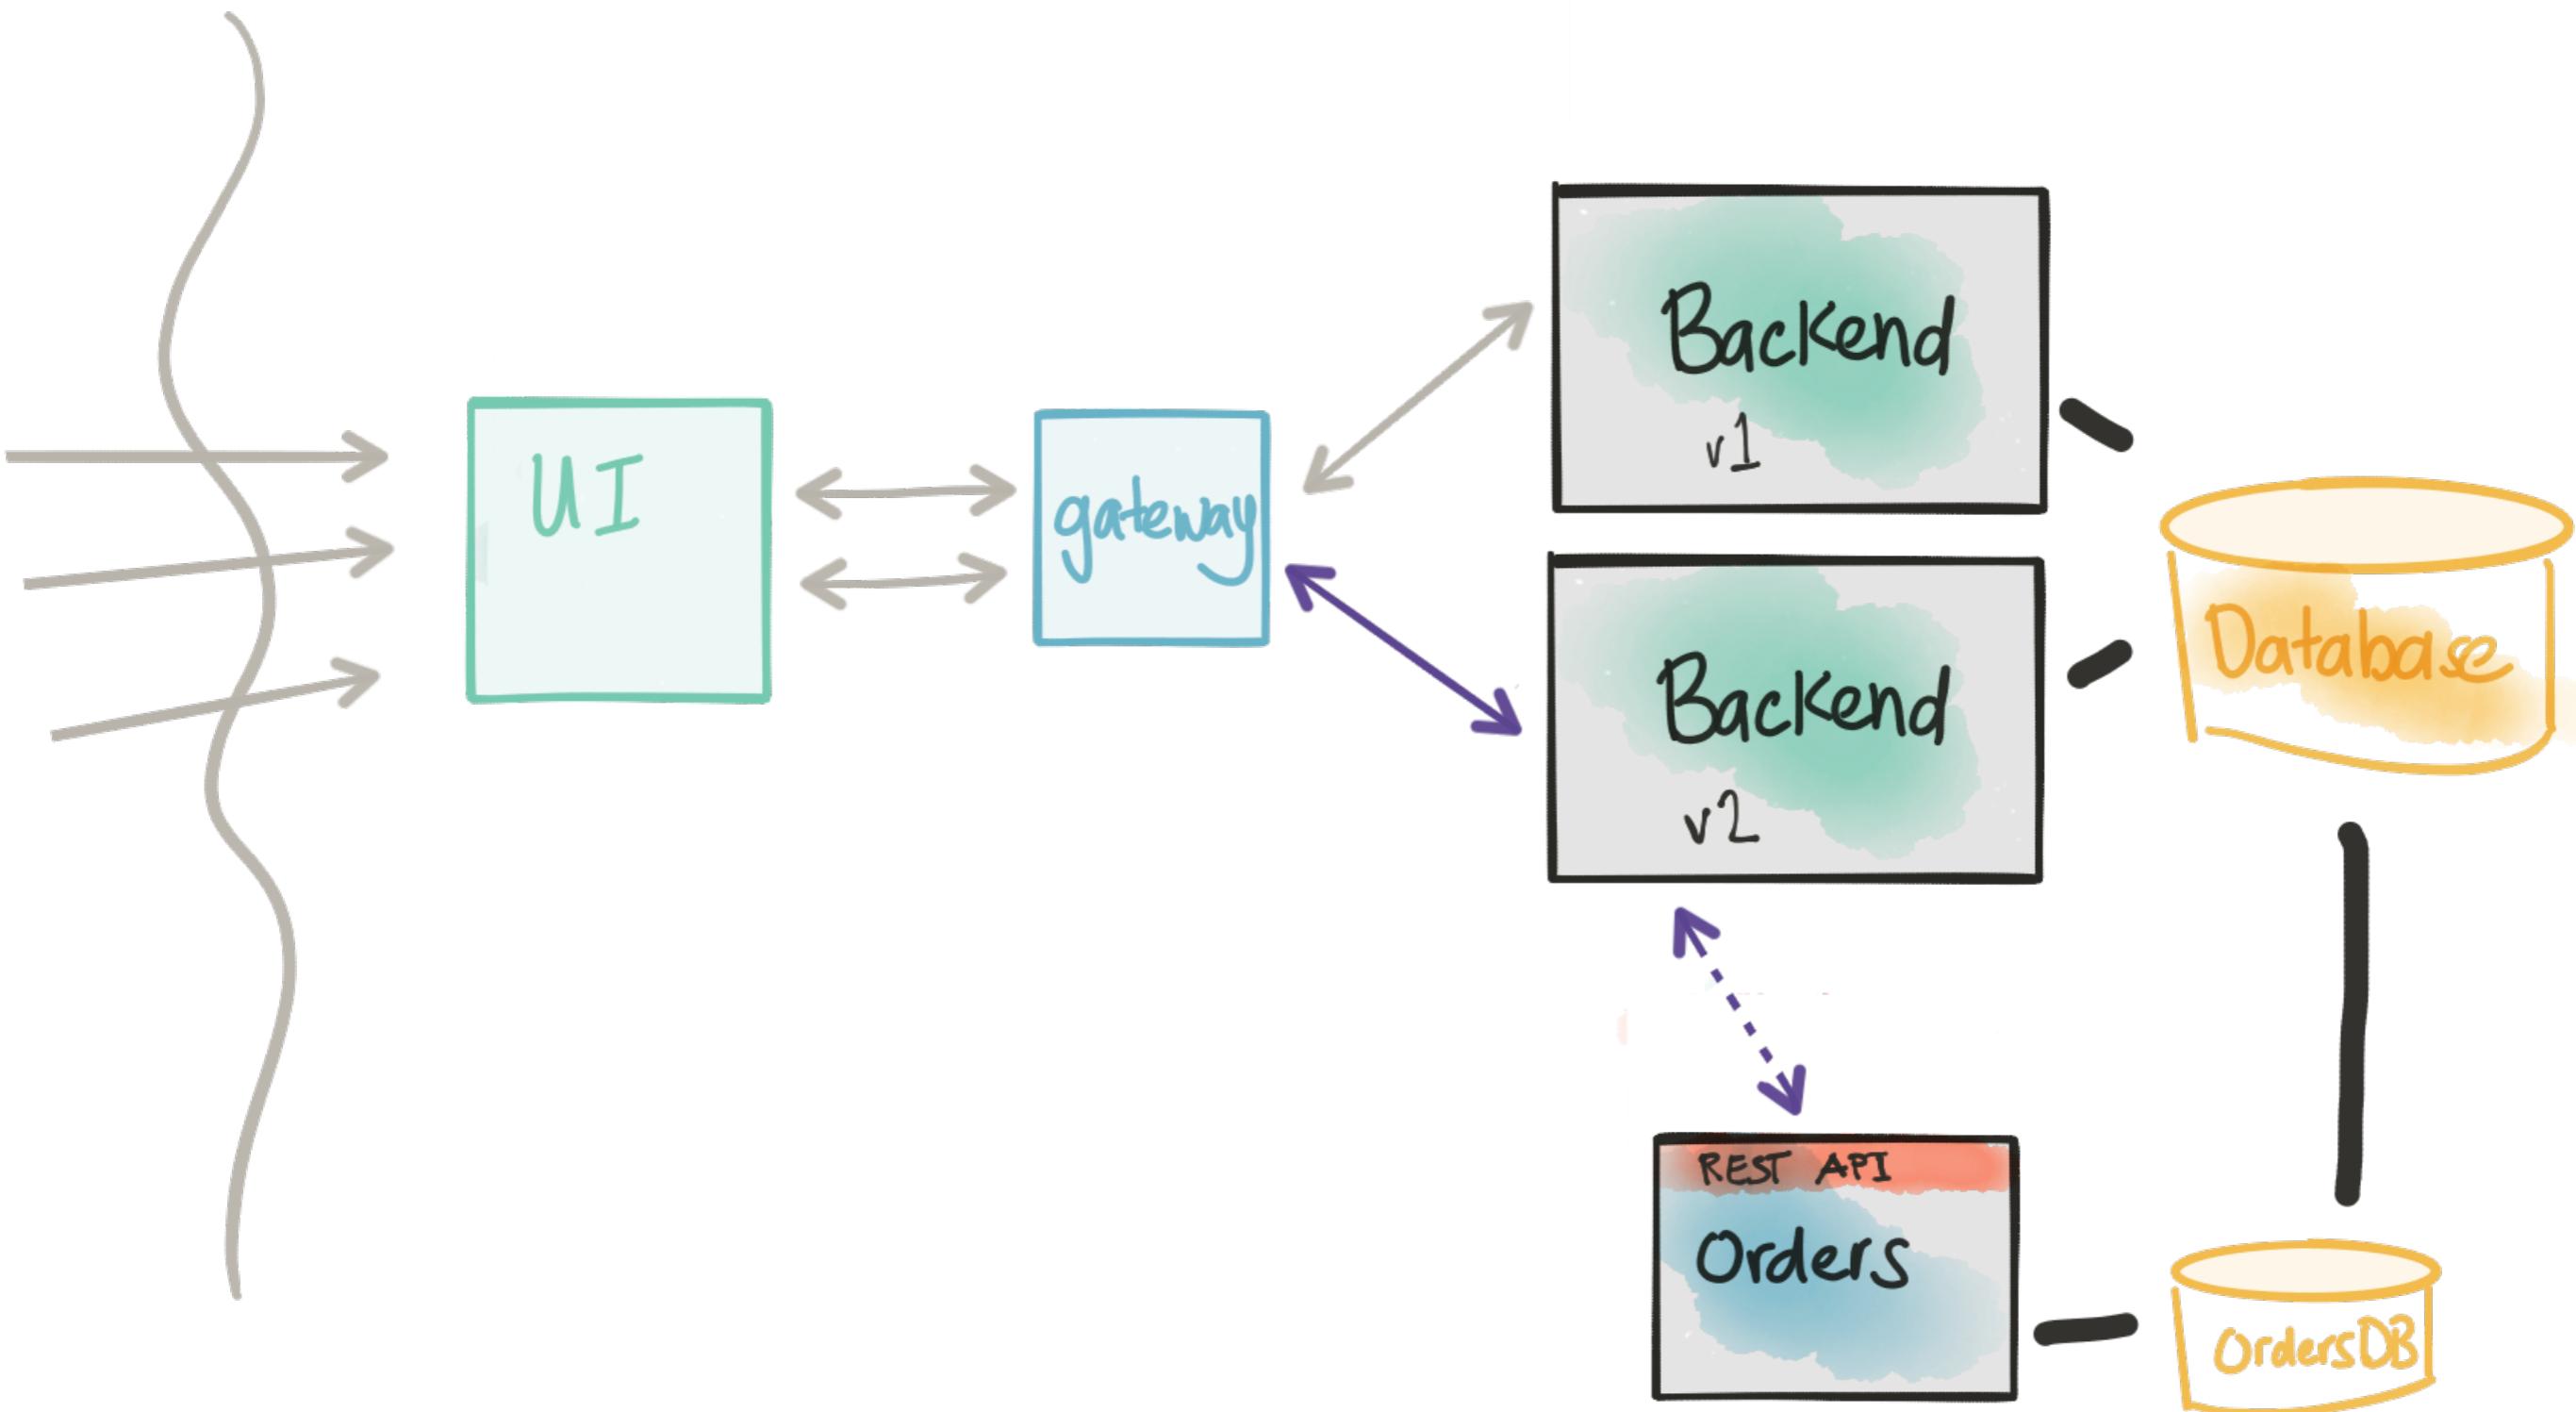
Connect the API with an implementation



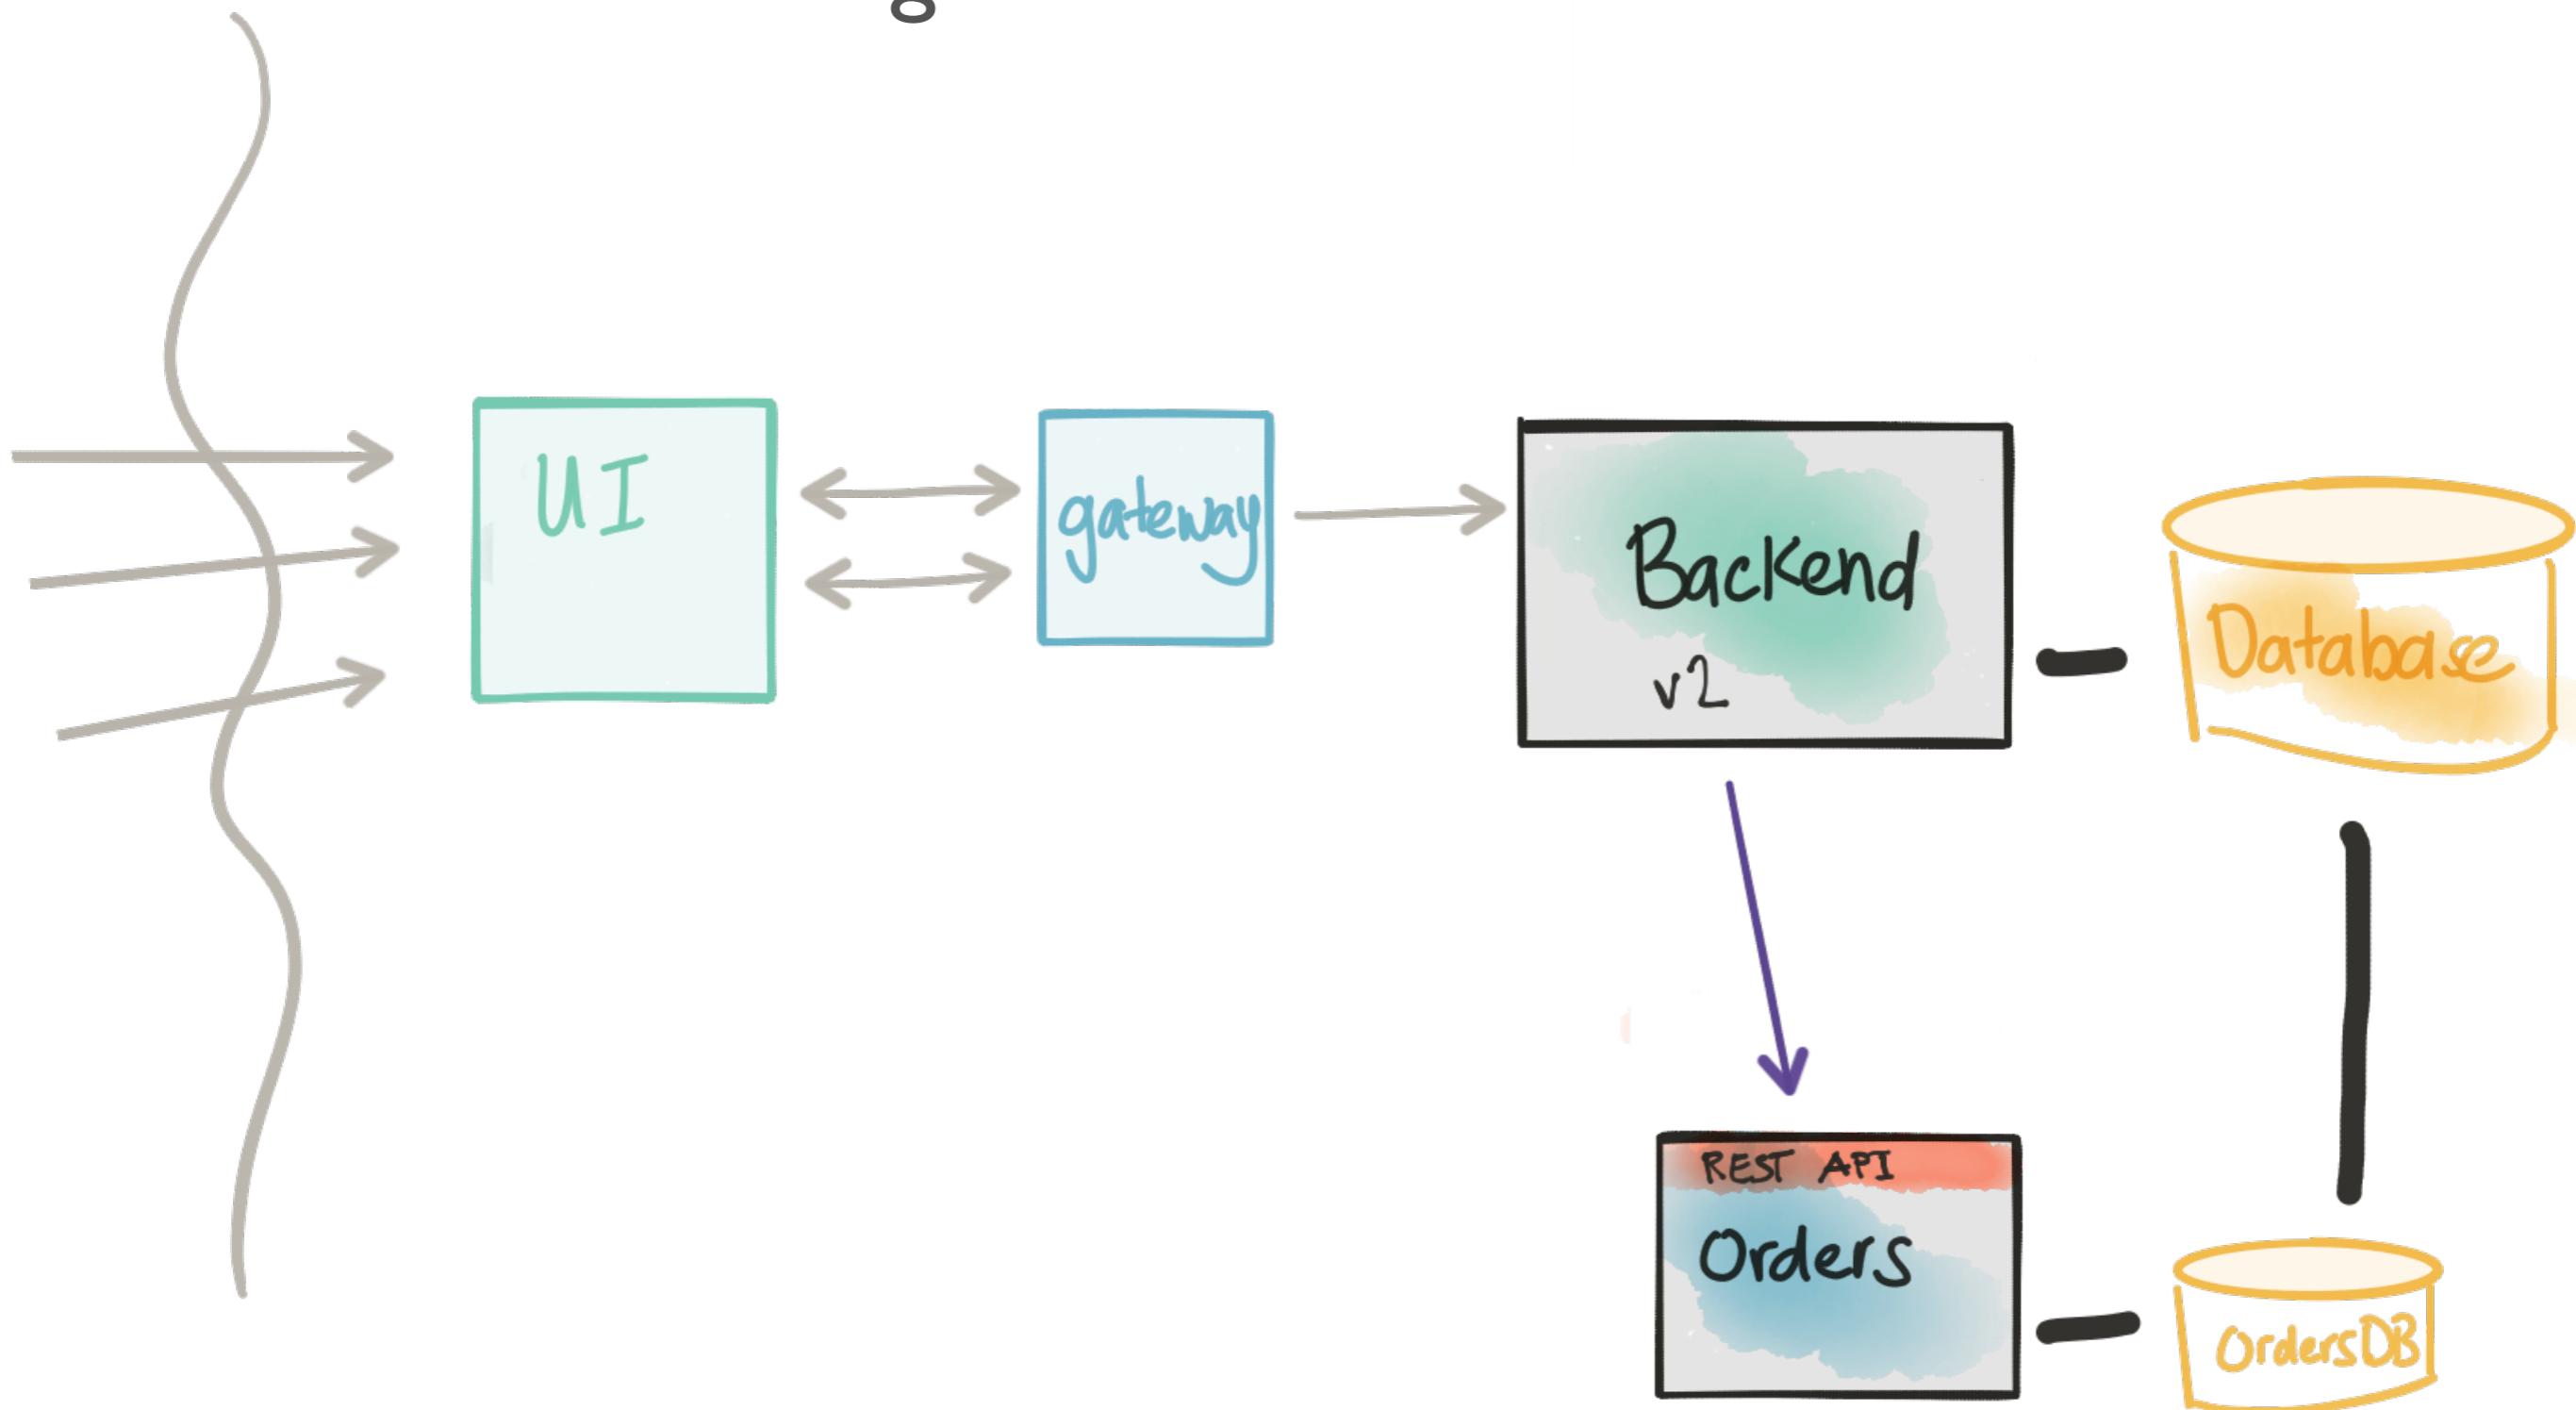
Start sending shadow traffic to the new service

dark launch

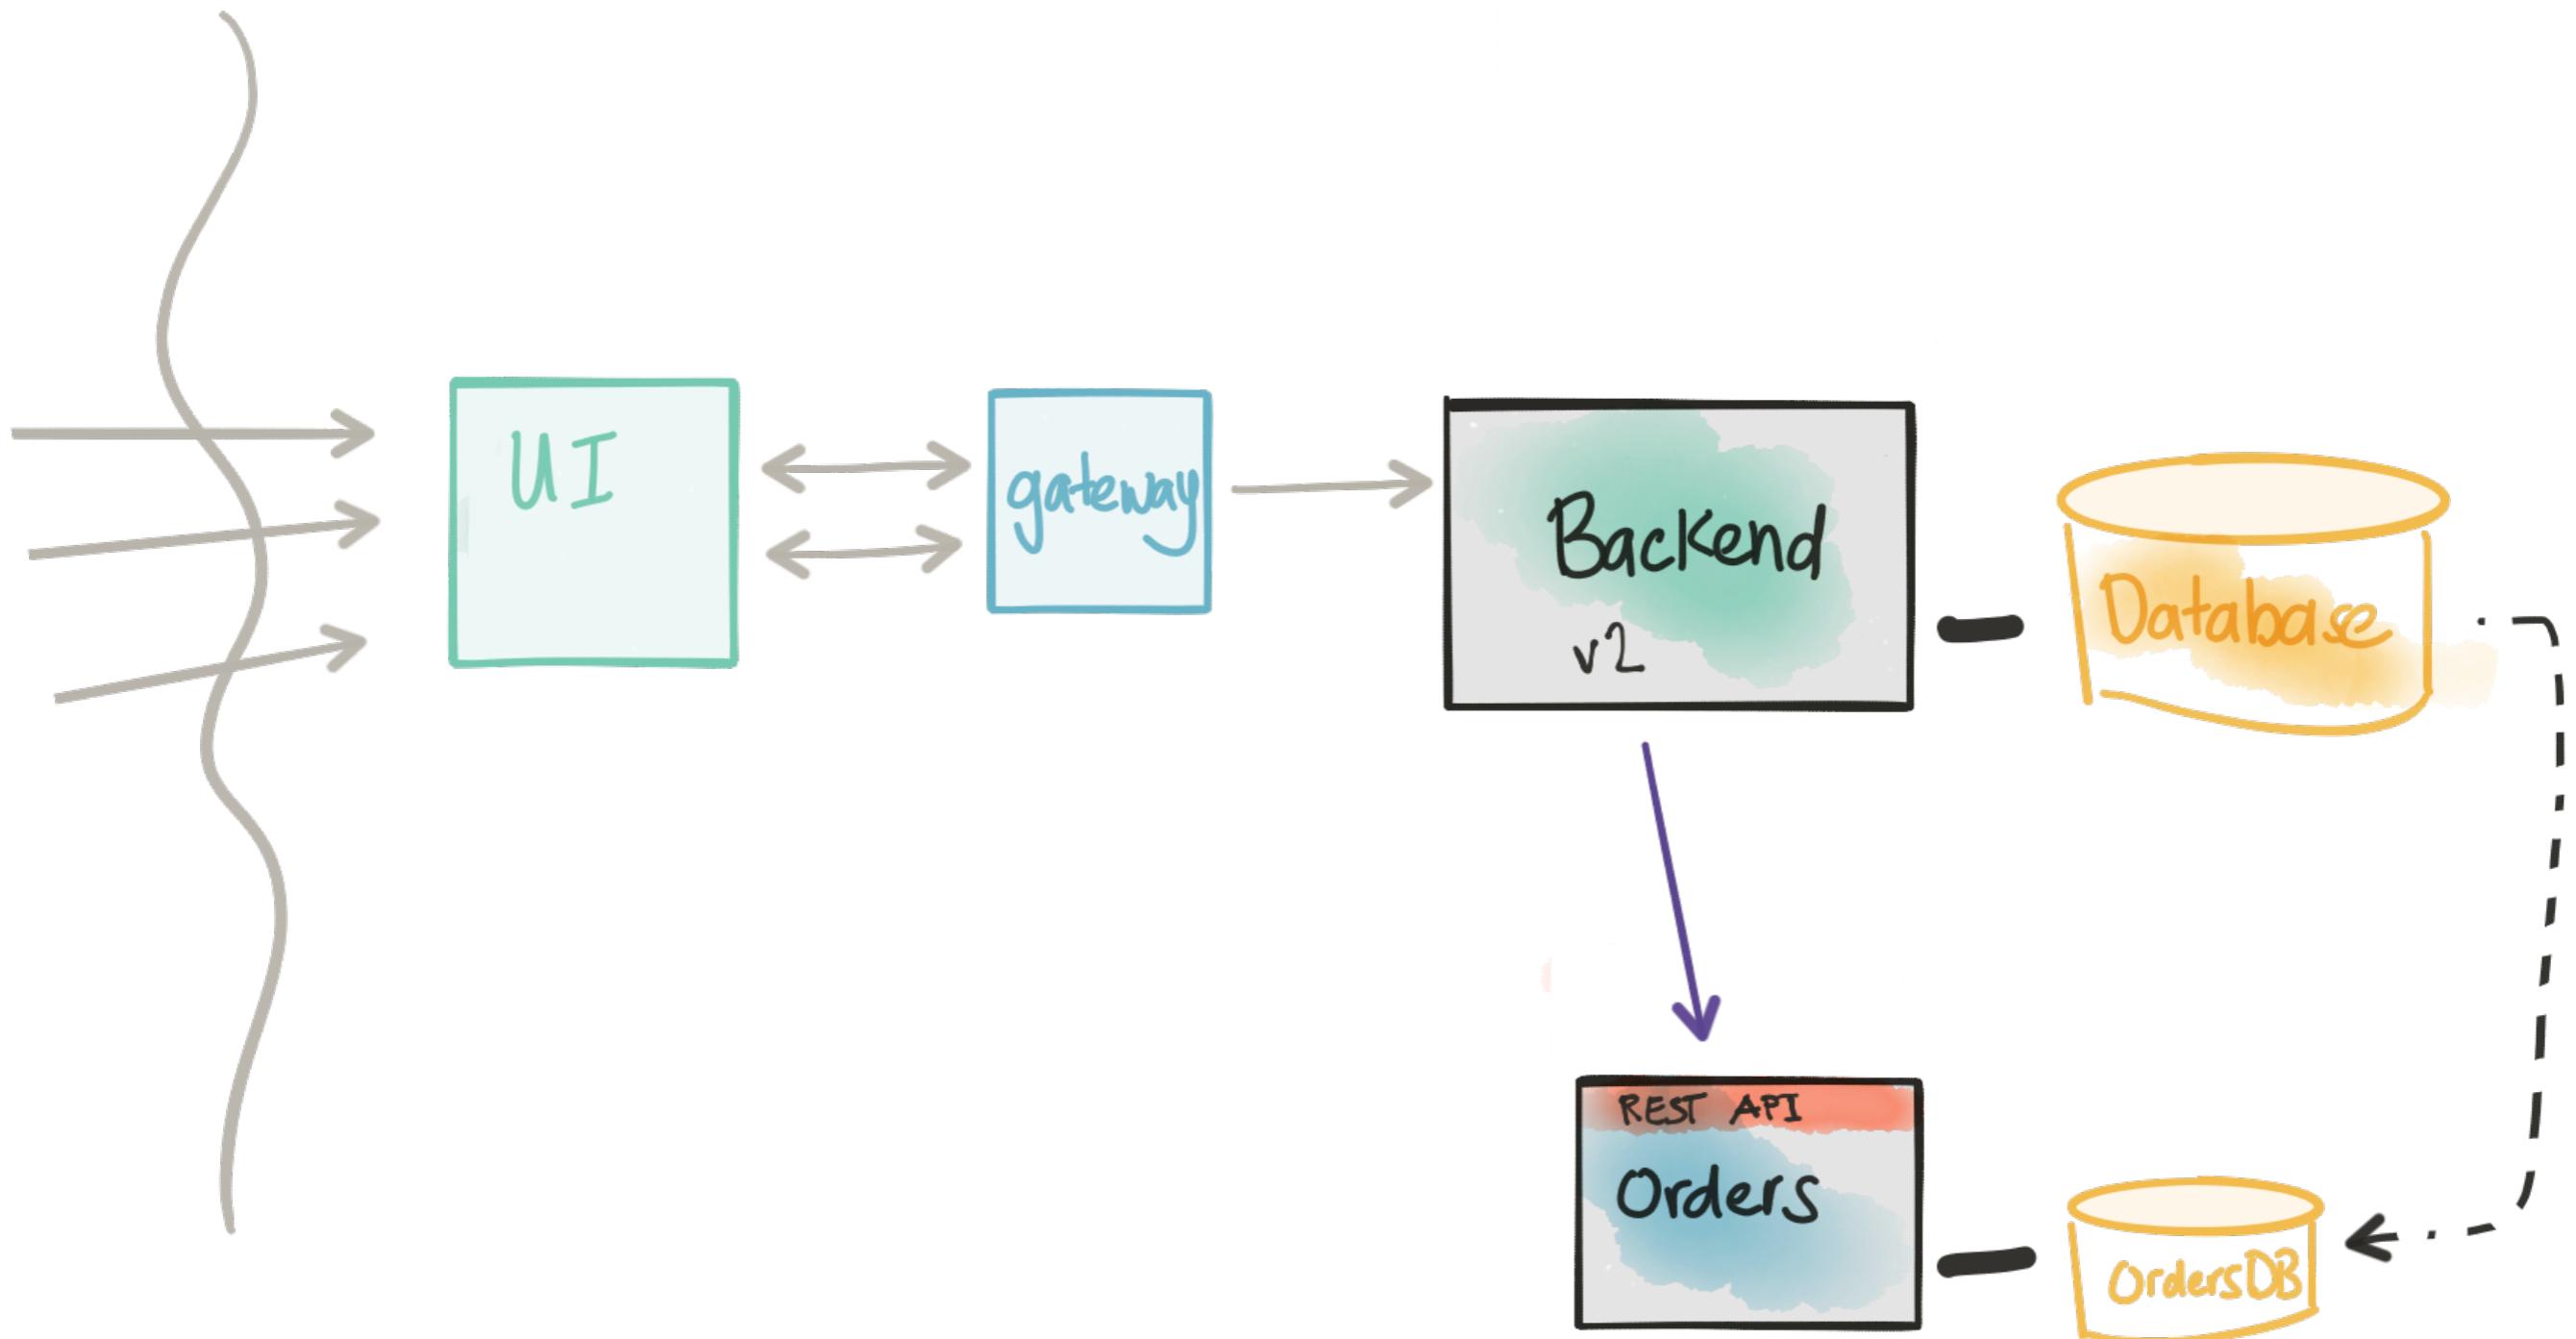




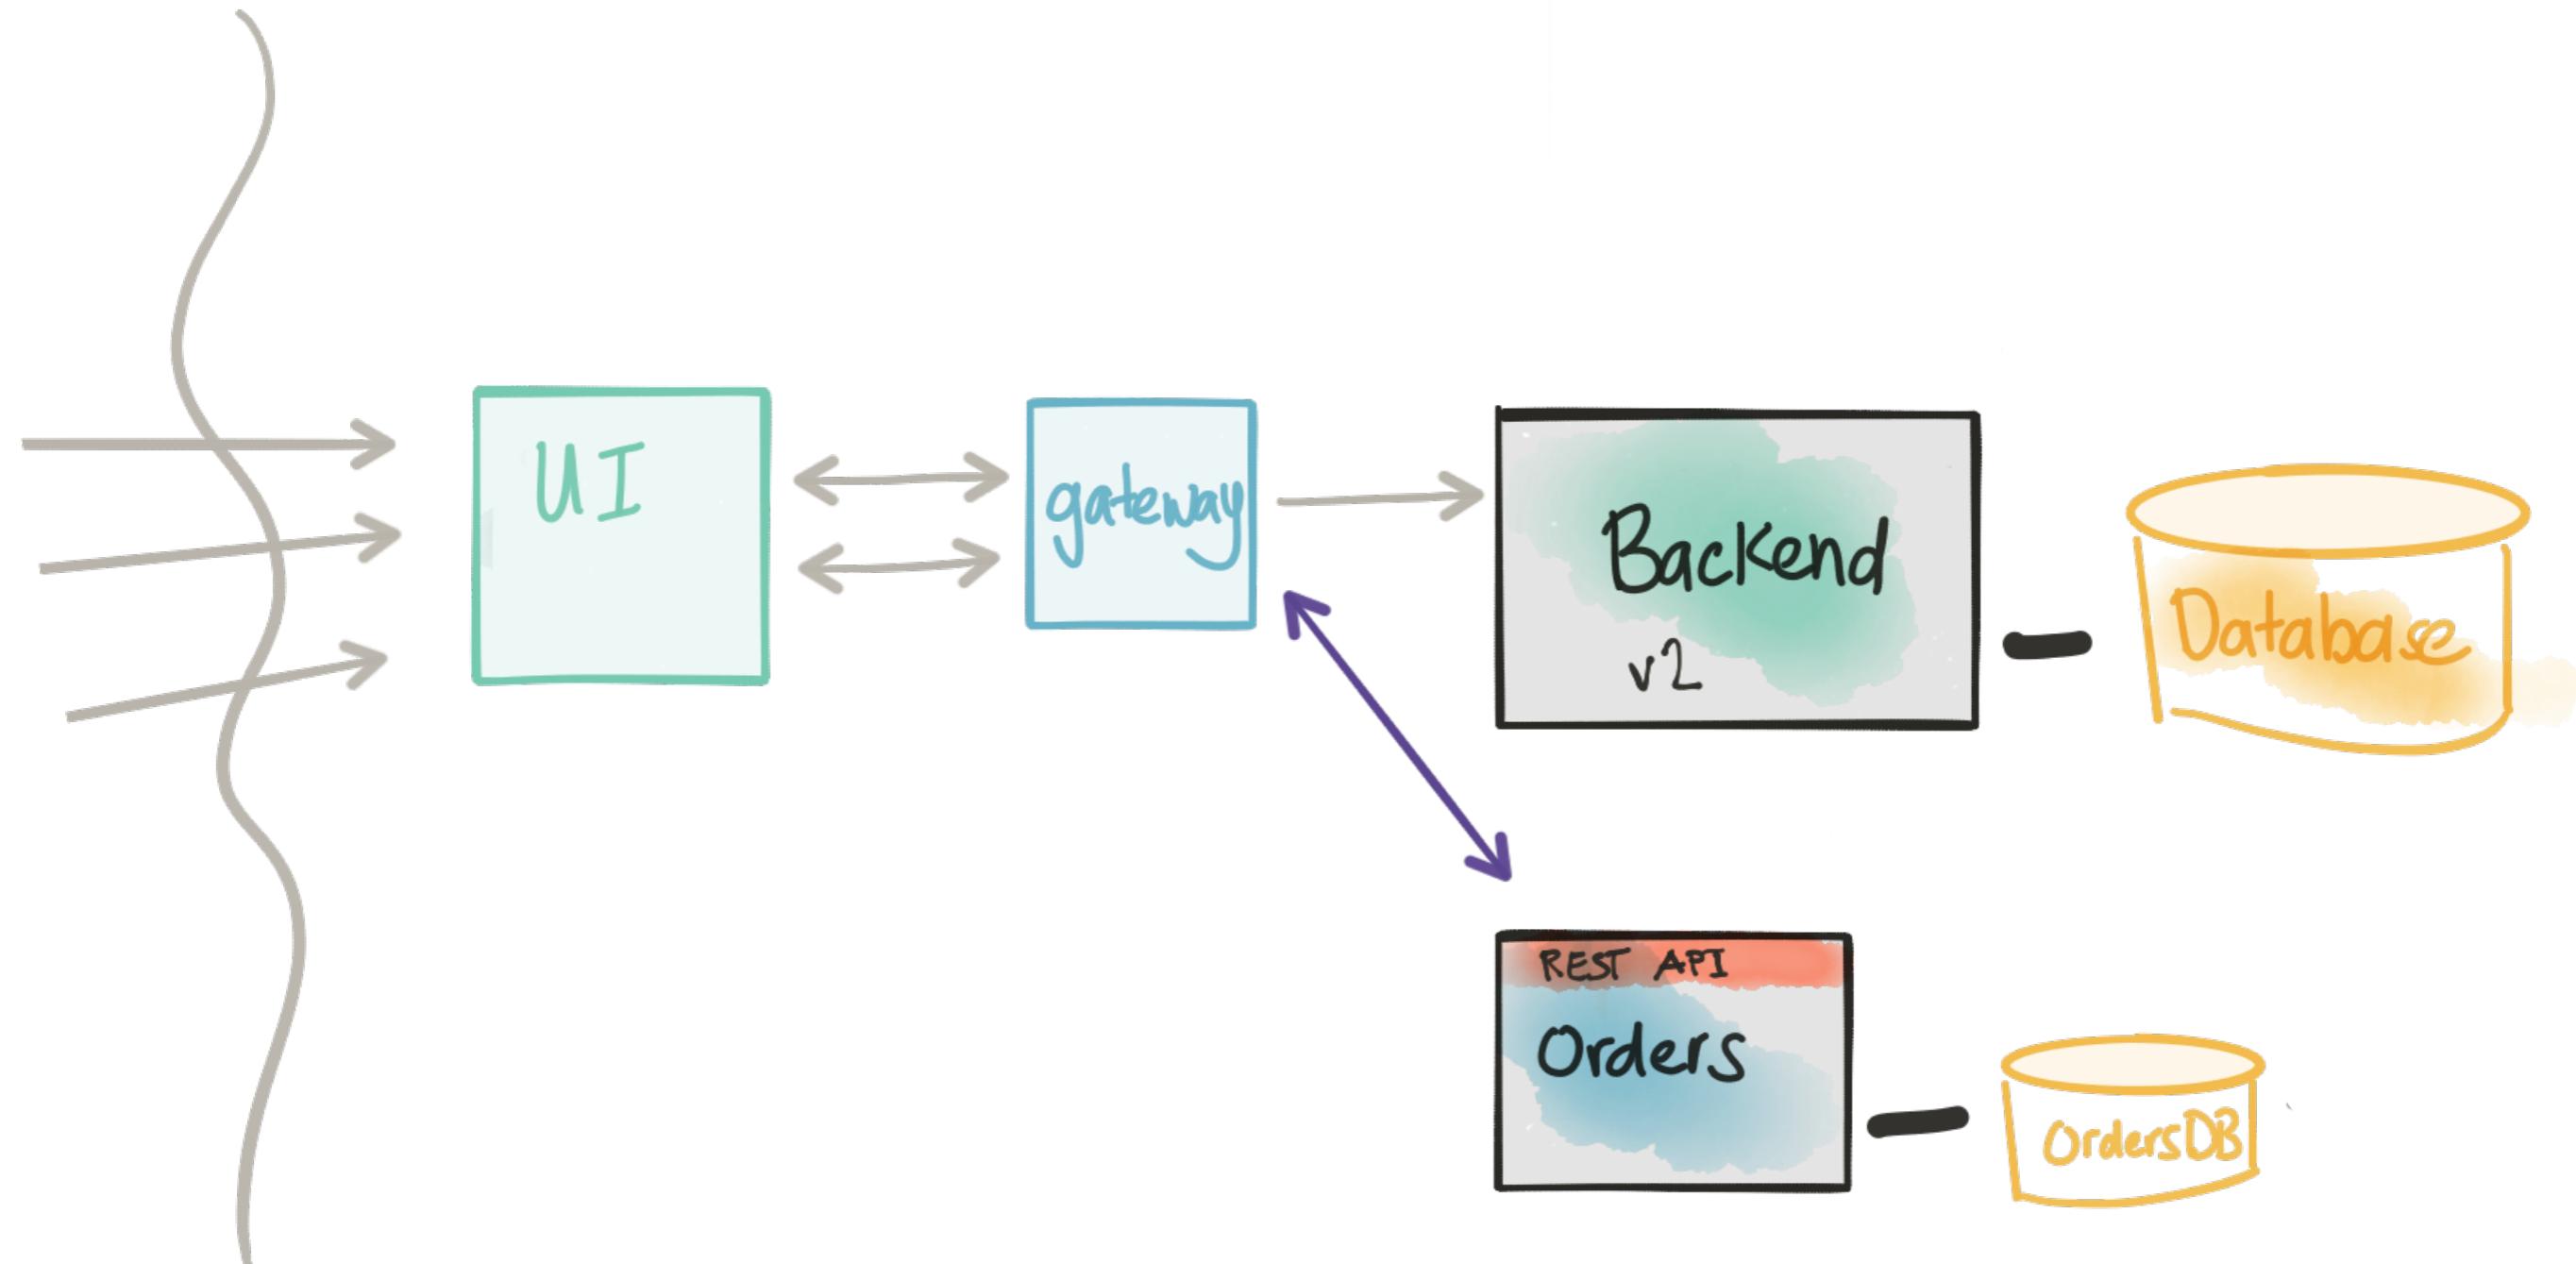
Rolling release to new service



Offline data ETL/migration



Disconnect/decouple the datastore





kubernetes

Hands on lab

Kubernetes

- Pods
- Services



kubernetes

PODs

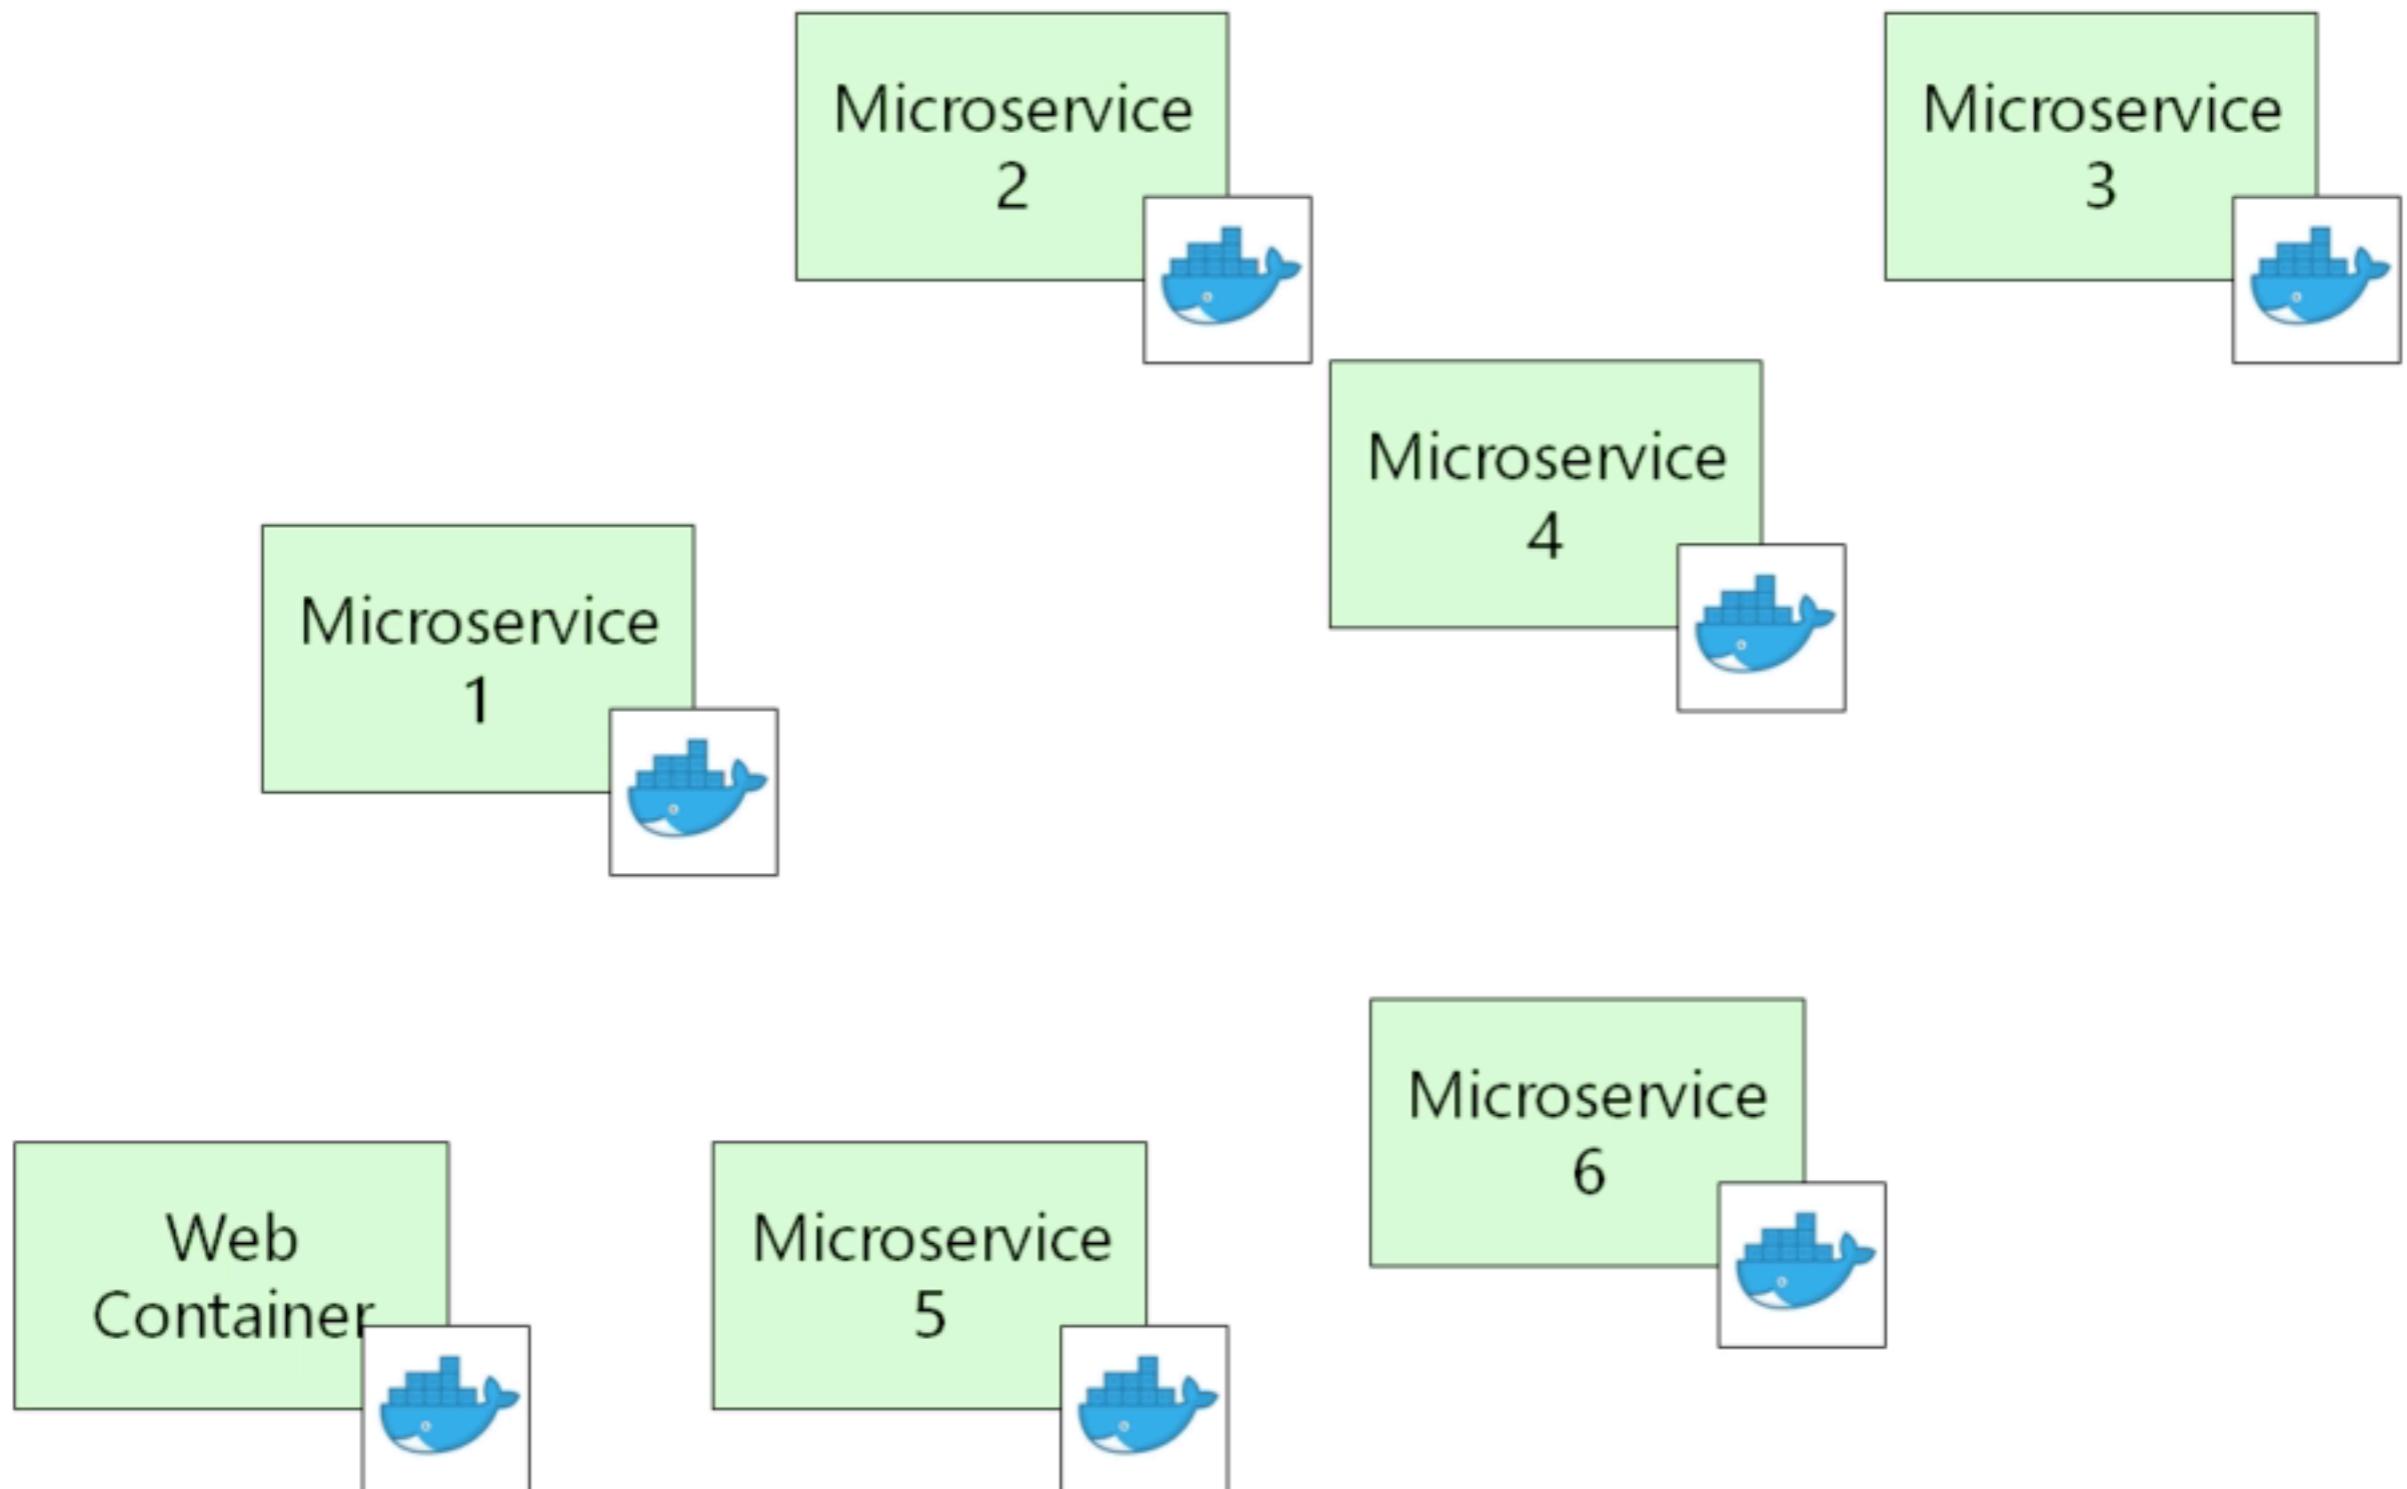
► What is a Pod?



<https://kubernetes.io/docs/concepts/workloads/pods/pod-overview/>

Pods are the most basic concept on kubernetes

► What are we planning to do?



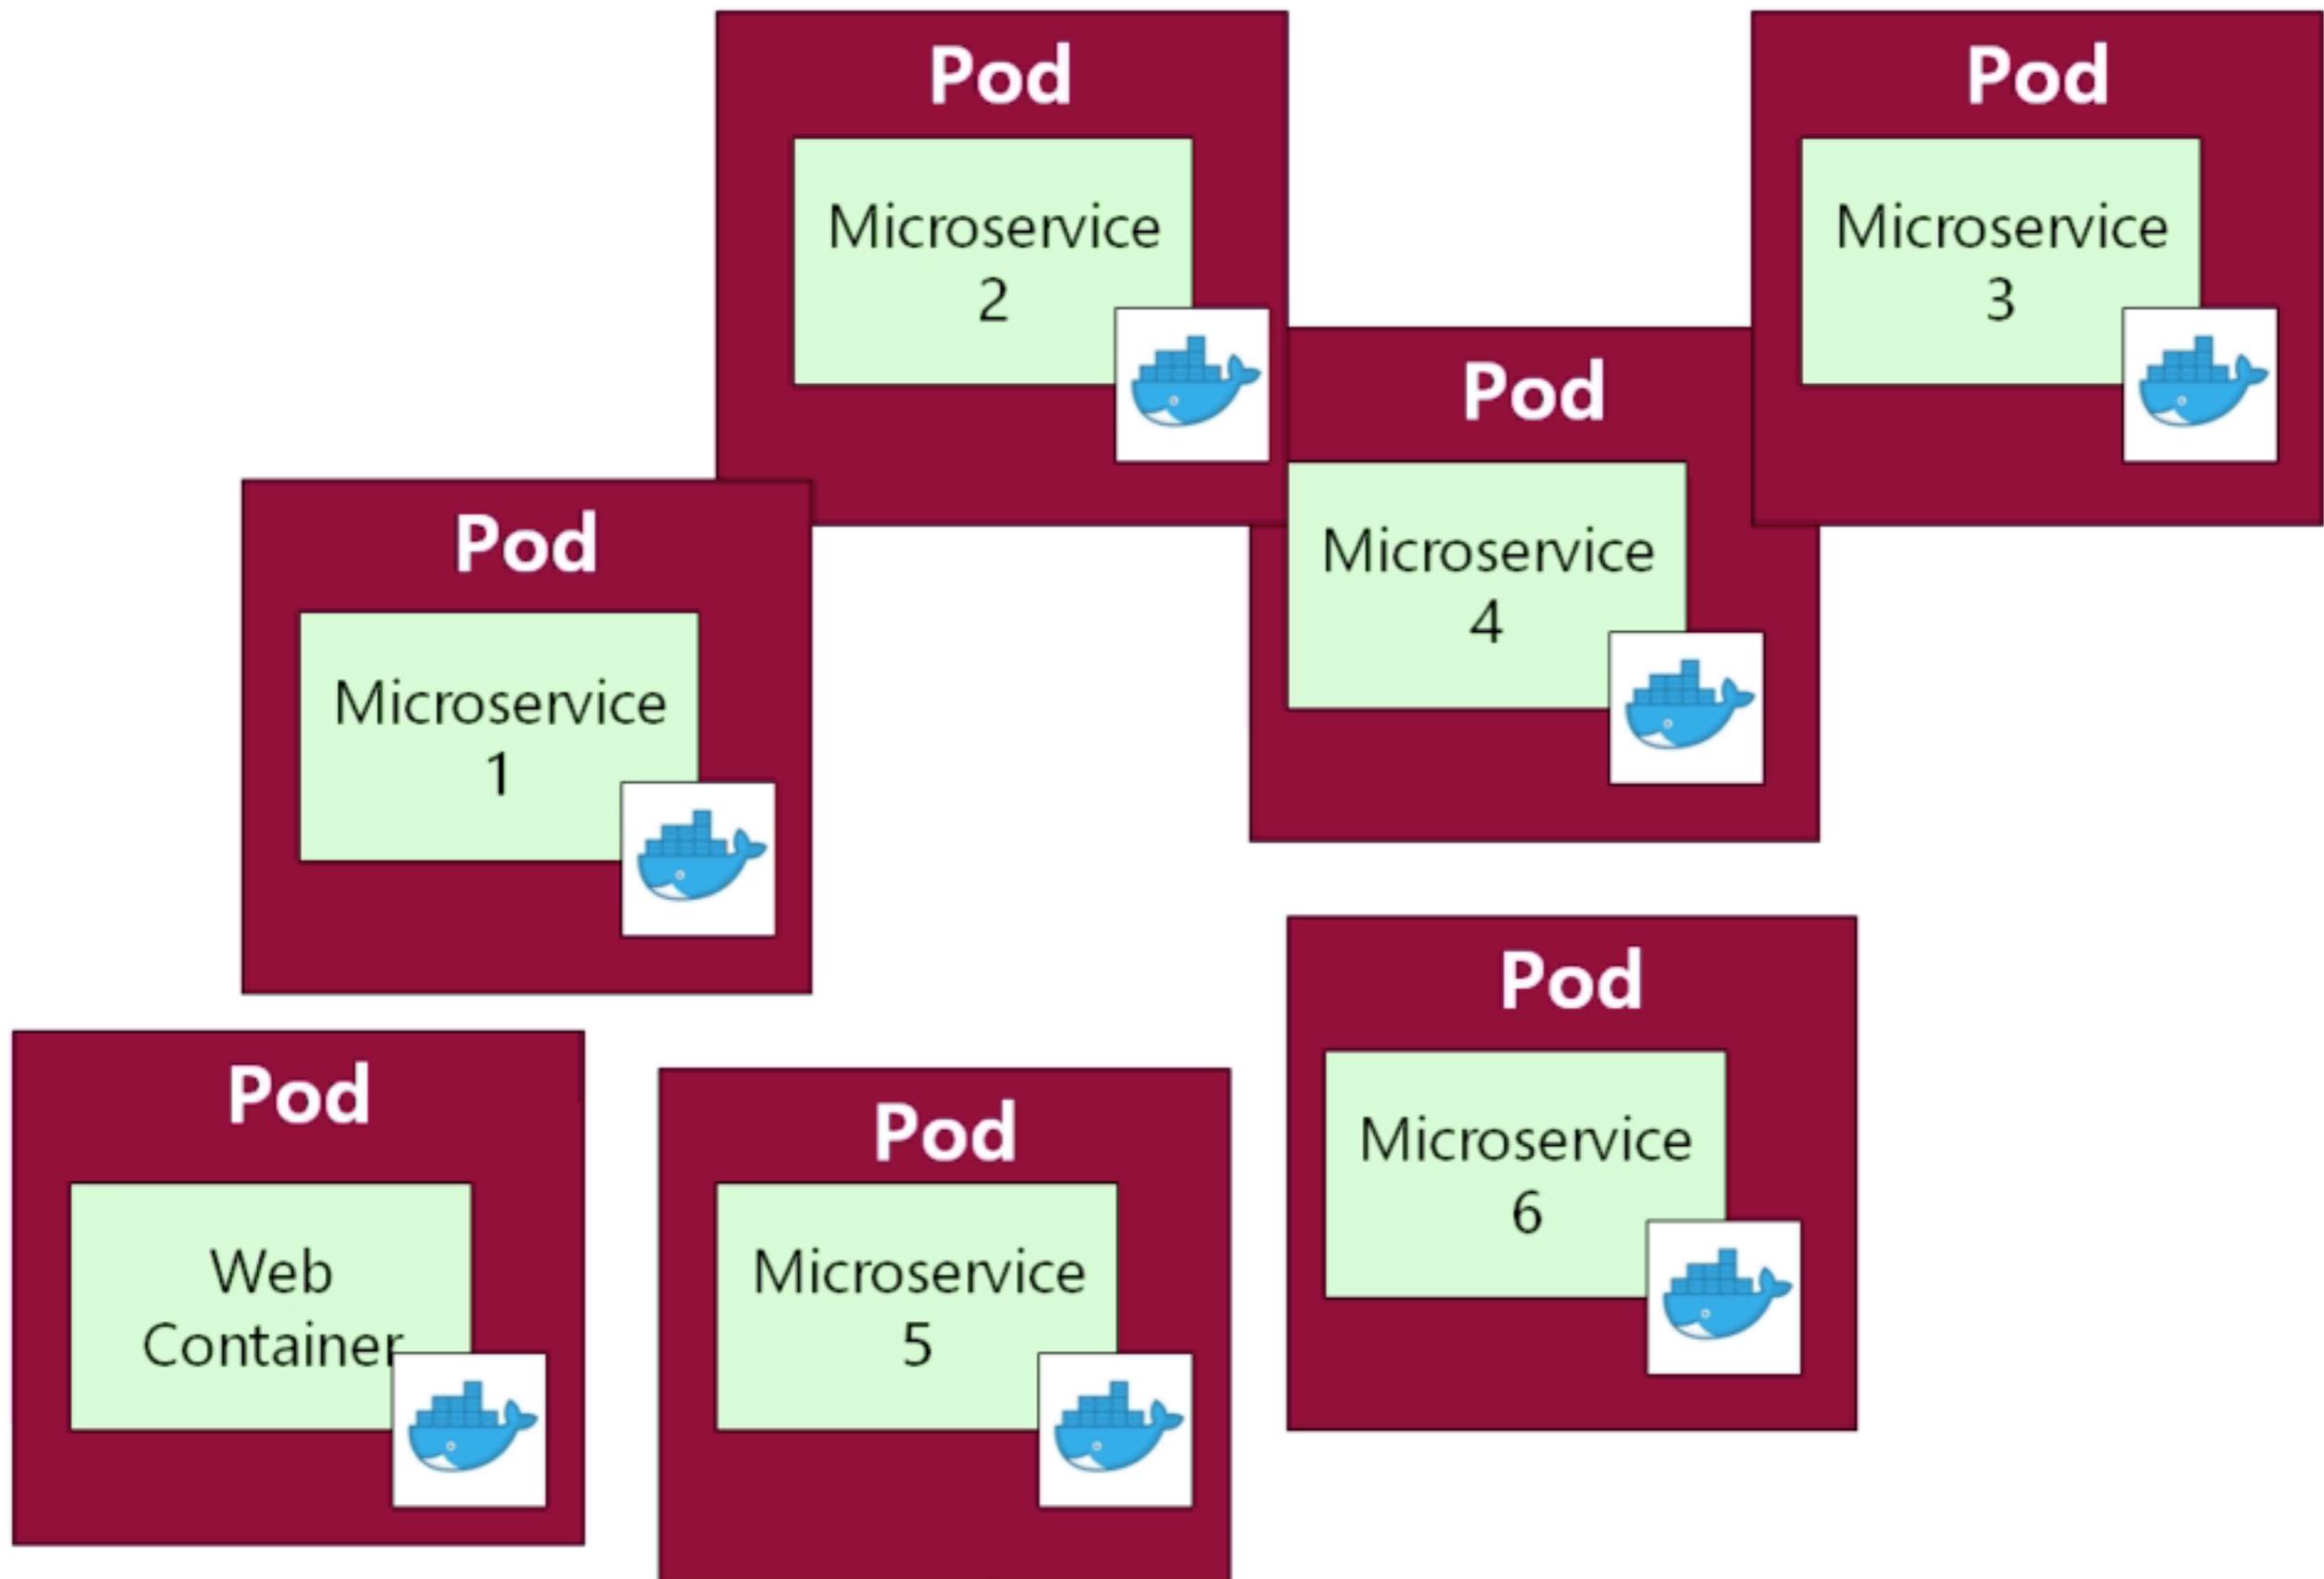
- ▶ Our job is to deploy a micro service architecture to the cloud
- ▶ Can we manage it ourselves?

We could log on to an Amazon AWS node and deploy each Microservice

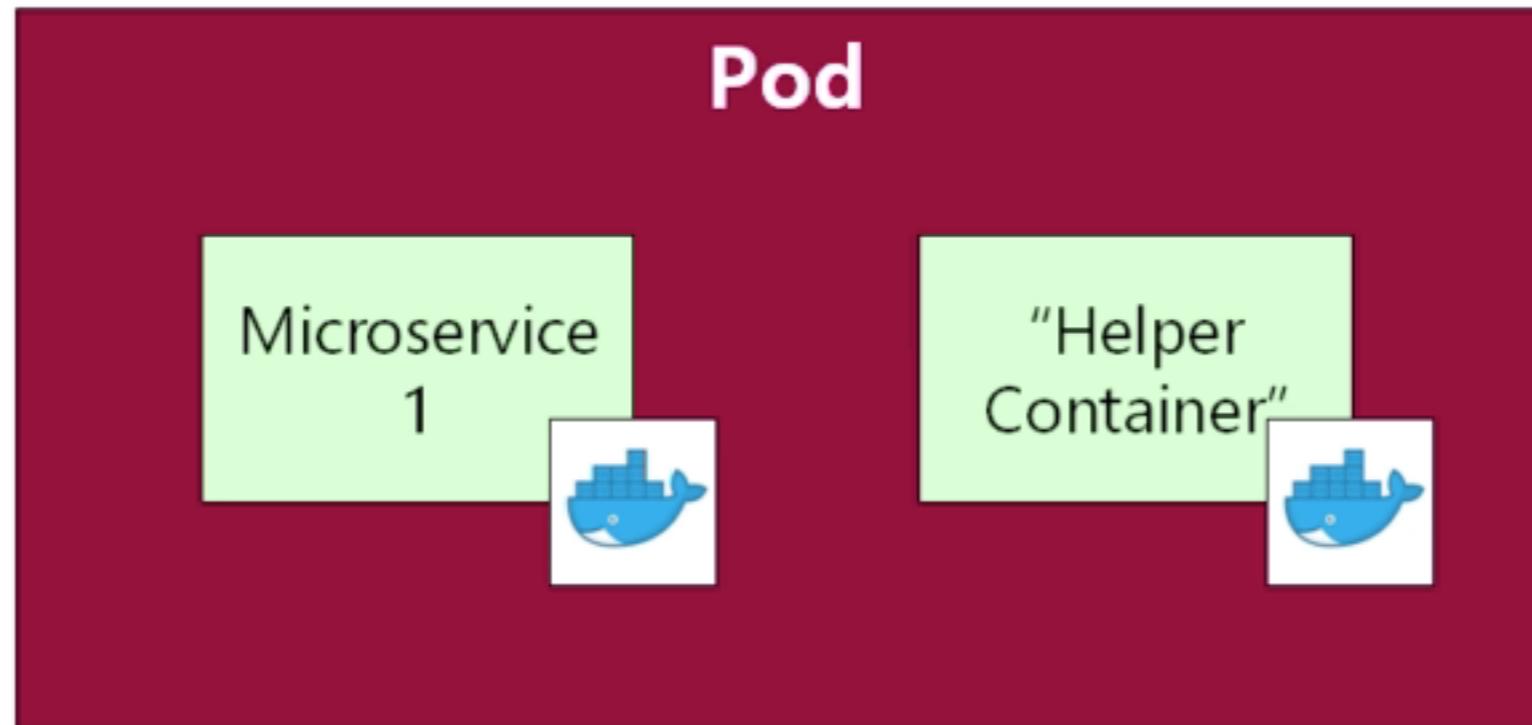
but who manages it?

It will be too much work...

- ▶ We will use kubernetes to orchestrate this system
- ▶ Responsible for starting and stopping these containers

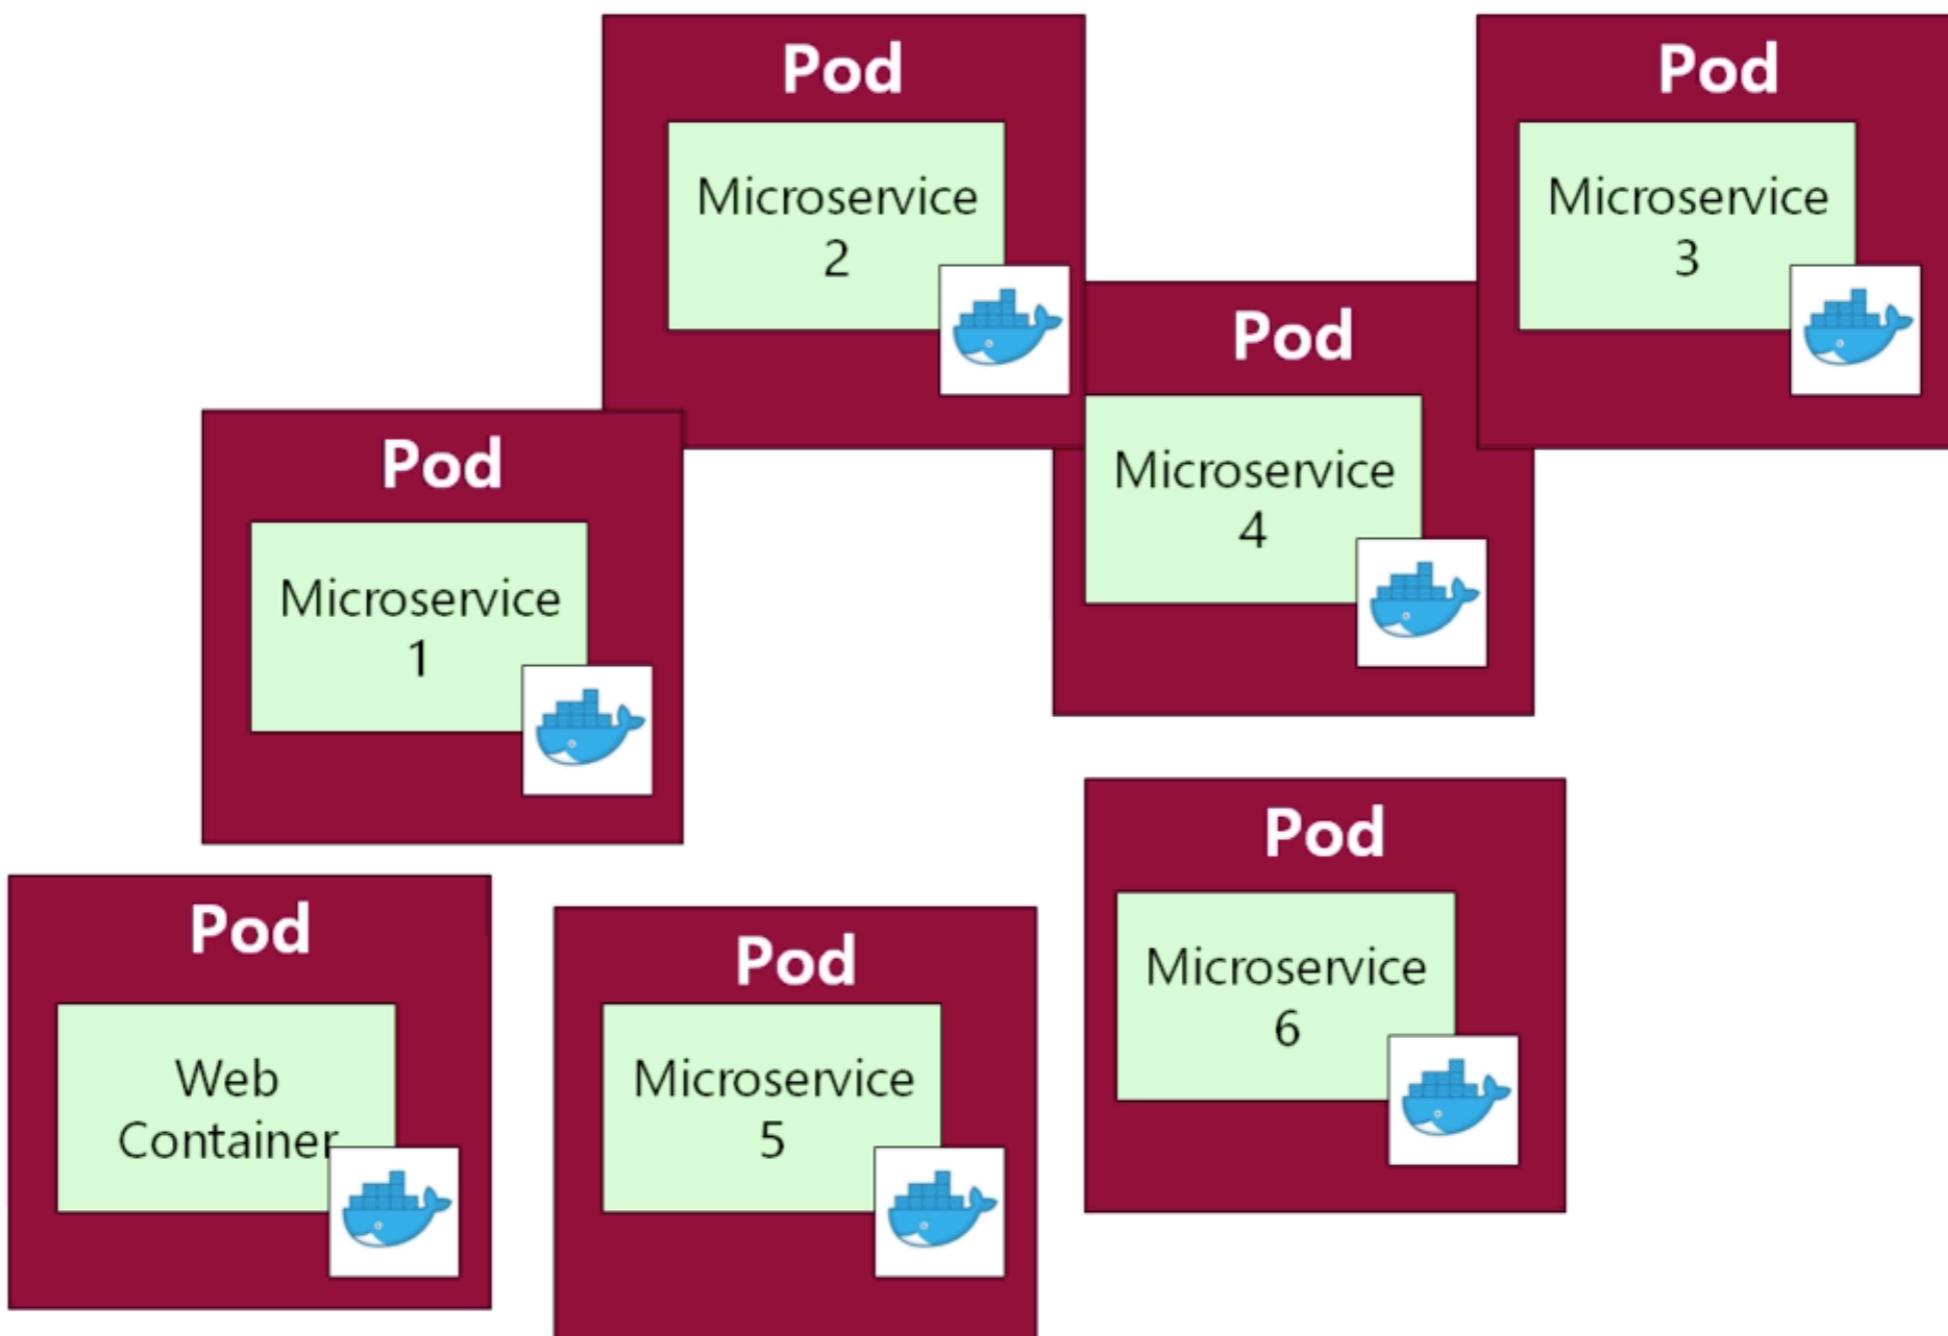


*each pod can be seen as a
wrapper to the container*



- ▶ We can have more than one container per pod
 - ▶ but it's not recommended
 - ▶ It's rare to see an architecture like that

- ▶ Kubernetes will manage these pods



- ▶ Pods are the most basic unit of deployment on kubernetes

Now let's write our first pod!

release0 9.76 MBLast updated **a year ago** by [richardchesterwood](#)

DIGEST	ARCHITECTURE	OS	SIZE
9b98fec20772	amd64	linux	9.76 MB

release0-5 9.76 MBLast updated **a year ago** by [richardchesterwood](#)

DIGEST	ARCHITECTURE	OS	SIZE
636f57ab20c3	amd64	linux	9.76 MB

release2 25.07 MBLast updated **a year ago** by [richardchesterwood](#)

DIGEST	ARCHITECTURE	OS	SIZE
ed7d720878ac	amd64	linux	25.07 MB

release1 25.07 MBLast updated **a year ago** by [richardchesterwood](#)

DIGEST	ARCHITECTURE	OS	SIZE
46bef951fc3c	amd64	linux	25.07 MB

How do we make a pod to deploy this image?

We want *kubernetes* to manage it!

- ▶ Release0 image is the first attempt to show a map

- Go to the documentation of a Pod

API Object

Pod is a top-level resource in the Kubernetes REST API. The [Pod API object](#) definition describes the object in detail.

Pod v1 core

[kubectl example](#)[curl example](#)

Pod Config to print "Hello World".

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
spec:
  containers:
  - name: ubuntu
    image: ubuntu:trusty
    command: ["echo"]
    args: ["Hello World"]
```

**on the
API page**

▶ YAML file for a basic pod

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
spec:
  containers:
    - name: ubuntu
      image: ubuntu:trusty
      command: ["echo"]
      args: ["Hello World"]
```

give the pod a name - think in a good name

name of the container in docker

image that is going to be build from optional

- ▶ First you need to see if minikube is up and running - **minikube status**

```
~ $ minikube status
host: Running
kubelet: Running
apiserver: Running
kubectl: Correctly Configured: pointing to minikube-vm at 192.168.99.100
```

- ▶ Create a YAML text file for the pod - **pod.yaml**

! pod.yaml

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: webapp
5    labels:
6      app: webapp
7      release: "0"
8  spec:
9    containers:
10   - name: webapp
11     image: richardchesterwood/k8s-fleetman-webapp-angular:release0
12
```

- ▶ `ls` on the folder of the file

```
~/mtsds ➔  
$ ls  
pod.yaml
```

- ▶ `kubectl get all`

```
~/mtsds ➔  
$ kubectl get all  
NAME                      READY   STATUS    RESTARTS   AGE  
service/kubernetes        ClusterIP  10.96.0.1  <none>    443/TCP
```

- ▶ `kubectl apply -f pod.yaml`

```
~/mtstsds
$ kubectl apply -f pod.yaml
pod/webapp created
```

- ▶ `kubectl get all`

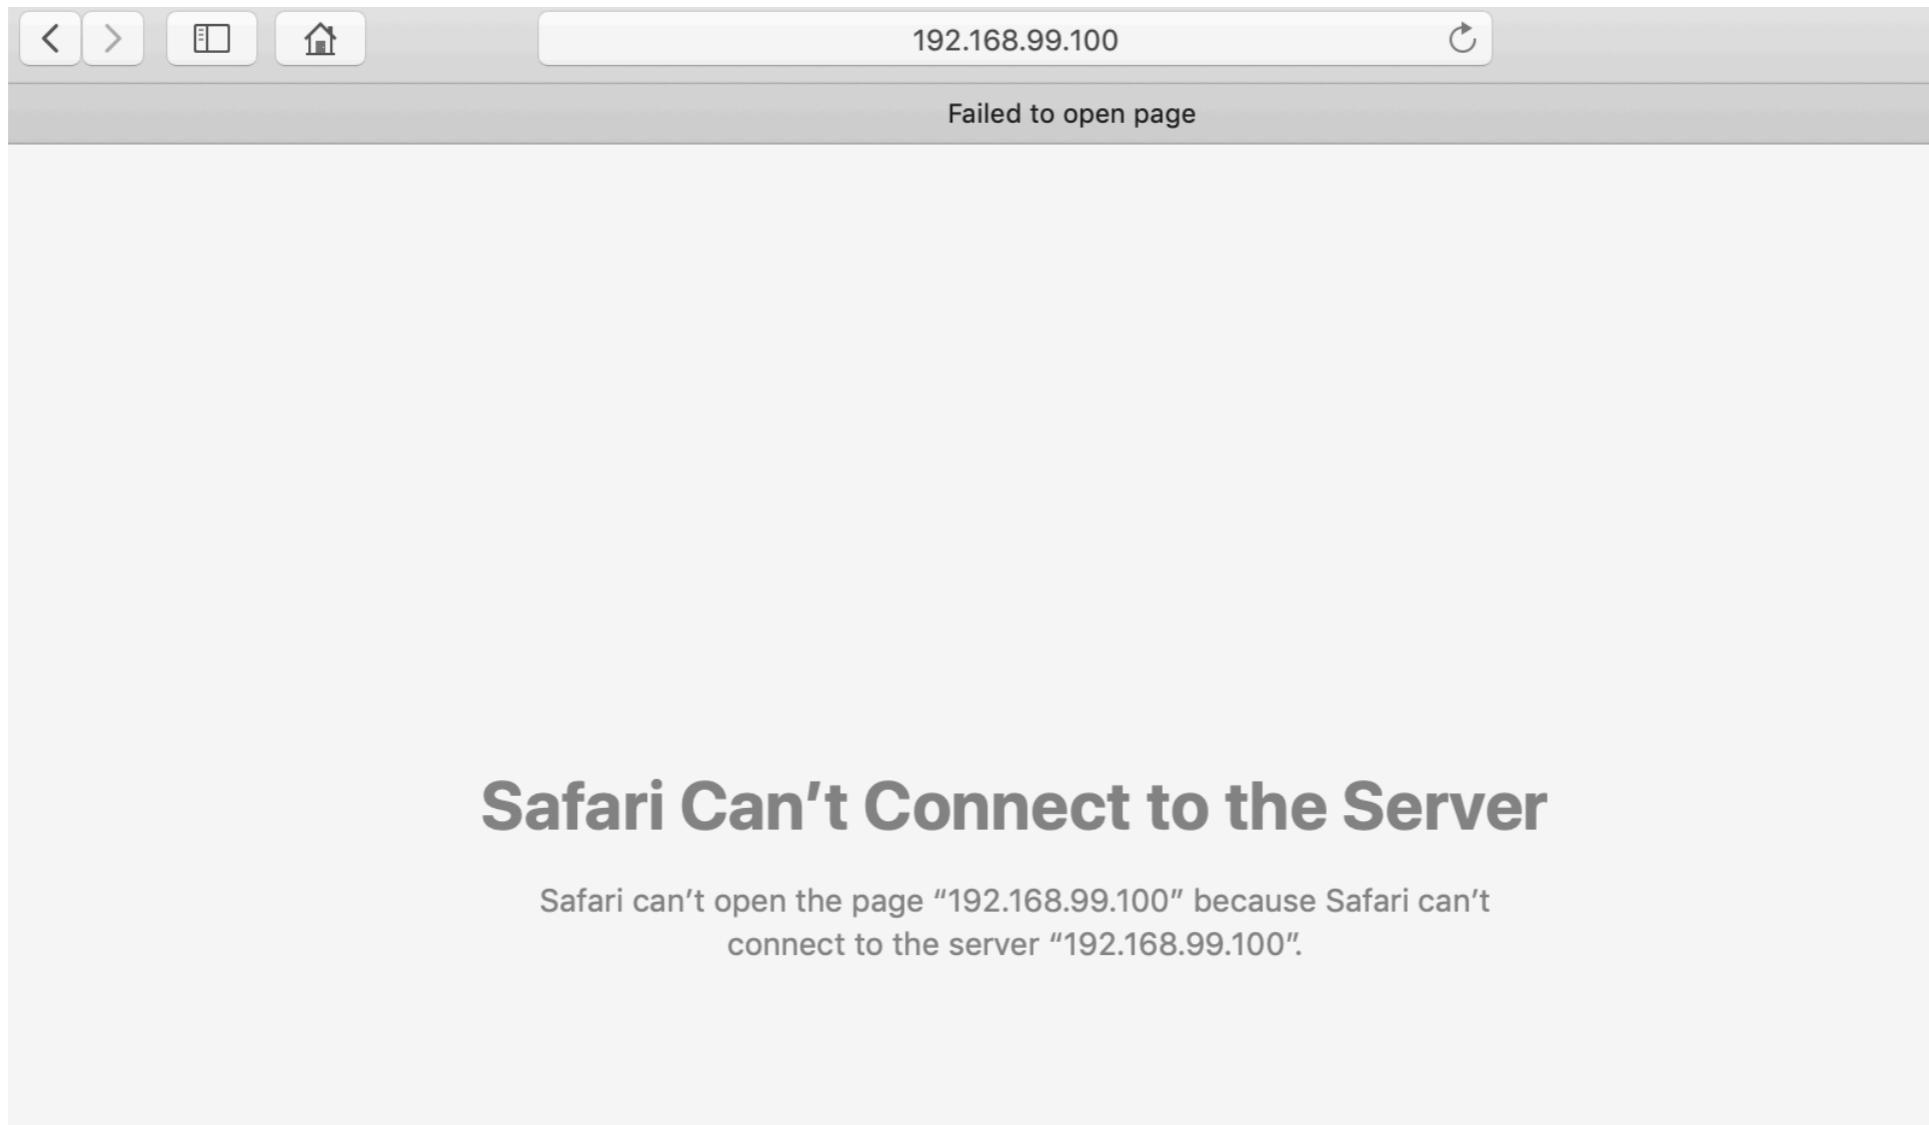
```
~/mtstsds
$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/webapp	1/1	Running	0	58s
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP

*Now we have our first
Kubernetes pod running on
a cluster!*

► Let's try it - minikube ip

```
~/mtstsds ➔  
$ ➔ minikube ip  
192.168.99.100
```



Kubernetes Cluster

Kubernetes Cluster

Pod

Web
Container
:80

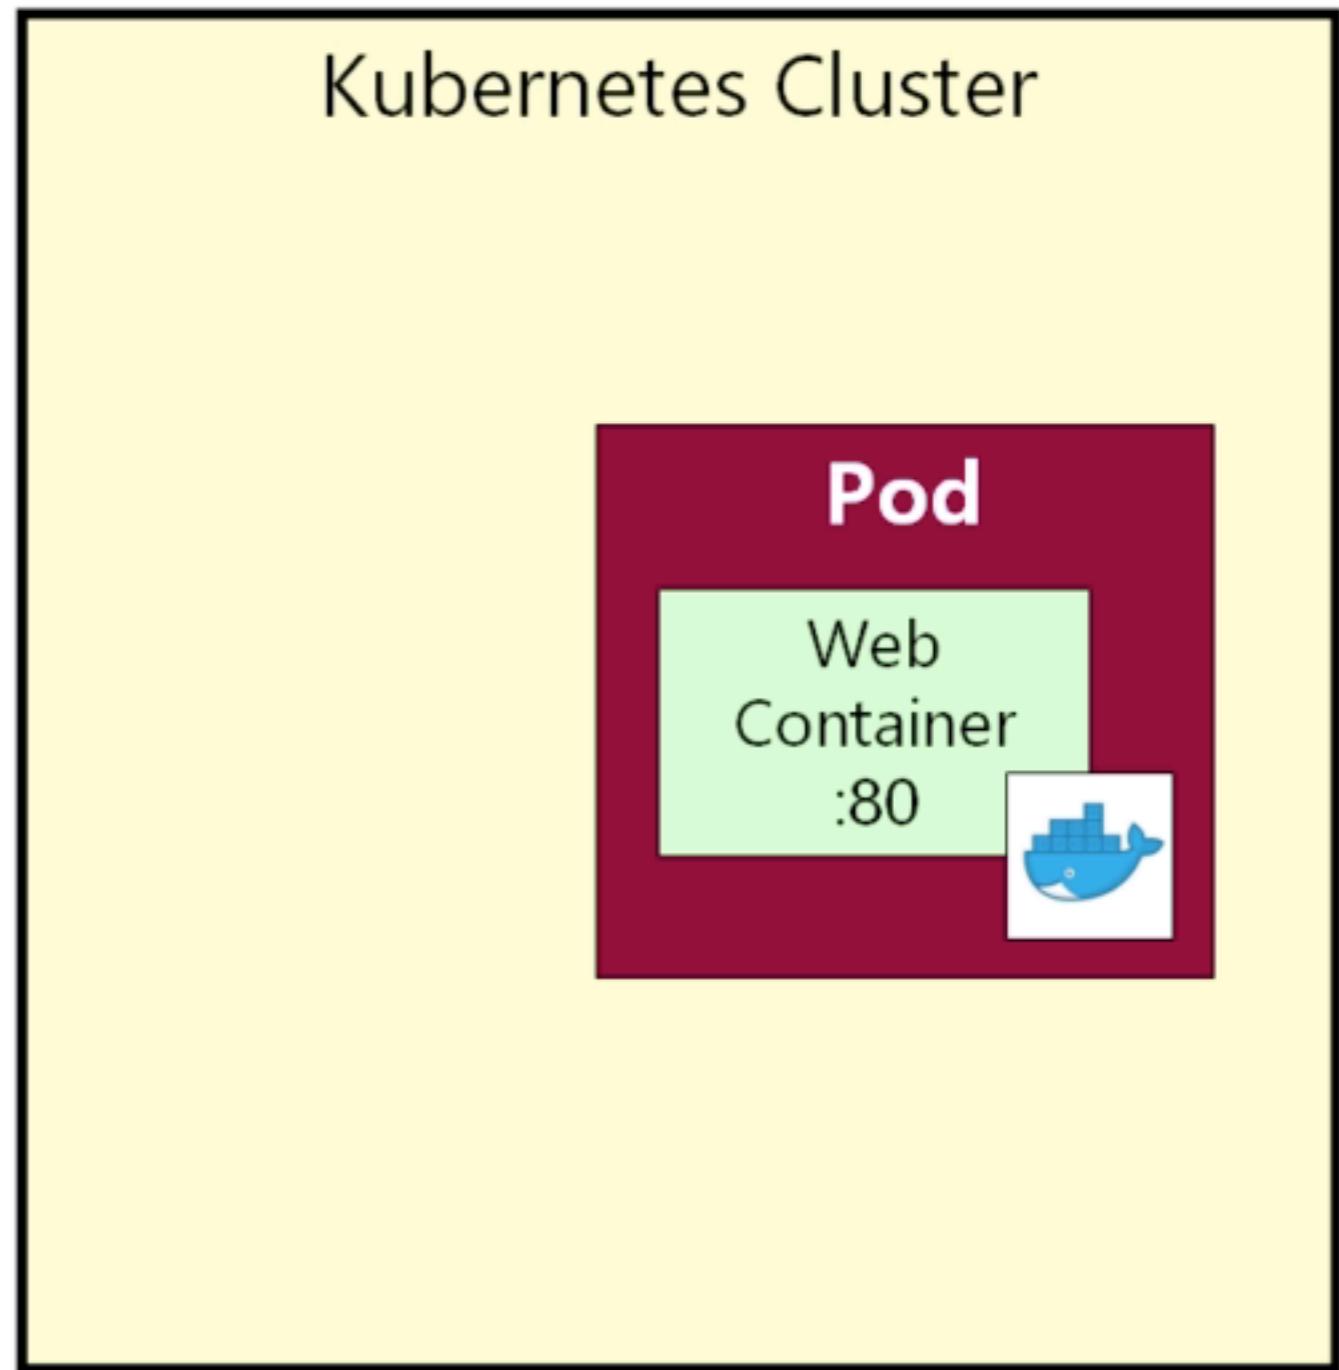
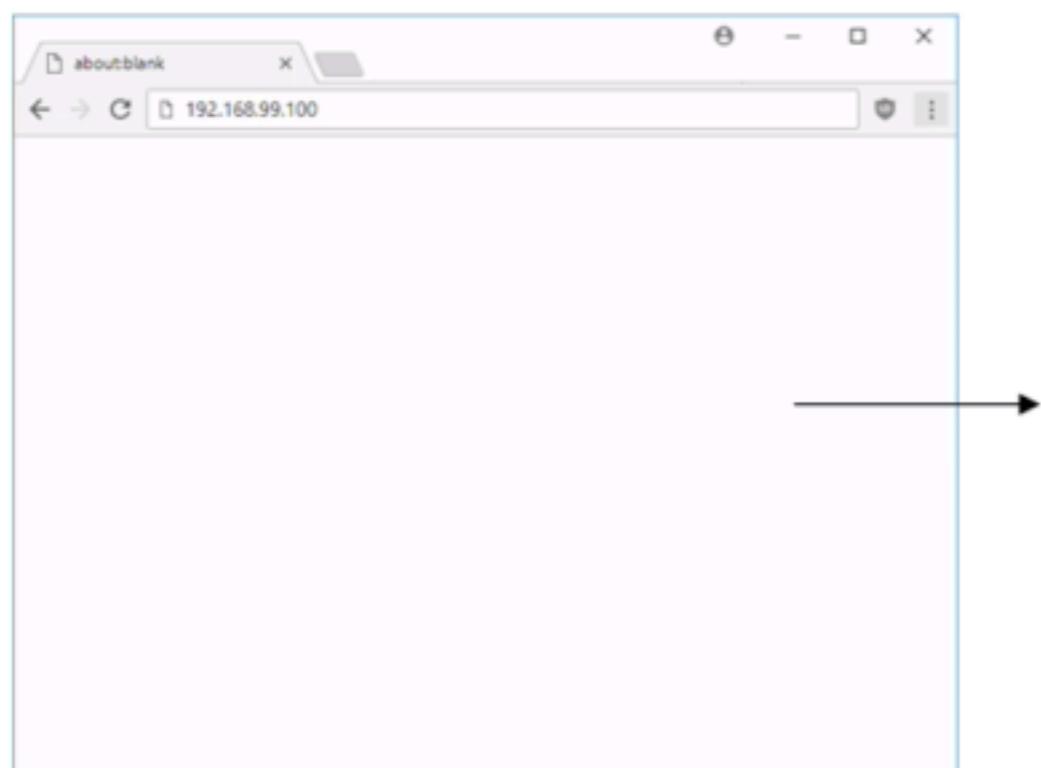


Kubernetes Cluster

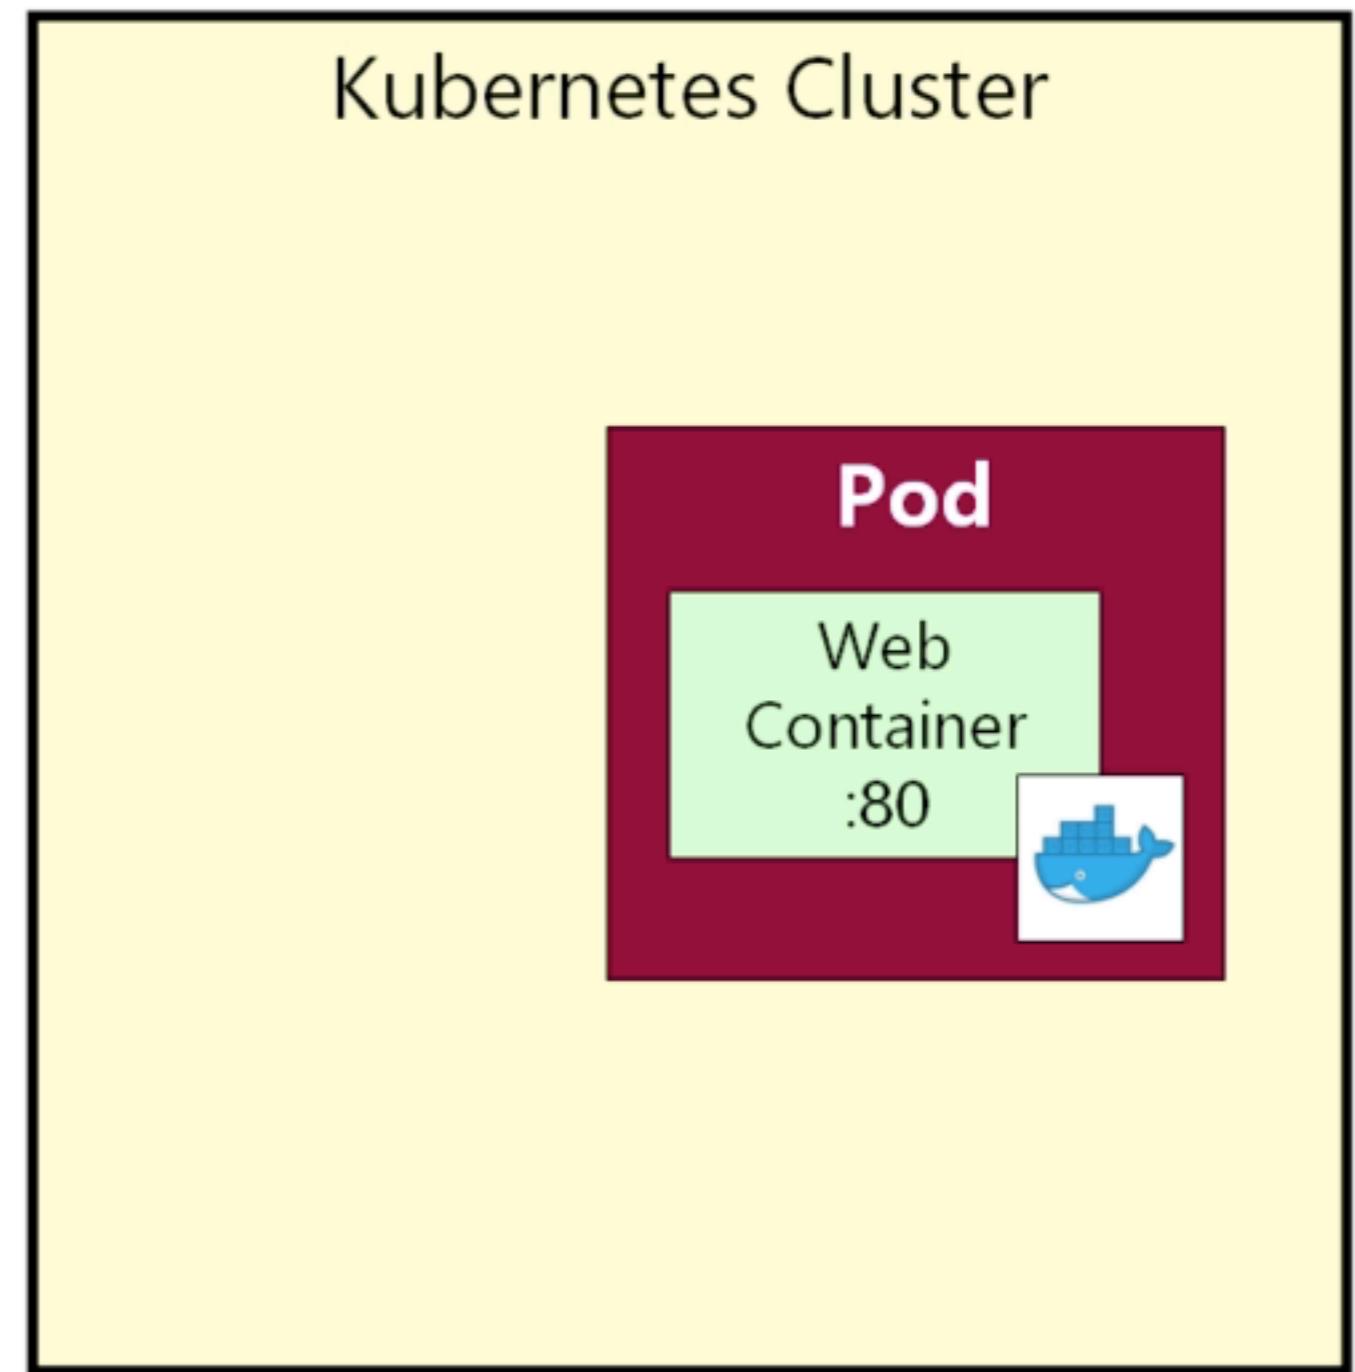
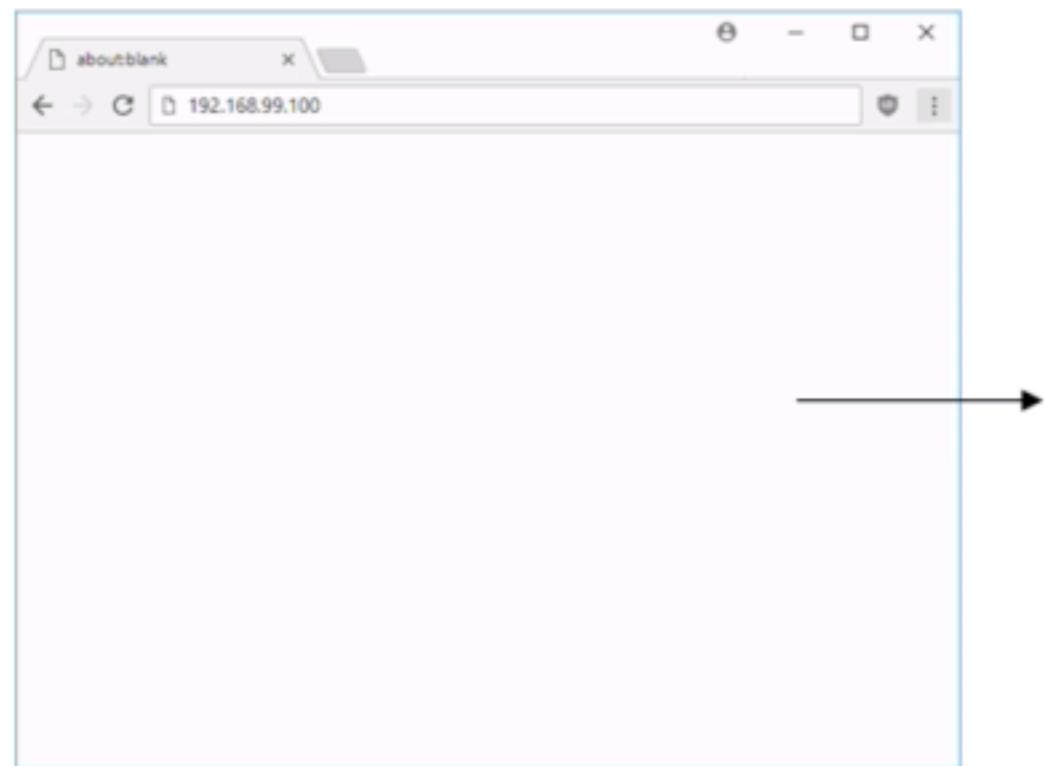
Pod

Web
Container
:80





Pods are not visible outside the cluster!



Pods are not visible outside the cluster!

- ▶ Part of the design of pods is that they are completely isolated and only accessible from inside the cluster itself
- ▶ We can't visit this pod from a web browser

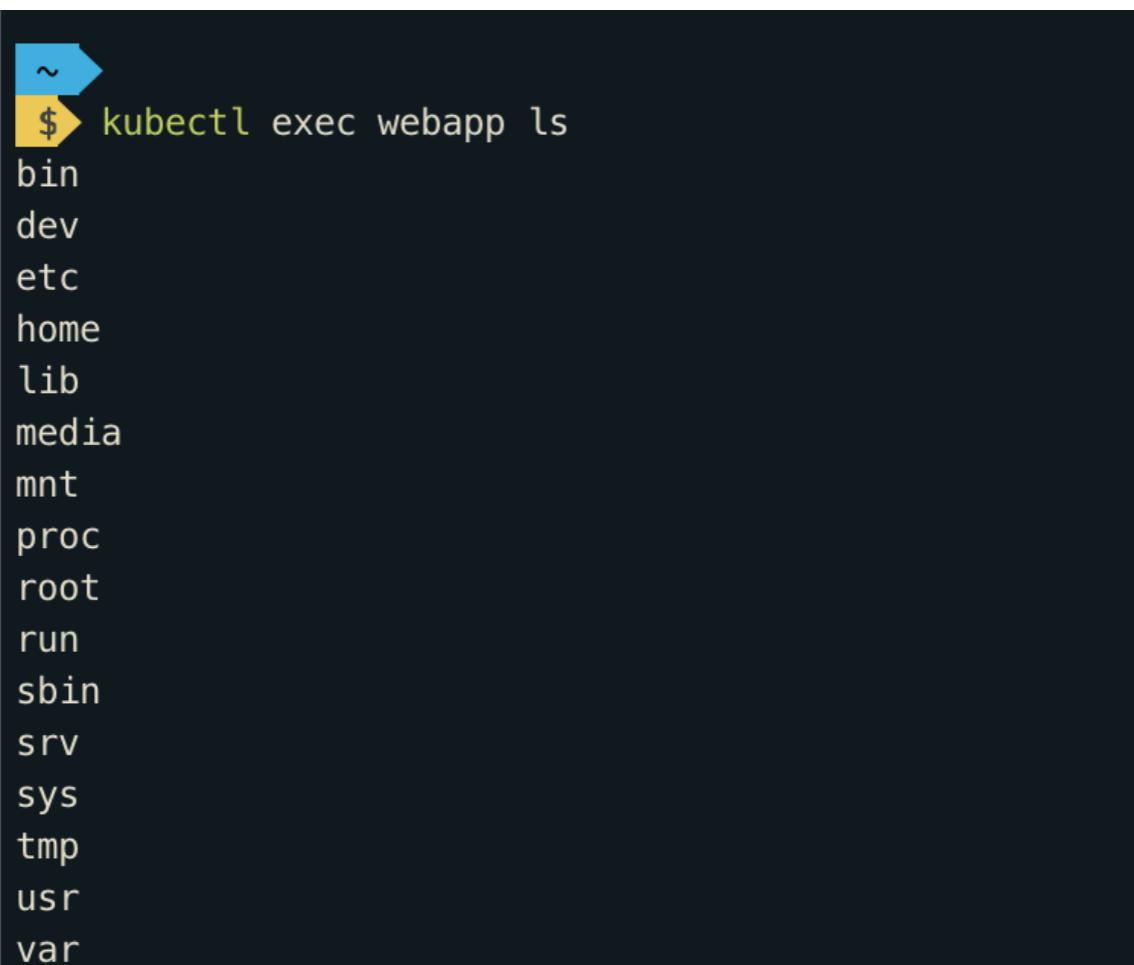
*Let's try some useful
commands*

► kubectl describe pod webapp

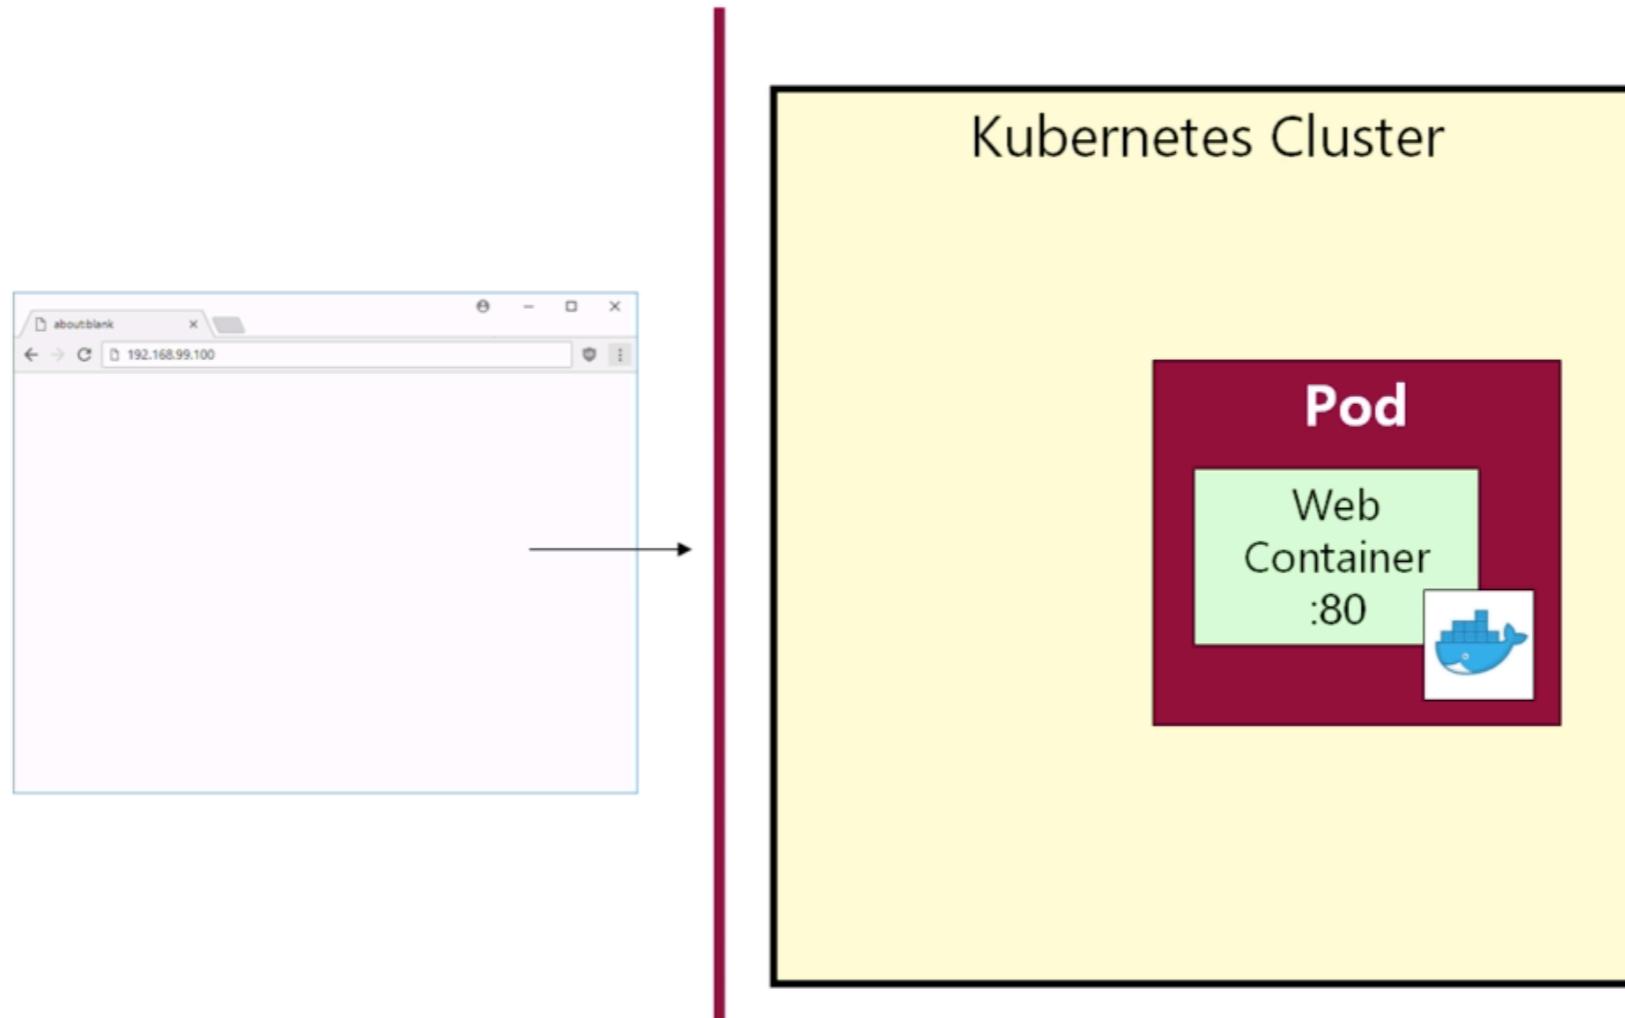
```
~ $ kubectl describe pod webapp
Name:           webapp
Namespace:      default
Priority:       0
Node:           minikube/10.0.2.15
Start Time:     Tue, 22 Oct 2019 11:35:21 +0100
Labels:          app=webapp
                  release=0
Annotations:    kubectl.kubernetes.io/last-applied-configuration:
                  {"apiVersion":"v1","kind":"Pod","metadata":{"annotations":{},"labels":{"webapp","namespace":"default"}}
Status:          Running
IP:             172.17.0.7
IPs:
  IP:  172.17.0.7
Containers:
```

node.kubernetes.io/unreachable: NOEXECUTE for 300s				
Events:				
Type	Reason	Age	From	Message
Normal	Scheduled	<unknown>	default-scheduler	Successfully assigned default/webapp to minikube
Normal	Pulled	5d23h	kubelet, minikube	Container image "richardchesterwood/k8s-node0" already present on machine
Normal	Created	5d23h	kubelet, minikube	Created container webapp
Normal	Started	5d23h	kubelet, minikube	Started container webapp
Warning	FailedMount	2m12s	kubelet, minikube	MountVolume.SetUp failed for volume "data" : failed to sync secret cache: timed out waiting for the condition
Normal	SandboxChanged	2m10s	kubelet, minikube	Pod sandbox changed, it will be killed
Normal	Pulled	2m9s	kubelet, minikube	Container image "richardchesterwood/k8s-node0" already present on machine
Normal	Created	2m8s	kubelet, minikube	Created container webapp
Normal	Started	2m7s	kubelet, minikube	Started container webapp

- ▶ We can't visit that pod with a browser
- ▶ But there is one thing we can do
- ▶ We can connect to the pod
- ▶ **kubectl exec webapp ls**



```
~ $ kubectl exec webapp ls
bin
dev
etc
home
lib
media
mnt
proc
root
run
sbin
srv
sys
tmp
usr
var
```



- ▶ We could get a shell to the container
- ▶ Do a curl to localhost port 80

- ▶ `kubectl -it exec webapp sh`

```
~ $ kubectl -it exec webapp sh  
/ #
```

- ▶ wget http://localhost:80

```
~ $ kubectl -it exec webapp sh  
/ # wget http://localhost:80  
Connecting to localhost:80 (127.0.0.1:80)  
index.html          100% |*****  
/ #
```

▶ cat index.html

```
~ $ kubectl -it exec webapp sh
/ # wget http://localhost:80
Connecting to localhost:80 (127.0.0.1:80)
index.html          100% |*****|=====
/ # cat index.html
<!doctype html>
<html lang="en">
<head>
  <meta charset="utf-8">
  <title>Fleet Management</title>
  <base href="/">

  <meta name="viewport" content="width=device-width, initial-scale=1">
  <link rel="icon" type="image/x-icon" href="favicon.png">
</head>
<body>
  <app-root></app-root>
<script type="text/javascript" src="runtime.js"></script><script type="text/javascript" src="styles.js"></script><script type="text/javascript" src="main.js"></script></body>
</html>
/ #
```

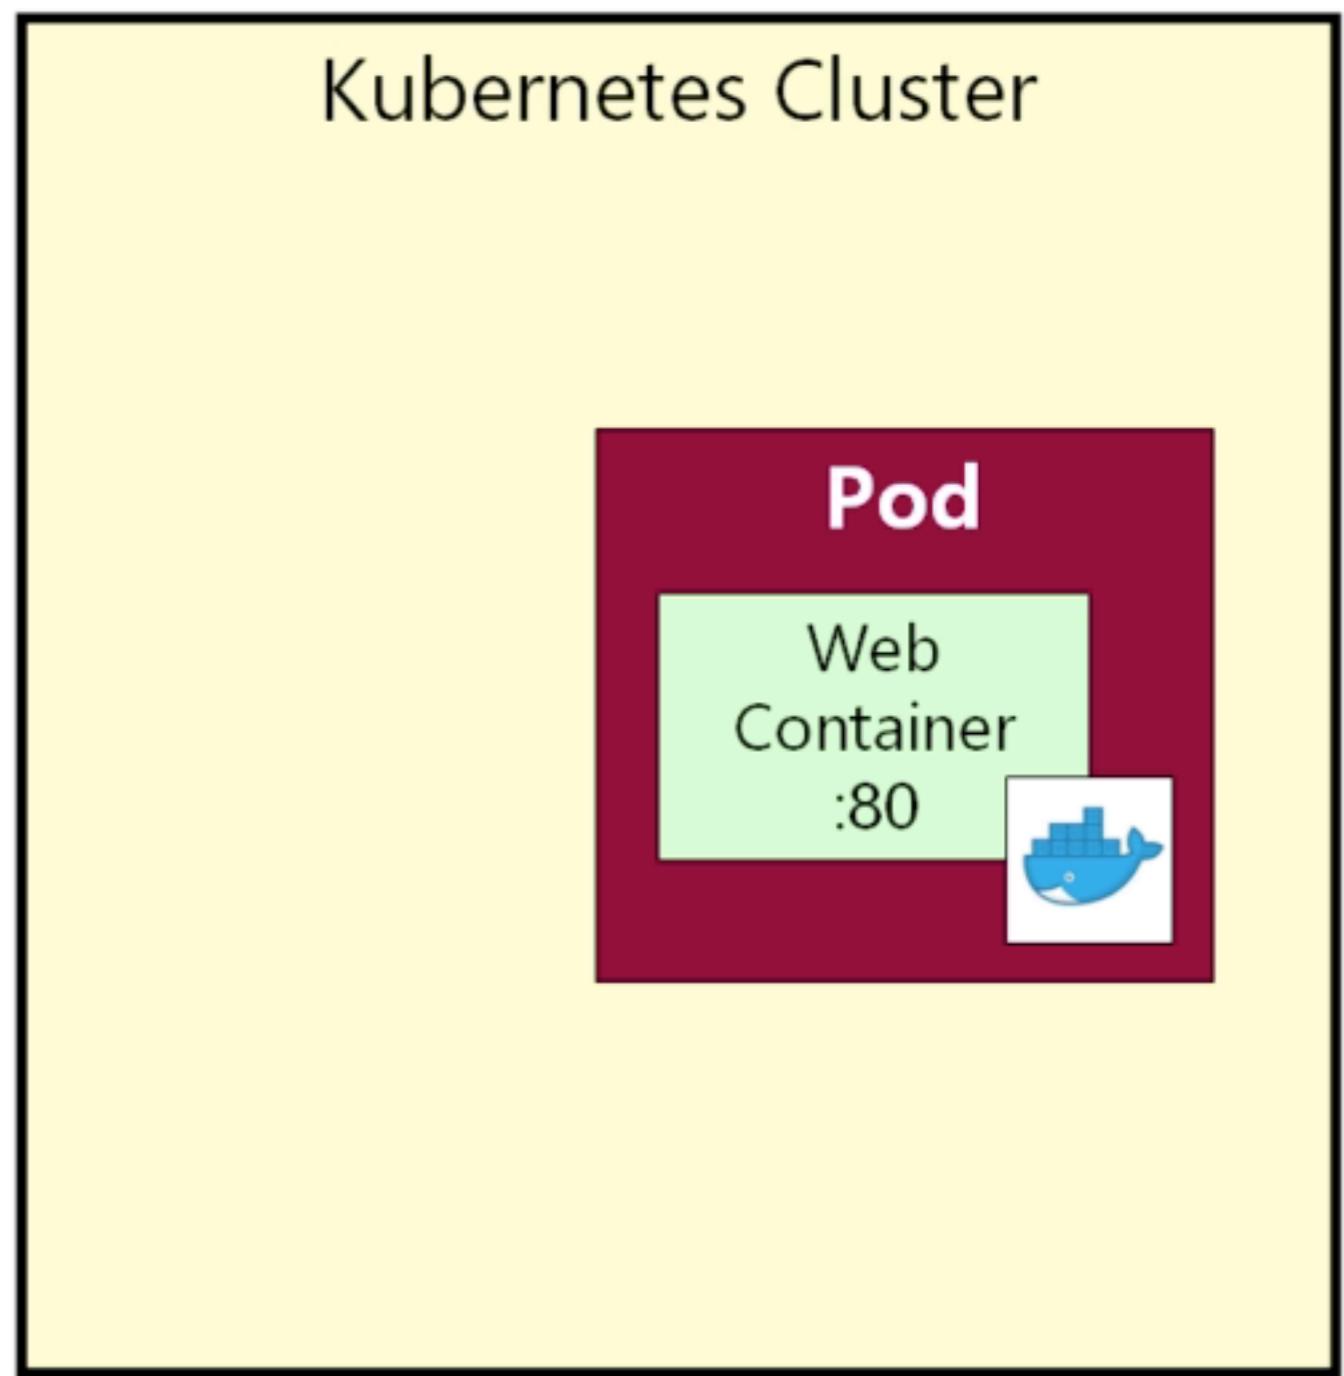
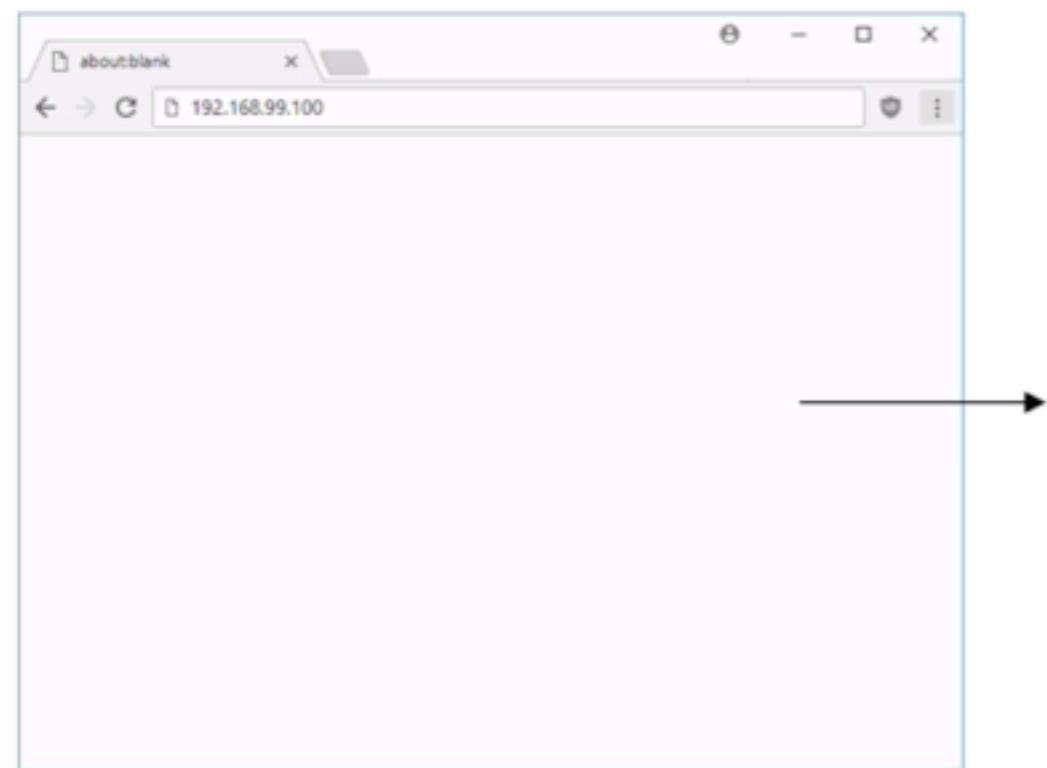
exit



kubernetes

Services

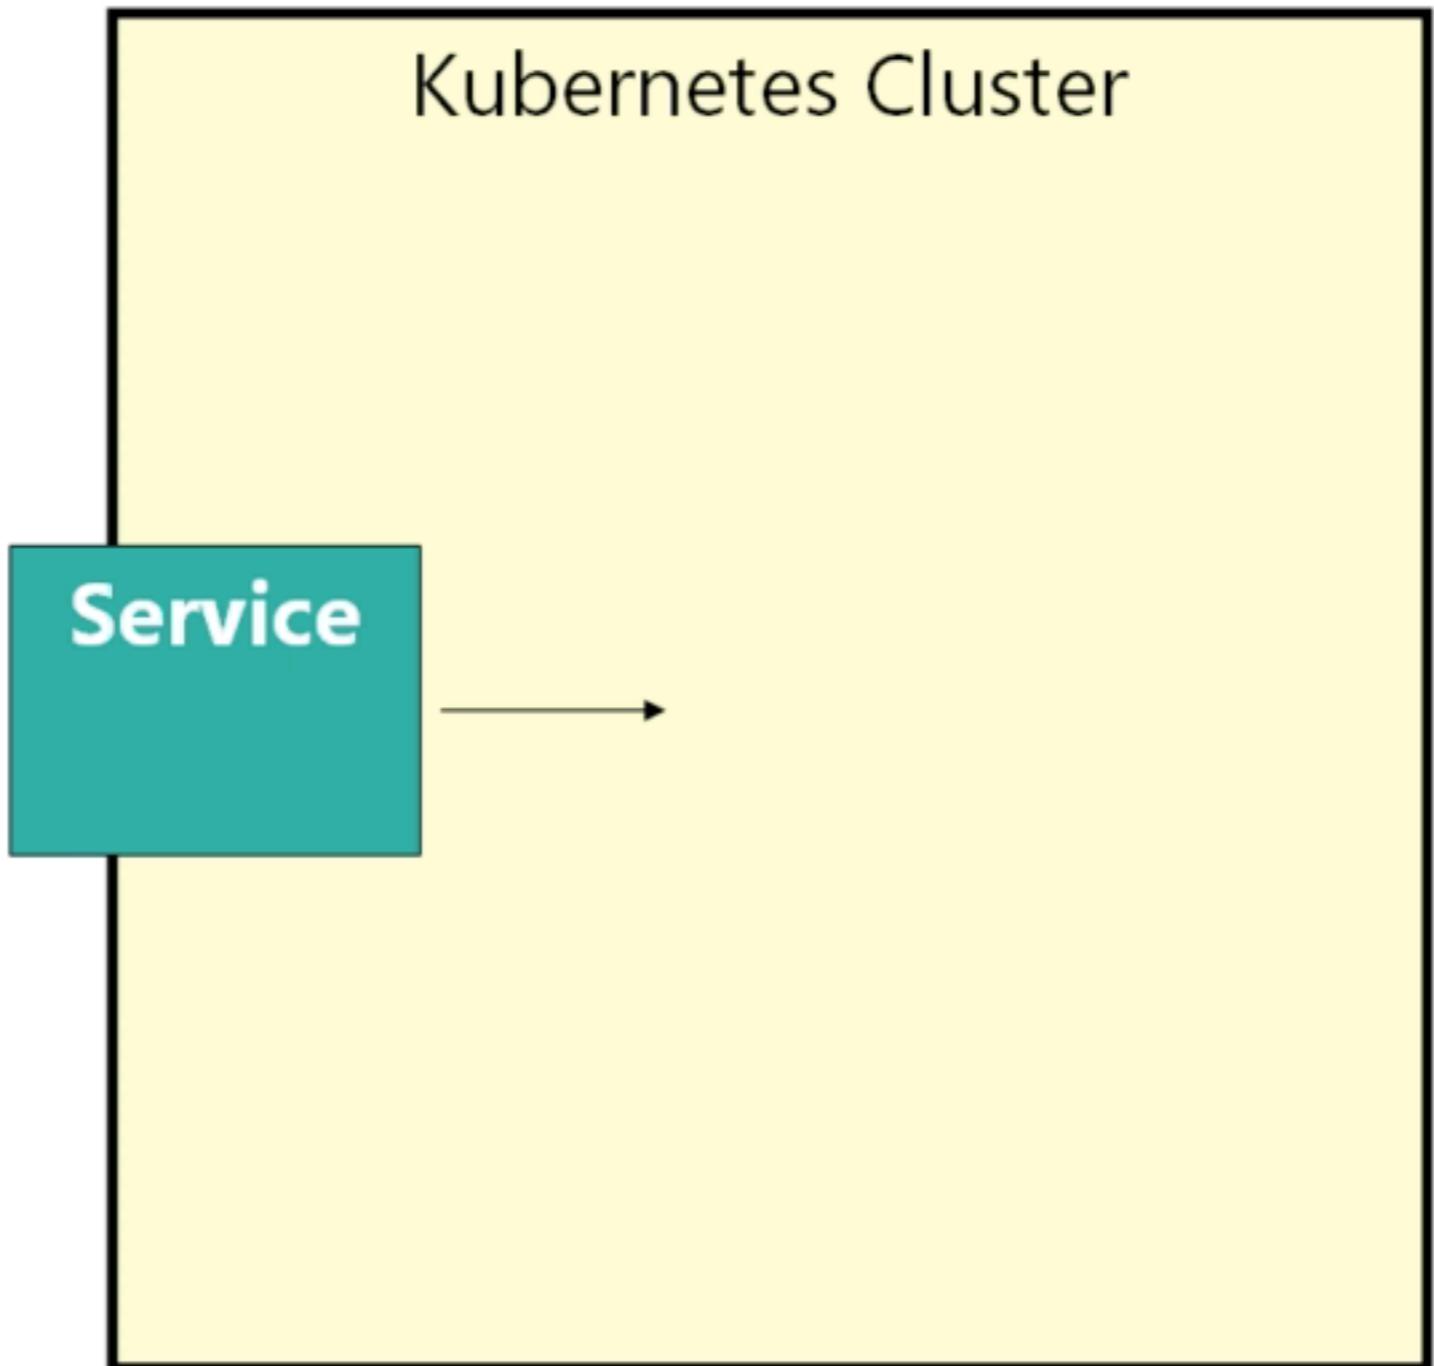
- ▶ We started our first pod
- ▶ It has a web container inside it
- ▶ We also know that the web container is running on port 80
- ▶ We tried to access the web container through a web browser - it failed



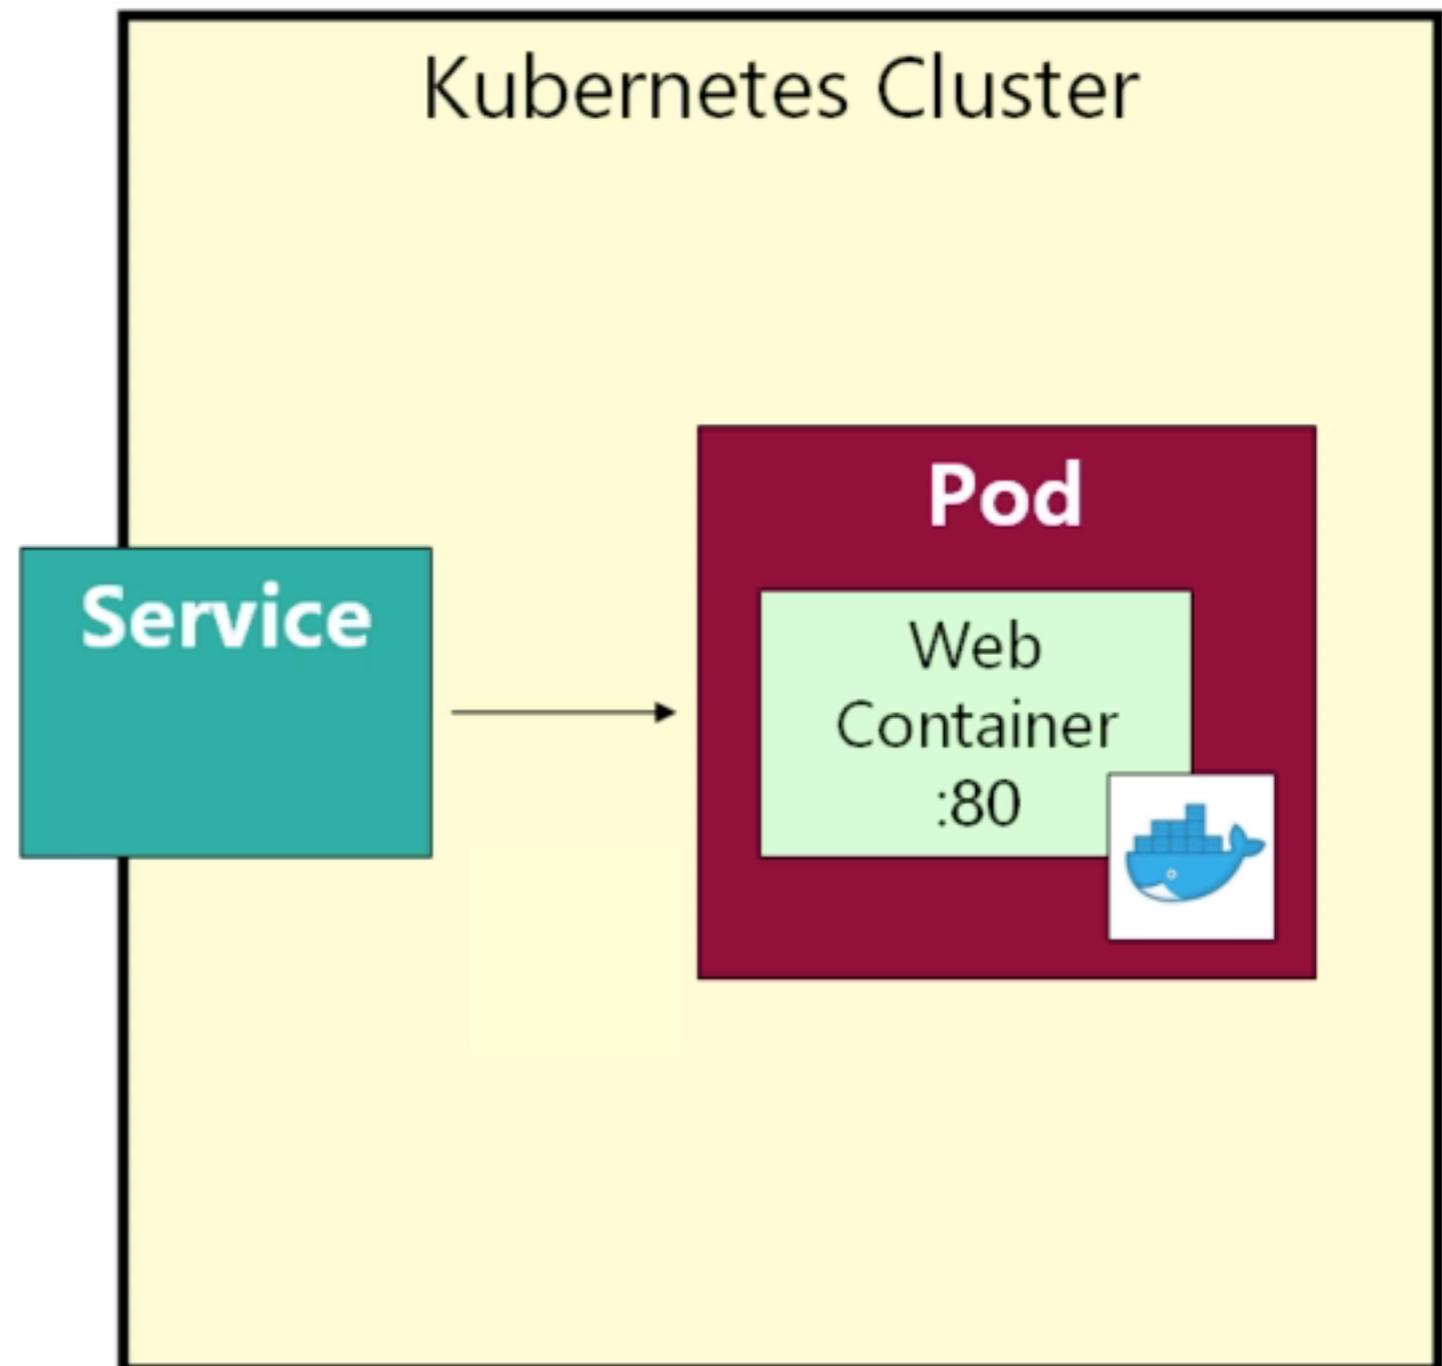
- ▶ Pods are designed to be disposable
- ▶ Pods have short live times
- ▶ Pods regularly die
- ▶ Pods are regularly recreated

We *treat pods like cattle*

*Kubernetes has a new
concept to solve this that is
a - Service*

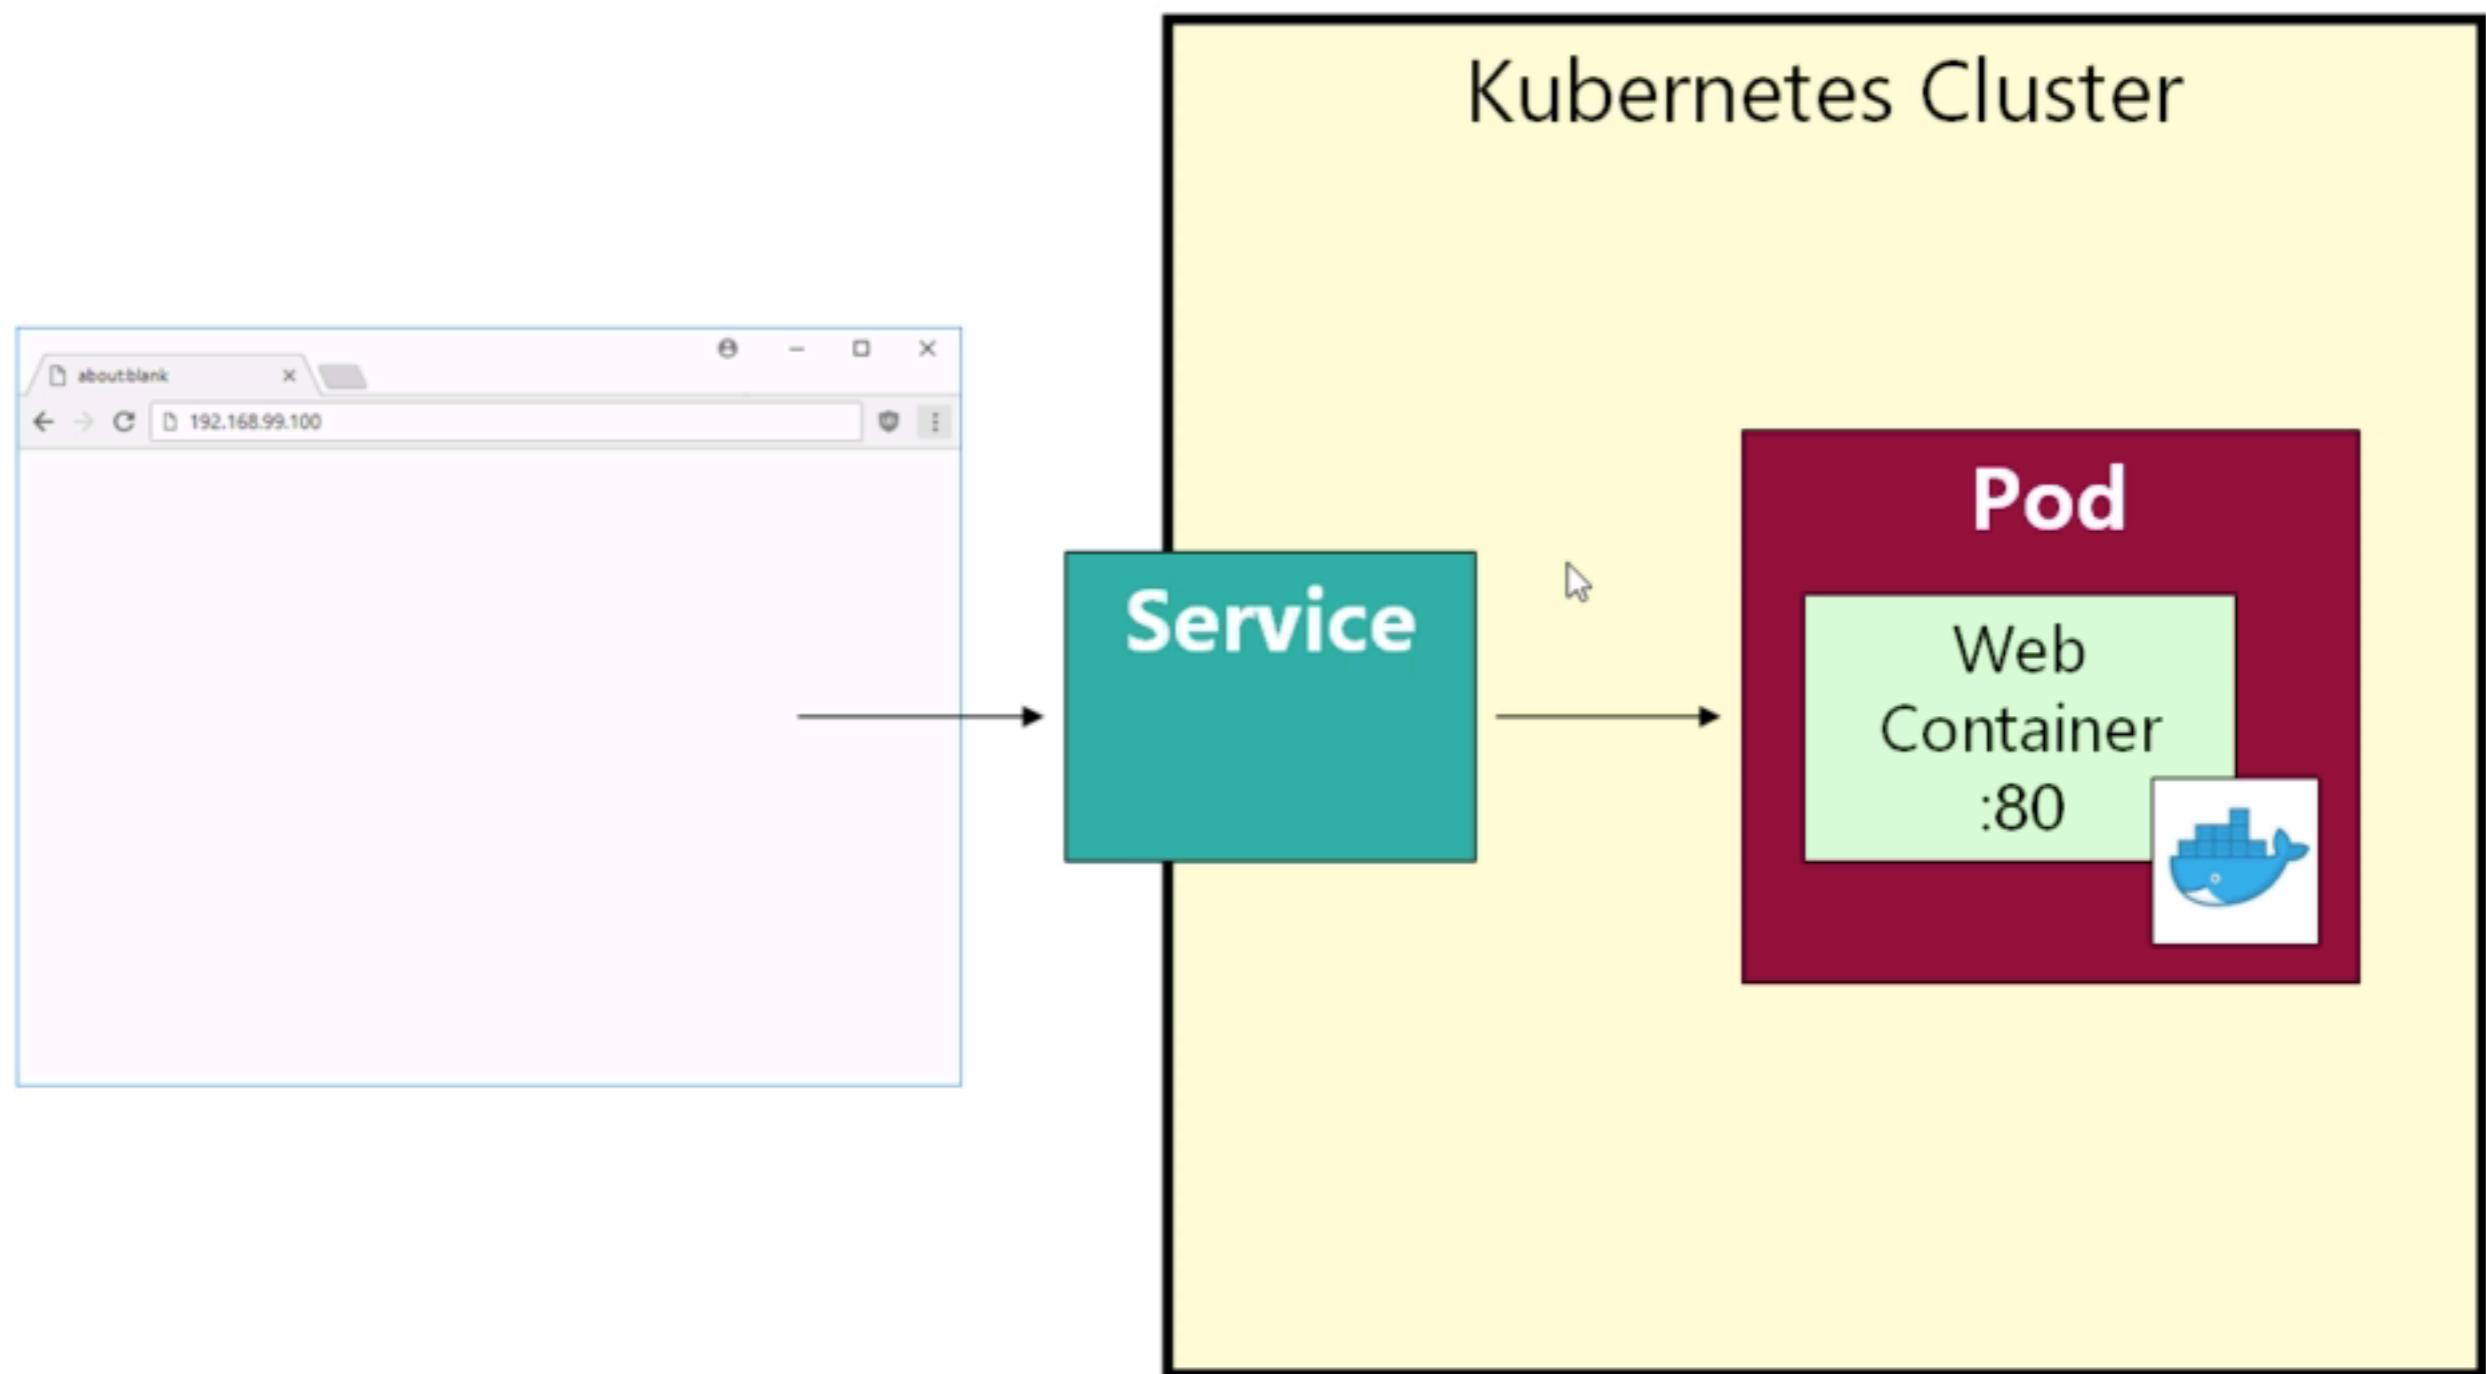


- ▶ A service is a long running object in Kubernetes
- ▶ A service will have an IP address
- ▶ And a service will have a stable fixed port

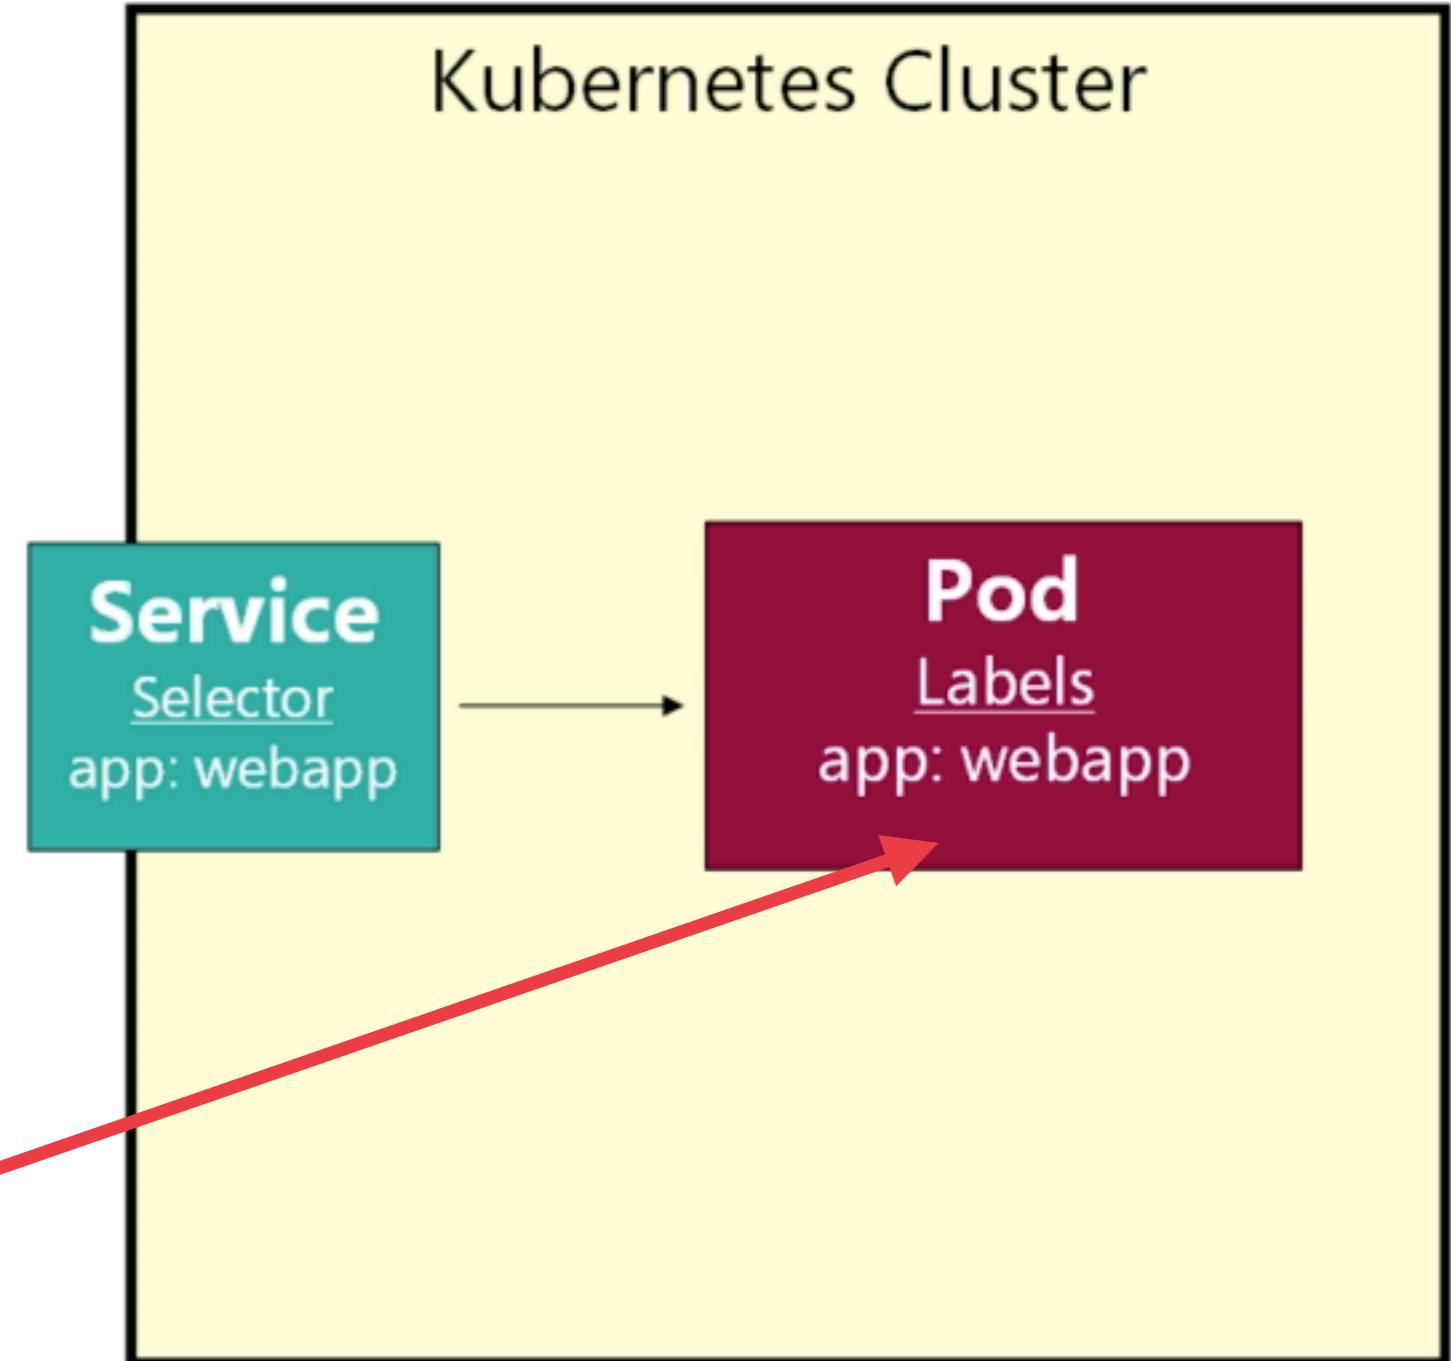


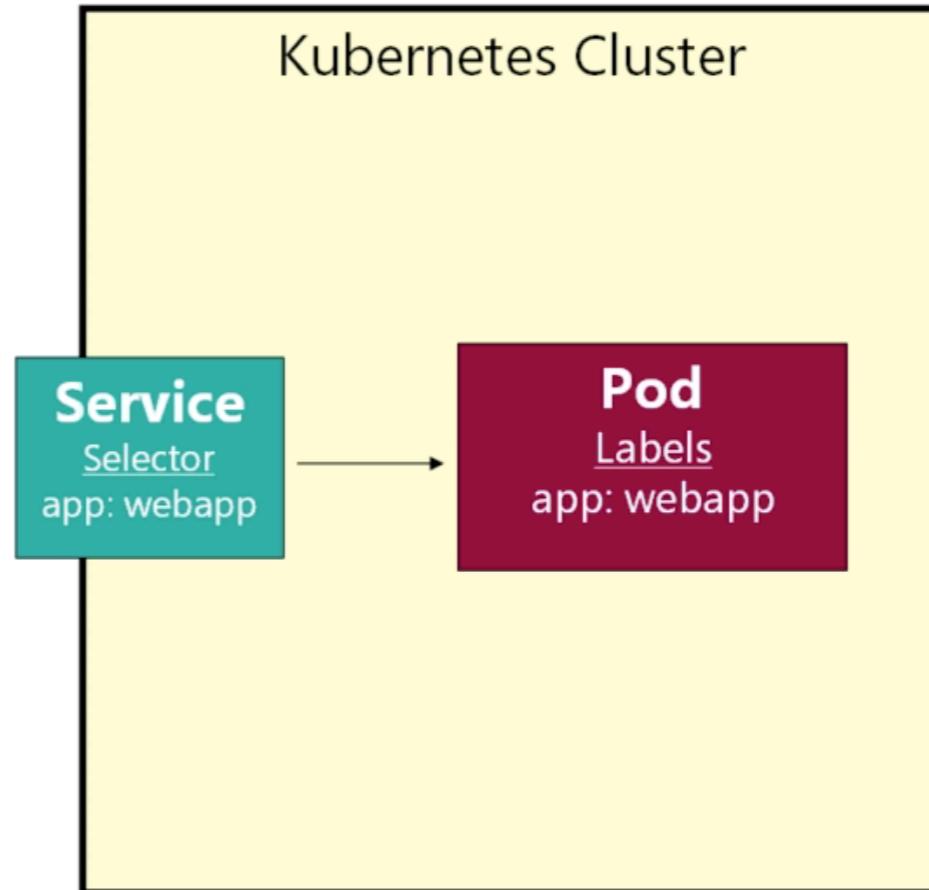
- We can attach services to pods

- ▶ With a service we can connect to the kubernetes cluster
- ▶ The service will find a suitable port to service that request



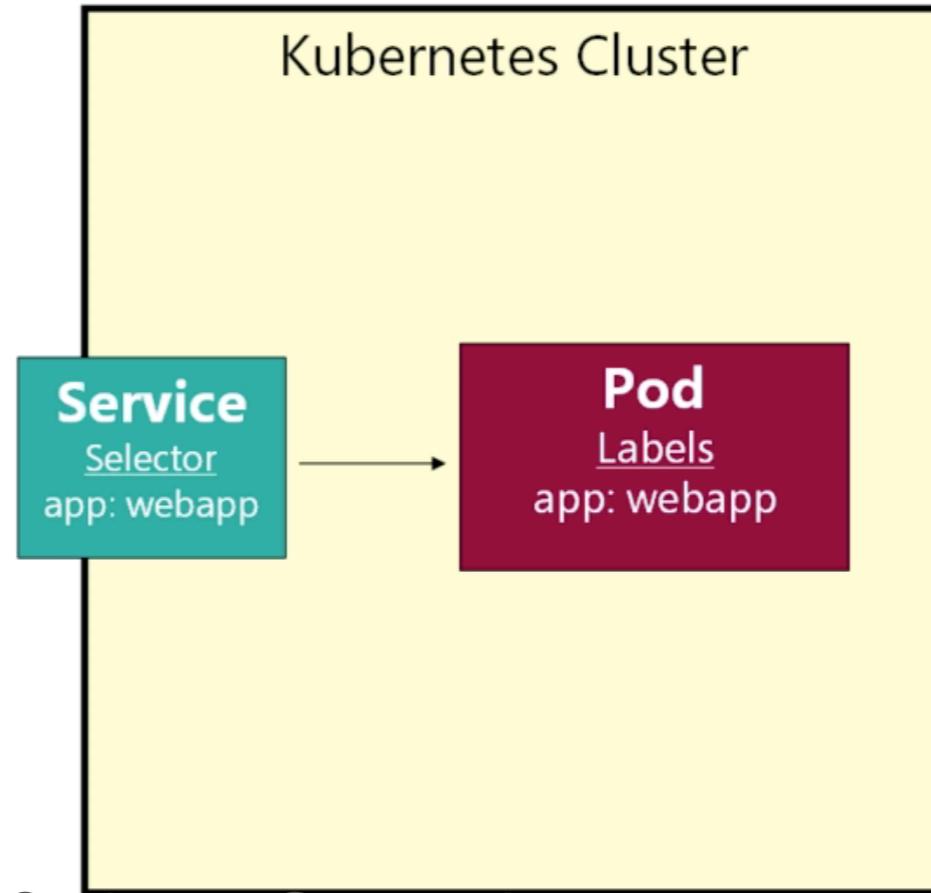
- ▶ The next important concept - Pod label
- ▶ We can setup a series of key-value pairs



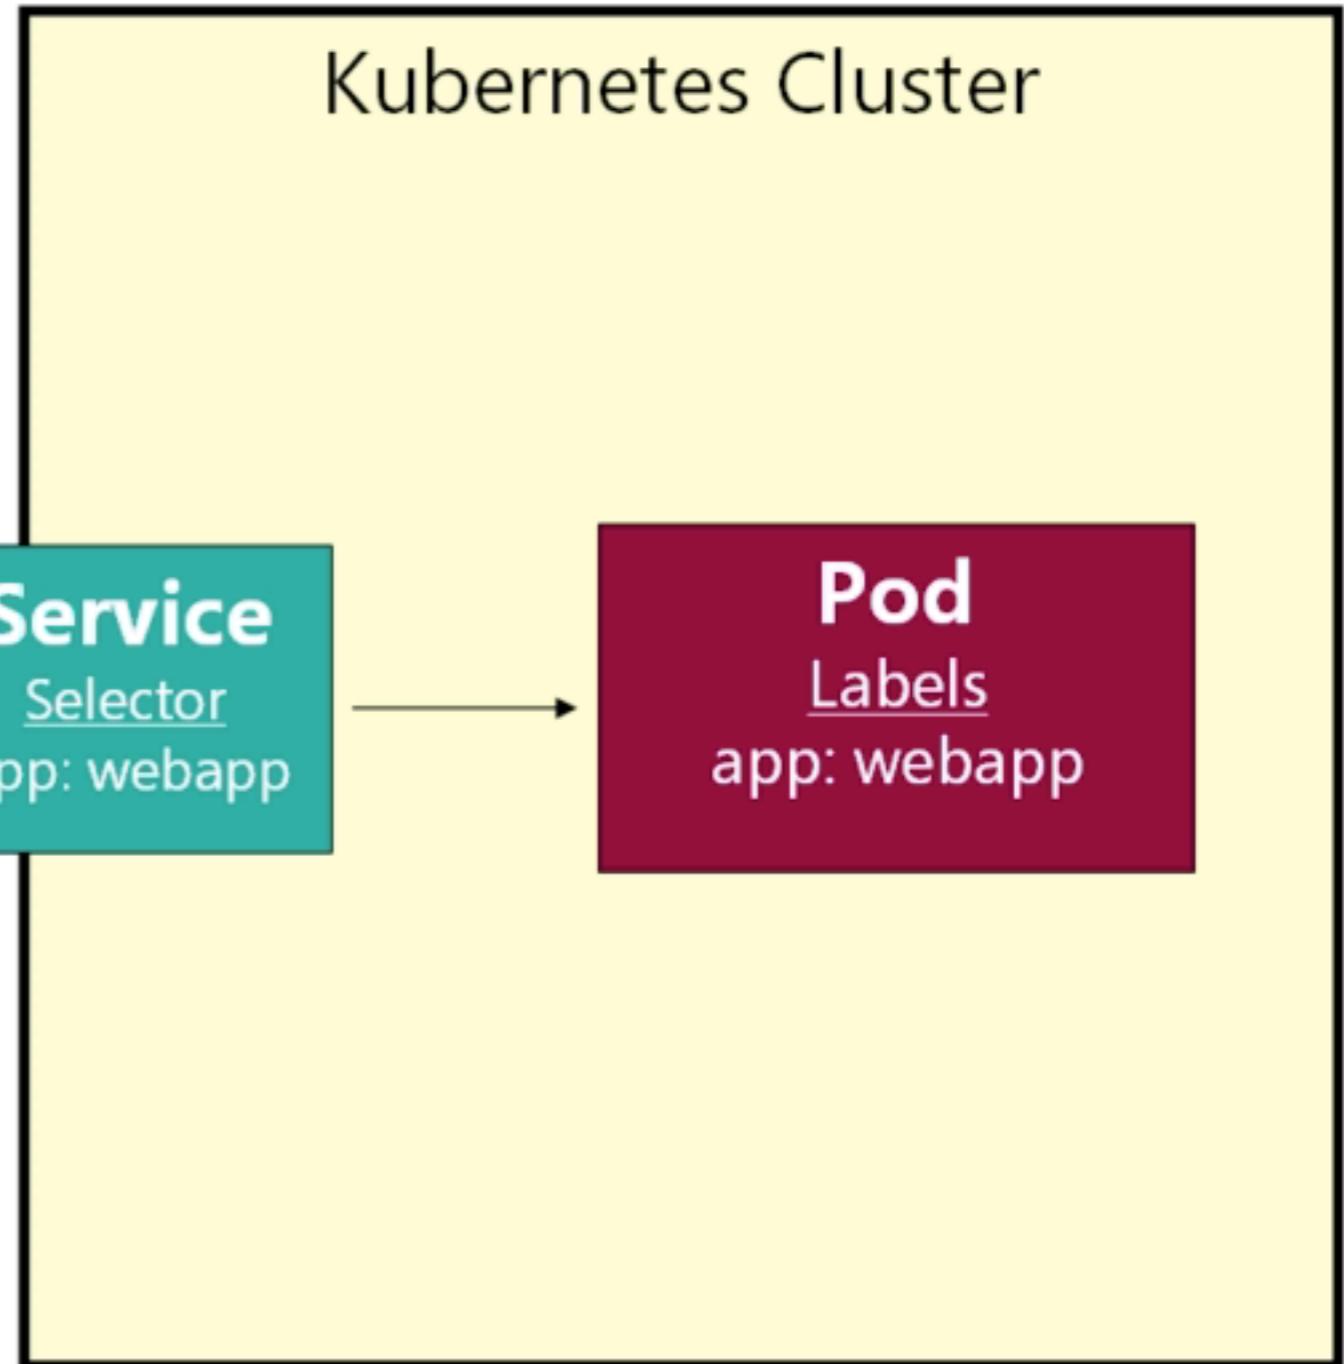
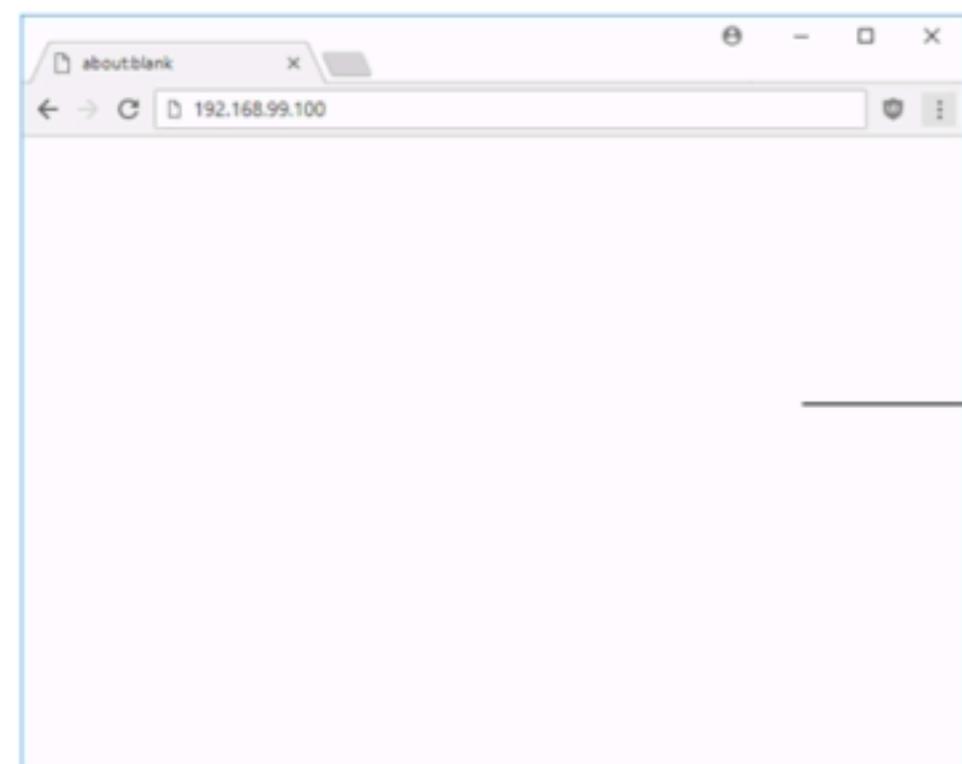


- ▶ In this example we have:
 - ▶ Key - app
 - ▶ Value - web app

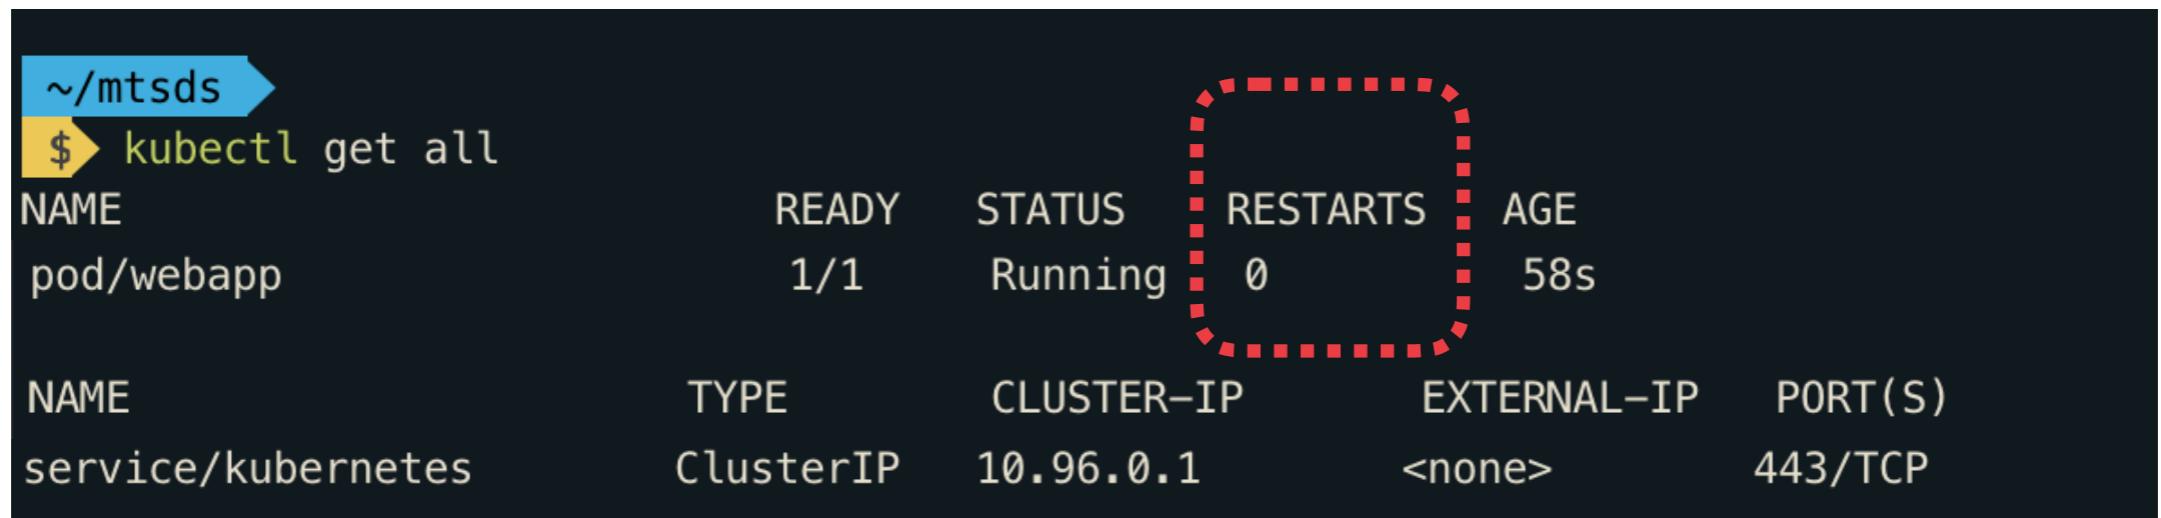
*This is the way we connect
Services to Pods*



- ▶ In the definition of the Service we provide a selector
 - ▶ Key - app
 - ▶ Value - web app



▶ kubectl get all



```
~/mtsds ➔ $ kubectl get all
NAME                 READY   STATUS    RESTARTS   AGE
pod/webapp           1/1     Running   0          58s
service/kubernetes   ClusterIP 10.96.0.1 <none>    443/TCP
```

NAME	READY	STATUS	RESTARTS	AGE
pod/webapp	1/1	Running	0	58s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP

- ▶ Create a service for that pod (**webapp**)
- ▶ Create a file **webapp-service.yaml**

API Overview - find the description for services

**SERVICE APIs**

Endpoints v1 core
EndpointSlice v1alpha1 discovery.k8s.io
Ingress v1beta1 networking.k8s.io
Service v1 core

CONFIG AND STORAGE APIs

ConfigMap v1 core
CSI Driver v1beta1 storage.k8s.io
CSI Node v1beta1 storage.k8s.io

Discovery & LB resources

Config & Storage resources

Cluster resources objects

Metadata resources are ol

Resource O

```
kind: Service
apiVersion: v1
metadata:
  # Unique key of the Service instance
  name: service-example
spec:
  ports:
    # Accept traffic sent to port 80
    - name: http
      port: 80
      targetPort: 80
  selector:
    # Loadbalance traffic across Pods matching
    # this label selector
    app: nginx
  # Create an HA proxy in the cloud provider
  # with an External IP address - *Only supported
  # by some cloud providers*
  type: LoadBalancer
```

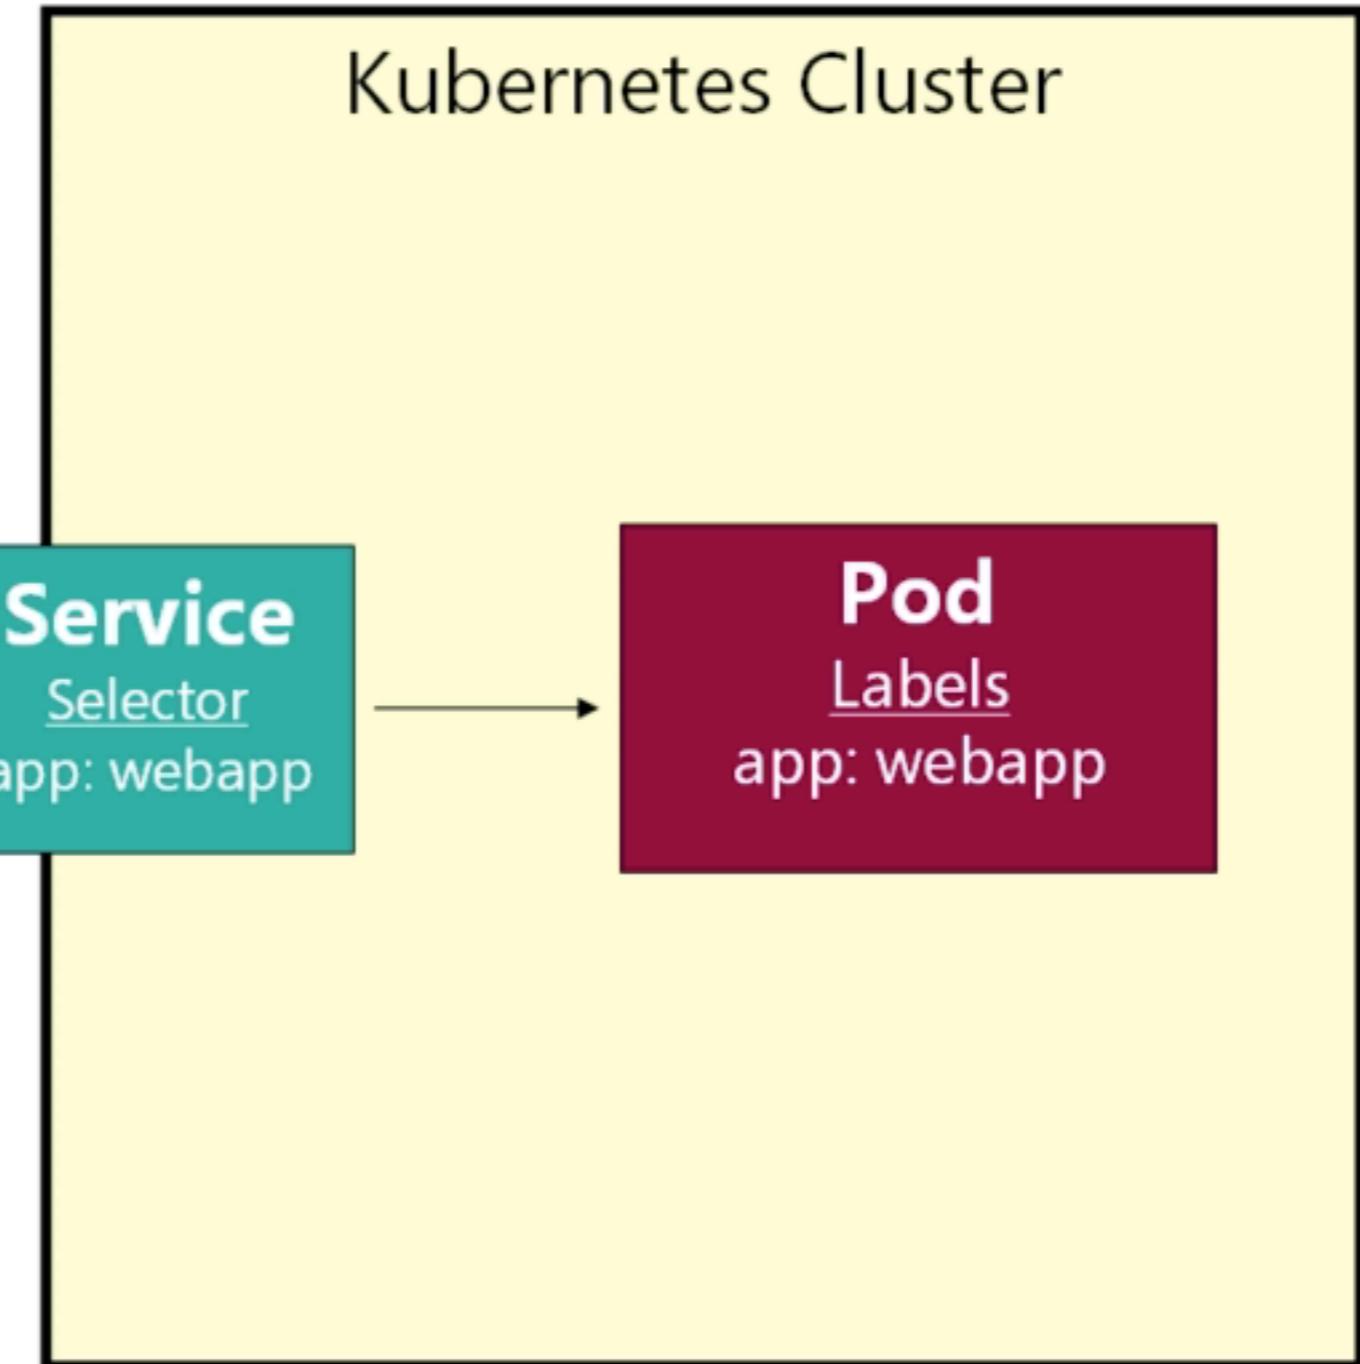
**give the service a name -
think in a good name**

later

**selector: key - app;
value - nginx**

later

- ▶ We are going to give to the pod a label app:webapp
- ▶ The service will select that key-value pair
 - app:webapp



```
! webapp-service.yaml
1   apiVersion: v1
2   kind: Service
3   metadata:
4     name: fleetman-webapp
5
6   spec:
7     selector:
8       app: webapp
9
10  ports:
11    - name: http
12      port: 80
13
14  type: ← ???
15
```

- ▶ Type options for Services:
 - ▶ **LoadBalancer** - we will learn that later and is not available on a local cluster
 - ▶ **ClusterIP** - This service will only be accessible from inside the cluster
 - ▶ We will only be using this if it is an internal service (private service)
 - ▶ **NodePort** - expose a port through the node
 - ▶ We can choose what port we want to expose to the outside world (>30000)

```
! webapp-service.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: fleetman-webapp
5
6  spec:
7    selector:
8      app: webapp
9
10   ports:
11     - name: http
12       port: 80
13       nodePort: 30080 ←
14
15   type: NodePort
16
```

- ▶ With this definition we are saying that we are going to have a service called fleetman-webapp
- ▶ It is going to be a network endpoint for any pod that finds in the cluster with the key value pair - app:webapp
- ▶ It is going to be talking to the port inside the pod (80) and it is going to expose to the outside world through the port 30080

- ▶ `kubectl apply -f webapp-service.yaml`

```
~ $ kubectl apply -f webapp-service.yaml
```

- ▶ `kubectl get all`

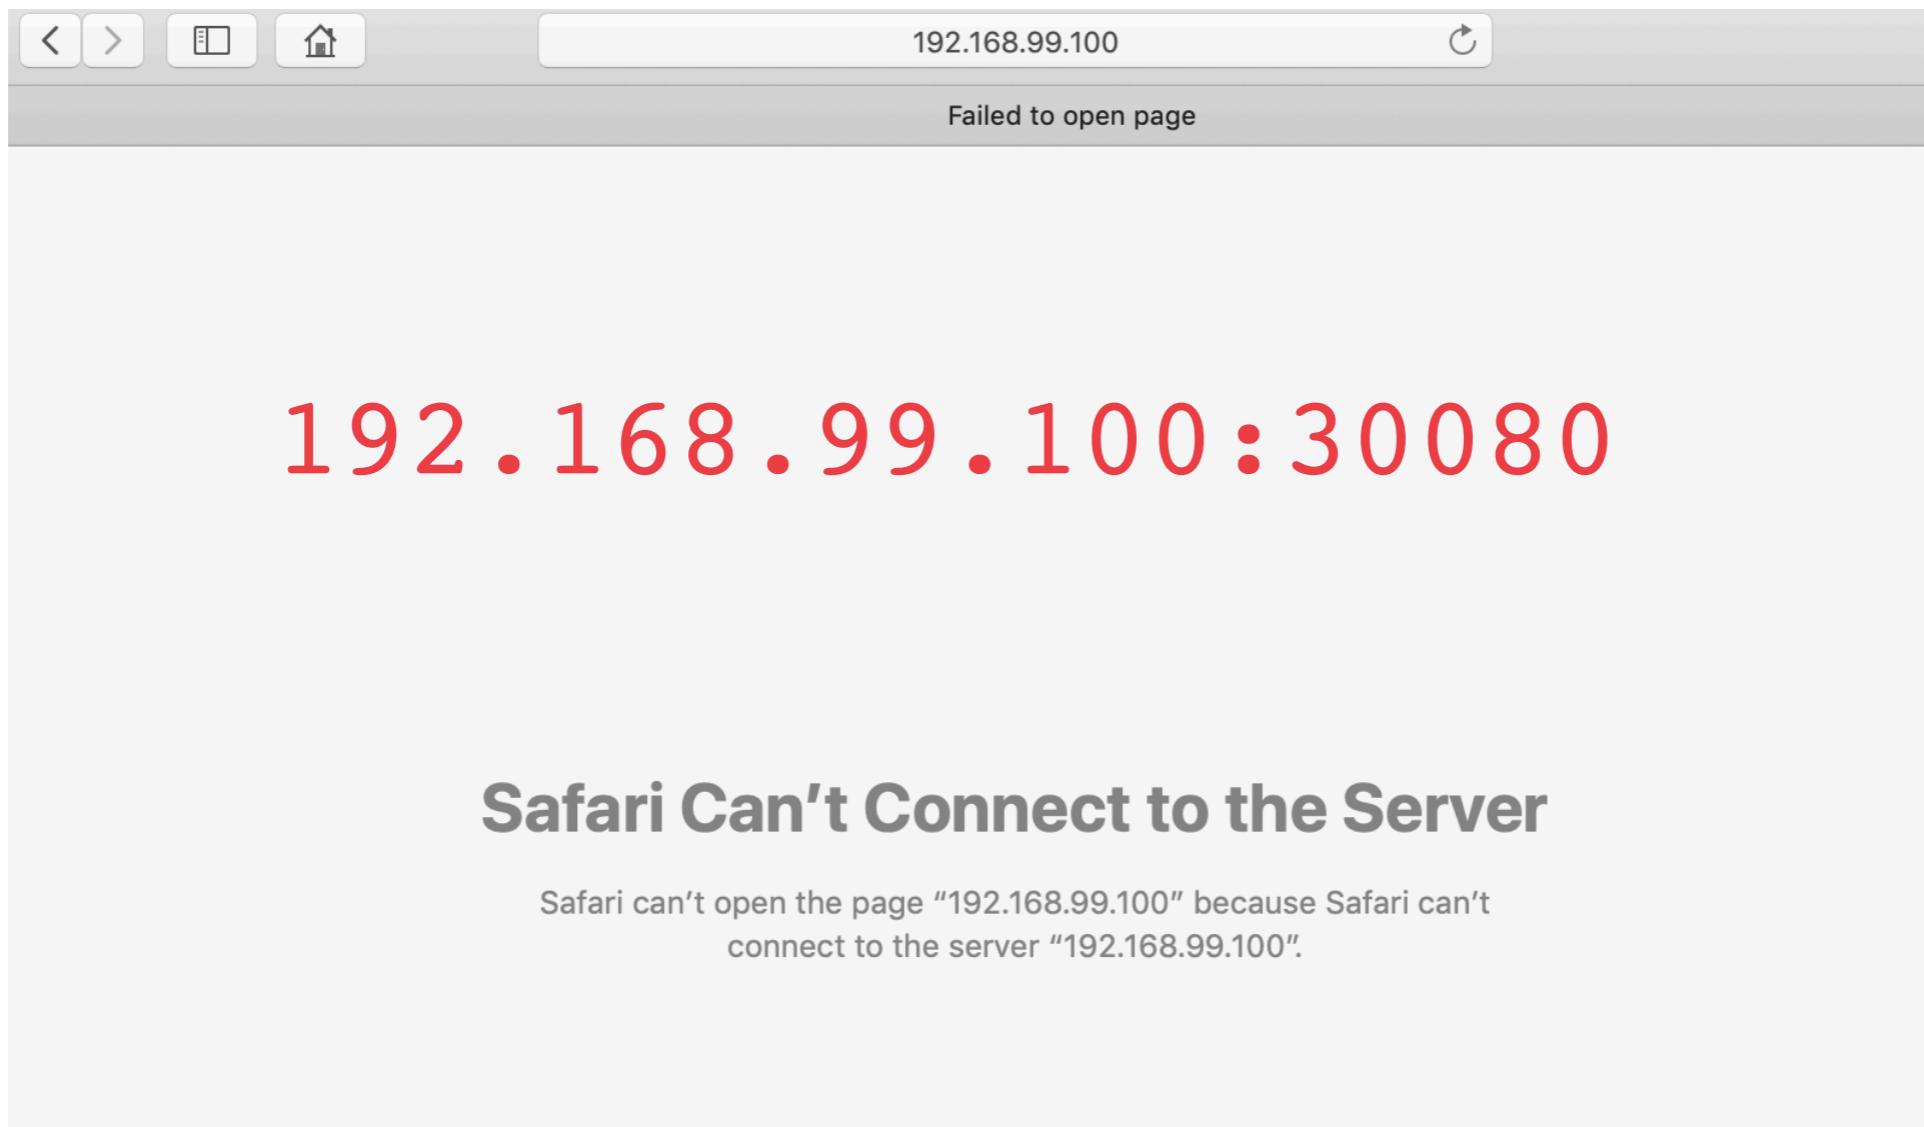
```
~ $ kubectl get all
```

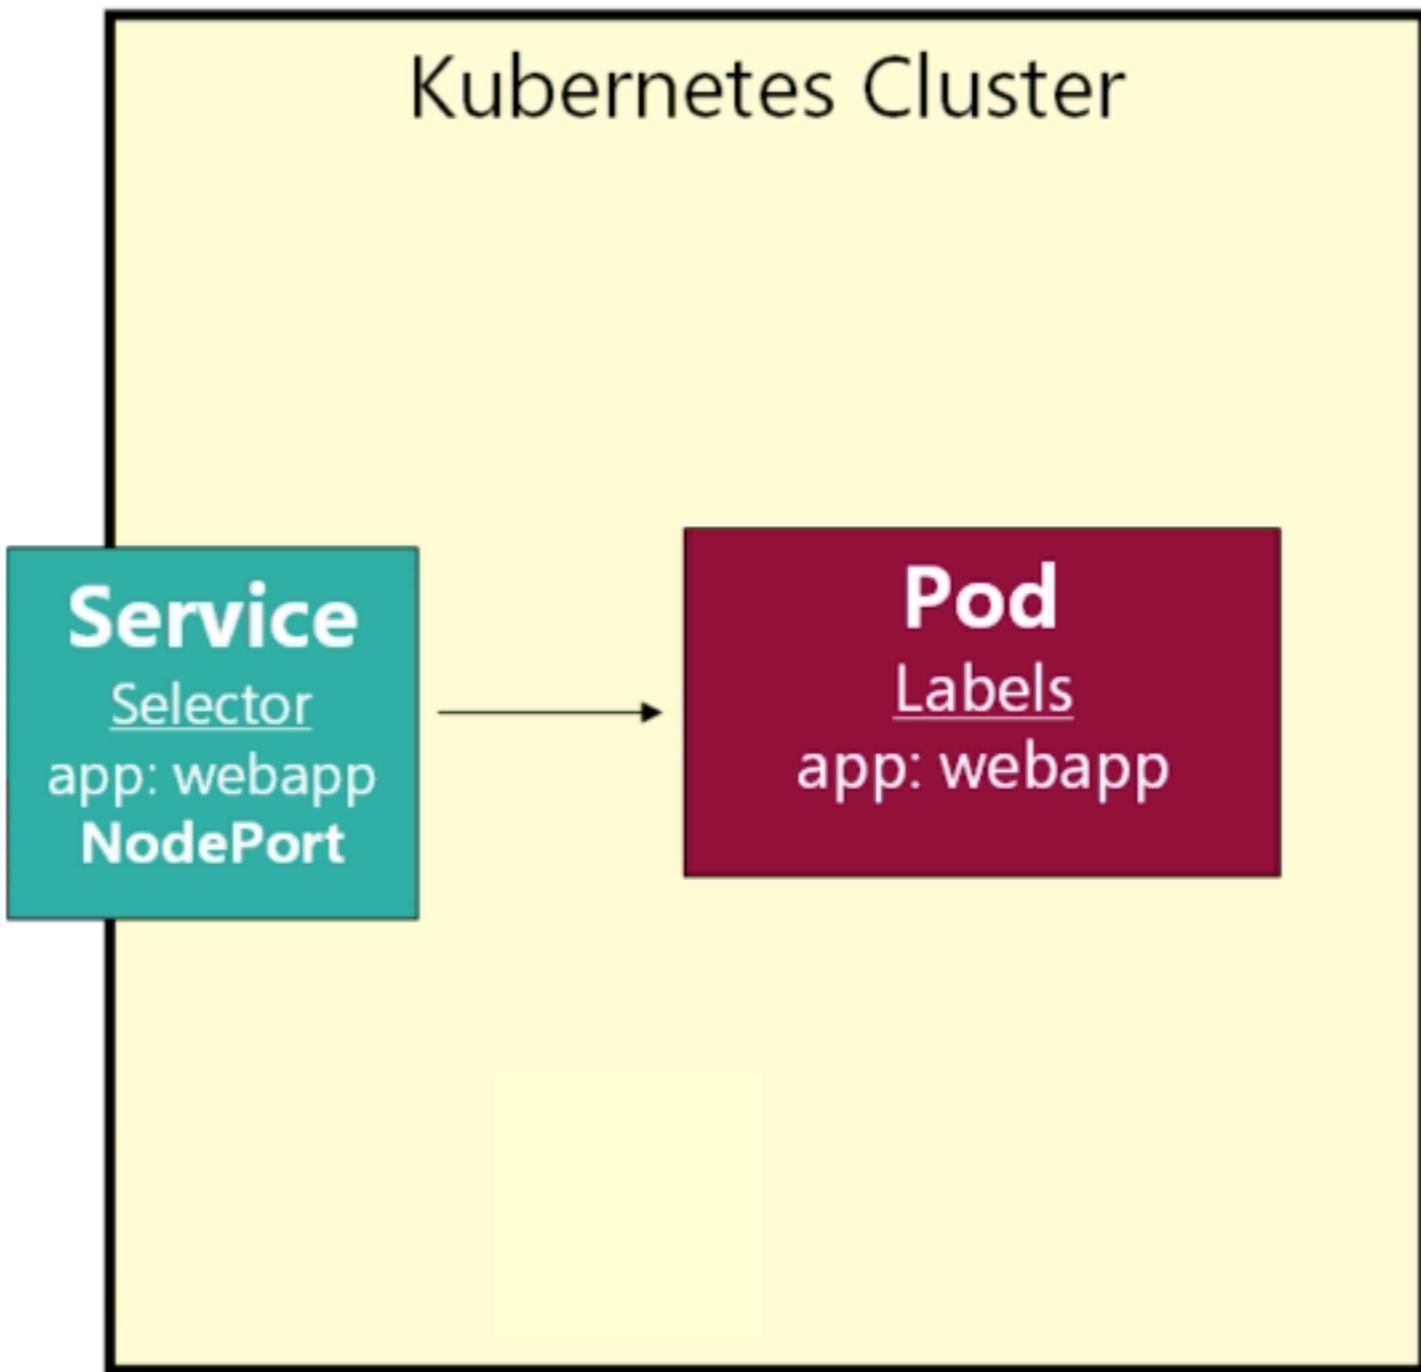
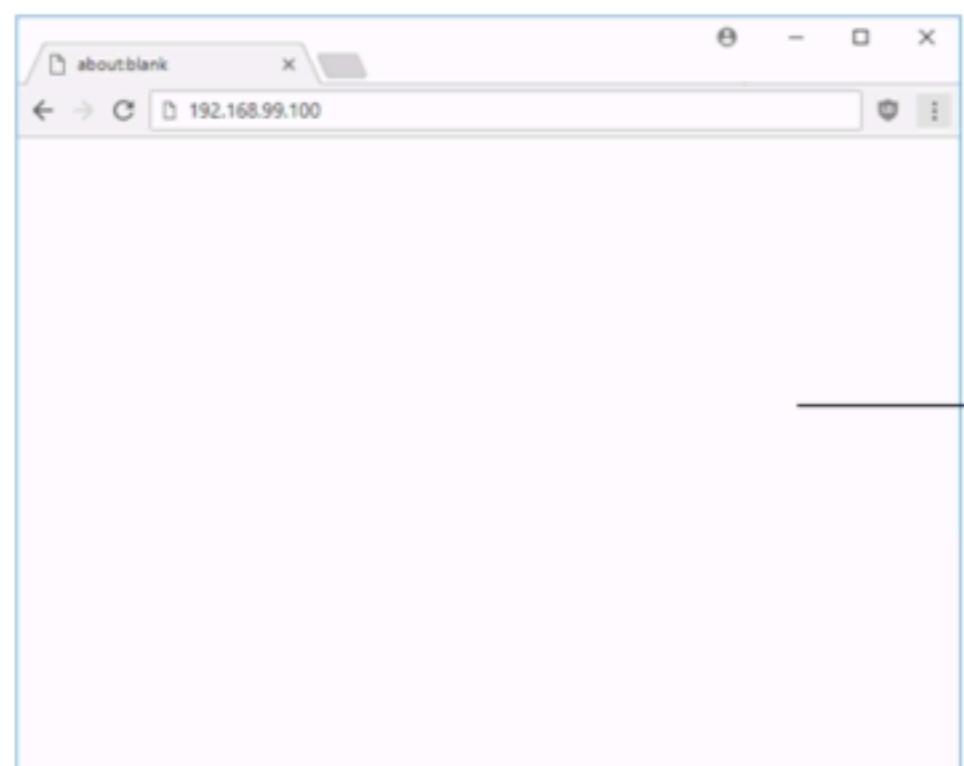
NAME	READY	STATUS	RESTARTS	AGE
pod/webapp	1/1	Running	1	6d5h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/fleetman-webapp	NodePort	10.106.54.9	<none>	80:30080/TCP
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP

► Let's try it - minikube ip

```
~/mtsd$ ➤  
$ ➤ minikube ip  
192.168.99.100
```





- ▶ The service looks for any pod that finds in the cluster with the key value pair - app:webapp

*We didn't define any label
on our Pod*

```
! pod.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: webapp
5    labels:
6      app: webapp
7
8  spec:
9    containers:
10   - name: webapp
11     image: richardchesterwood/k8s-fleetman-webapp-angular:release0
12
```



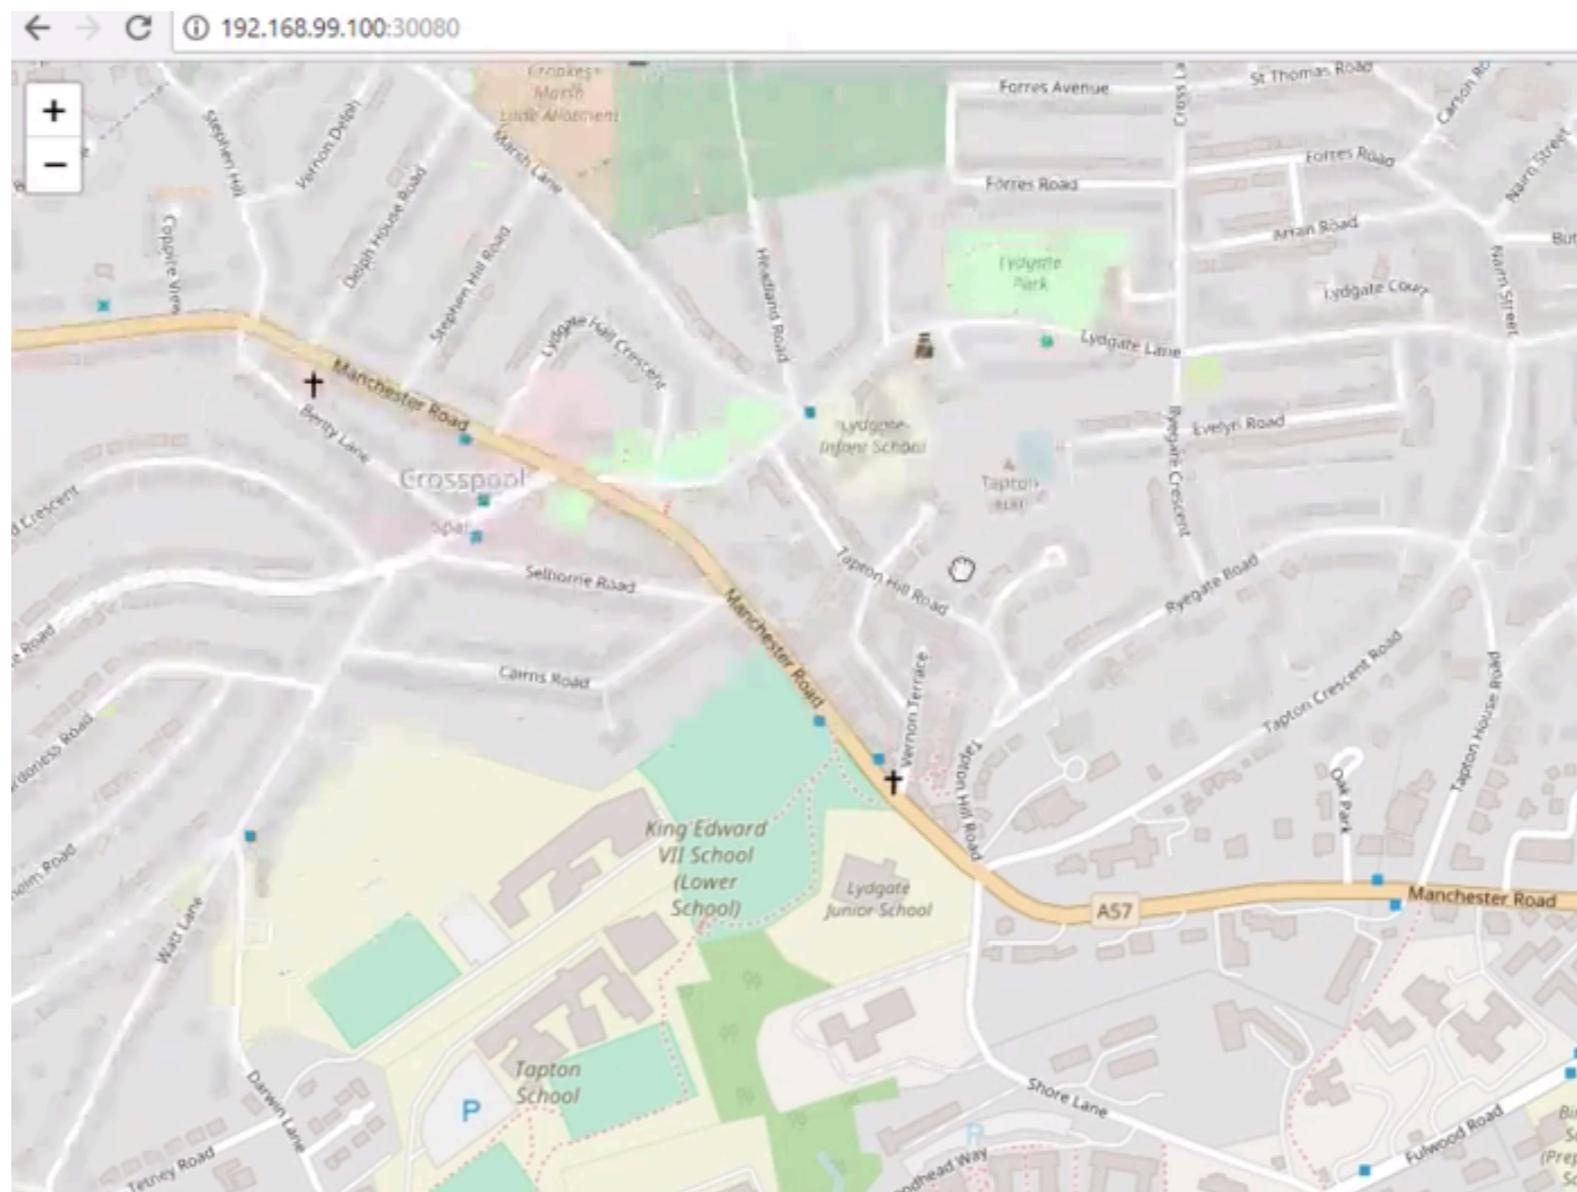
- ▶ `kubectl apply -f pod.yaml`

```
~$ kubectl apply -f pod.yaml
```

- ▶ `kubectl apply -f webapp-service.yaml`

```
~$ kubectl apply -f webapp-service.yaml
```

Let's try it now!



- ▶ Now we have been told by the developer that a new version is available - release0-5

Deploy the new version!

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: webapp
5    labels:
6      app: webapp
7
8  spec:
9    containers:
10   - name: webapp
11     image: richardchesterwood/k8s-fleetman-webapp-angular:release0
12
```



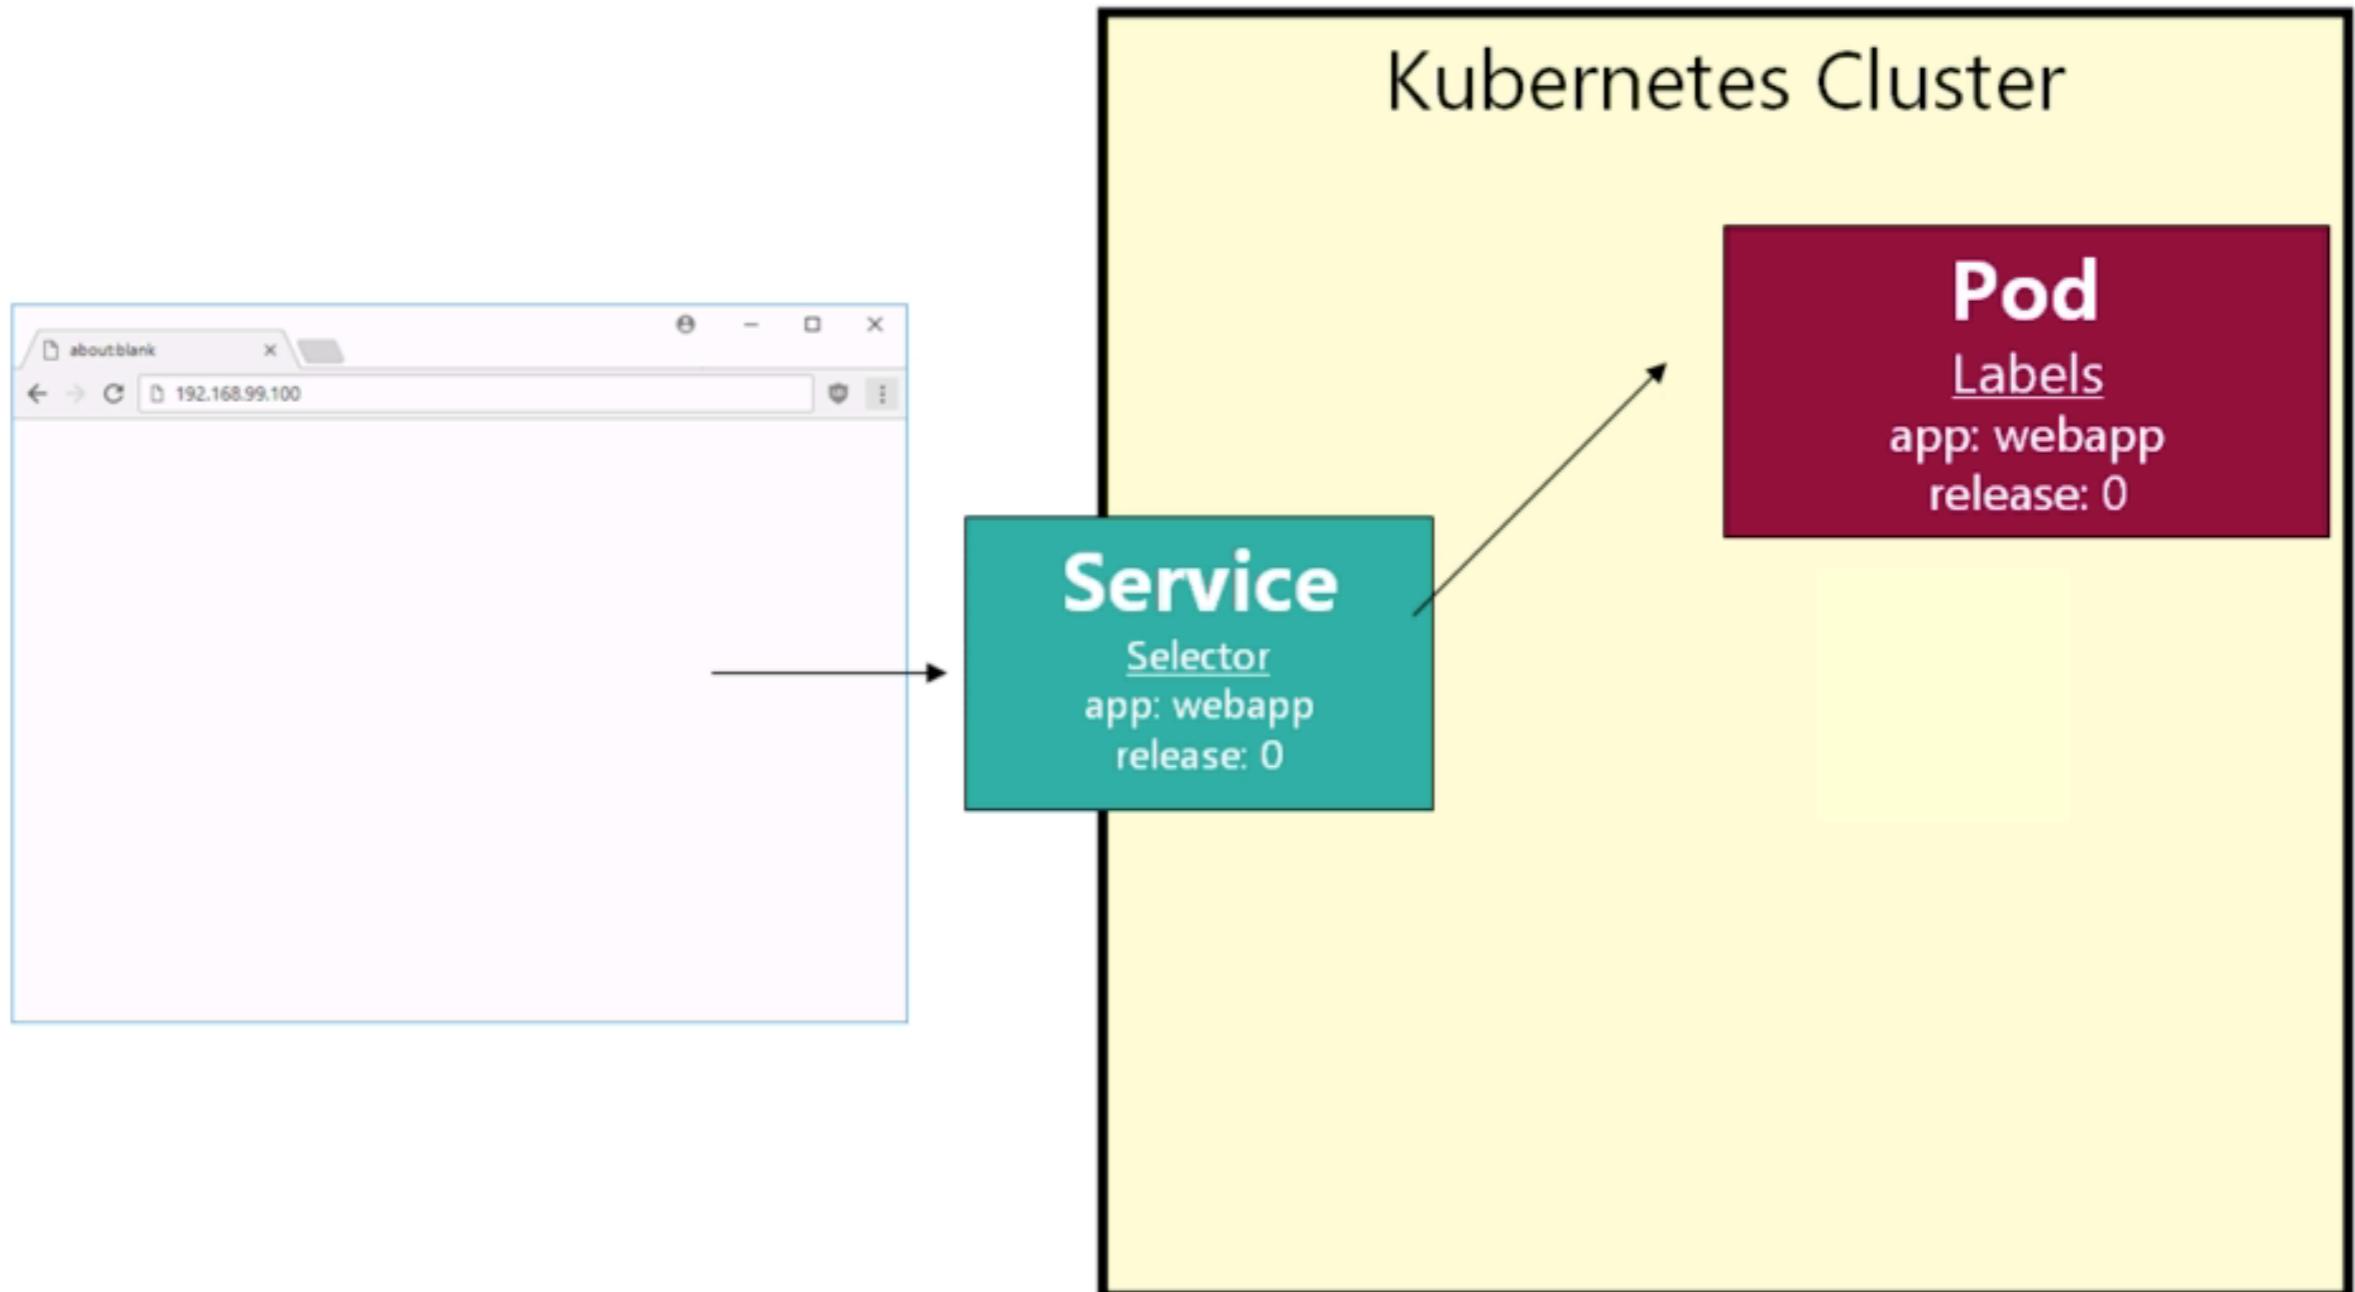
Should we change this?

release0 → release0-5

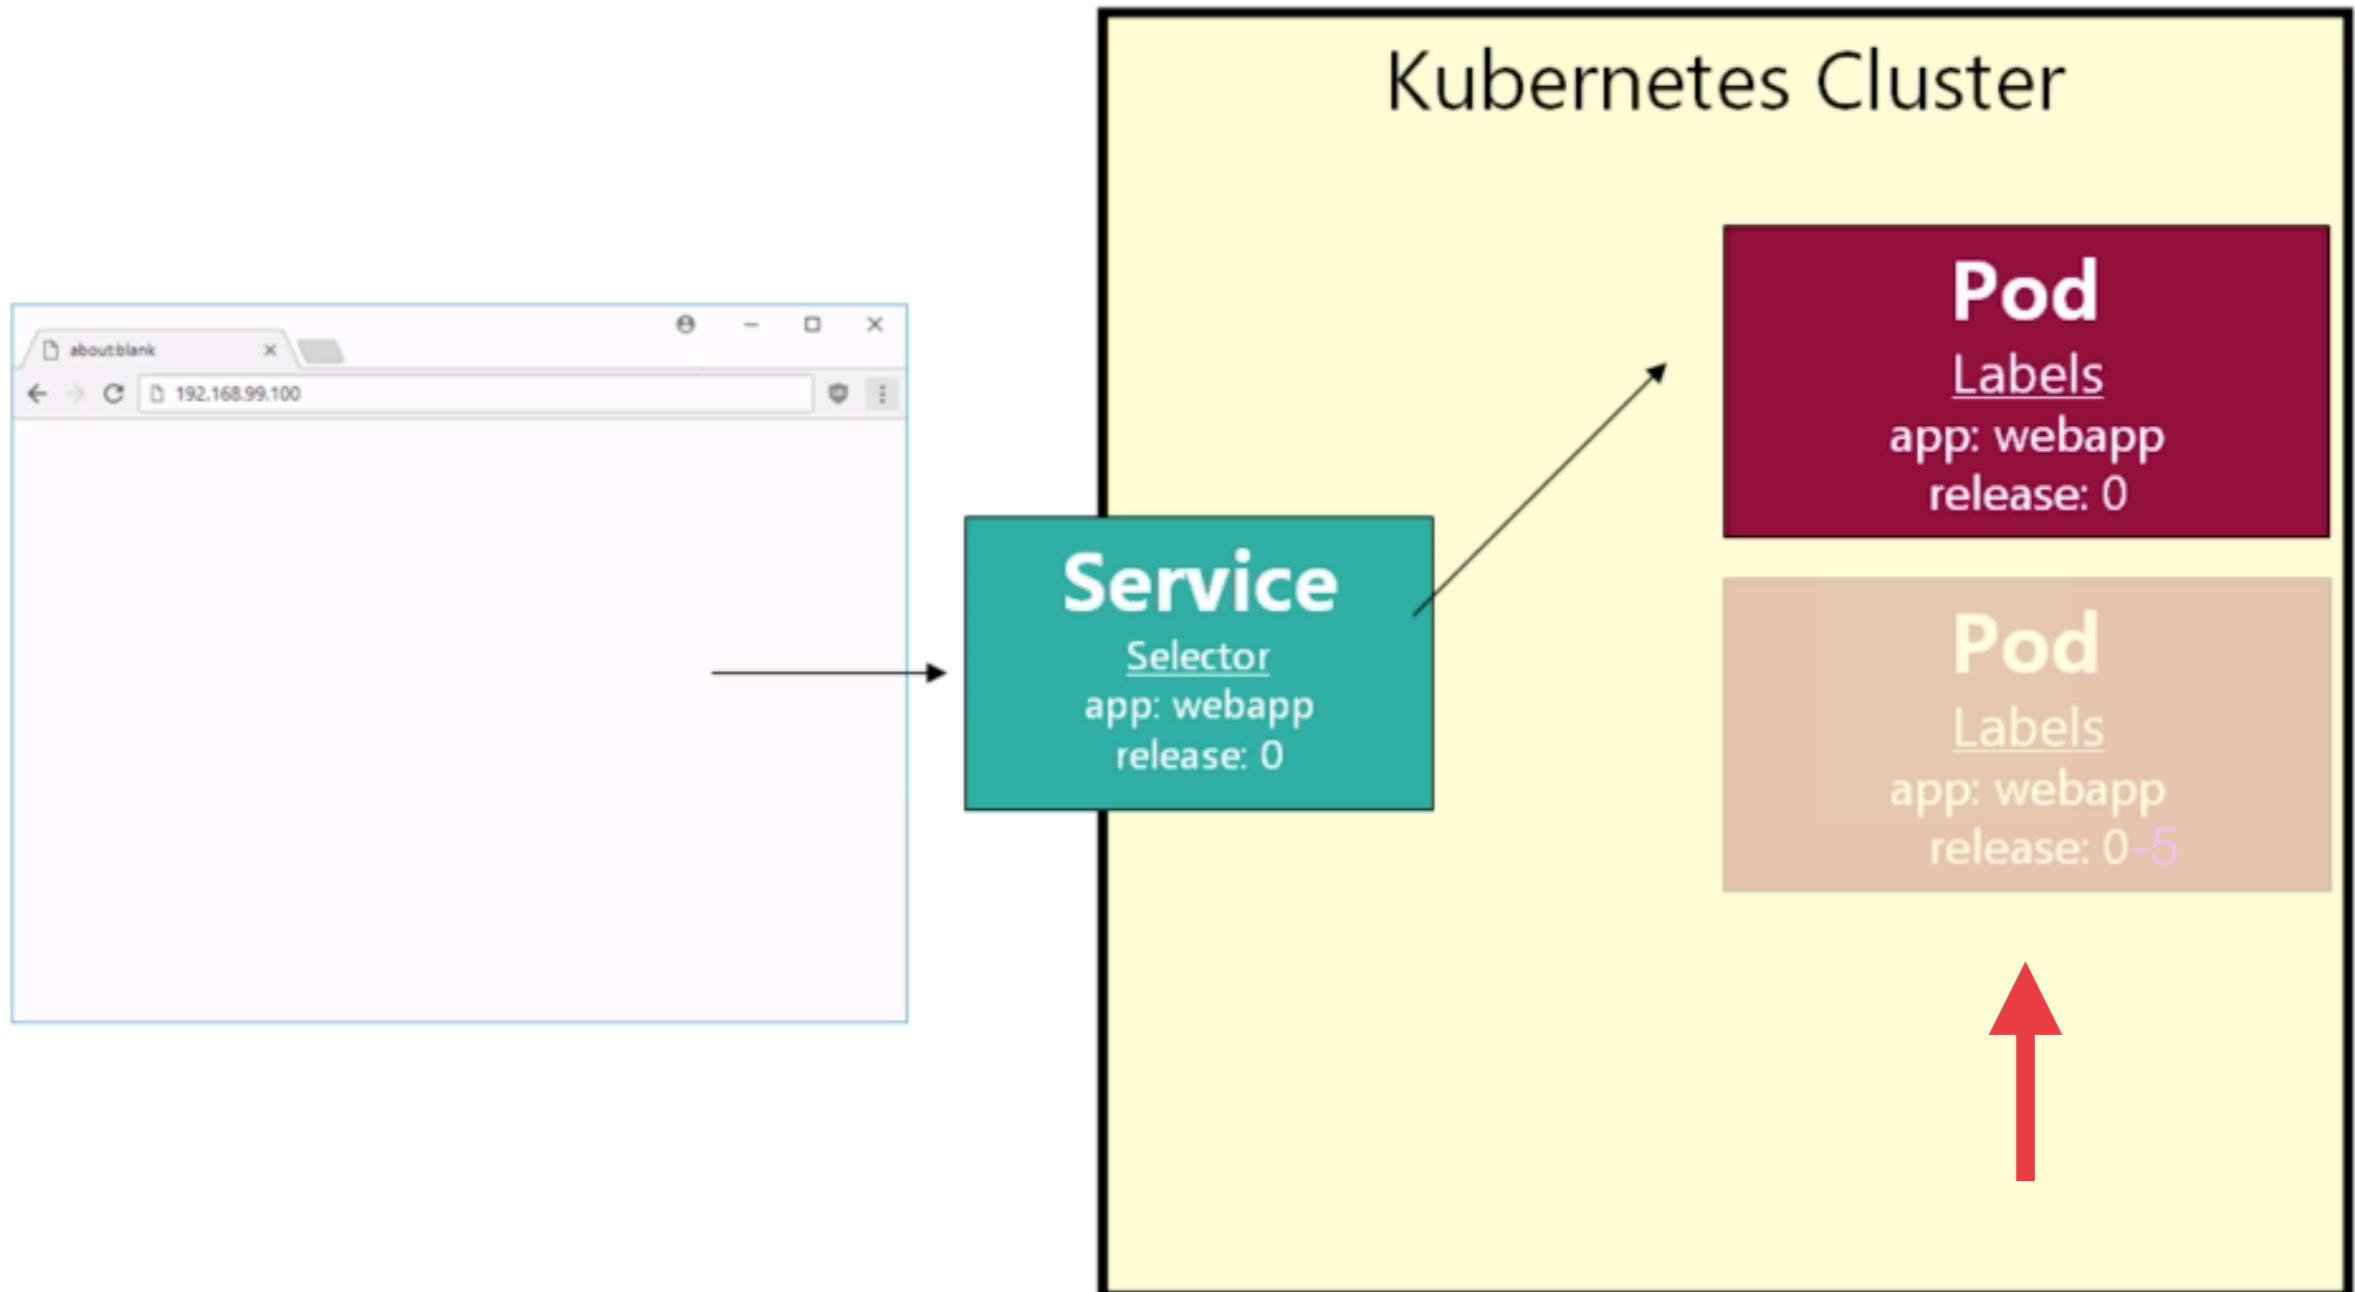
- ▶ That is not a good idea:
 - ▶ The pod is only a wrapper to a container; Kubernetes will have to stop this pod, start a new pod
 - ▶ Pull a new container image
 - ▶ Start that pod

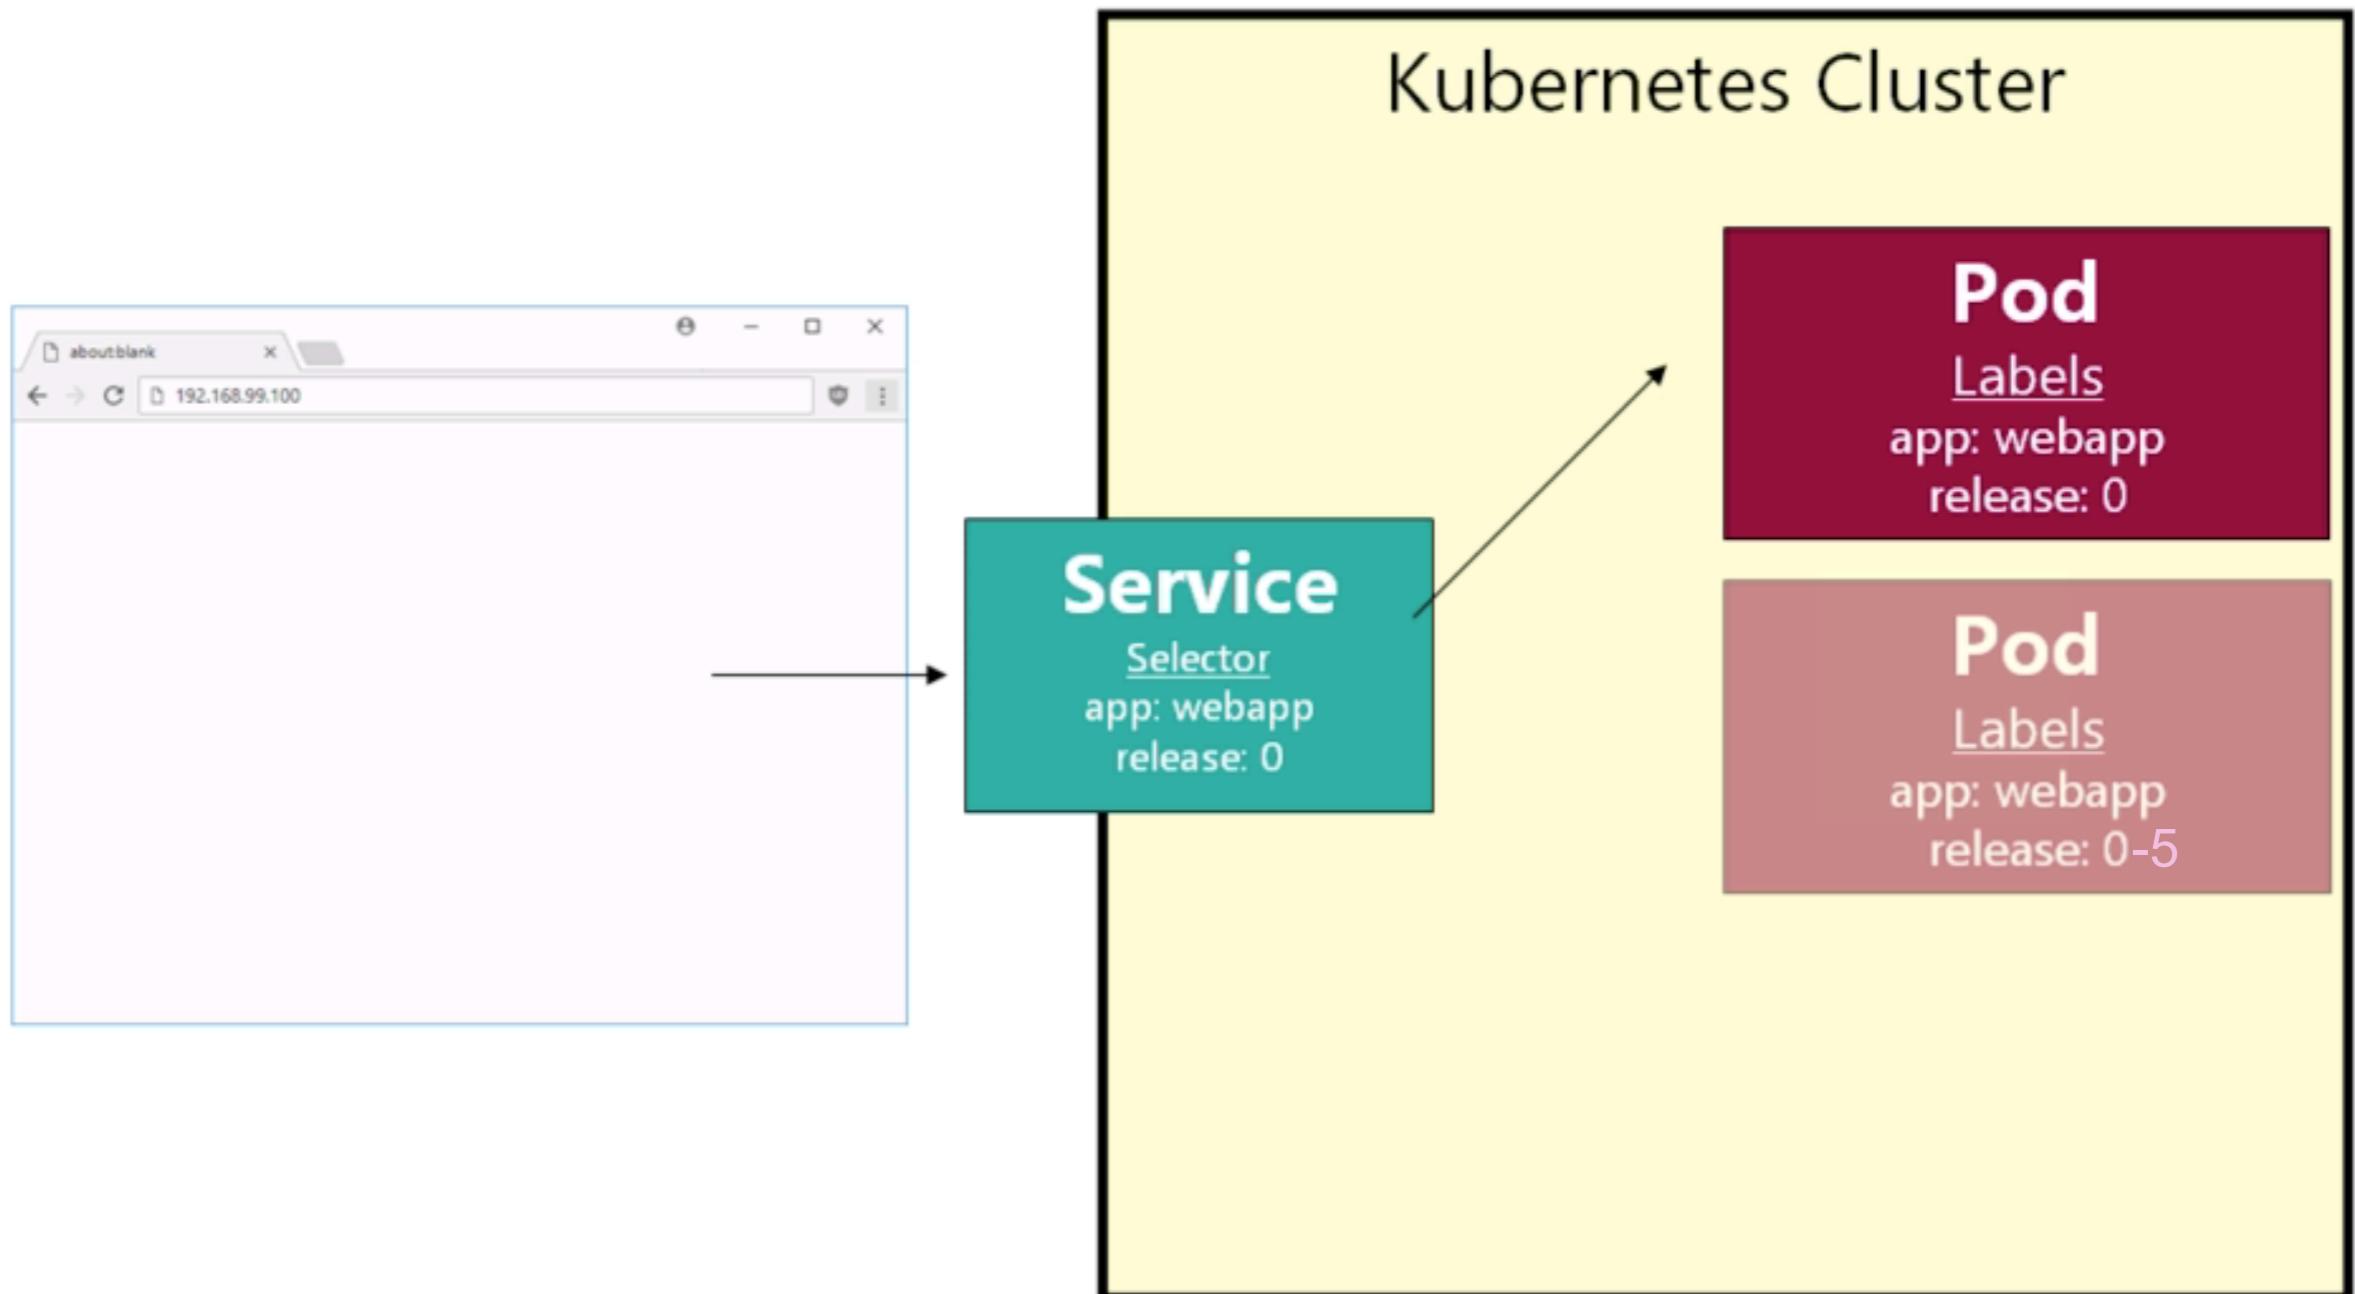
*We are going to have
some downtime! :(*

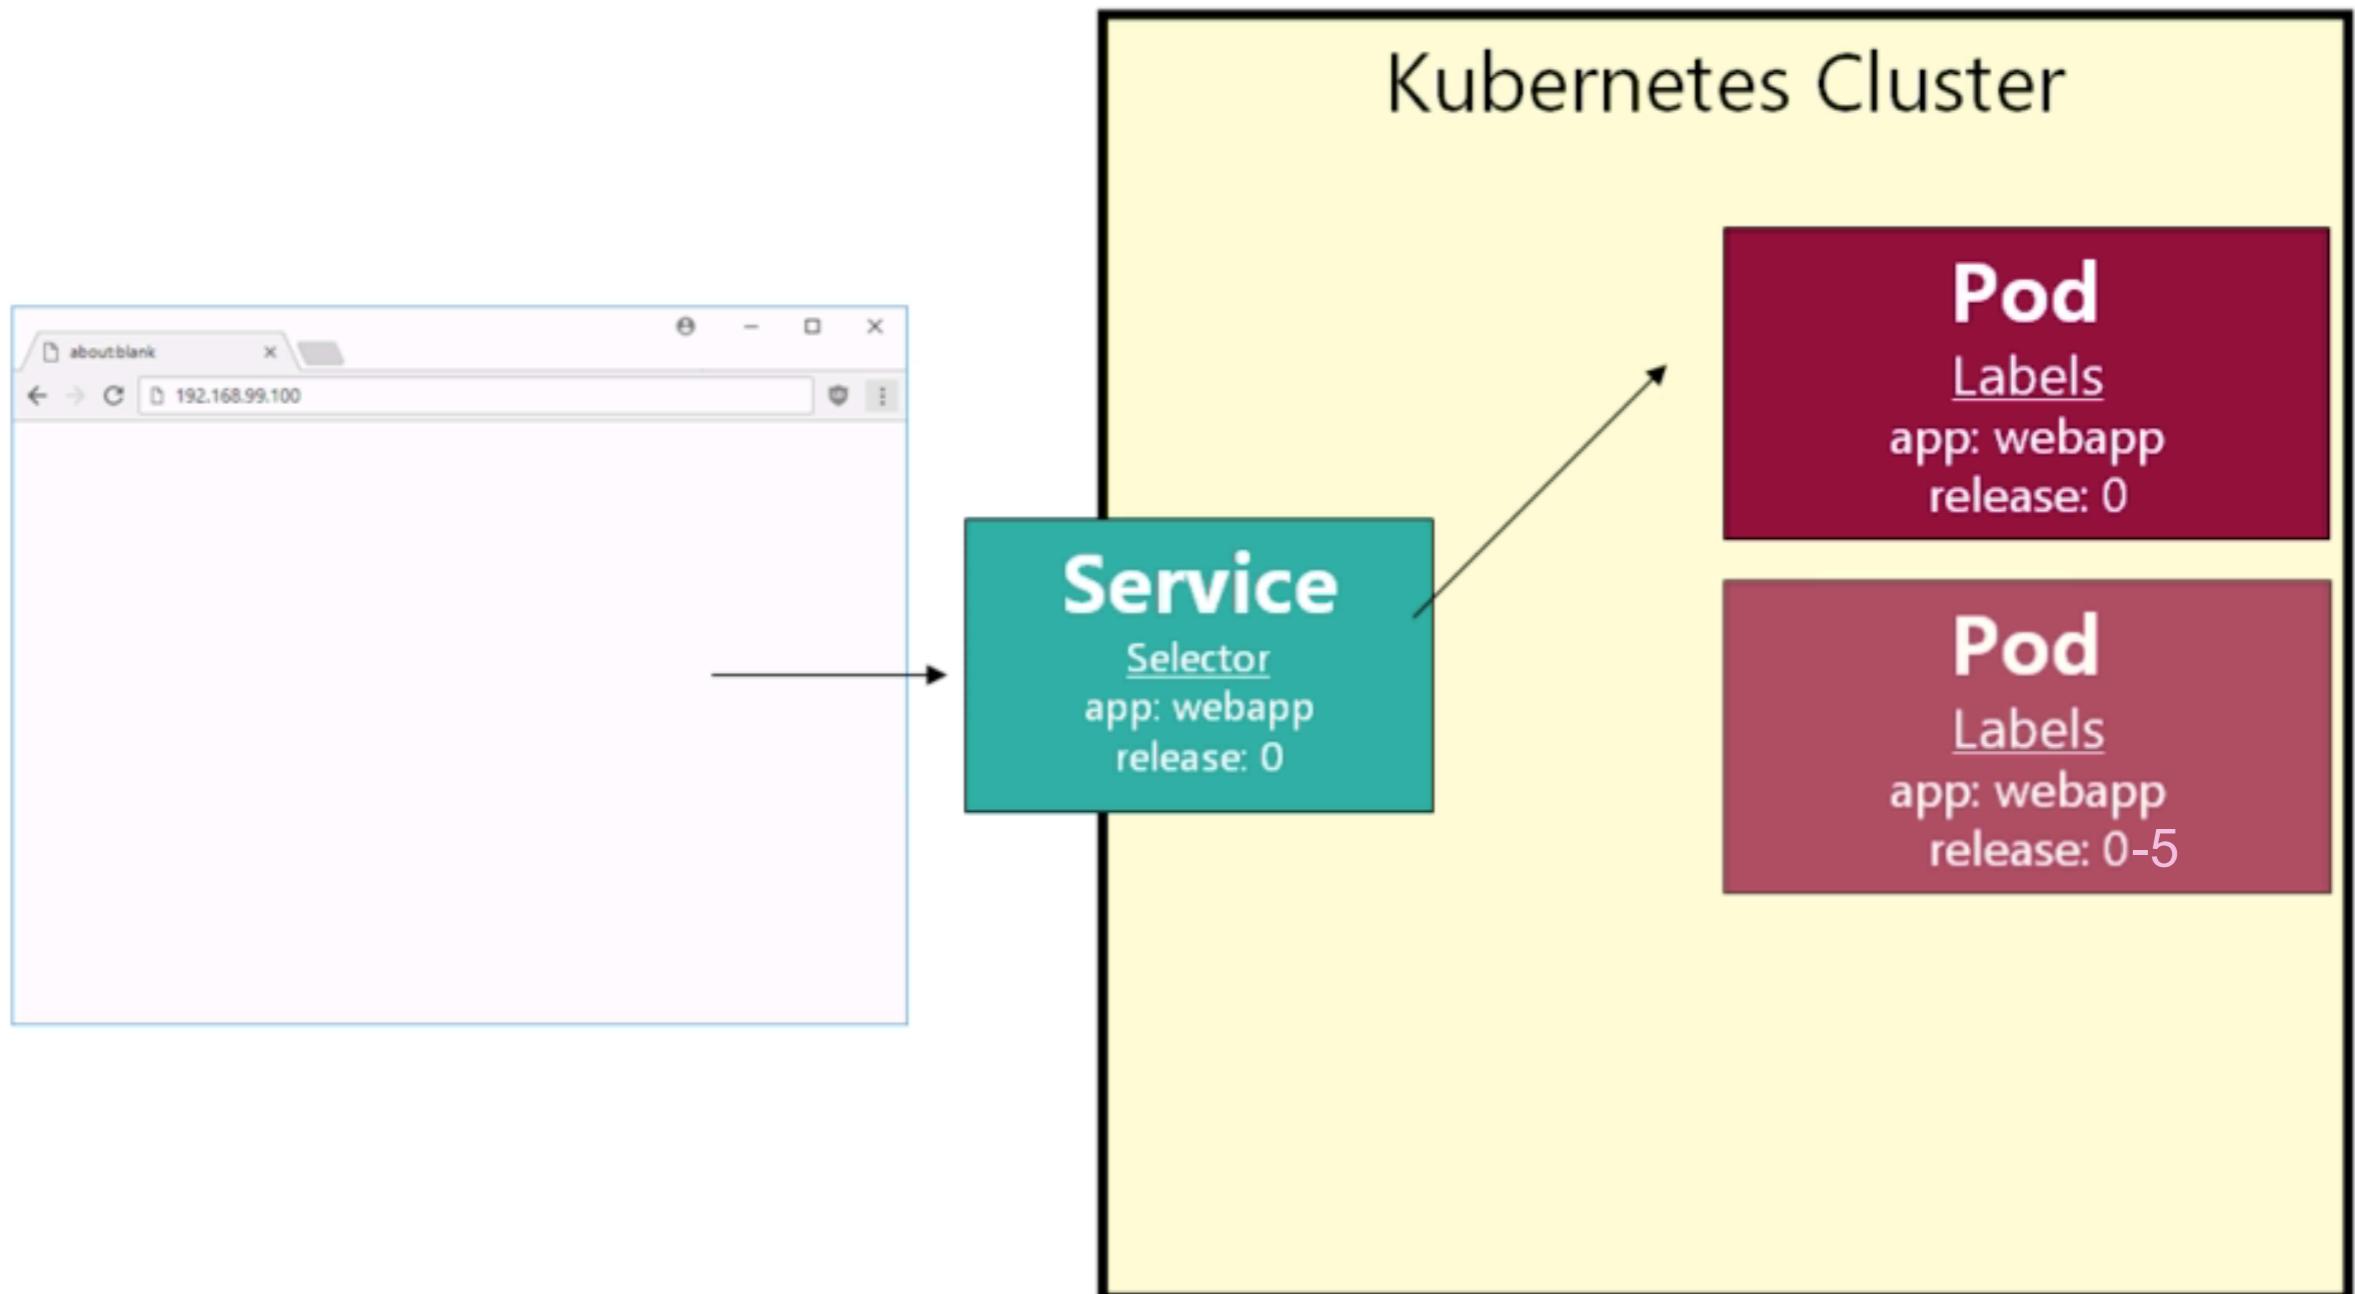
*In the next lessons we are
going to learn how to
deploy with zero downtime*

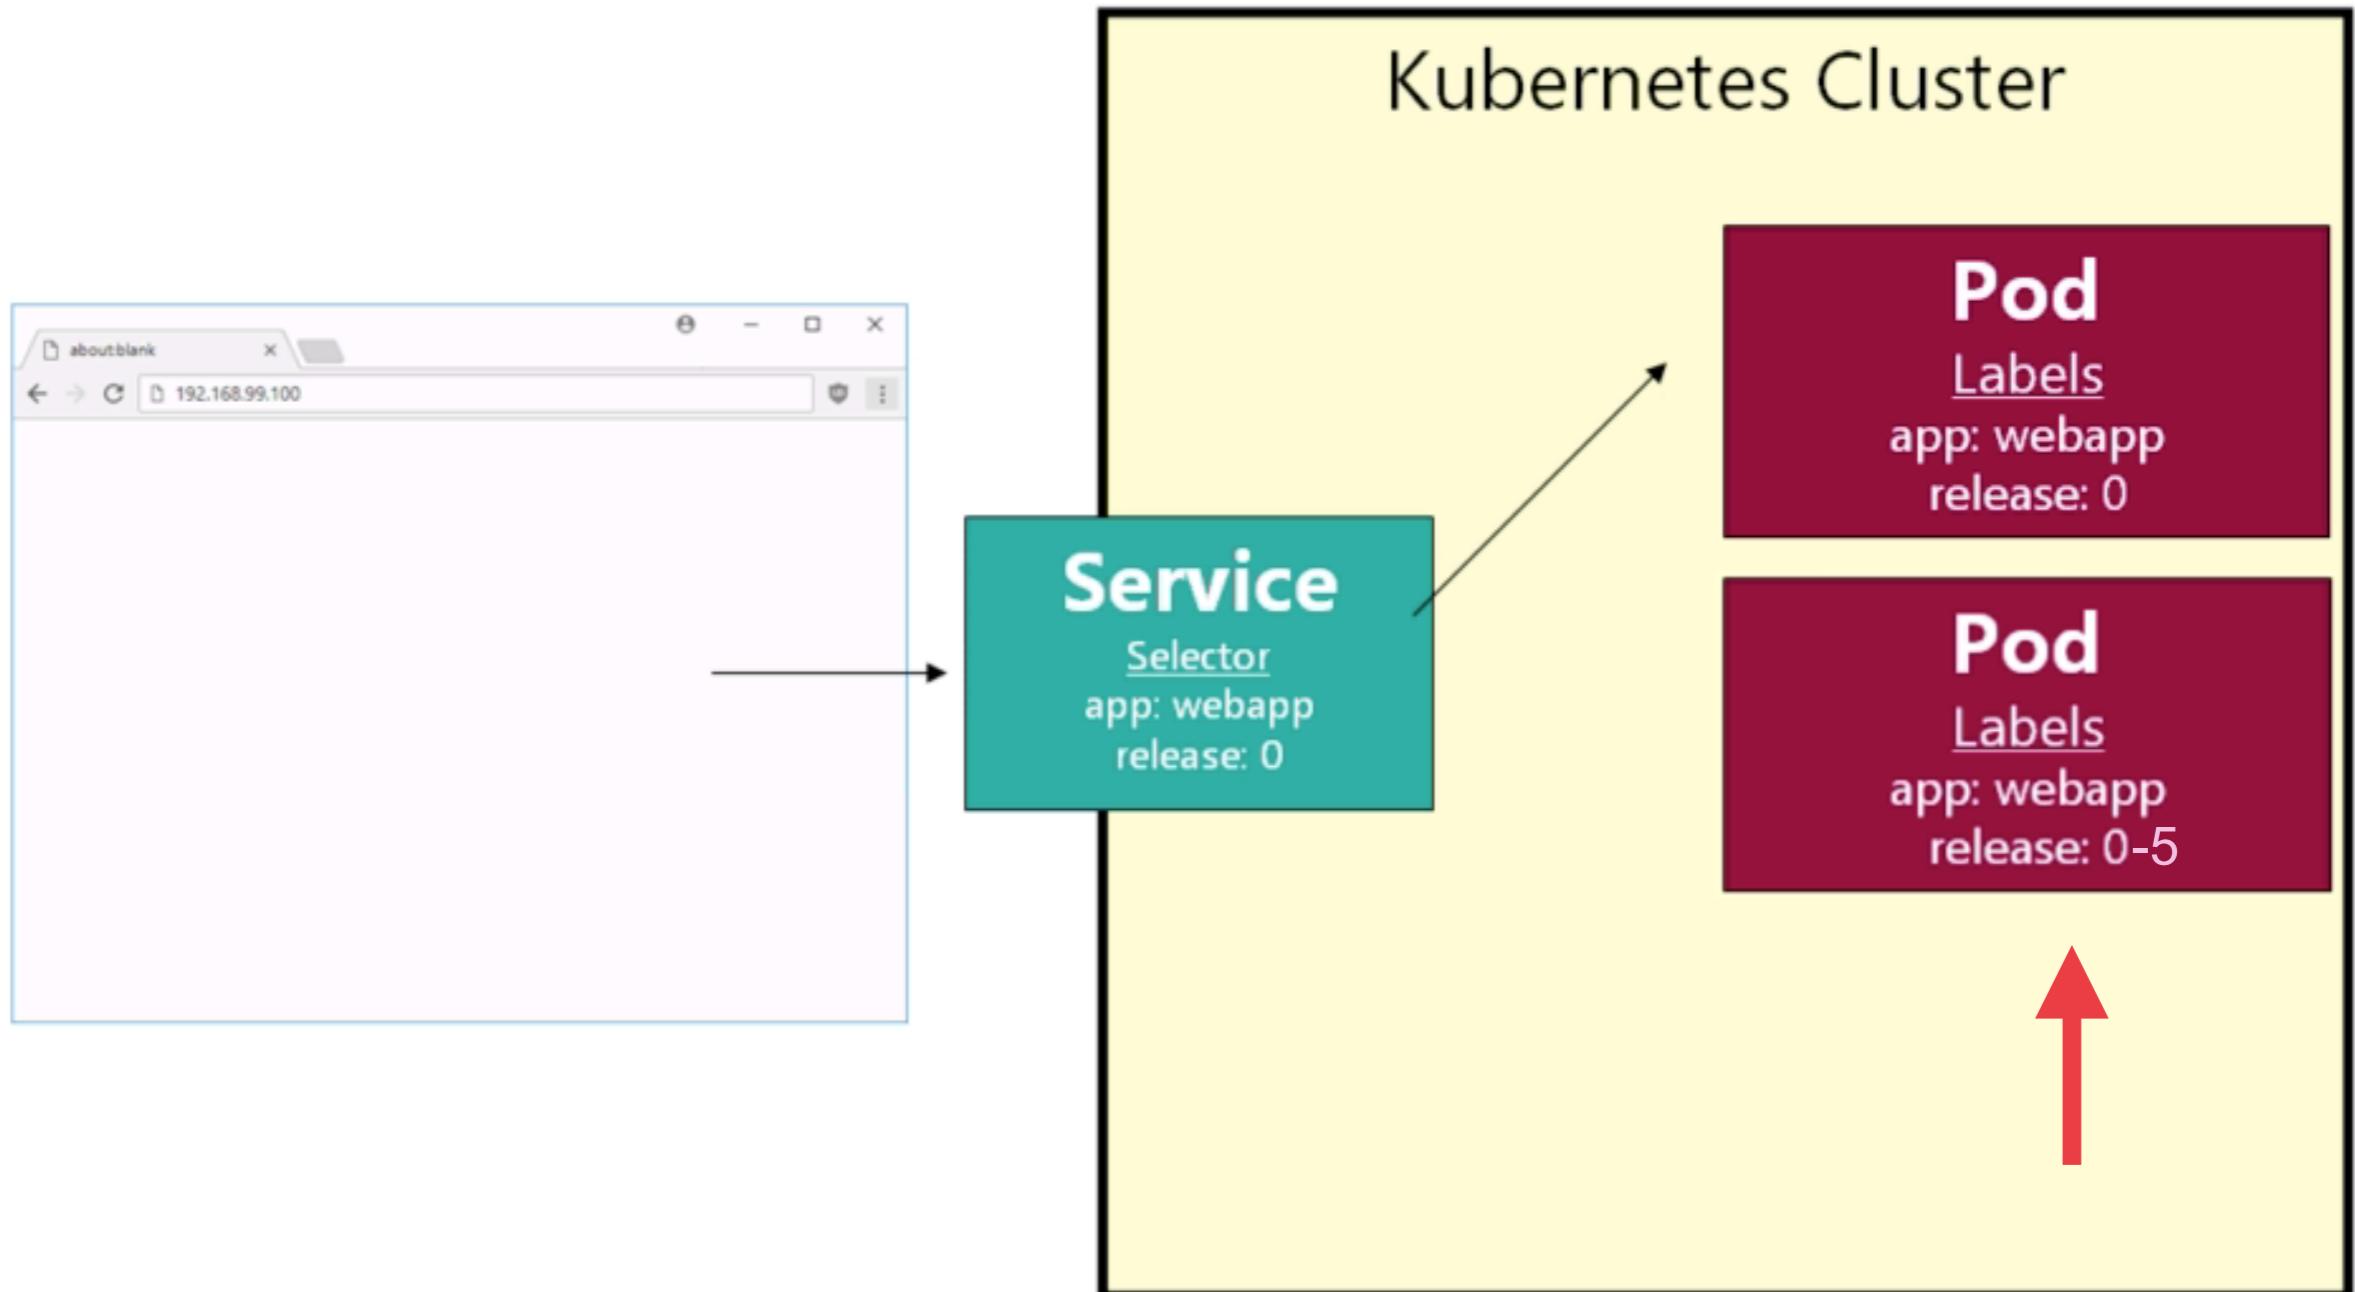


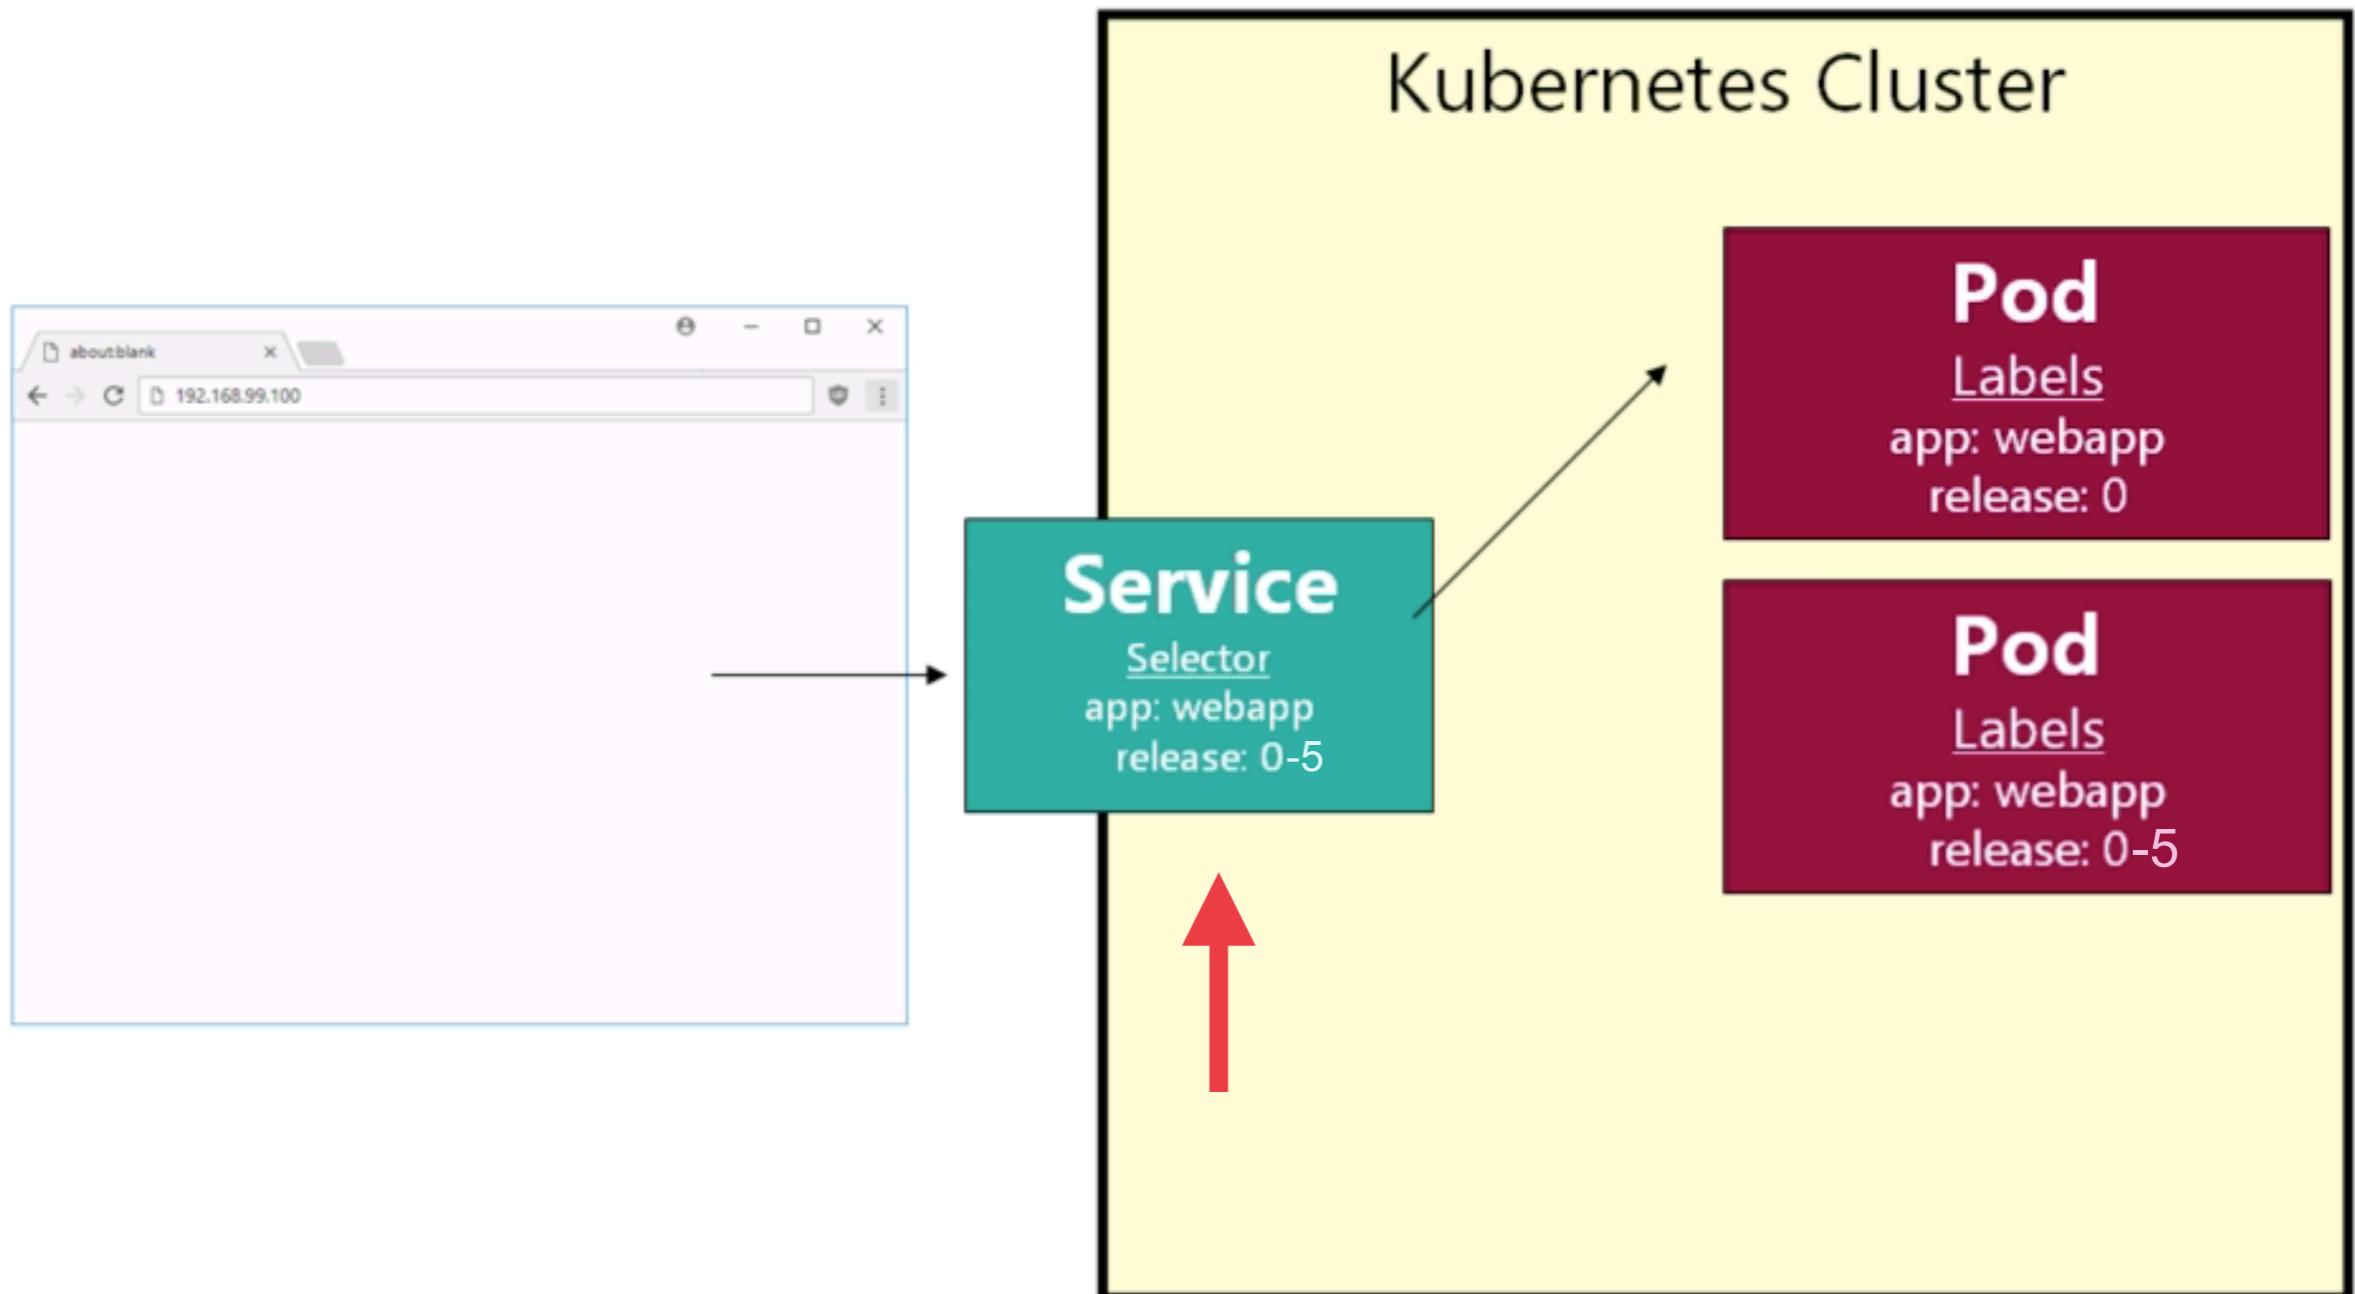
- ▶ Until we learn new concepts we can deploy with zero downtime using labels

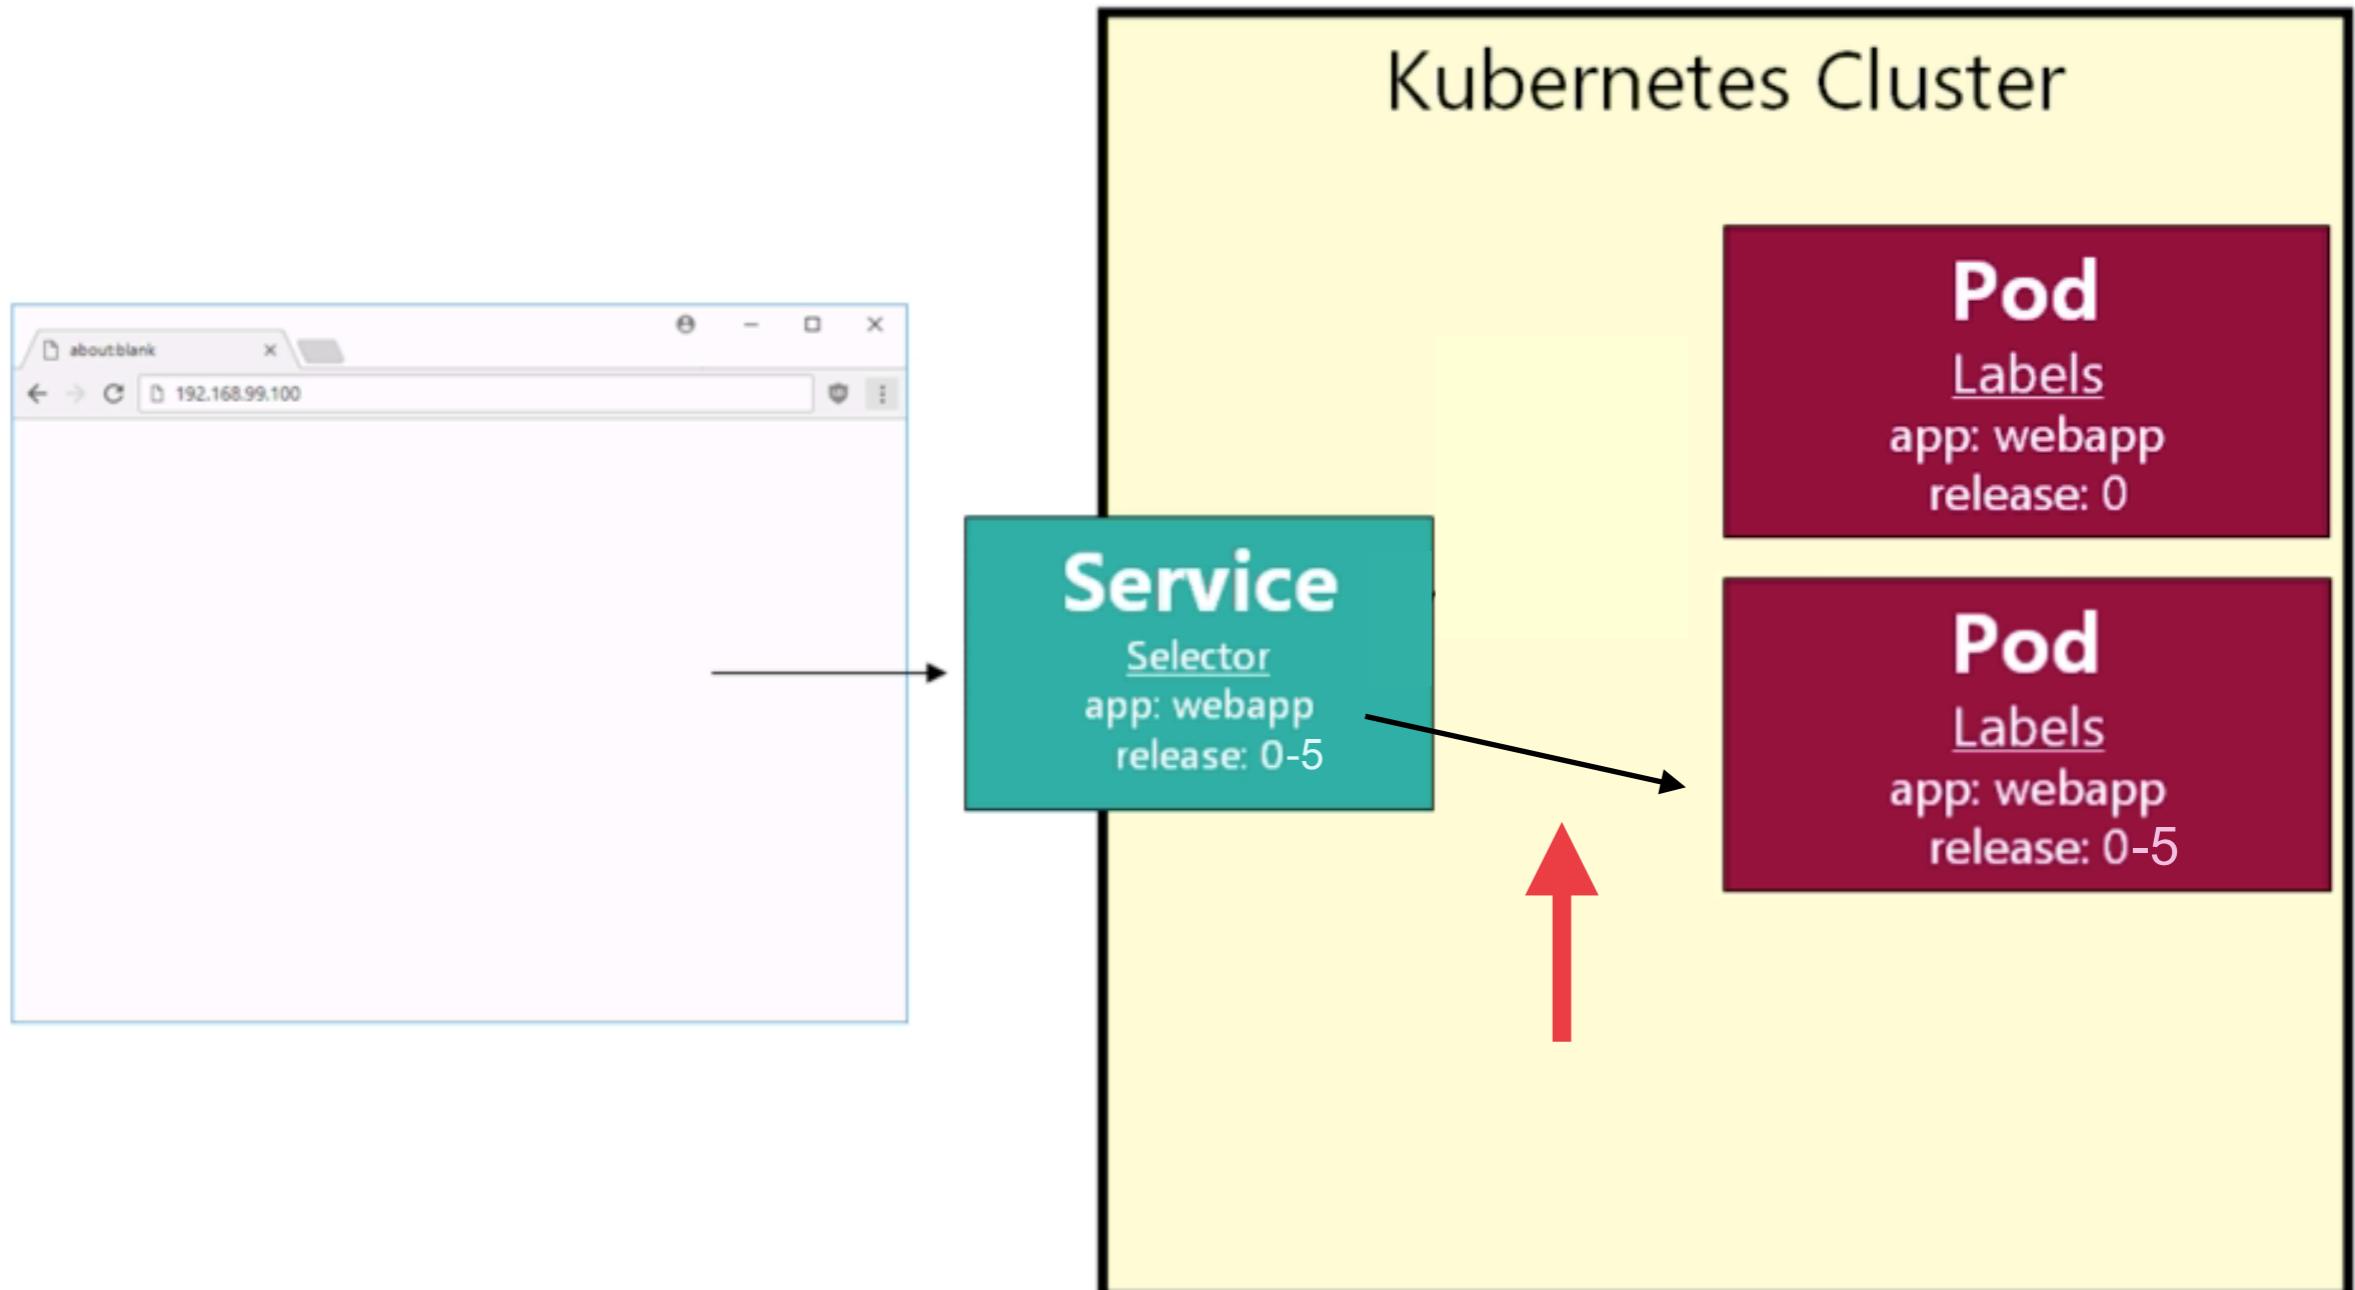












- ▶ This is a very fast operation the pod is already started

- ▶ We can switch between release0 and release0-5 at will

- ▶ We can start thinking that is possible support a production and development release on the same cluster (it is not the right way to do it or the more elegant one, but it is possible)

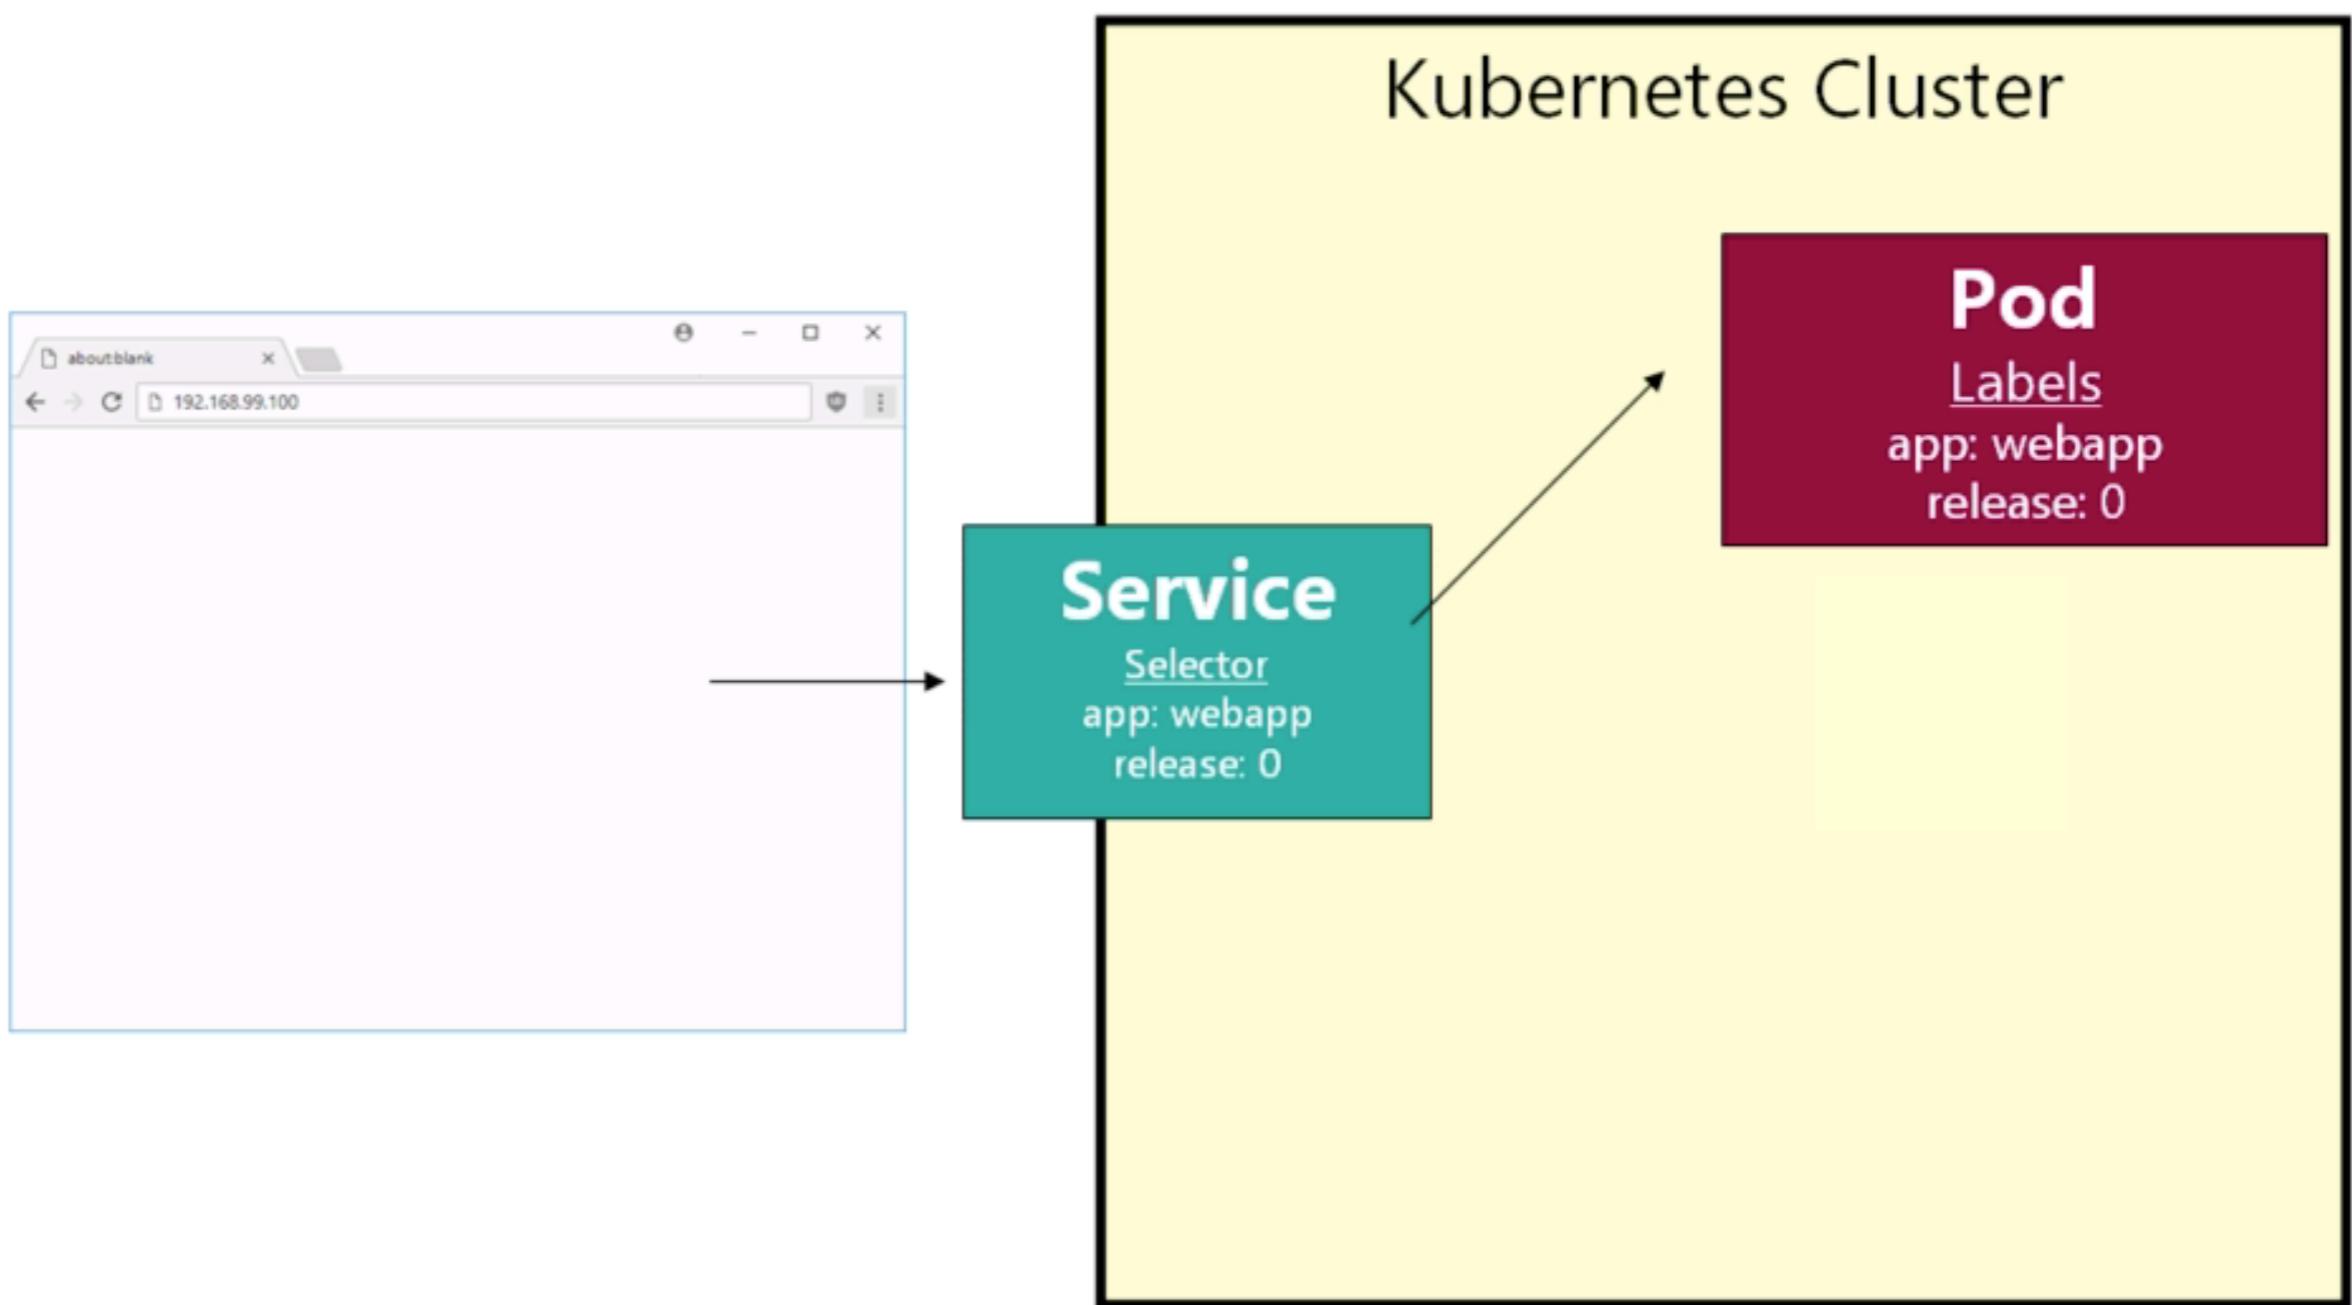
```
! pod.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: webapp
5    labels:
6      app: webapp
7      release: "0" ←
8
9  spec:
10   containers:
11     - name: webapp
12       image: richardchesterwood/k8s-fleetman-webapp-angular:release0
13
14 --- ←
15
```

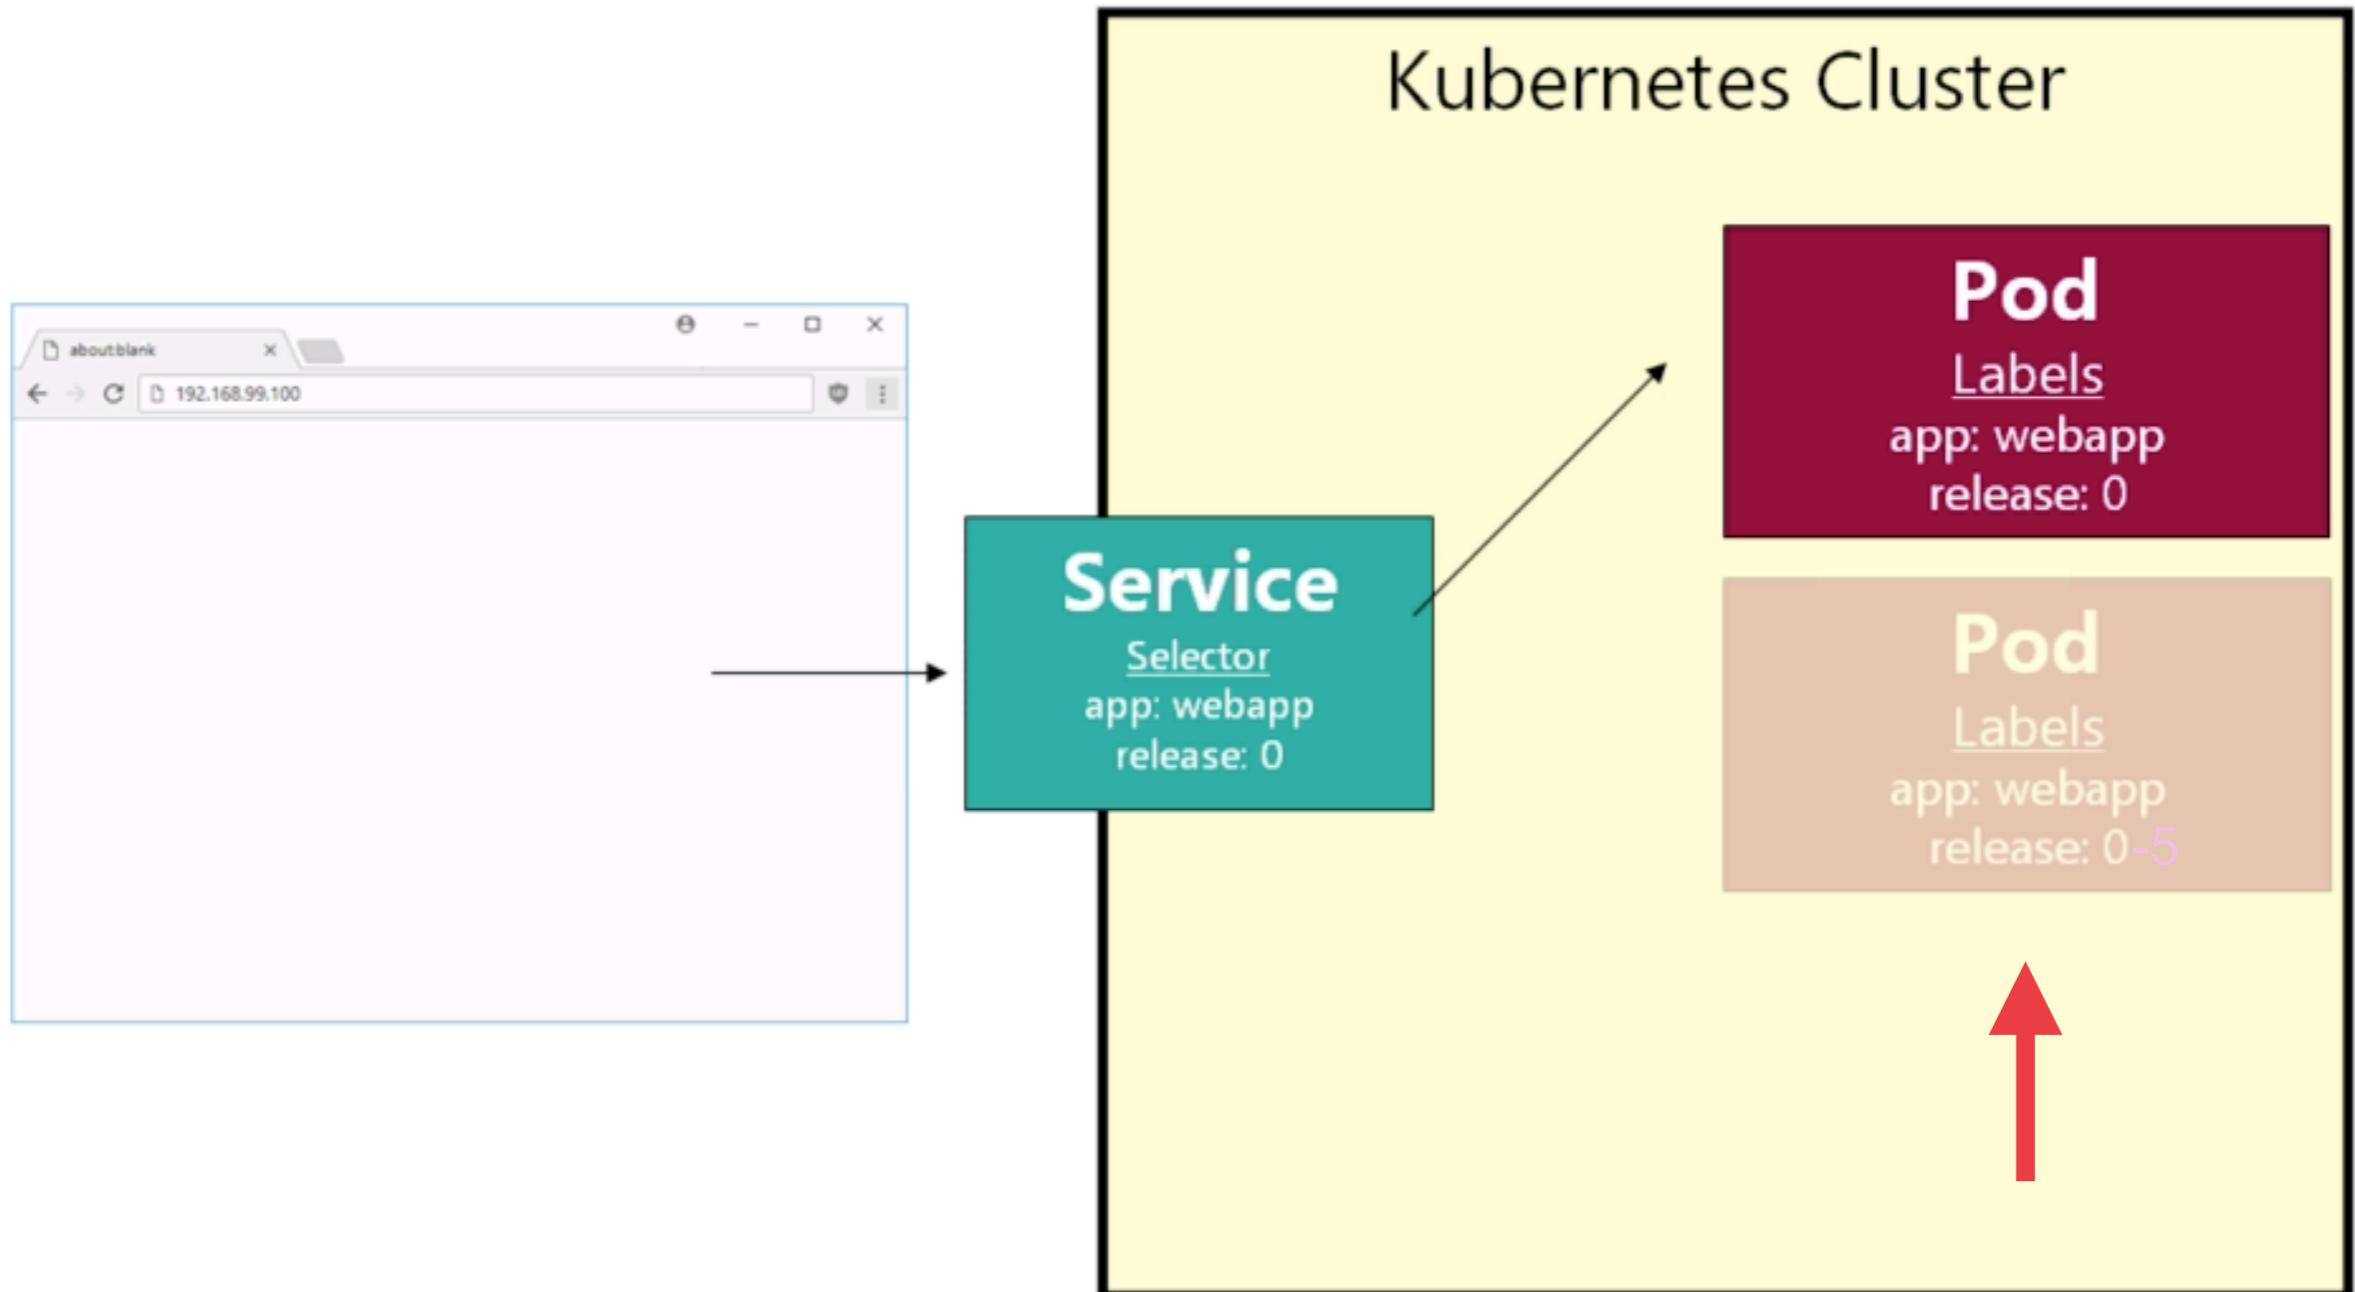
yaml document separator

```
16    apiVersion: v1
17    kind: Pod
18    metadata:
19      name: webapp-release-0-5 ←
20      labels:
21        app: webapp
22        release: "0-5" ←
23
24    spec:
25      containers:
26        - name: webapp
27          image: richardchesterwood/k8s-fleetman-webapp-angular:release0-5
28
```

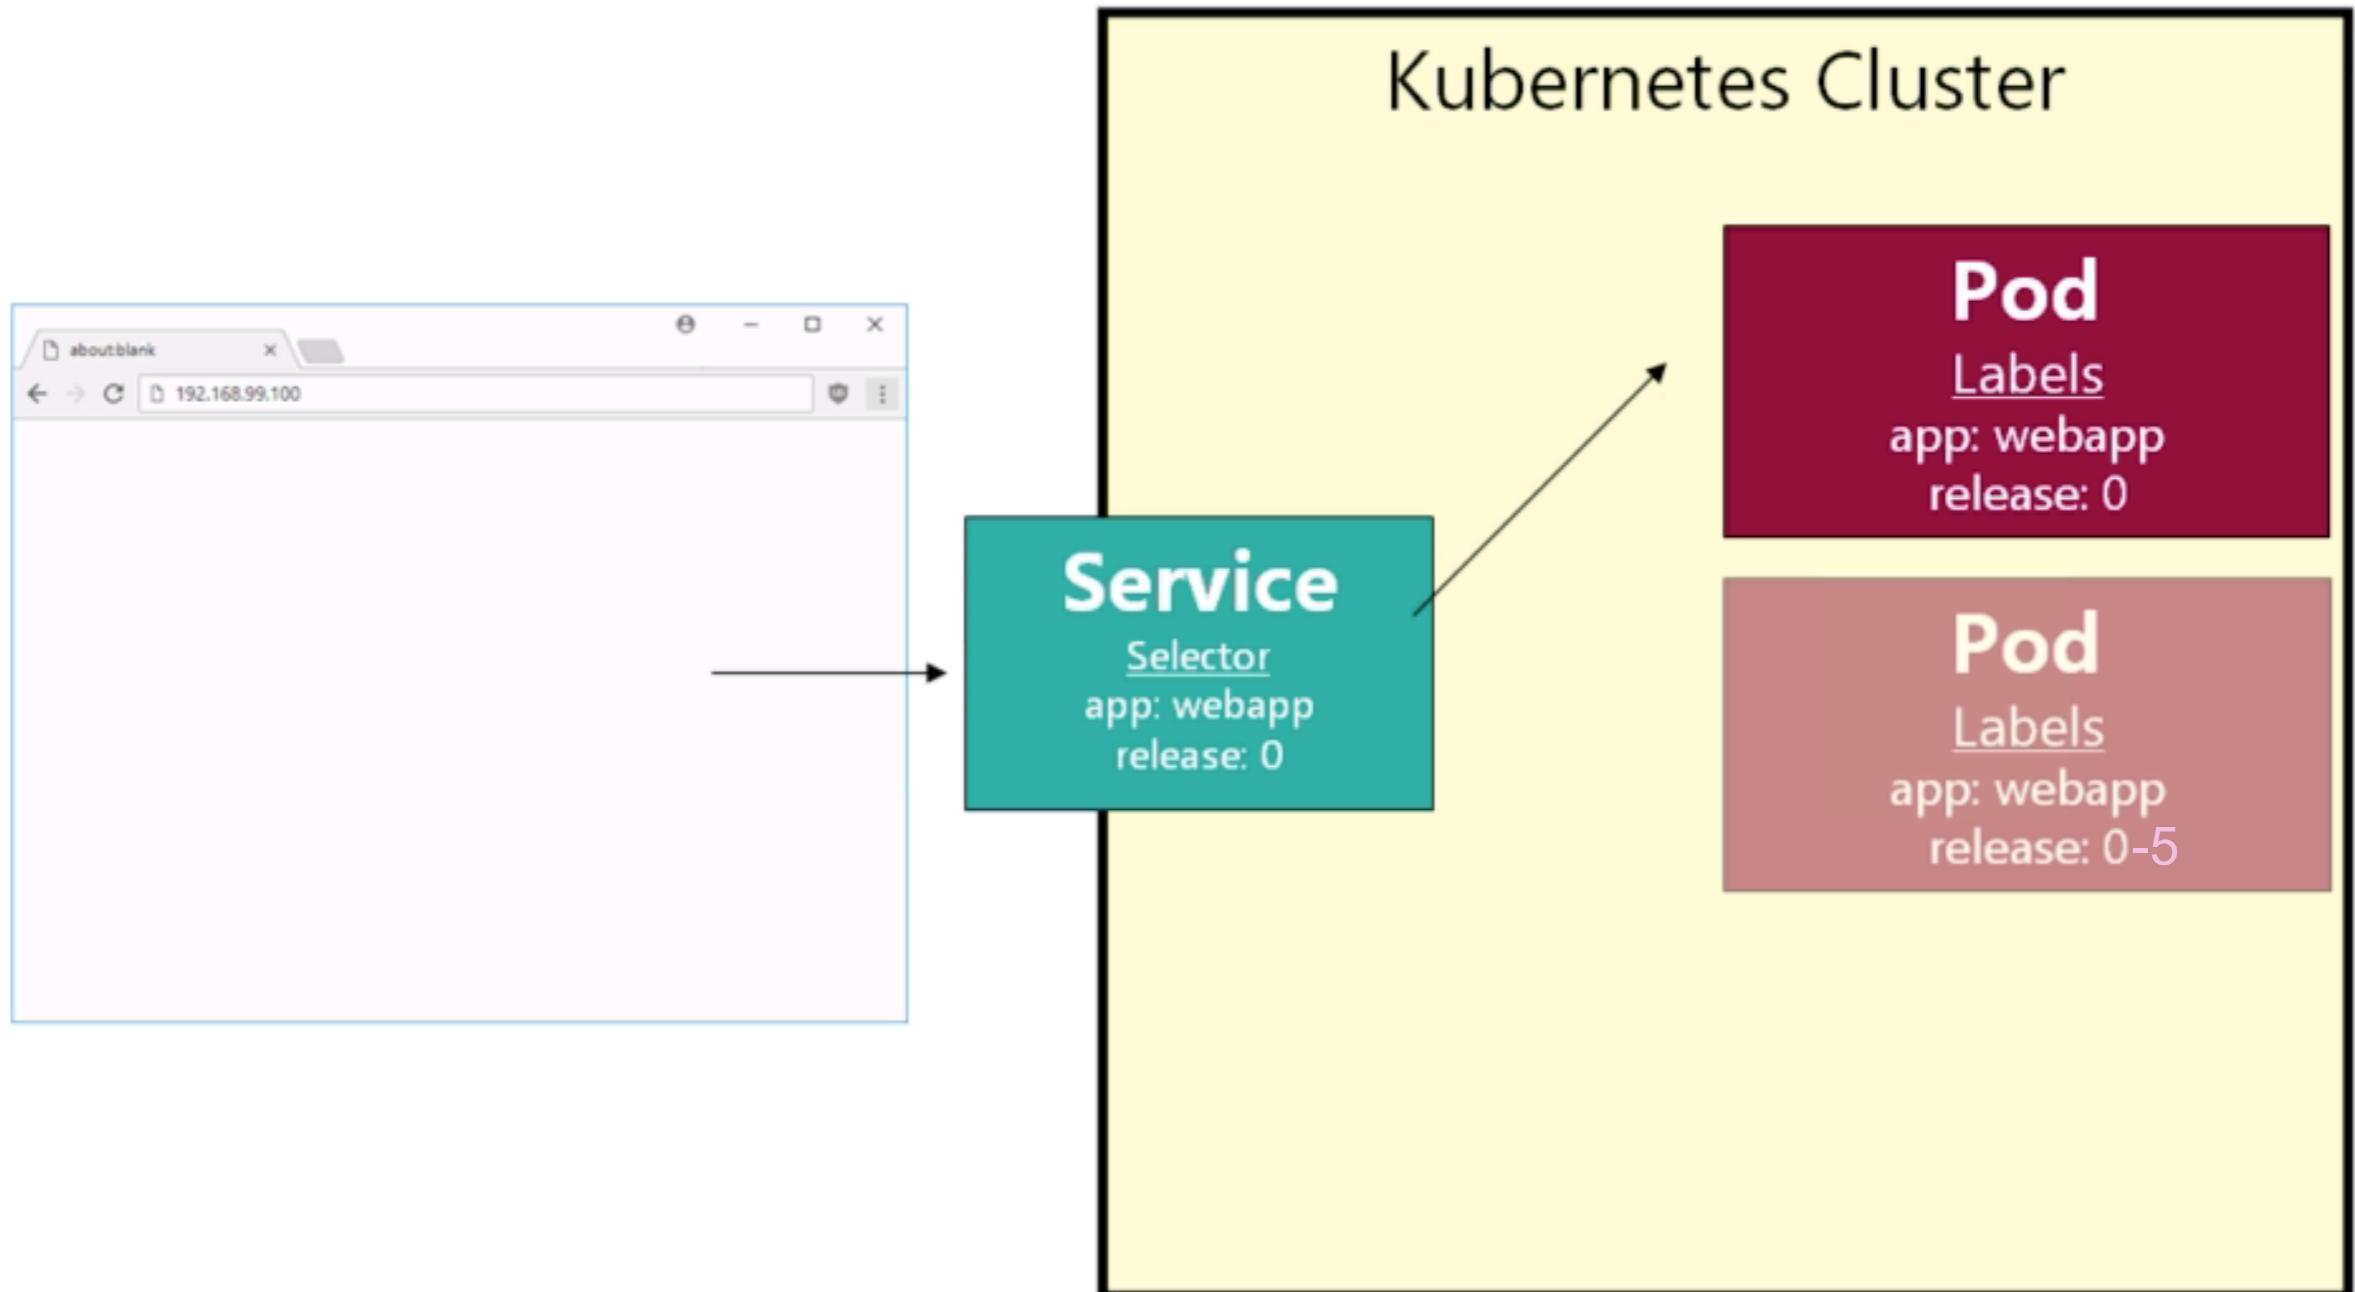
```
! webapp-service.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: fleetman-webapp
5
6  spec:
7    selector:
8      app: webapp
9      release: "0" ←
10
11  ports:
12    - name: http
13      port: 80
14      nodePort: 30080
15
16  type: NodePort
17
```

This is our scenario

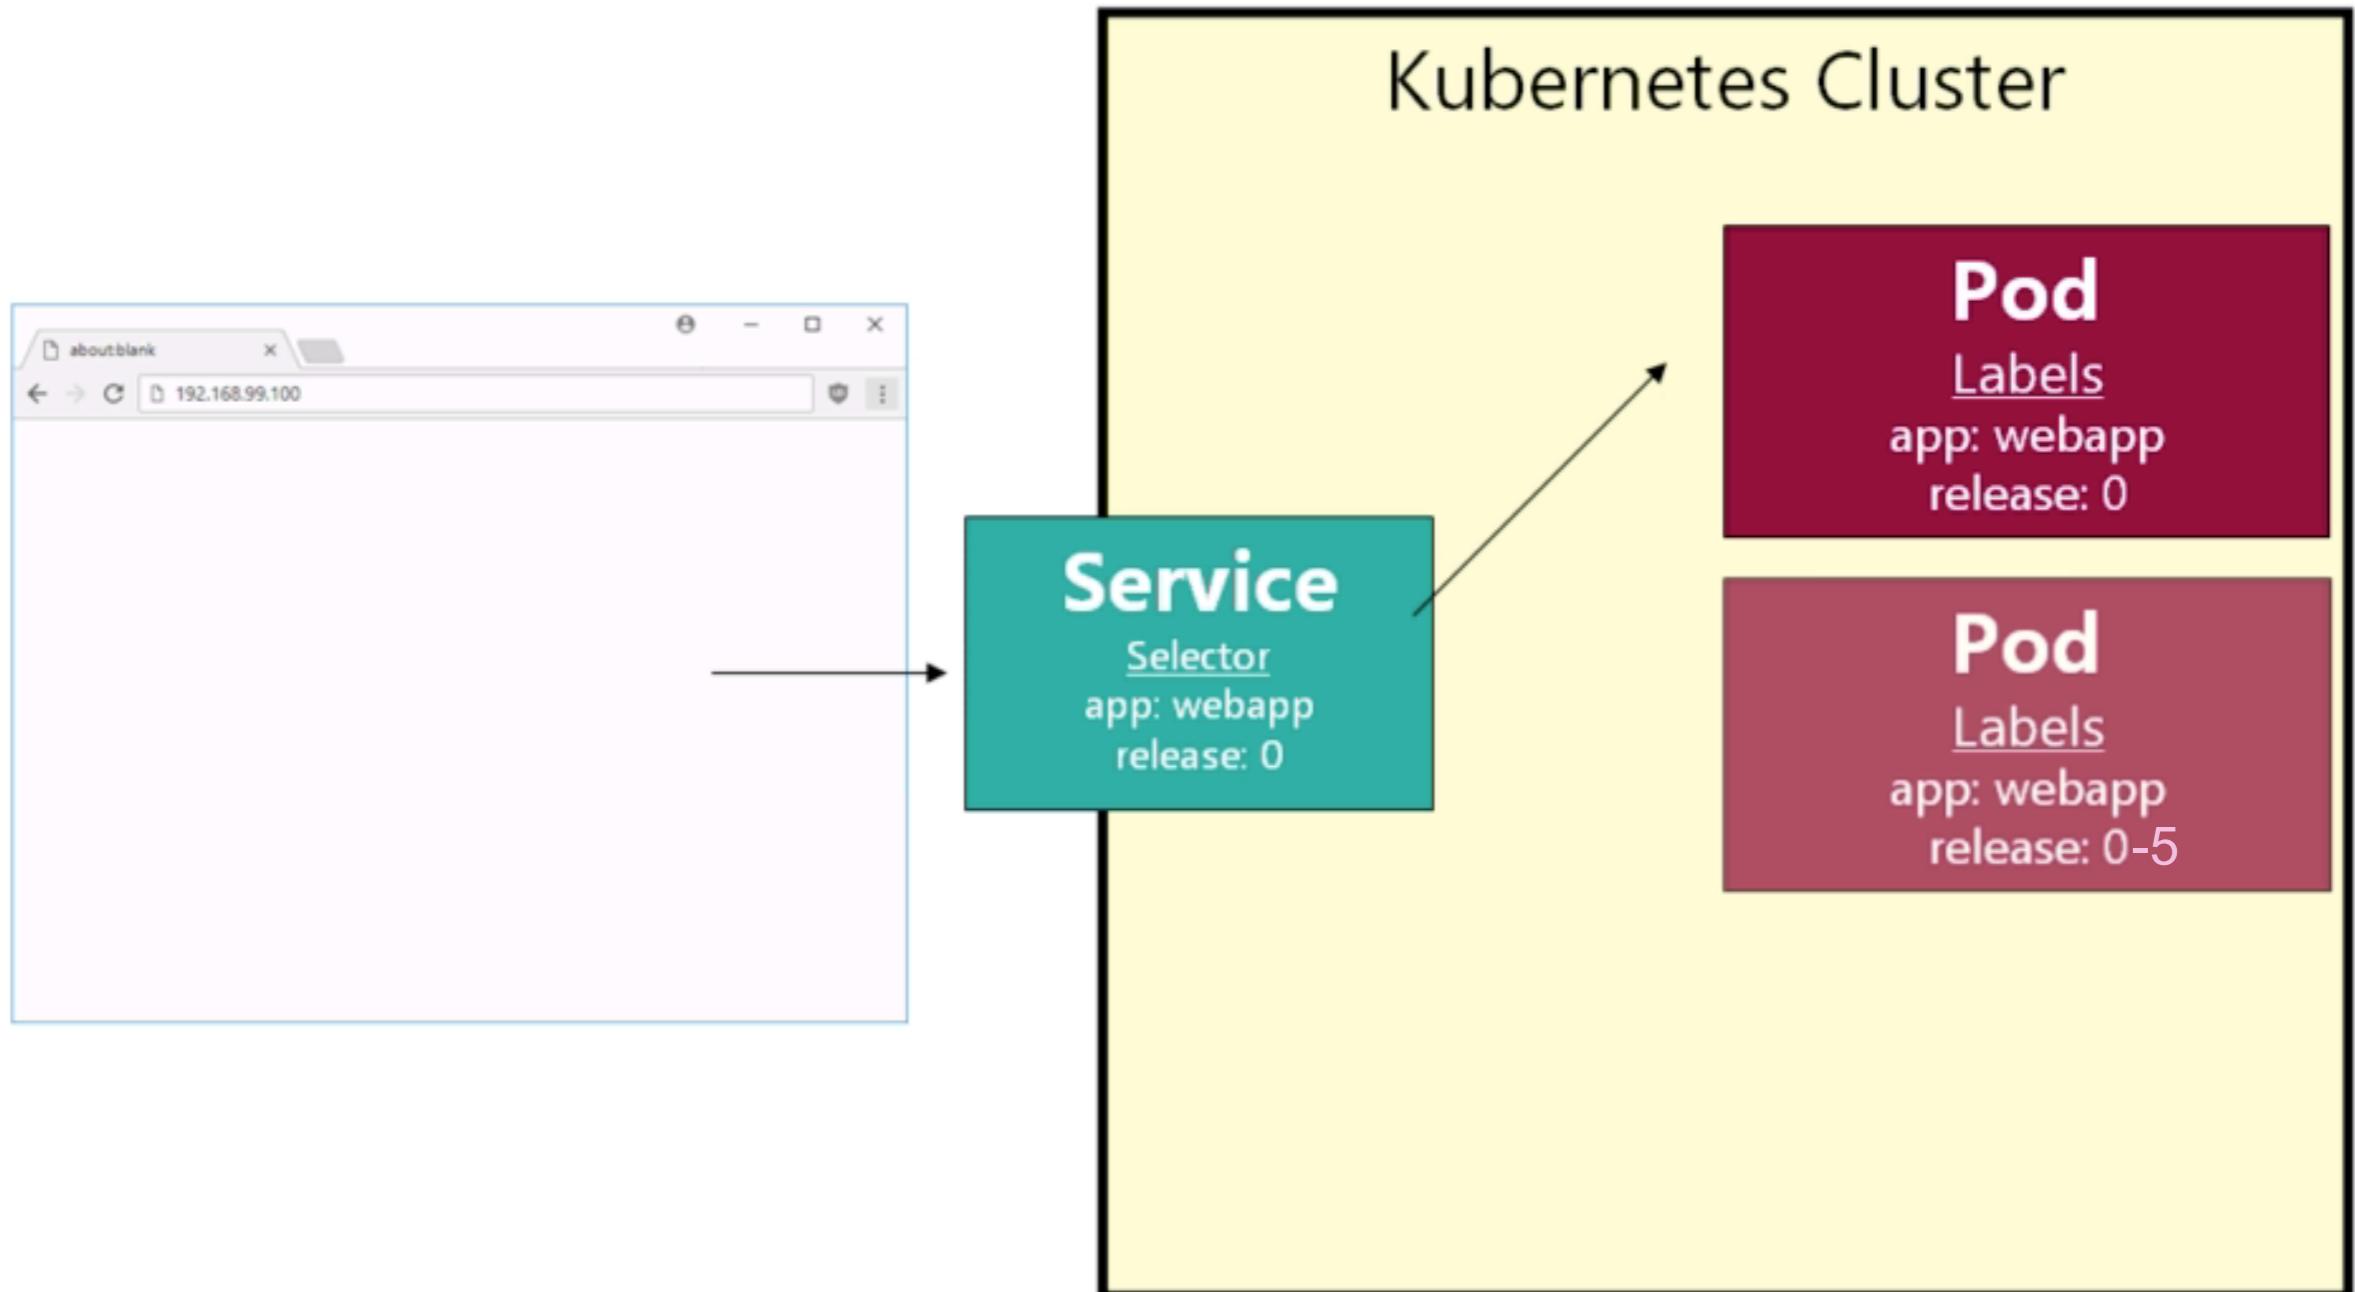




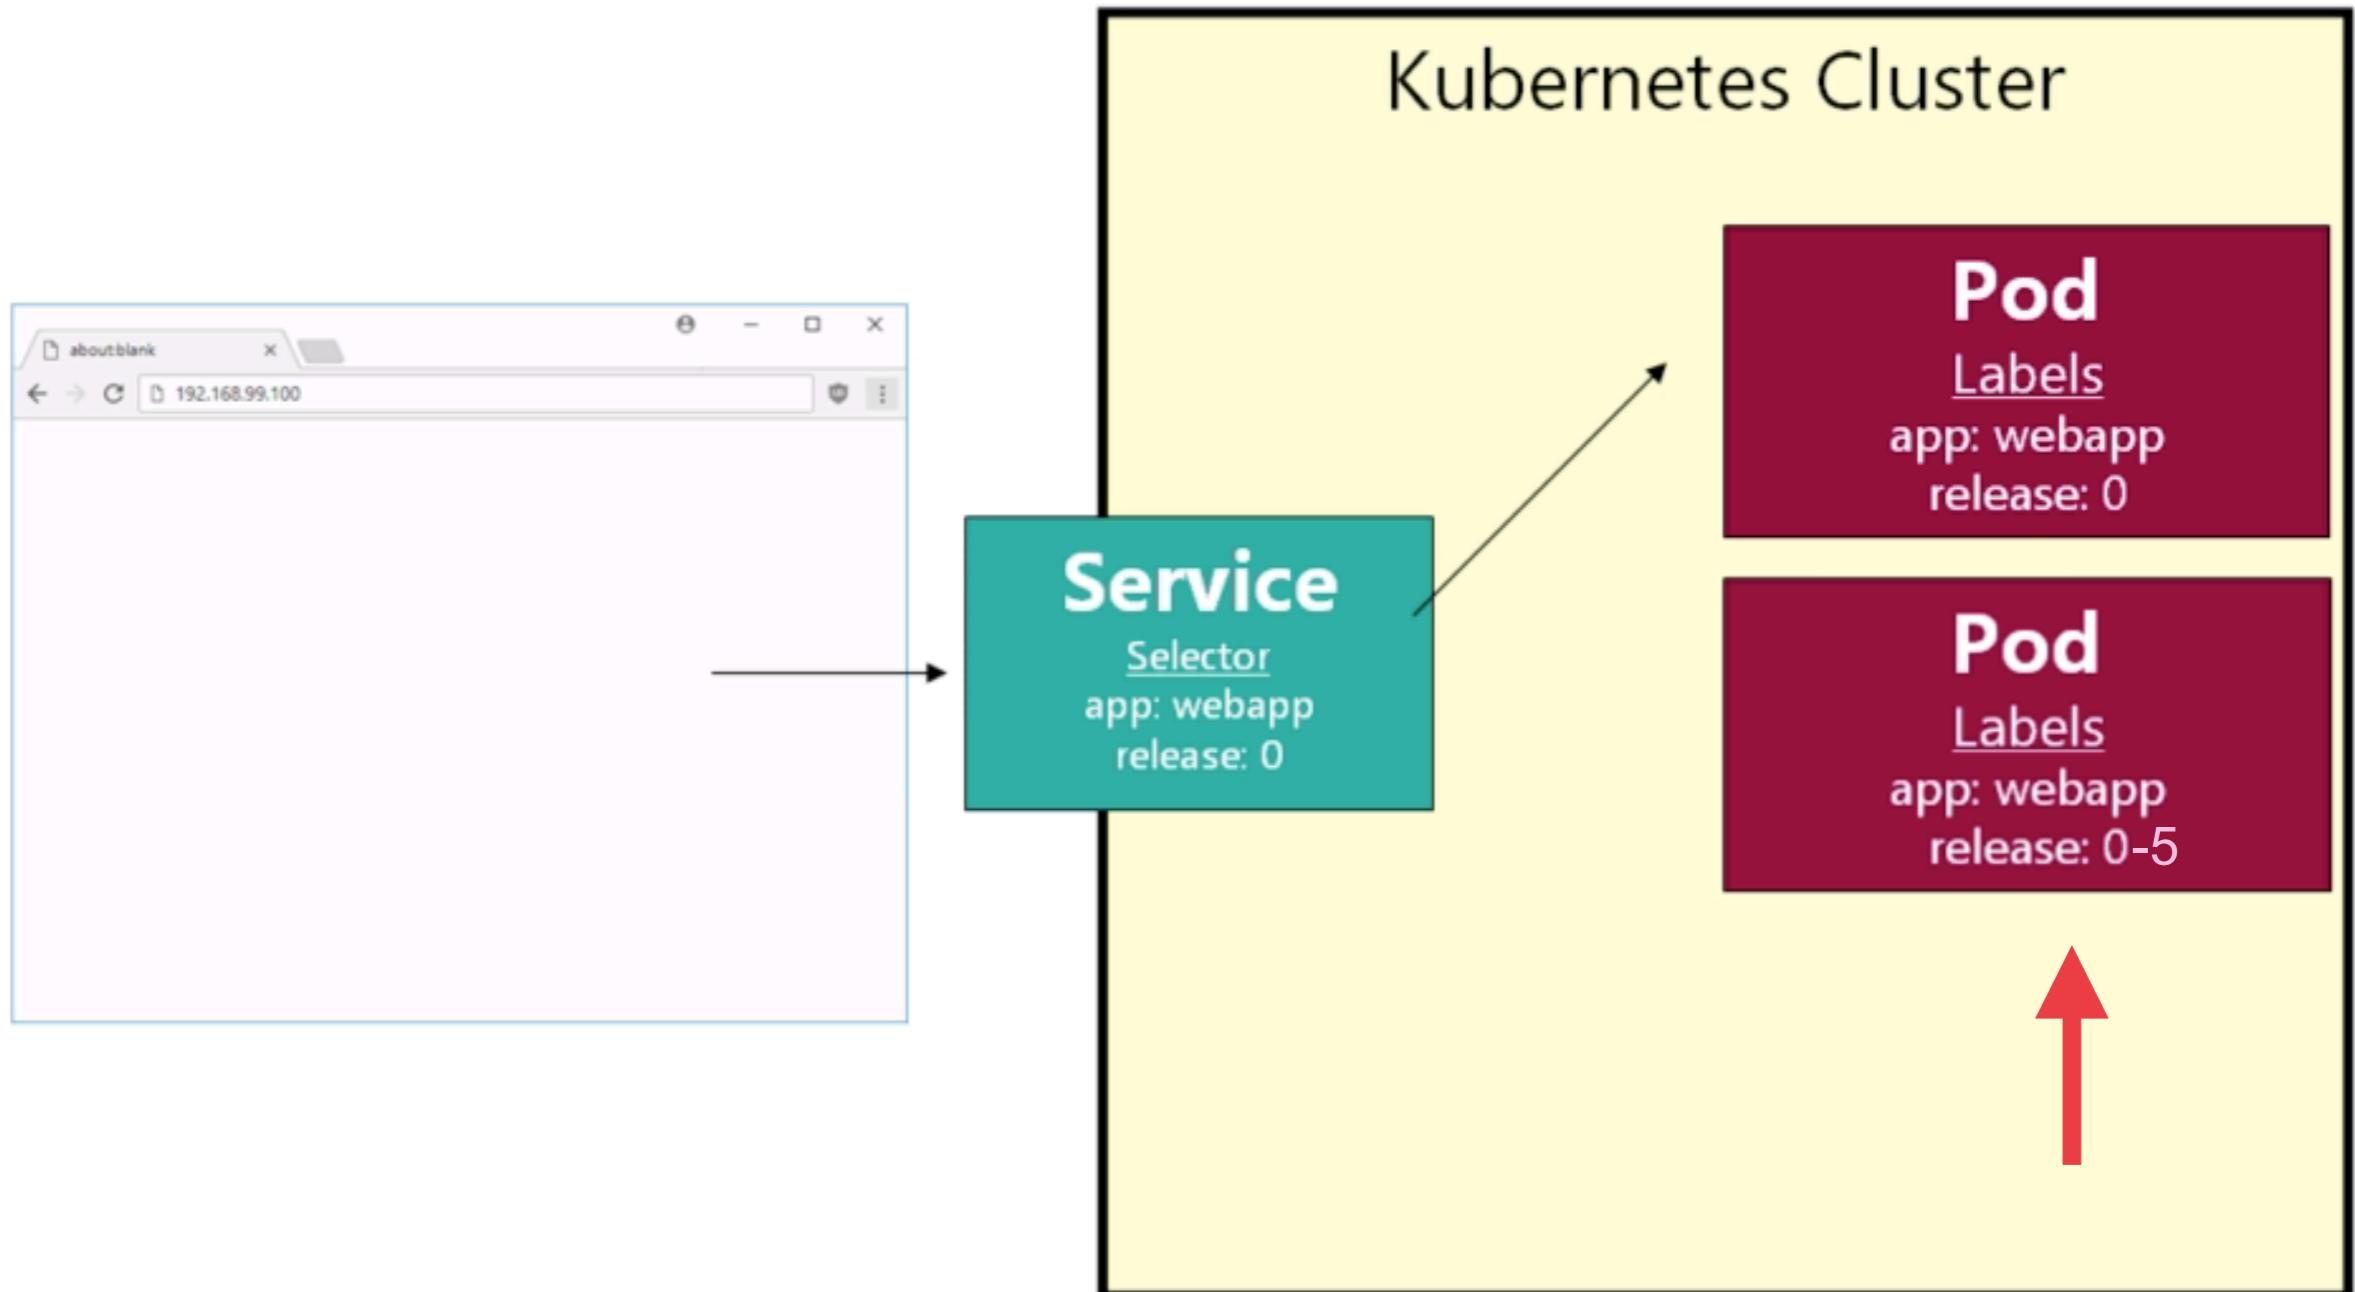
Pod deploying...



Pod deploying...



Pod deploying...



Pod deployed!

- ▶ `kubectl apply -f pod.yaml`



```
$ kubectl apply -f pod.yaml
```

- ▶ `kubectl apply -f webapp-service.yaml`



```
$ kubectl apply -f webapp-service.yaml
```

▶ **kubectl get all**

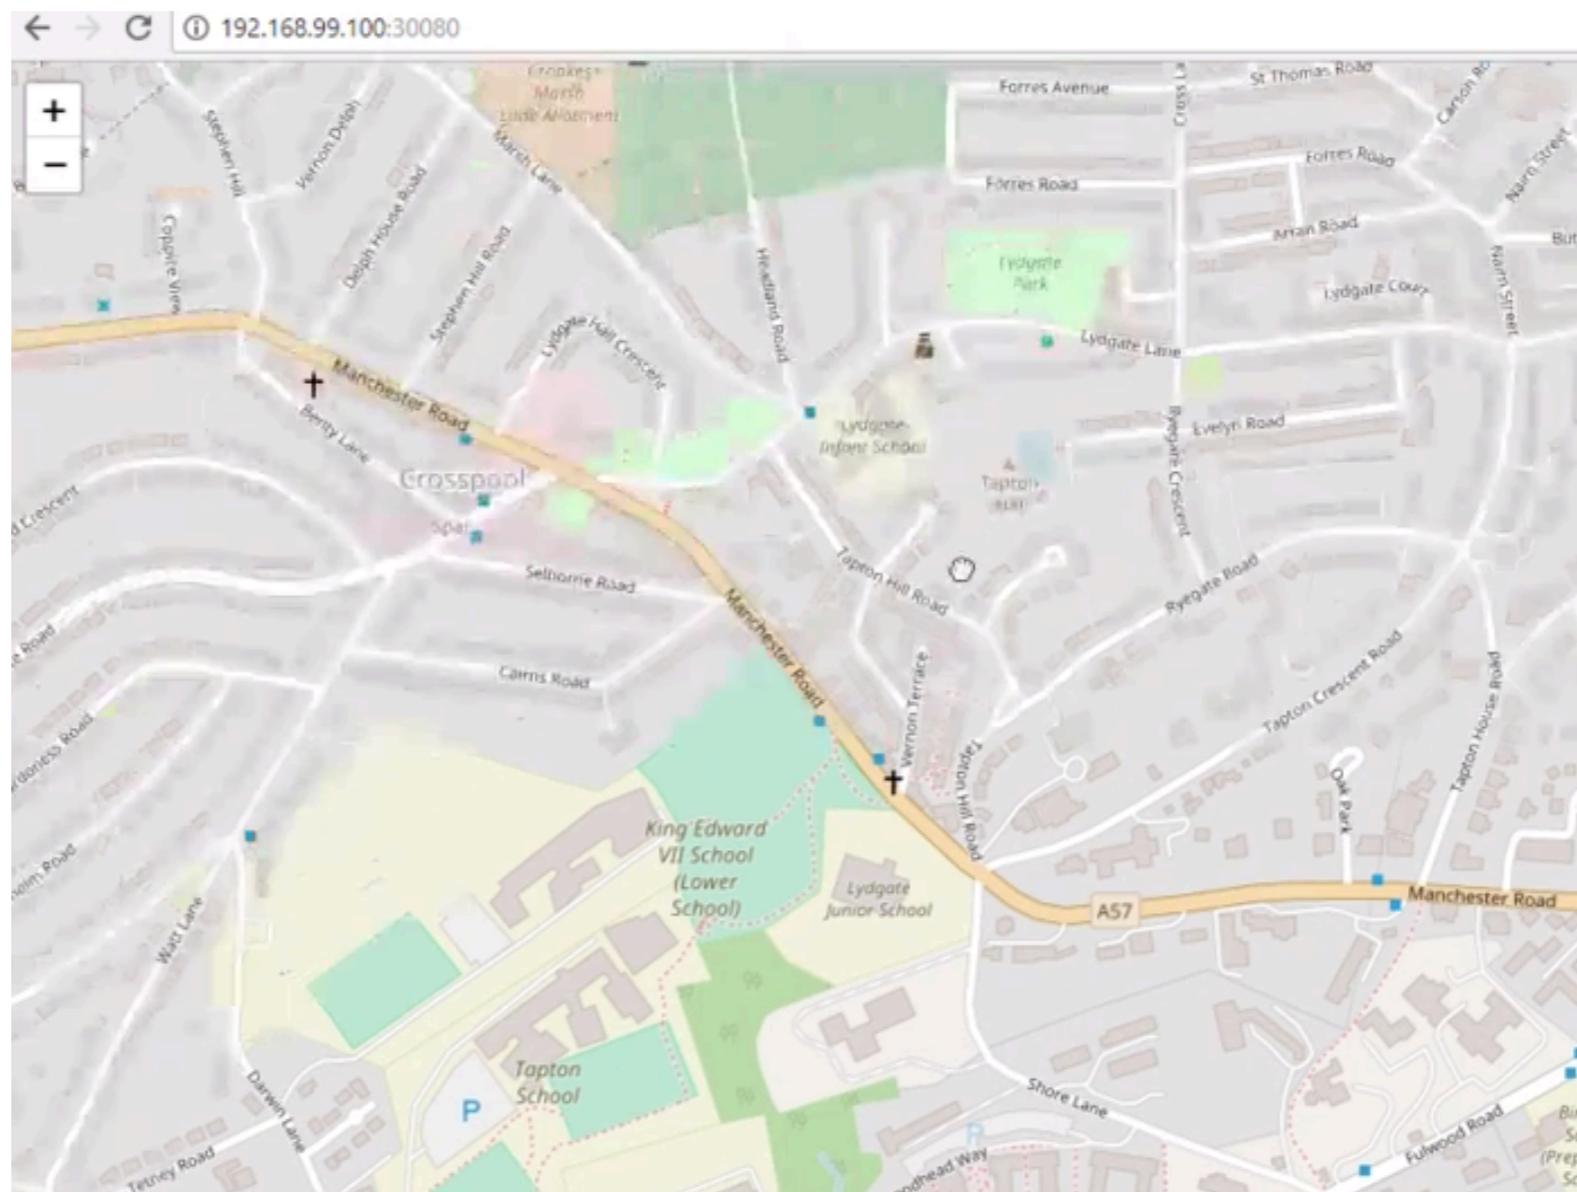
NAME	READY	STATUS	RESTARTS	AGE
pod/webapp	1/1	Running	1	20h
pod/webapp-release-0-5	1/1	Running	0	1m
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/fleetman-webapp	NodePort	10.107.11.35	<none>	80:30080/TCP
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP

- ▶ `kubectl describe service fleetman-webapp`

```
Name:                      fleetman-webapp
Namespace:                 default
Labels:                    <none>
Annotations:               kubectl.kubernetes.io/last-applied-configuration=
                           {"apiVersion": "v1", "kind": "Service", "metadata": {"name": "fleetman-webapp"}, "spec": {"ports": [{"port": 80, "targetPort": 80, "type": "NodePort"}], "selector": {"app": "webapp", "release": "0"}, "targetPort": 30080}, "status": {"loadBalancer": {}, "ports": [{"ip": "10.107.11.35", "nodePort": 30080, "port": 80, "protocol": "TCP"}]}}
Selector:                  app=webapp,release=0
Type:                      NodePort
IP:                        10.107.11.35
Port:                      http  80/TCP
TargetPort:                80/TCP
NodePort:                  http  30080/TCP
Endpoints:                 172.17.0.4:80
Session Affinity:          None
External Traffic Policy:   Cluster
Events:                    <none>
```



Let's try it now!



*Now let's change to
relesase0-5*

```
! webapp-service.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: fleetman-webapp
5
6  spec:
7    selector:
8      app: webapp
9      release: "0-5" ←
10
11  ports:
12    - name: http
13      port: 80
14      nodePort: 30080
15
16  type: NodePort
17
```

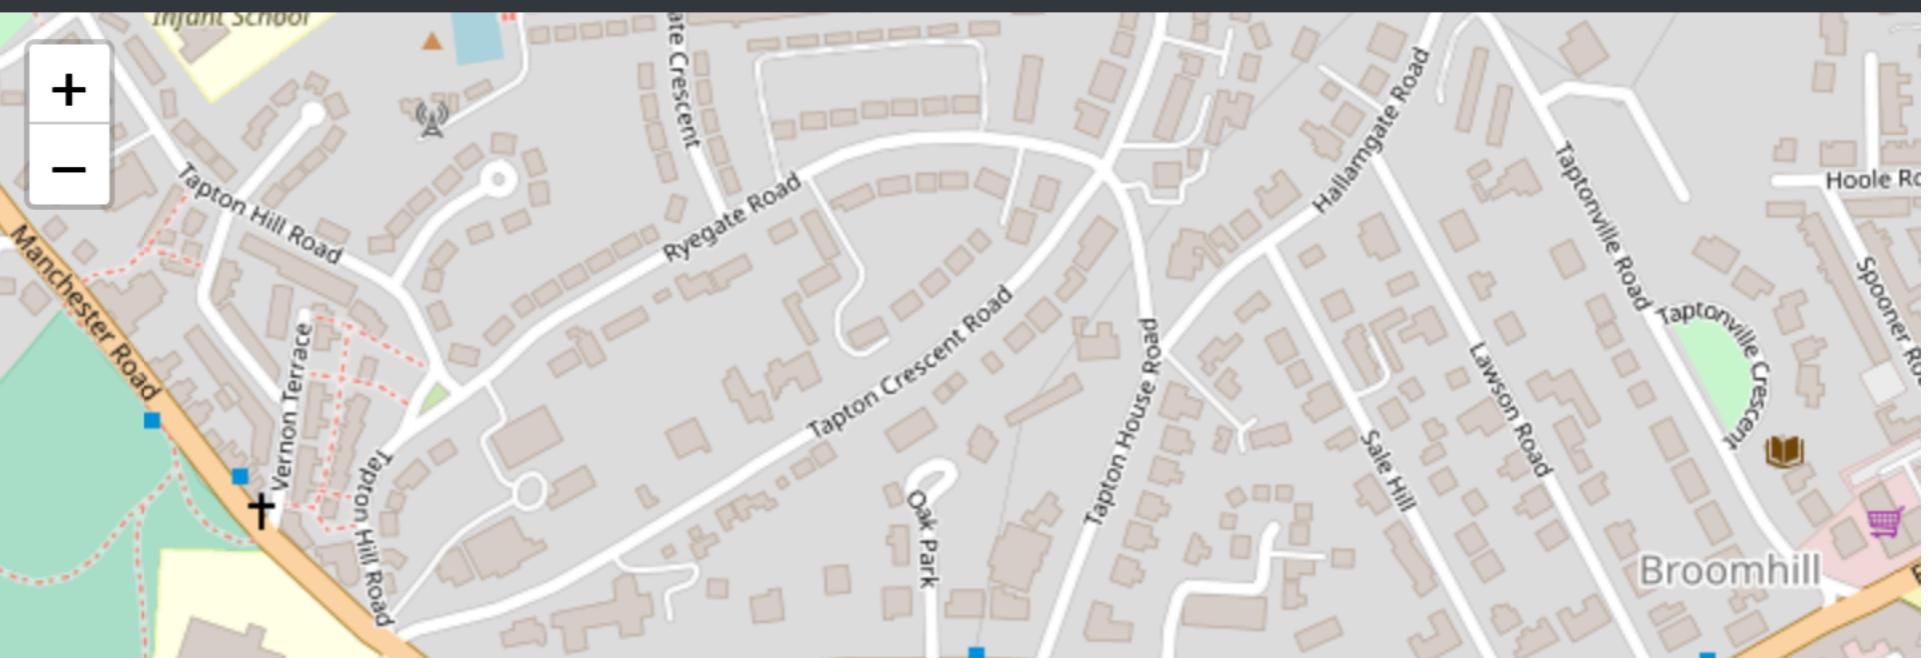
- ▶ kubectl apply -f webapp-service.yaml

```
~ $ kubectl apply -f webapp-service.yaml
```

Fleet Management System PROTOTYPE. Release 0.5

Name	Last seen	Speed mph
------	-----------	-----------

Live vehicle updates will appear here.
Once we've implemented it!



Some useful commands

- ▶ kubectl get all
- ▶ Kubectl get pods
- ▶ Kubectl get po
- ▶ Kubectl get po --show-labels
- ▶ Kubectl get po --show-labels -l release=0

- <https://www.martinfowler.com/bliki/StranglerApplication.html>
- <http://microservices.io/patterns/monolithic.html>
- <http://microservices.io/patterns/microservices.html>
- Richardson, Ch. 2016. Refactoring a Monolith into Microservices
- Vivek Juneja, 2016. From Monoliths To Microservices: An Architectural Strategy

P.PORTO

Polytechnic of
Porto

ESTG - School of
Management and Technology