

P.PORTO

Kubernetes Aula 6

Methods and Techniques for
Software Development

2019/2020

P.PORTO

Kubernetes Lesson 6

Methods and Techniques for
Software Development

2019/2020



kubernetes


Hands on lab

Kubernetes

- Namespaces
- Service Discovery
- Fully Qualified Domain Name

*How do we network
containers together?*

- ▶ How will we do that?



- ▶ Lets imagine I deployed some kind of application container
- ▶ I have it here as a java application




- ▶ I want this application to be able to store its data to a database
- ▶ I already said that is a bad idea to incorporate a database as well as the application in the same container


- ▶ docker containers are designed to be a single service
- ▶ It is possible to expose multiple services in one container but its not something that we want to do
- ▶ We want to have a separate container to deploy the image with the database



- ▶ We need to have these two containers networking with each other



- ▶ In kubernetes if you decide to deploy these two containers into a single Pod its very easy because containers can see each other using localhost

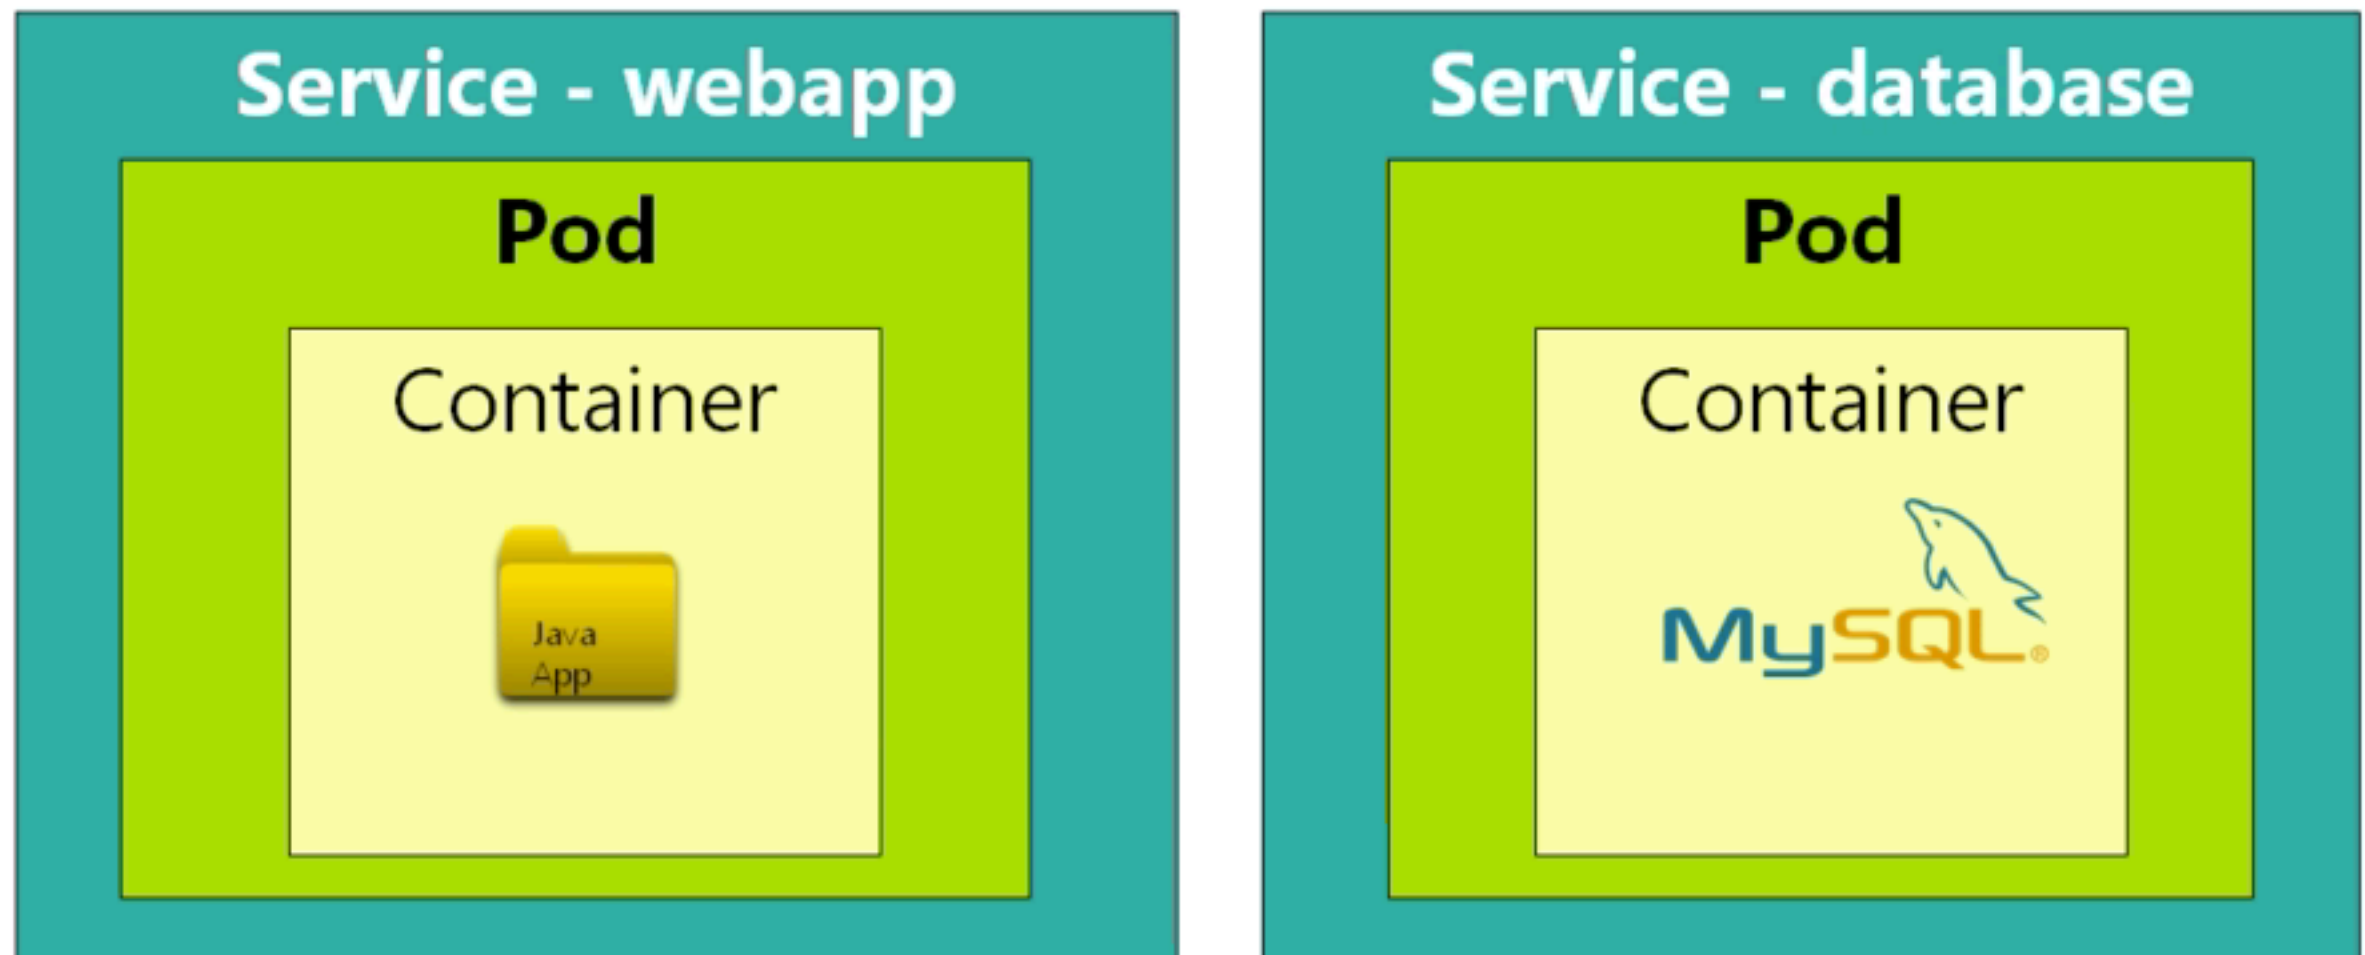


- ▶ If I were to write any code in this java application I could just do a lookup of the address **localhost:3306** and I would see the data in mysql container database

- ▶ It is acceptable to have multiple containers in a single pod in kubernetes
- ▶ But is not a very good idea.... it is not recommended
- ▶ Having an application and database container running in the same pod would make the pod more complicated to manage

- ▶ If the pods fails we need to find out if it failed because of the application container or because of the database container - it makes things more complicated

- ▶ If we want to deploy the application as well as the database in not only different containers but also in separate pods



- ▶ The problem is, we are not going to know the IP address of that service because it is allocated dynamically by kubernetes
- ▶ The next time we run the kubernetes cluster there will be different IP addresses allocated to the services

What is the solution?

*Kubernetes maintains its
own private DNS service*

Service - webapp

Pod

Container



Service - database

Pod

Container



Service – kube-dns

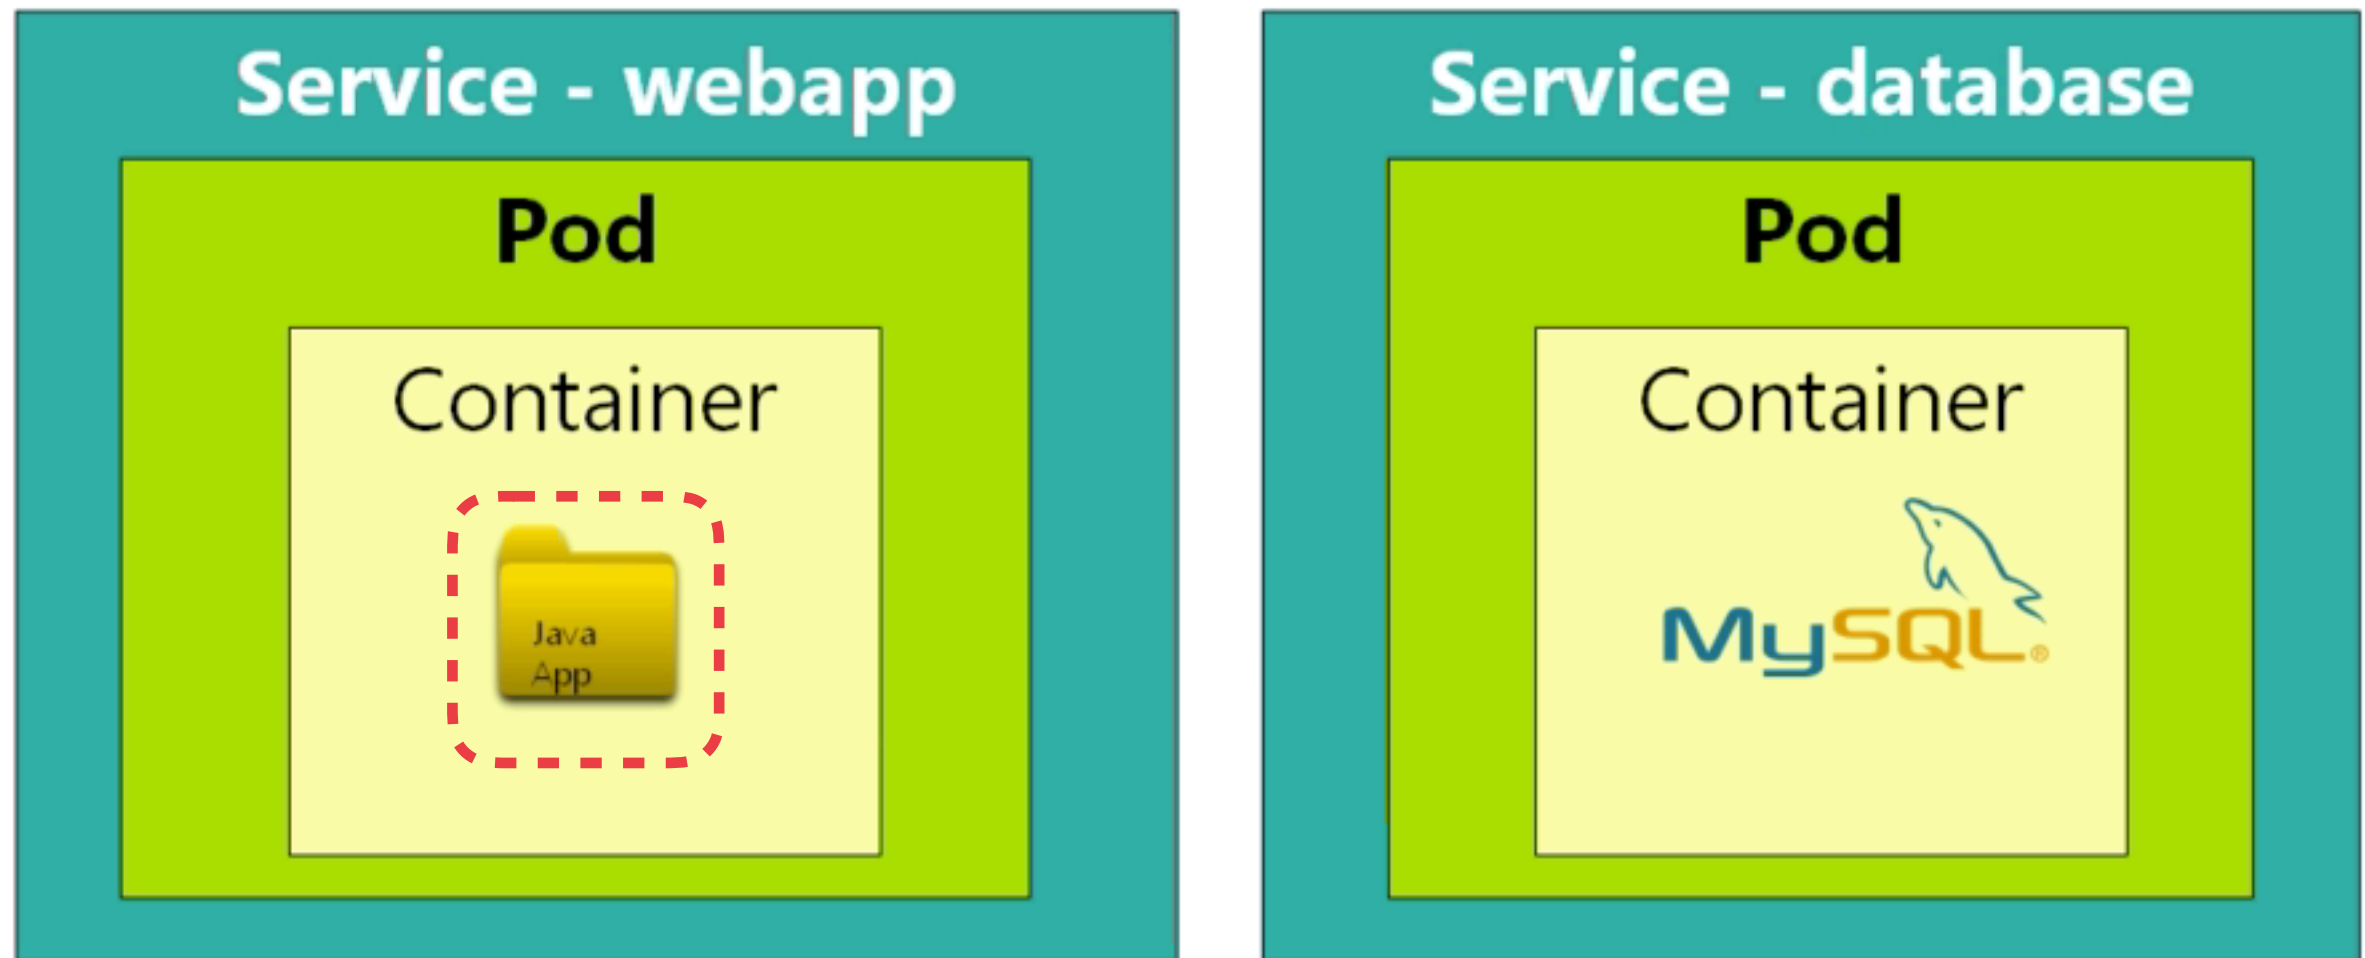
Pod

DNS Container

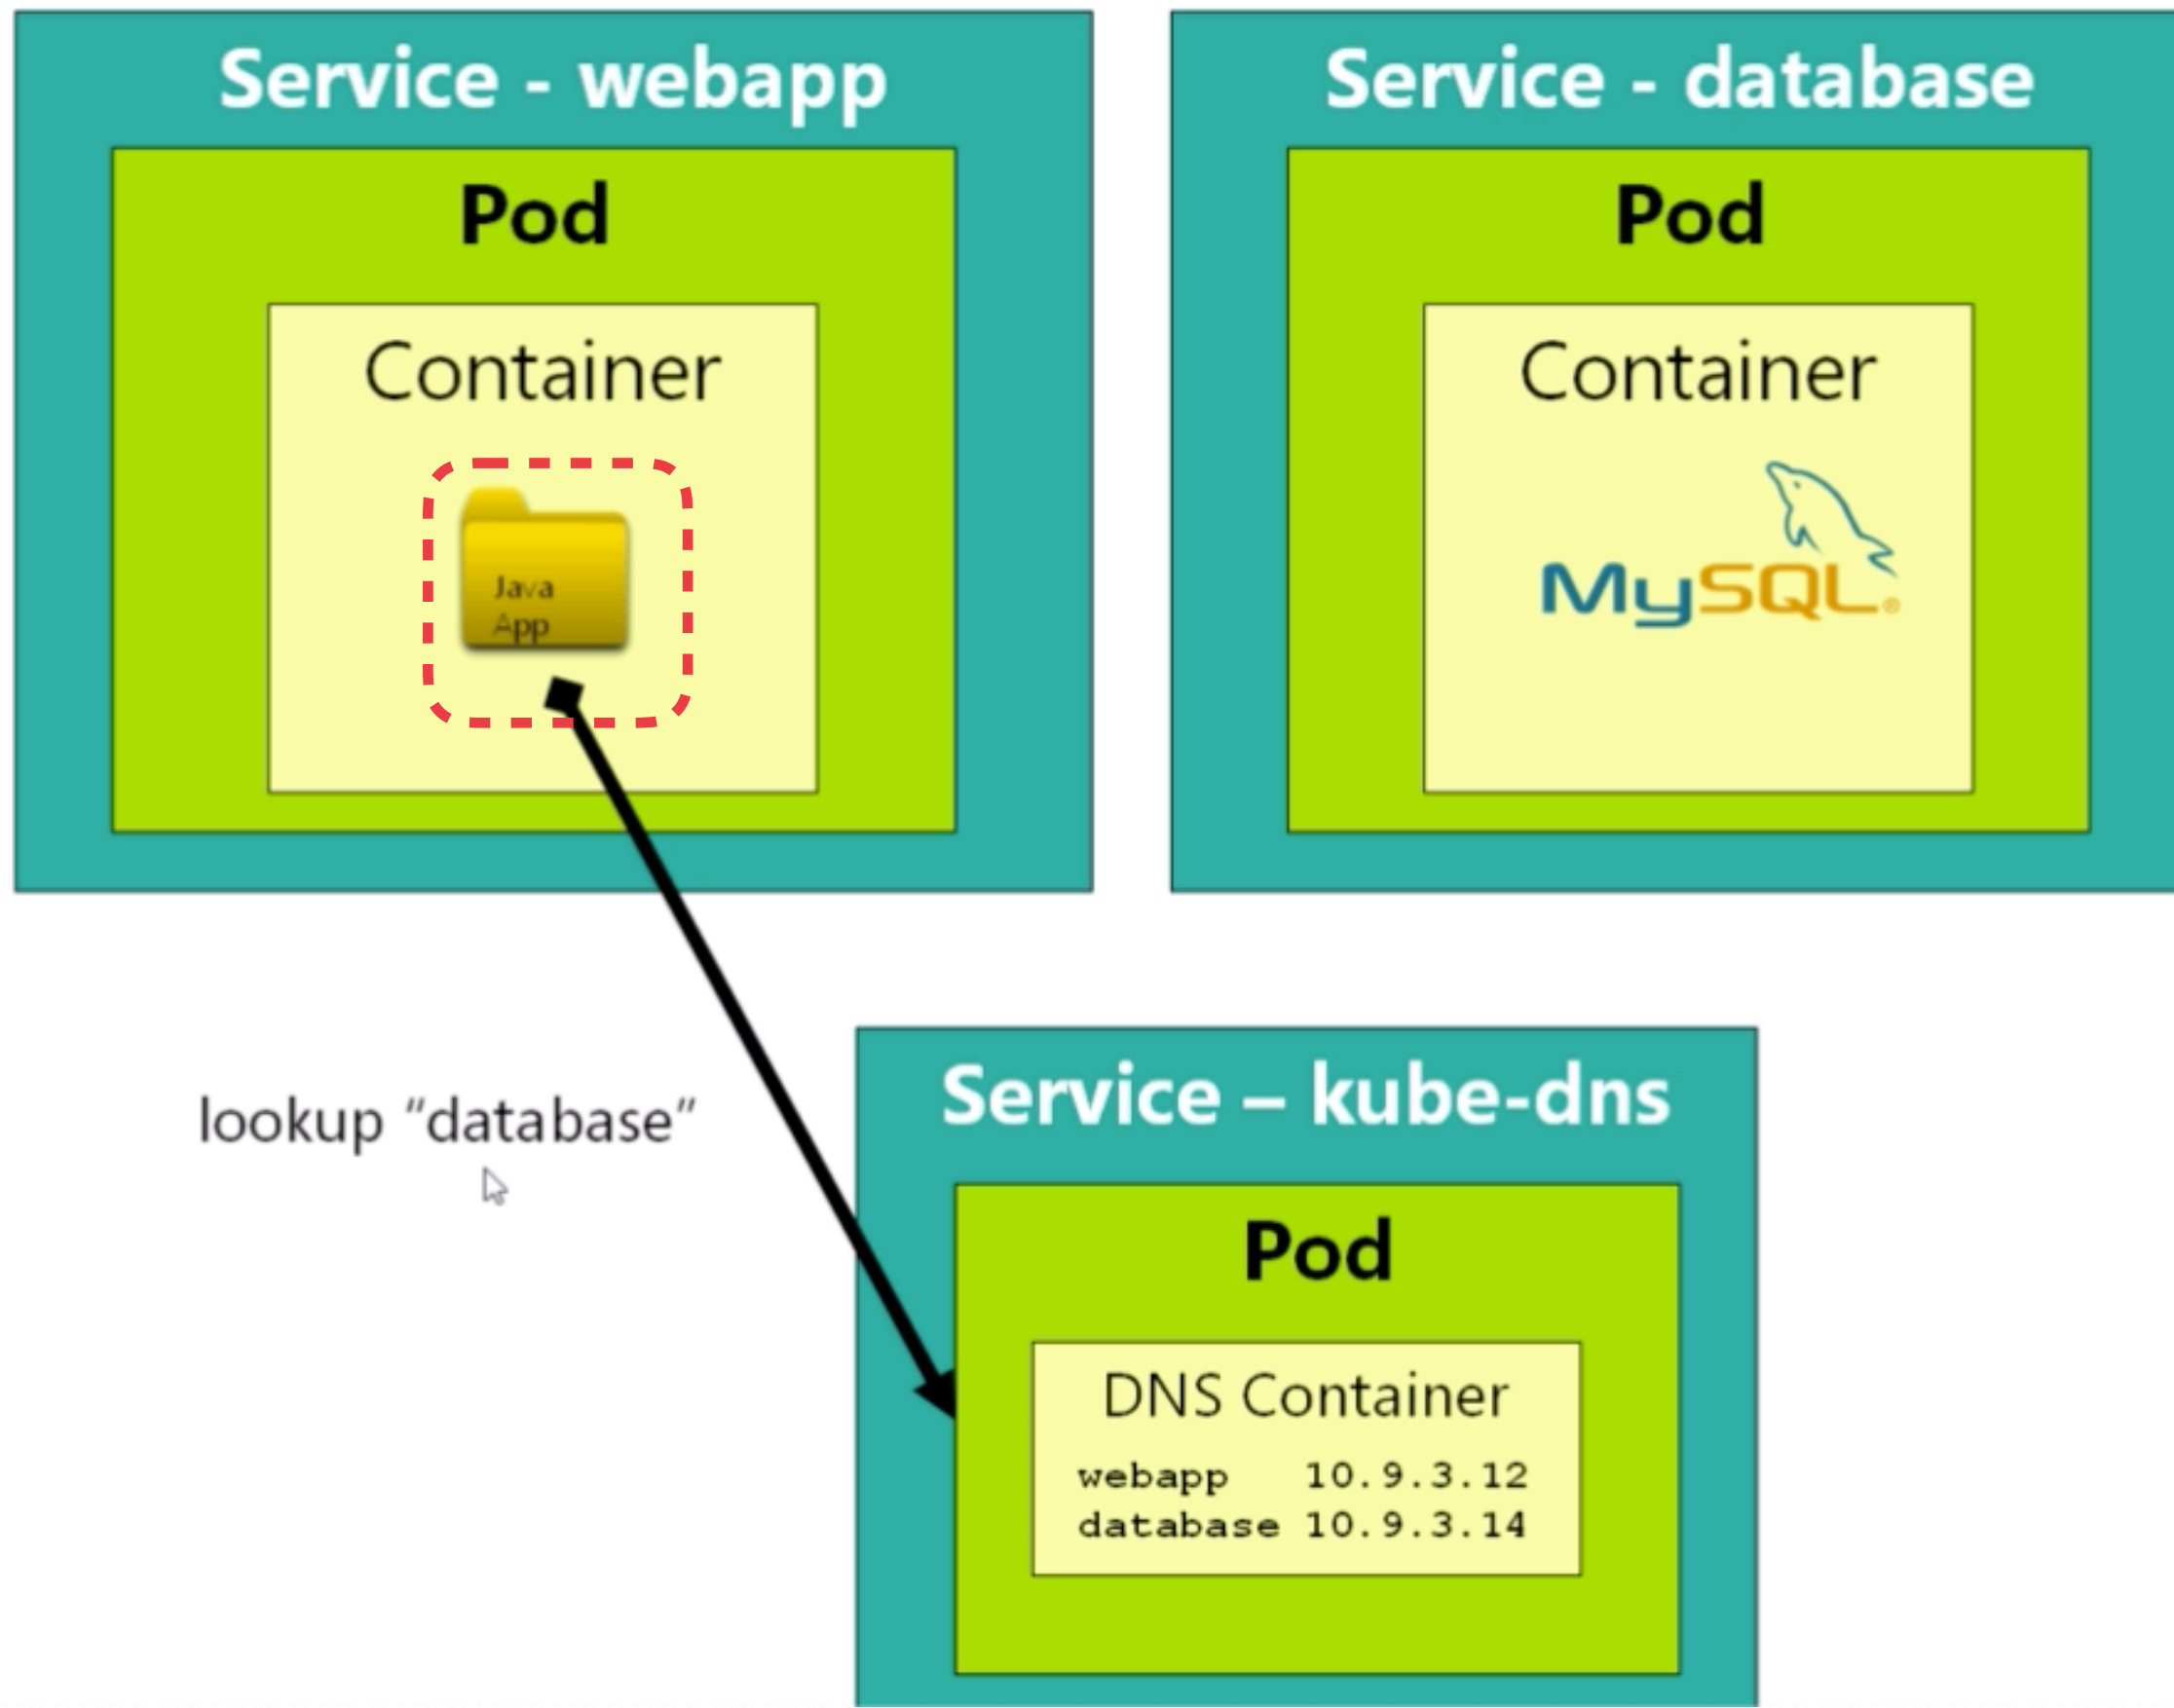
webapp	10.9.3.12
database	10.9.3.14

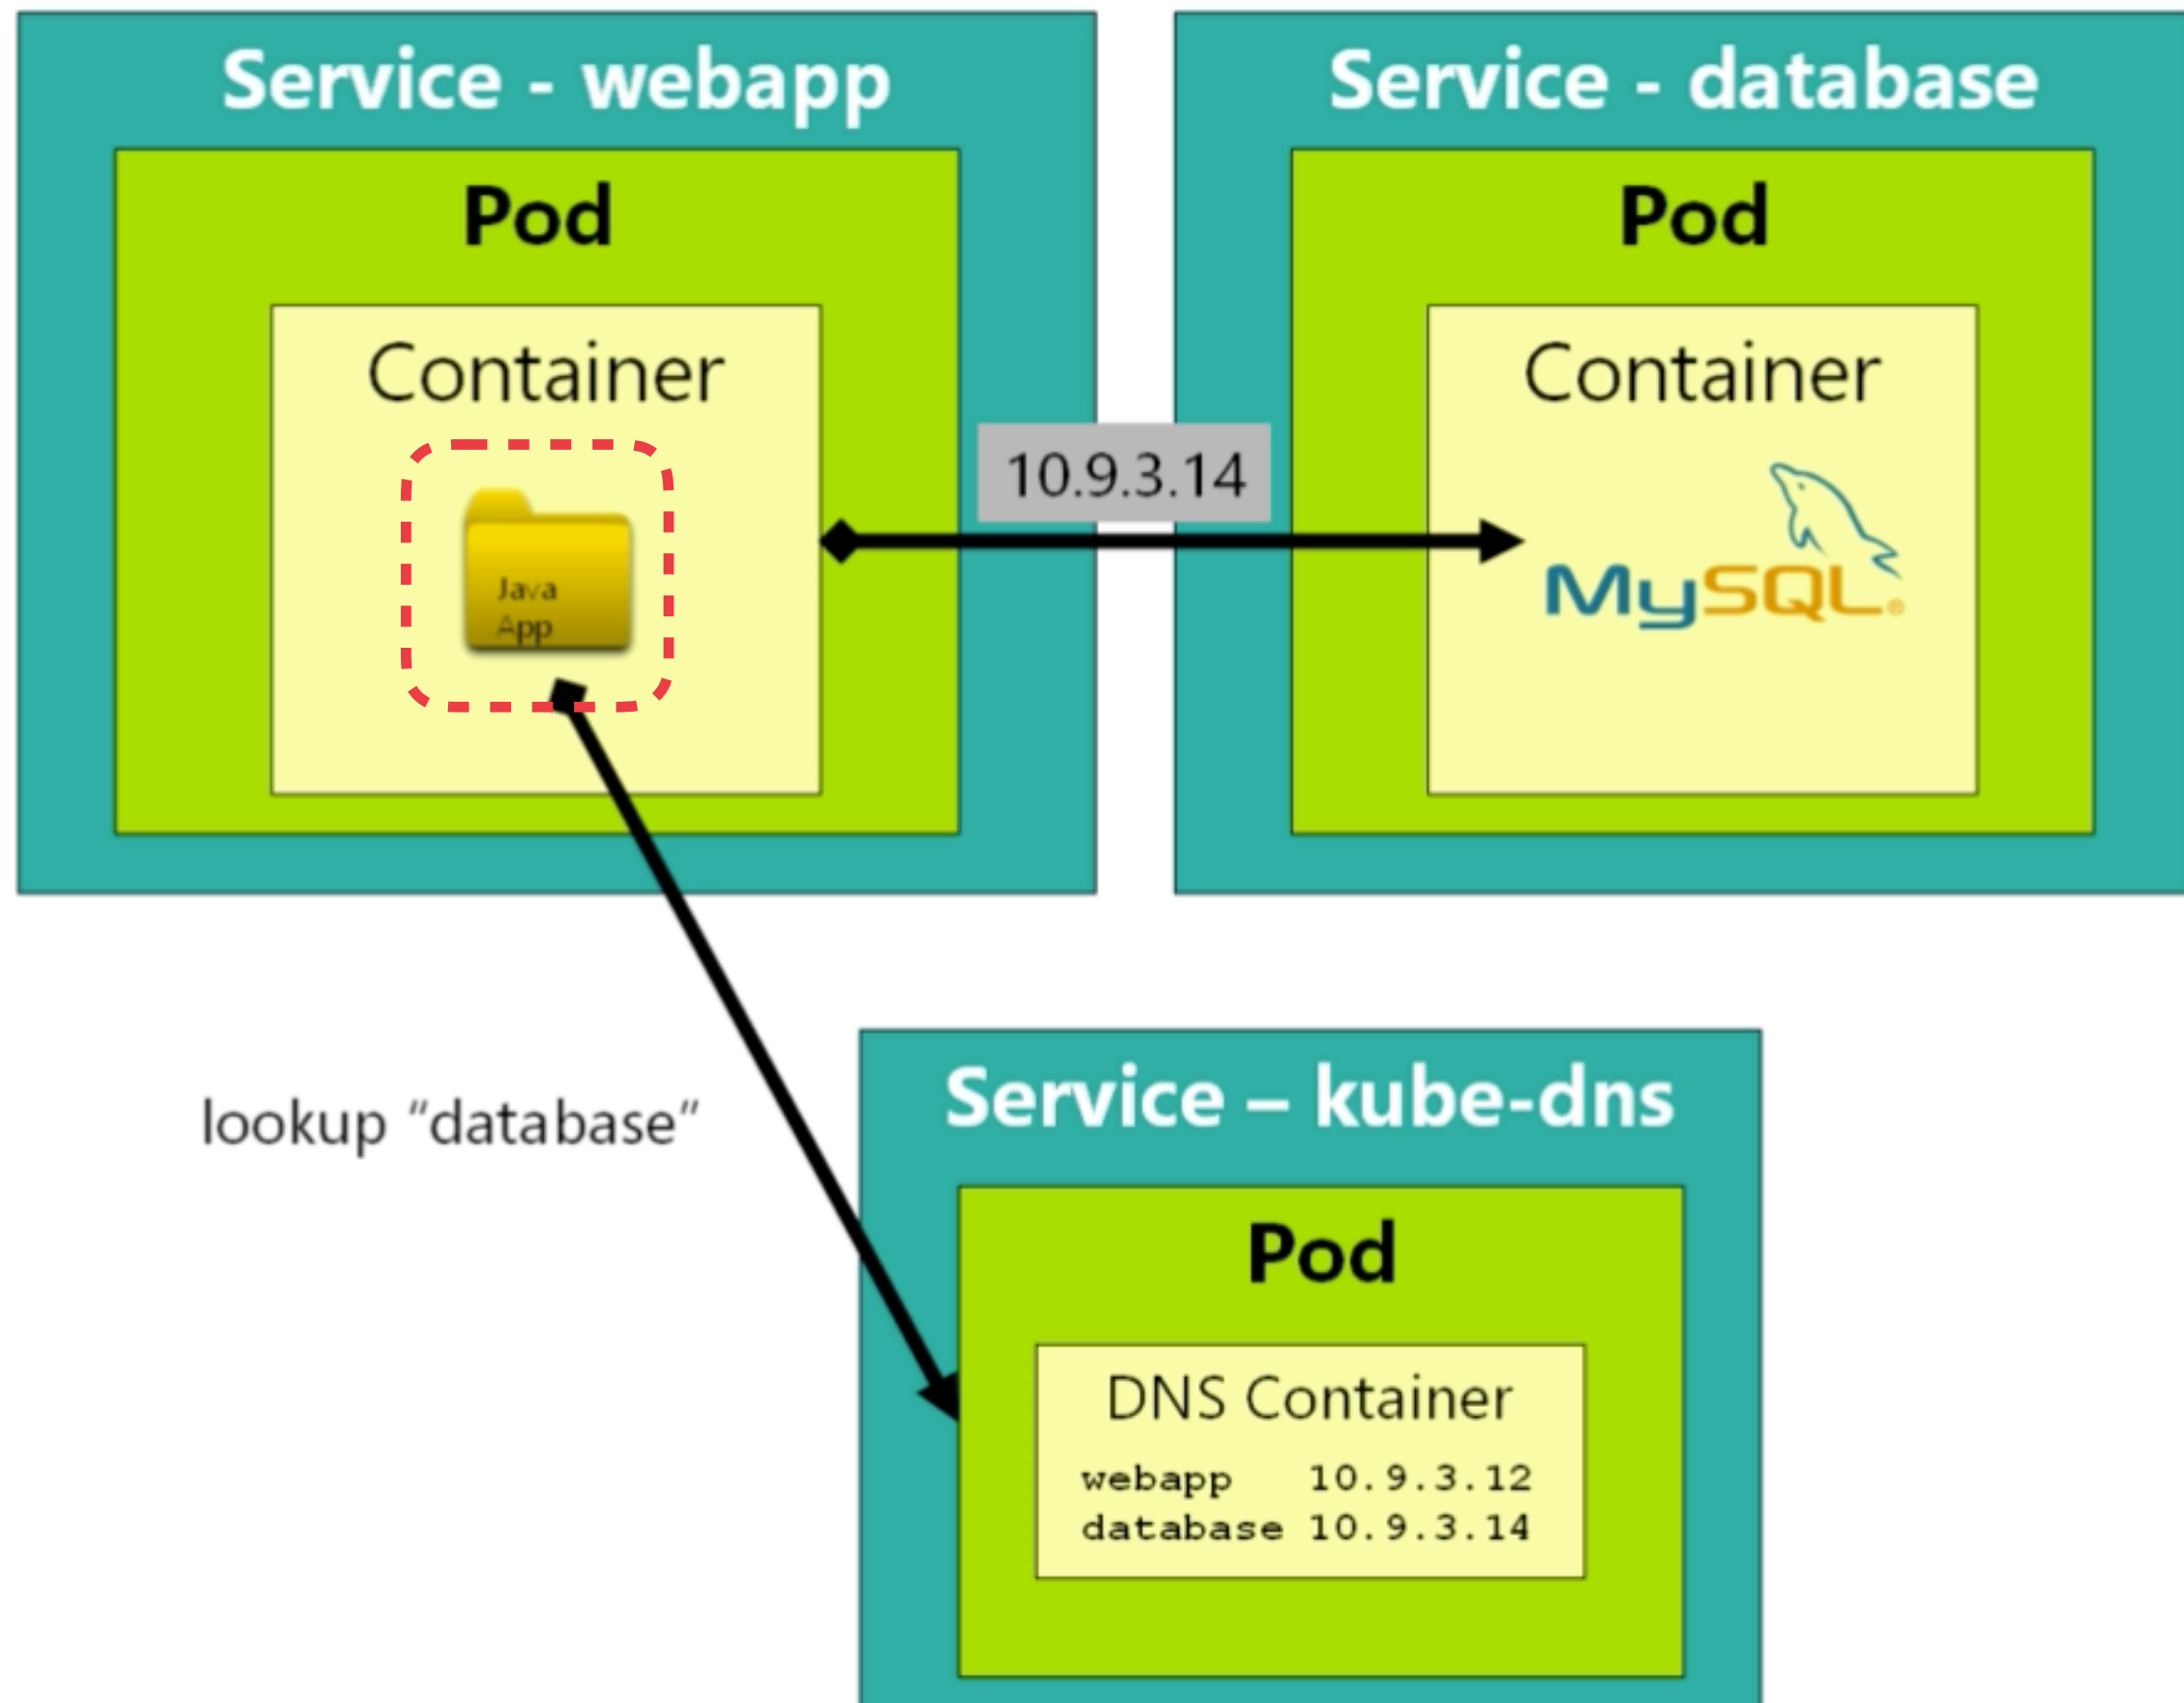
- ▶ The service is called kube-dns
- ▶ It's not needed to configure this service, it has been running automatically in background
- ▶ This service is wrapping a pod named kube-dns

- ▶ The DNS services consists in a key-value database
- ▶ The keys are labels - the names of the Kubernetes services
- ▶ The values are the IP addresses of those services
- ▶ Kubernetes has the full responsibility to maintain this service



- ▶ In our code if we need to refer to the database we could just use the string "database"





*Why didn't we see this
kube-dns before?*

► kubectl get all

```
$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/queue	1/1	Running	0	35m
pod/webapp-ccb5c74c9-mdzrb	1/1	Running	0	35m
pod/webapp-ccb5c74c9-wjsnn	1/1	Running	0	35m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/fleetman-queue	NodePort	10.104.209.185	<none>	8161:30010/TCP	35m
service/fleetman-webapp	NodePort	10.101.253.163	<none>	80:30080/TCP	35m
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	36m

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/webapp	2	2	2	2	35m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/webapp-ccb5c74c9	2	2	2	35m

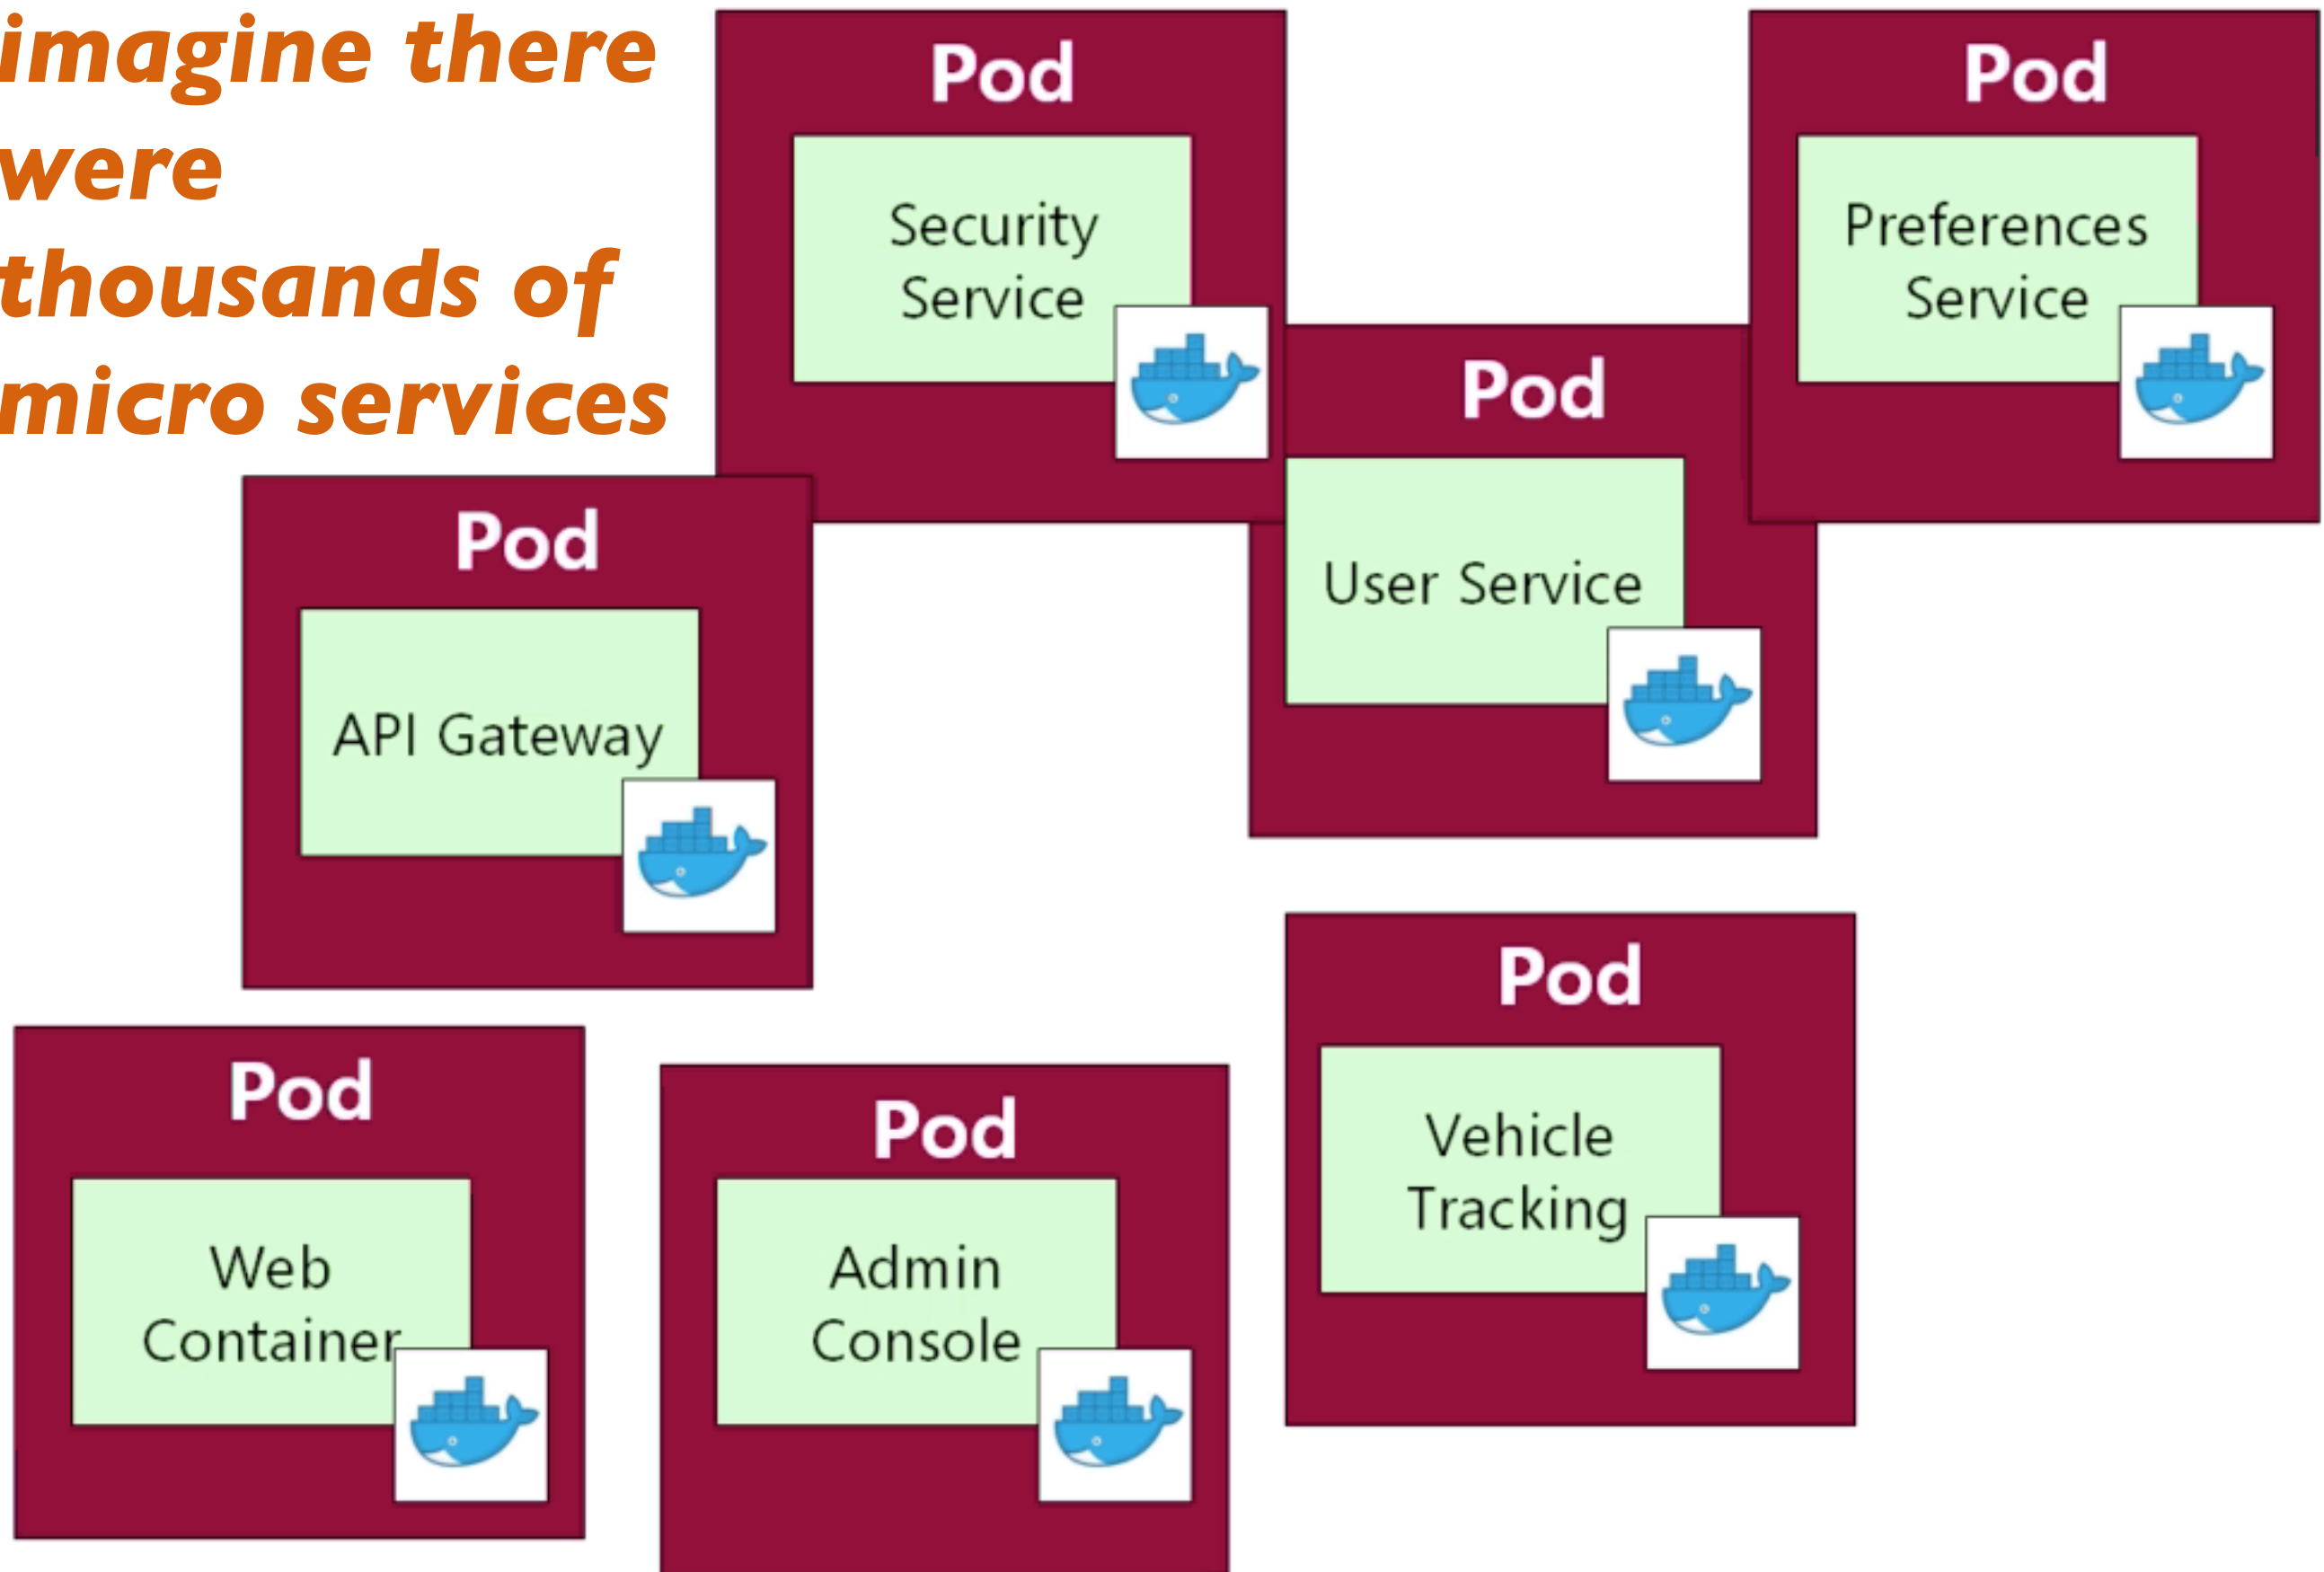
► There's not any reference to DNS...

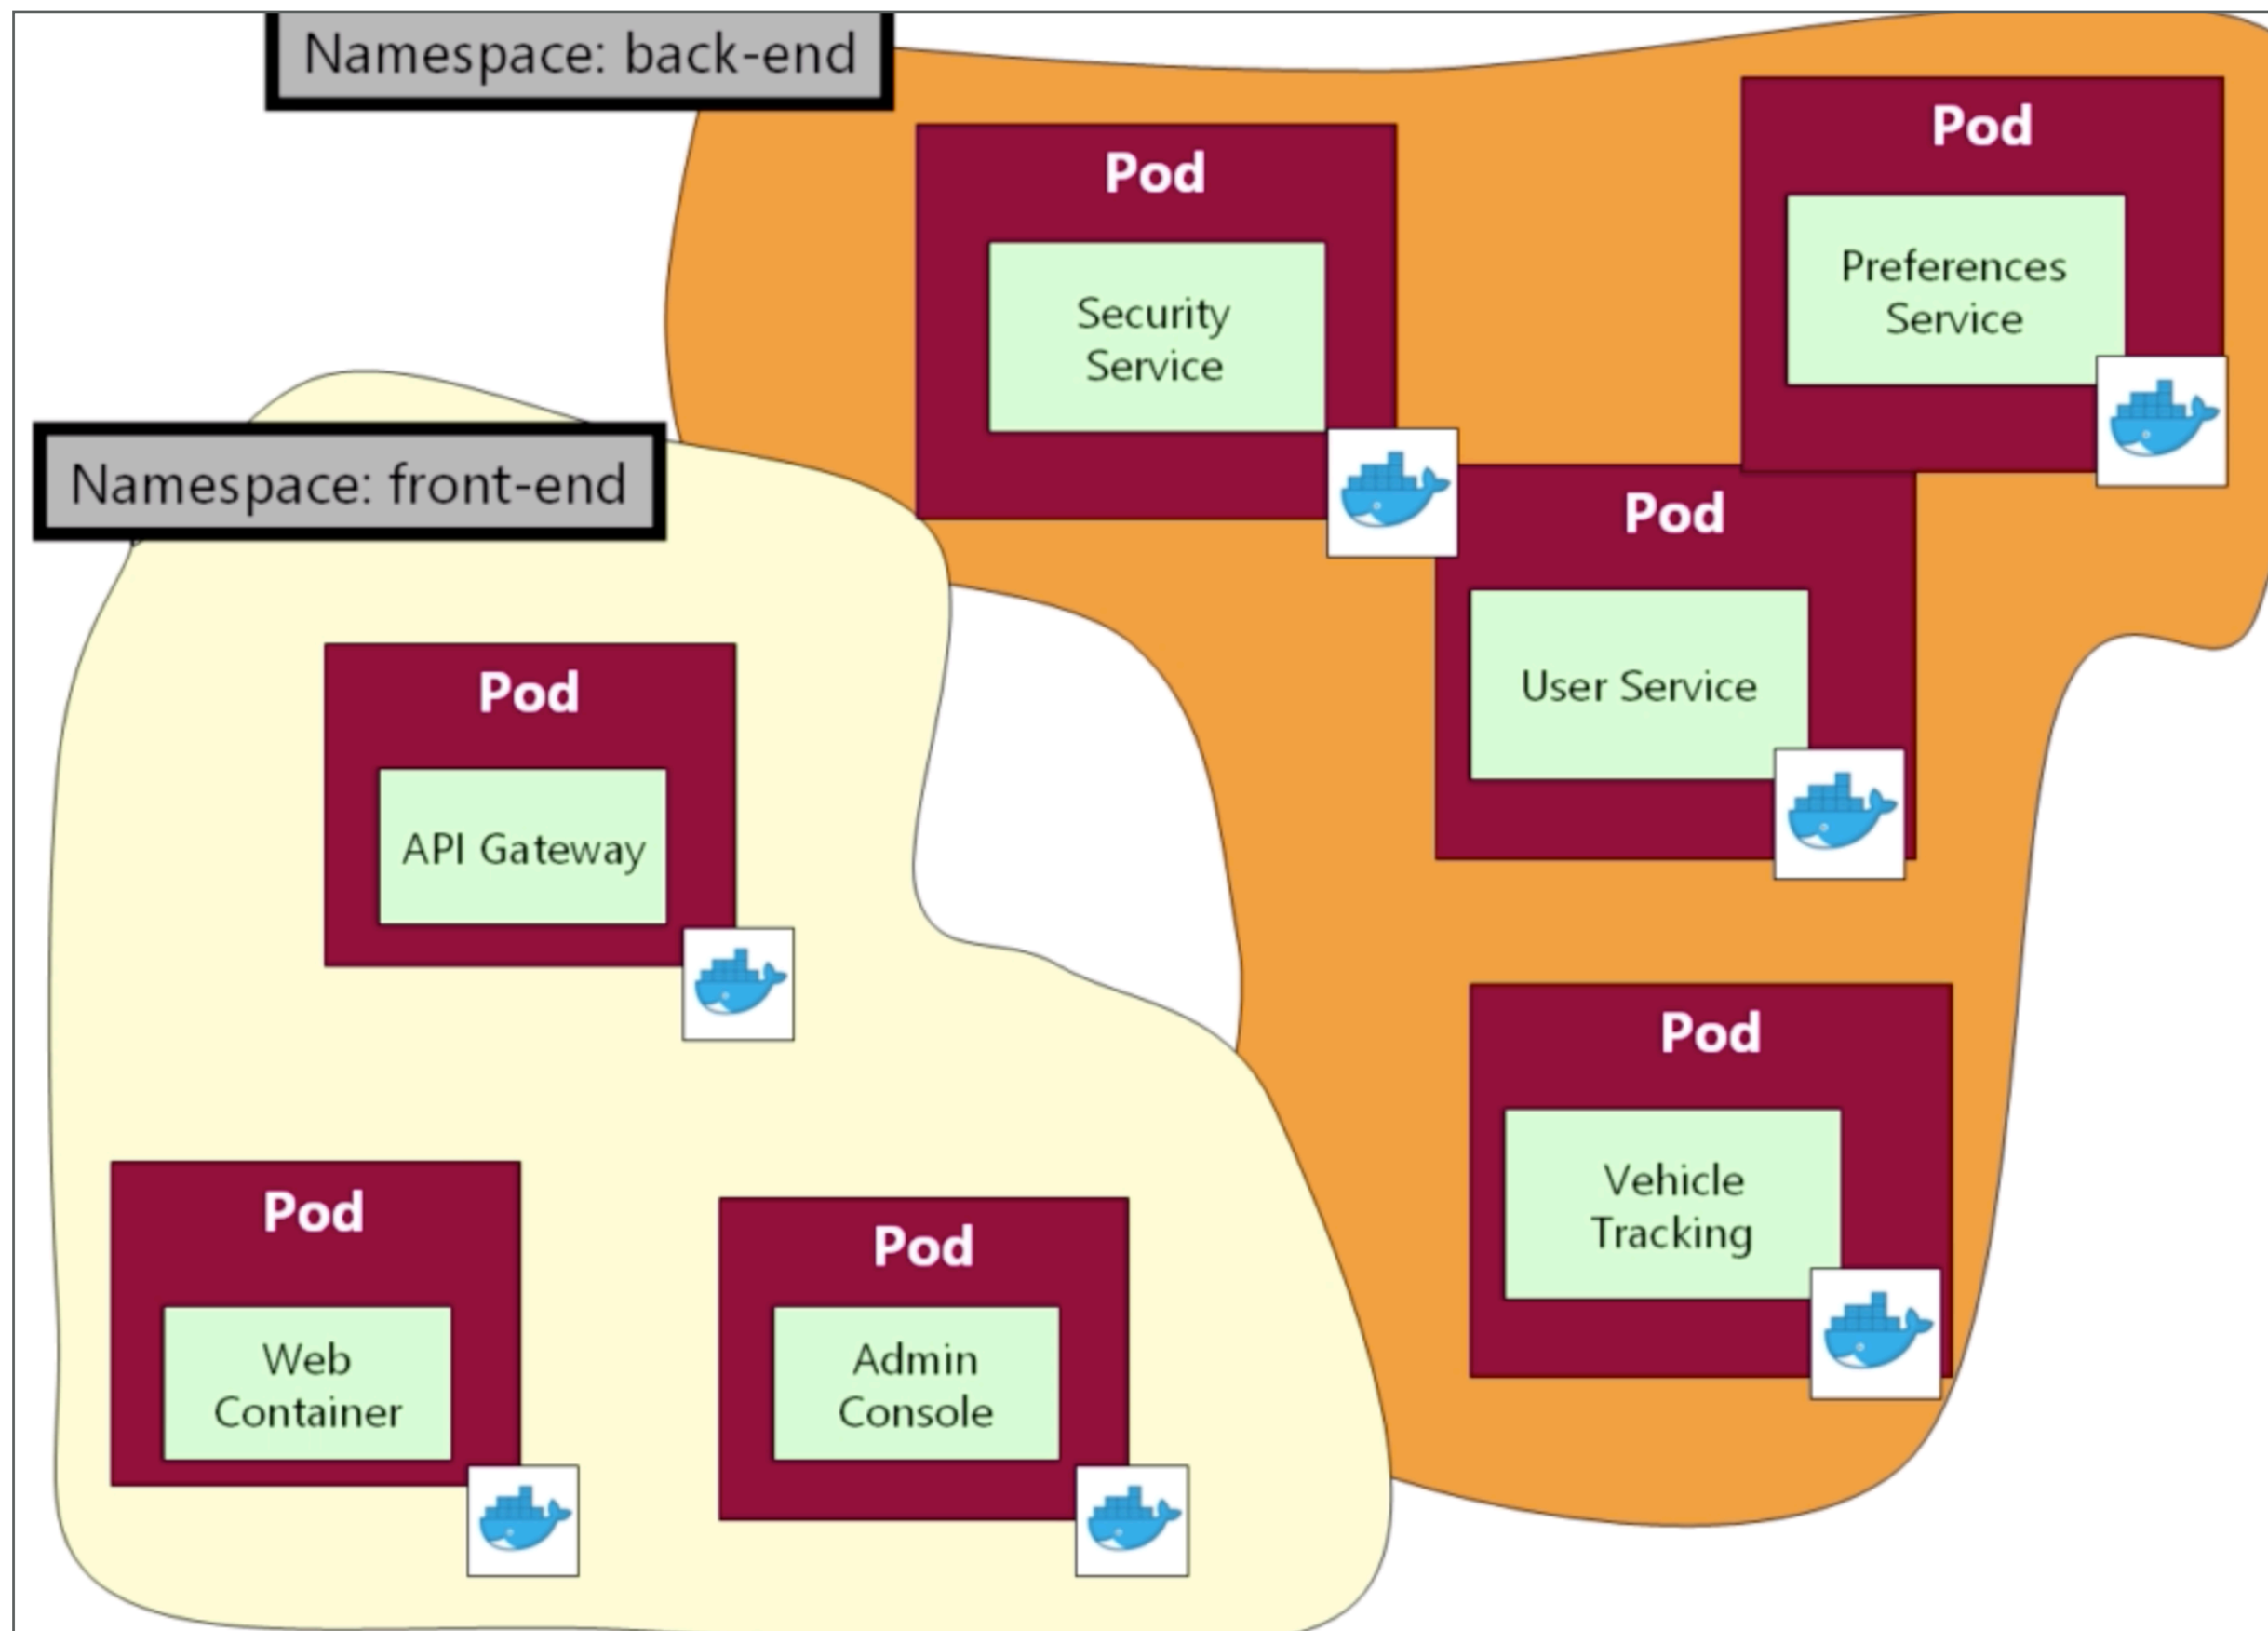
*There is a new concept we
didn't talk about...*

namespace

- ▶ Namespaces are a way of partitioning resources in kubernetes into separate areas

*imagine there
were
thousands of
micro services*






- ▶ We just created pods and services
- ▶ What happens when you don't specify a namespace?
 - ▶ That resource is put in the default namespace
- ▶ In the same way, when we don't specify the namespace we only see the resources on the default namespace
- ▶ Kubernetes comes with additional namespaces that we never seen

Lets verify them!

- ▶ `kubectl get namespaces`
- ▶ `kubectl get ns`

```
$ kubectl get ns
NAME           STATUS    AGE
default        Active    6d
kube-public    Active    6d
kube-system    Active    6d
```



- ▶ These are all the namespaces defined in the system
- ▶ Let's see the other namespaces

- ▶ `kubectl get po` *default ns*

```
$ kubectl get po
```

NAME	READY	STATUS	RESTARTS	AGE
queue	1/1	Running	0	1h
webapp-ccb5c74c9-mdzrb	1/1	Running	0	1h
webapp-ccb5c74c9-wjsnn	1/1	Running	0	1h

- ▶ `kubectl get po -n kube-system`

```
$ kubectl get po -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
etcd-minikube	1/1	Running	0	2h
kube-addon-manager-minikube	1/1	Running	1	6d
kube-apiserver-minikube	1/1	Running	0	2h
kube-controller-manager-minikube	1/1	Running	0	2h
kube-dns-86f4d74b45-f5kxg	3/3	Running	4	6d
kube-proxy-jbwq5	1/1	Running	0	2h
kube-scheduler-minikube	1/1	Running	0	2h
kubernetes-dashboard-5498ccf677-z1k8p	1/1	Running	3	6d
storage-provisioner	1/1	Running	3	6d

► `kubectl get all -n kube-system`

```
$ kubectl get all -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
pod/etcd-minikube	1/1	Running	0	2h
pod/kube-addon-manager-minikube	1/1	Running	1	6d
pod/kube-apiserver-minikube	1/1	Running	0	2h
pod/kube-controller-manager-minikube	1/1	Running	0	2h
pod/kube-dns-86f4d74b45-f5kxg	3/3	Running	4	6d
pod/kube-proxy-jbwq5	1/1	Running	0	2h
pod/kube-scheduler-minikube	1/1	Running	0	2h
pod/kubernetes-dashboard-5498ccf677-zlk8p	1/1	Running	3	6d
pod/storage-provisioner	1/1	Running	3	6d

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP	6d
service/kubernetes-dashboard	NodePort	10.105.32.177	<none>	80:30000/TCP	6d

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
daemonset.apps/kube-proxy	1	1	1	1	1	<none>	6d

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/kube-dns	1	1	1	1	6d
deployment.apps/kubernetes-dashboard	1	1	1	1	6d

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/kube-dns-86f4d74b45	1	1	1	6d
replicaset.apps/kubernetes-dashboard-5498ccf677	1	1	1	6d

Exercise

Describe the service kube-dns

► `kubectl describe svc kube-dns`

```
$ kubectl describe svc kube-dns  
Error from server (NotFound): services "kube-dns" not found
```

??

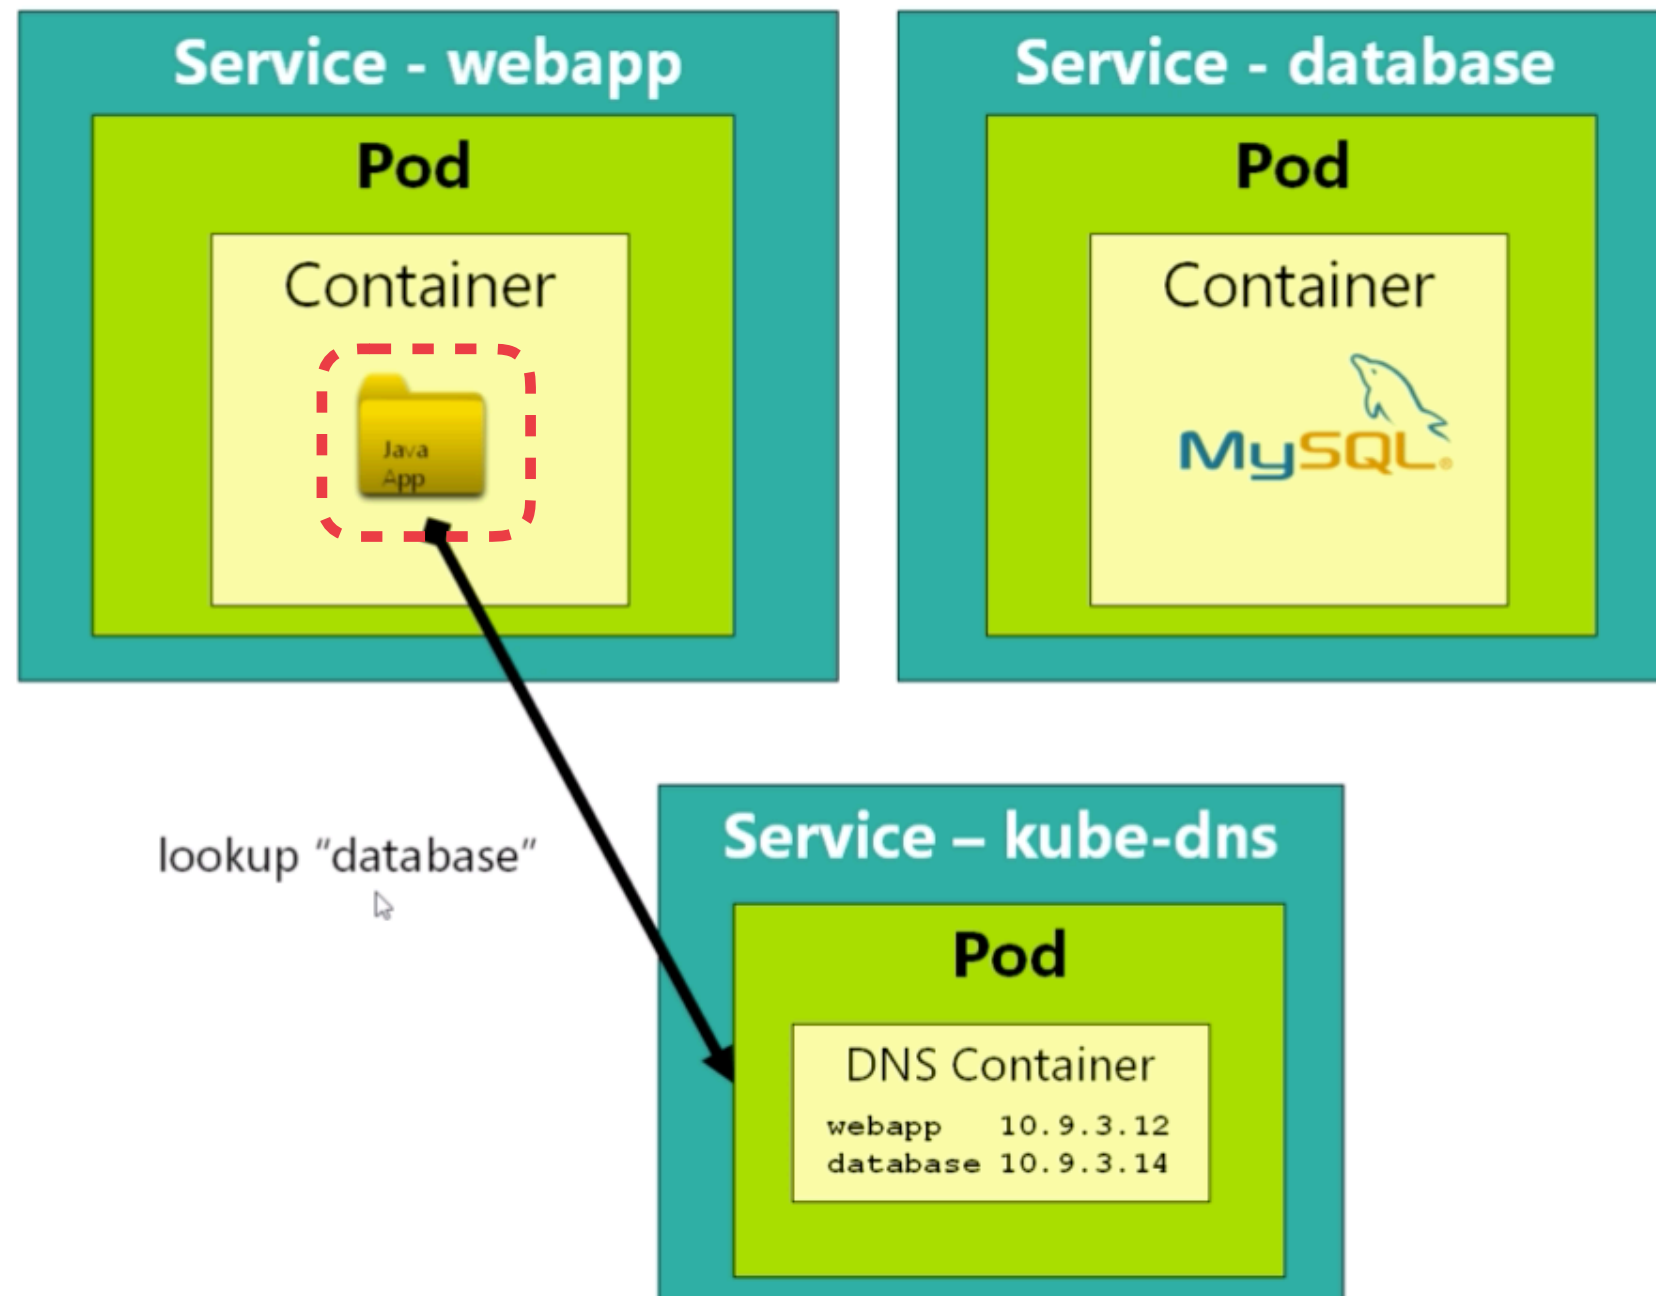
- ▶ `kubectl describe svc kube-dns -n kube-system`

```
$ kubectl describe svc kube-dns -n kube-system
Name: kube-dns
Namespace: kube-system
Labels: k8s-app=kube-dns
        kubernetes.io/cluster-service=true
        kubernetes.io/name=KubeDNS
Annotations: <none>
Selector: k8s-app=kube-dns
Type: ClusterIP
IP: 10.96.0.10
Port: dns 53/UDP
TargetPort: 53/UDP
Endpoints: 172.17.0.2:53
Port: dns-tcp 53/TCP
TargetPort: 53/TCP
Endpoints: 172.17.0.2:53
Session Affinity: None
Events: <none>
```

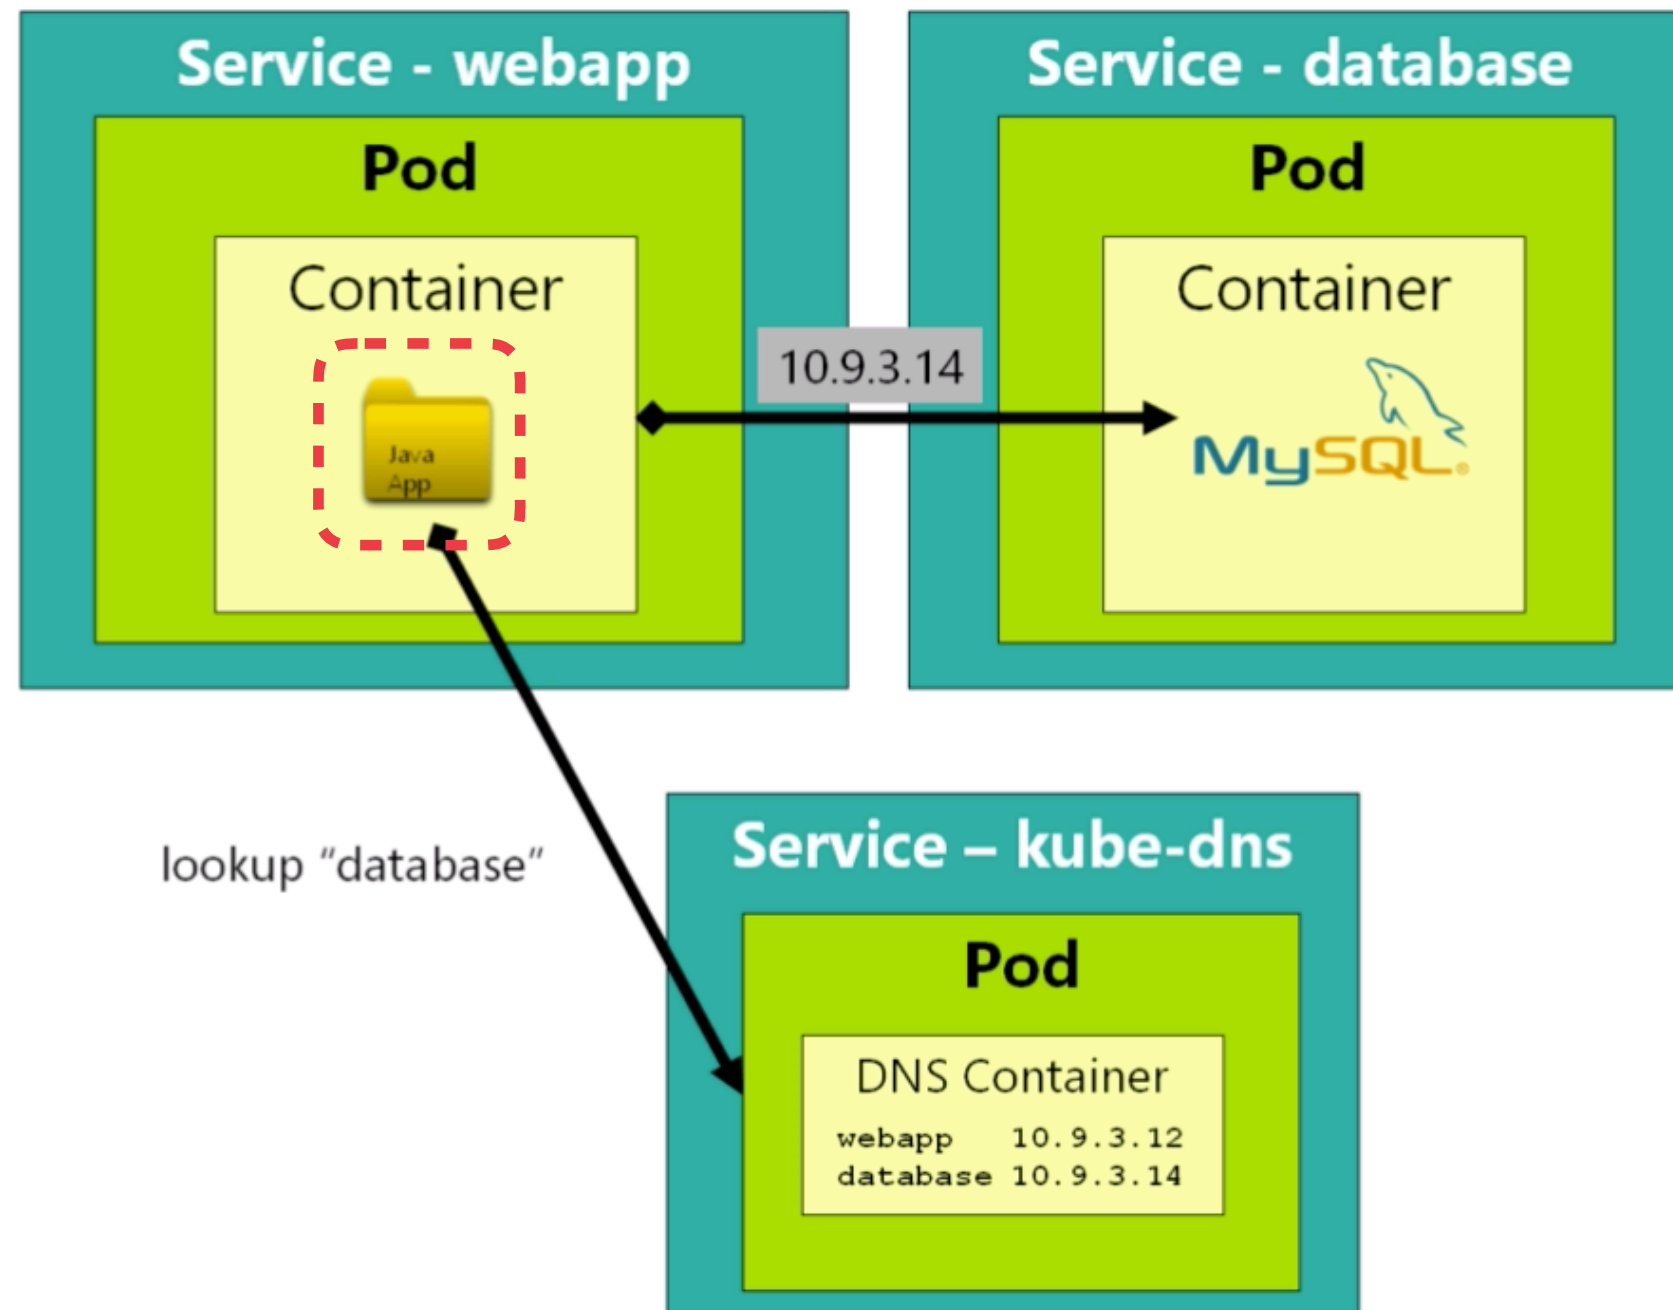
- ▶ We already understand now why by default we can't see the kube-dns service

Back to the work we were doing

- ▶ We want from our webapp application to reference a database container in a different pod



- Now we know that its as easy as doing a lookup for "database" in the kube-dns service



- ▶ That will give us the IP address of the target service

- ▶ We already have the webapp container, but we don't have the database container in our
- ▶ Let's create a new MySQL container in its own Pod and a Service

► networking-tests.yaml


```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: mysql
5    labels:
6      app: mysql
7  spec:
8    containers:
9      - name: mysql
10        image: mysql:5
11        env:
12          # Use secret in real life
13          - name: MYSQL_ROOT_PASSWORD
14            value: password
15          - name: MYSQL_DATABASE
16            value: fleetman
17  ---
```

←
Pod

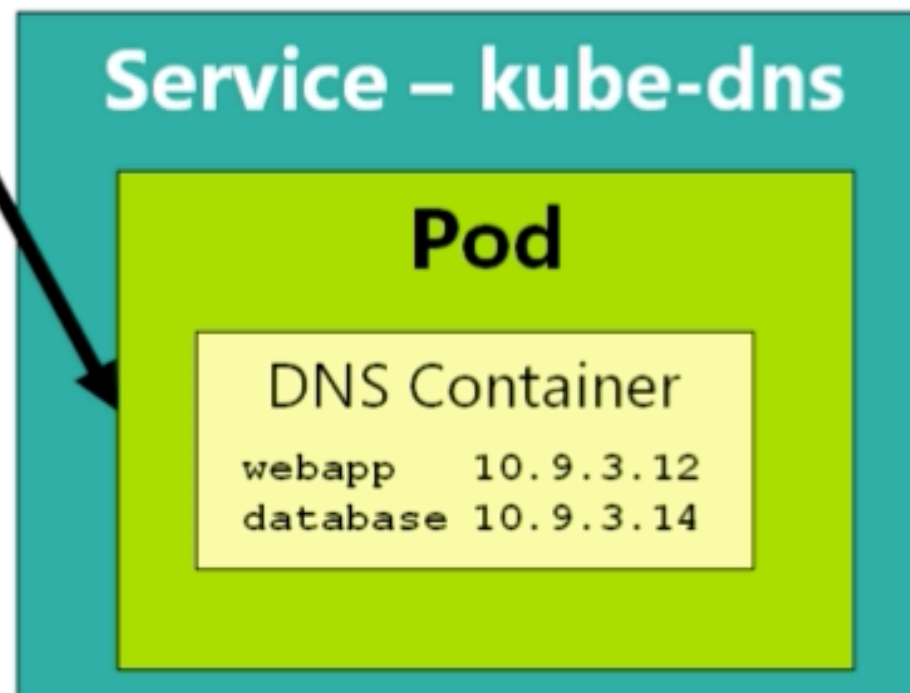

```
17 ---
18 kind: Service
19 apiVersion: v1
20 metadata:
21 | name: database
22 spec:
23 | selector:
24 |   app: mysql
25 | ports:
26 | - port: 3306
27 |   type: ClusterIP
28
```



Service



lookup "database"



Nobody outside this cluster will be able to access this database service

► ls

```
$ ls  
networking-tests.yaml  pods.yaml  services.yaml
```

► kubectl apply -f networking-tests.yaml

```
$ kubectl apply -f networking-tests.yaml  
pod "mysql" created  
service "database" created
```

► kubectl get all

```
$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/mysql	1/1	Running	0	11s
pod/queue	1/1	Running	0	1h
pod/webapp-ccb5c74c9-mdzrb	1/1	Running	0	1h
pod/webapp-ccb5c74c9-wjsnn	1/1	Running	0	1h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/database	ClusterIP	10.109.75.215	<none>	3306/TCP	12s
service/fleetman-queue	NodePort	10.104.209.185	<none>	8161:30010/TCP	1h
service/fleetman-webapp	NodePort	10.101.253.163	<none>	80:30080/TCP	1h
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	1h

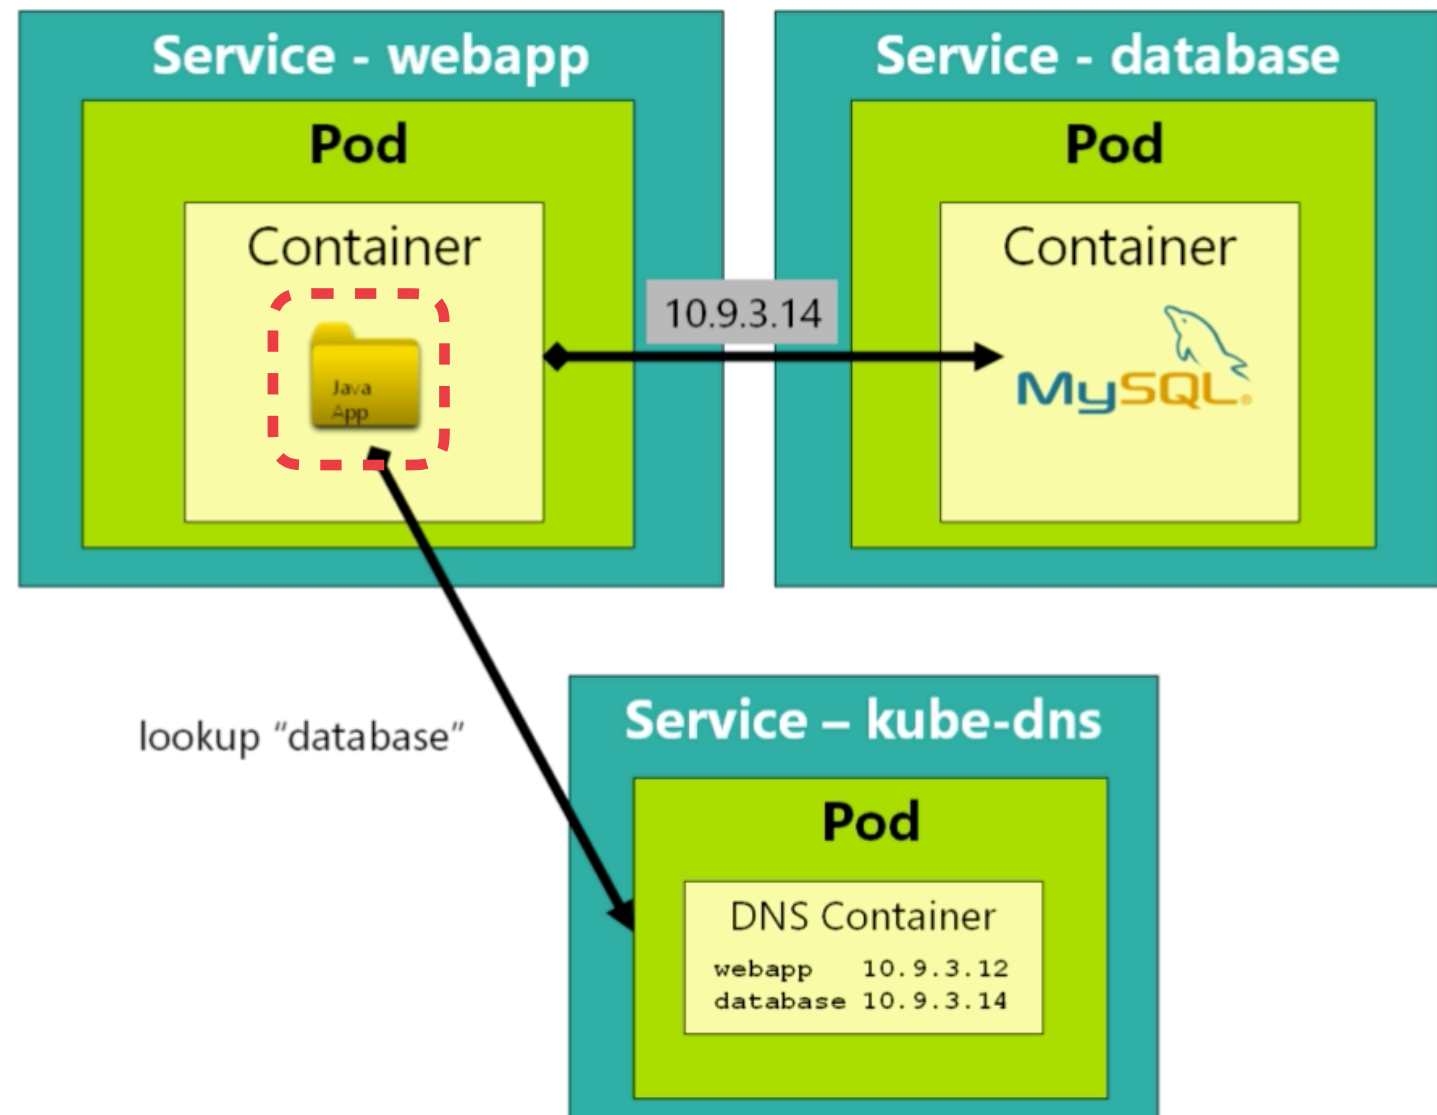
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/webapp	2	2	2	2	1h

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/webapp-ccb5c74c9	2	2	2	1h

- ▶ We are not going to write any code inside this webapp container that calls the database
- ▶ we are just going to do a very basic demonstration - we are going to simulate that happening

- ▶ `kubectl exec -it webapp-ccb5c74c9-mdzrb sh`
- ▶ `ls`

```
ls
bin
dev
etc
home
lib
media
mnt
proc
root
run
sbin
srv
sys
tmp
usr
var
```



- ▶ How this networking mechanism works?
- ▶ I want to be able to access the remote database container
- ▶ How does our webapp container know how to find the dns service?

- ▶ The answer is in a piece of automatic configuration that Kubernetes is doing behind the scenes for us in all of our containers
- ▶ Kubernetes will automatically do some management of the container and automatically configures the dns system

We can verify that

- ▶ Remember you are in the webapp container
- ▶ `cat /etc/resolv.conf`

```
/ # cat /etc/resolv.conf
nameserver 10.96.0.10
search default.svc.cluster.local svc.cluster.local cluster.local
options ndots:5
/ #
```

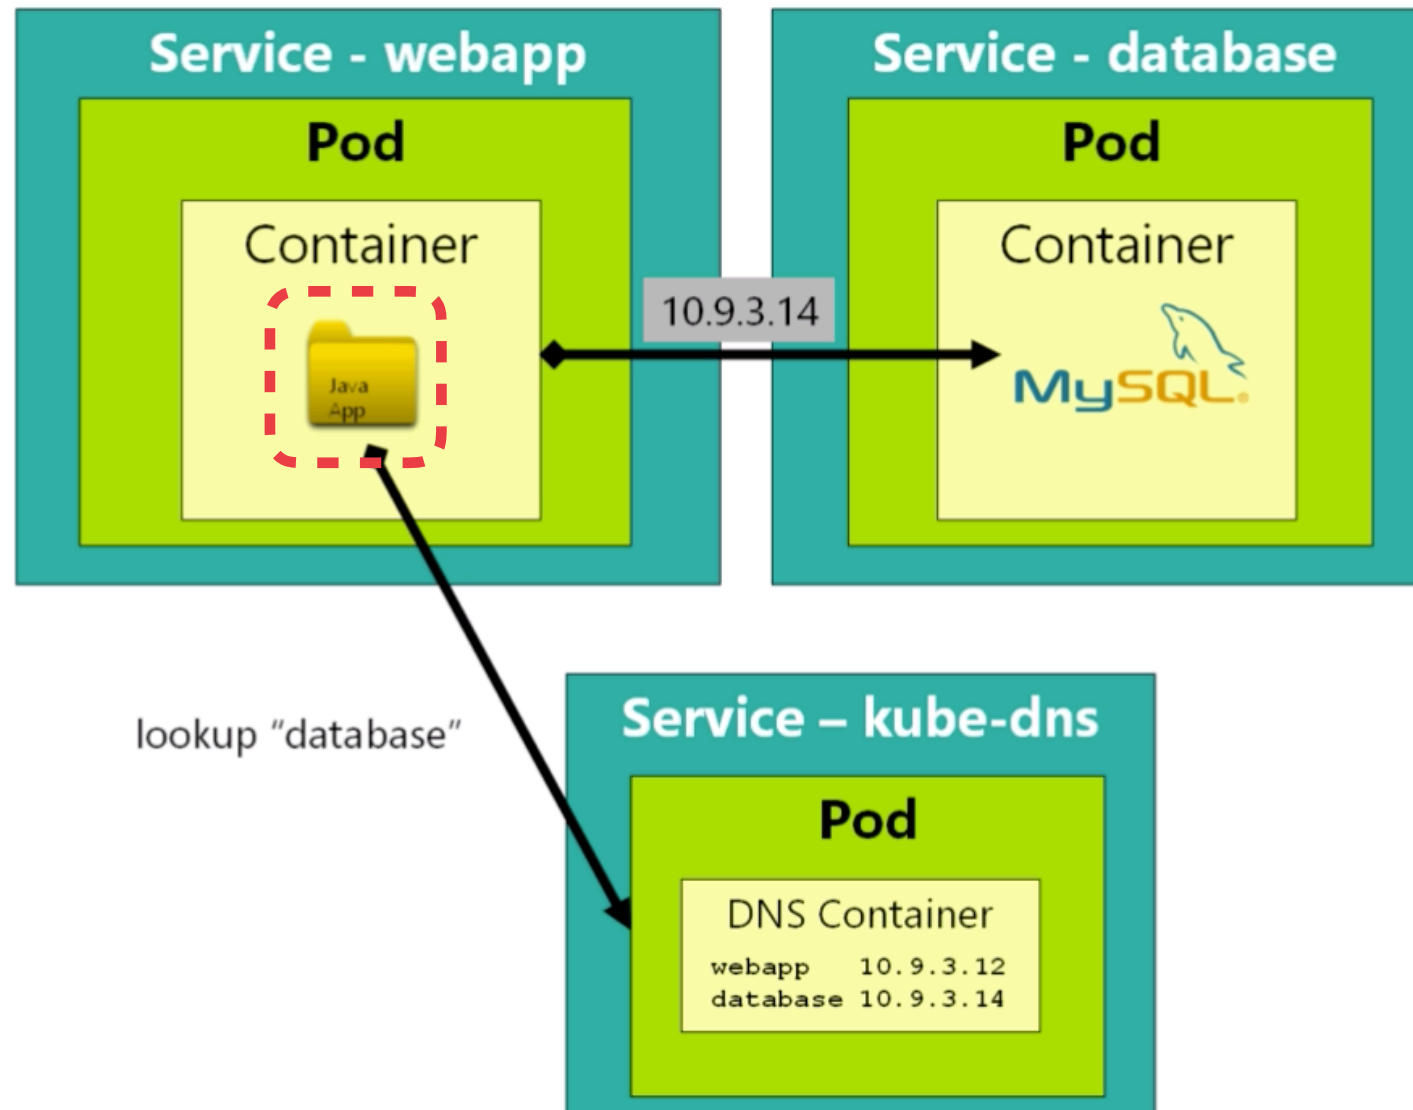
- ▶ This is the file that configures how the dns name resolution is going to work

- ▶ The first line tells the linux system, the container, when we make a reference to the domain name we need to find the dns server located in "10.96.0.10" (in my case)
- ▶ `kubectl get all -n kube-system`

```
$ kubectl get all -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
pod/etcd-minikube	1/1	Running	0	2h
pod/kube-addon-manager-minikube	1/1	Running	1	6d
pod/kube-apiserver-minikube	1/1	Running	0	2h
pod/kube-controller-manager-minikube	1/1	Running	0	2h
pod/kube-dns-86f4d74b45-f5kxg	3/3	Running	4	6d
pod/kube-proxy-jbwq5	1/1	Running	0	2h
pod/kube-scheduler-minikube	1/1	Running	0	2h
pod/kubernetes-dashboard-5498ccf677-zlk8p	1/1	Running	3	6d
pod/storage-provisioner	1/1	Running	3	6d

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP
service/kubernetes-dashboard	NodePort	10.105.32.177	<none>	80:30000/TCP



- Now we understand how it works (next slide)

1. We try to make a request to a URL with the domain database inside it - this container doesn't know what that means
2. Uses that resolv.conf configuration to tell where the DNS service is
3. Where it will lookup that name and it will respond with the correct IP address

- ▶ We can see this working by using the nslookup command
- ▶ nslookup google.com

```
/ # nslookup google.com
nslookup: can't resolve '(null)': Name does not resolve

Name:      google.com
Address 1:  216.58.204.78 1hr25s13-in-f14.1e100.net
Address 2:  2a00:1450:4009:814::200e 1hr25s13-in-x0e.1e100.net
```

- ▶ nslookup database

```
/ # nslookup database
nslookup: can't resolve '(null)': Name does not resolve

Name:      database
Address 1:  10.109.75.215 database.default.svc.cluster.local
/ # |
```


- ▶ exit of our shell - `exit`

```
/ # exit
```

- ▶ `kubectl get all`

```
$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/mysql	1/1	Running	0	19m
pod/queue	1/1	Running	0	1h
pod/webapp-ccb5c74c9-mdzrb	1/1	Running	0	1h
pod/webapp-ccb5c74c9-wjsnn	1/1	Running	0	1h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/database	ClusterIP	10.109.75.215	<none>	3306/TCP	19m
service/fleetman-queue	NodePort	10.104.209.185	<none>	8161:30010/TCP	1h
service/fleetman-webapp	NodePort	10.101.253.163	<none>	80:30080/TCP	1h
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	1h

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/webapp	2	2	2	2	1h

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/webapp-ccb5c74c9	2	2	2	1h

- ▶ go back to the shell
- ▶ `kubectl exec -it webapp-ccb5c74c9-mdzrb sh`
- ▶ We can't use mysql because it is not installed
- ▶ `mysql`

```
/ # mysql  
sh: mysql: not found  
/ #
```

- ▶ Let's install the mysql client

- ▶ `apk update` (this a is linux alpine)

```
/ # apk update
fetch http://dl-cdn.alpinelinux.org/alpine/v3.7/main/x86_64/APKINDEX.tar.gz
fetch http://dl-cdn.alpinelinux.org/alpine/v3.7/community/x86_64/APKINDEX.tar.gz
v3.7.0-214-g519be0a2d1 [http://dl-cdn.alpinelinux.org/alpine/v3.7/main]
v3.7.0-207-gac61833f9b [http://dl-cdn.alpinelinux.org/alpine/v3.7/community]
OK: 9055 distinct packages available
/ #
```

- ▶ `apk add mysql-client`

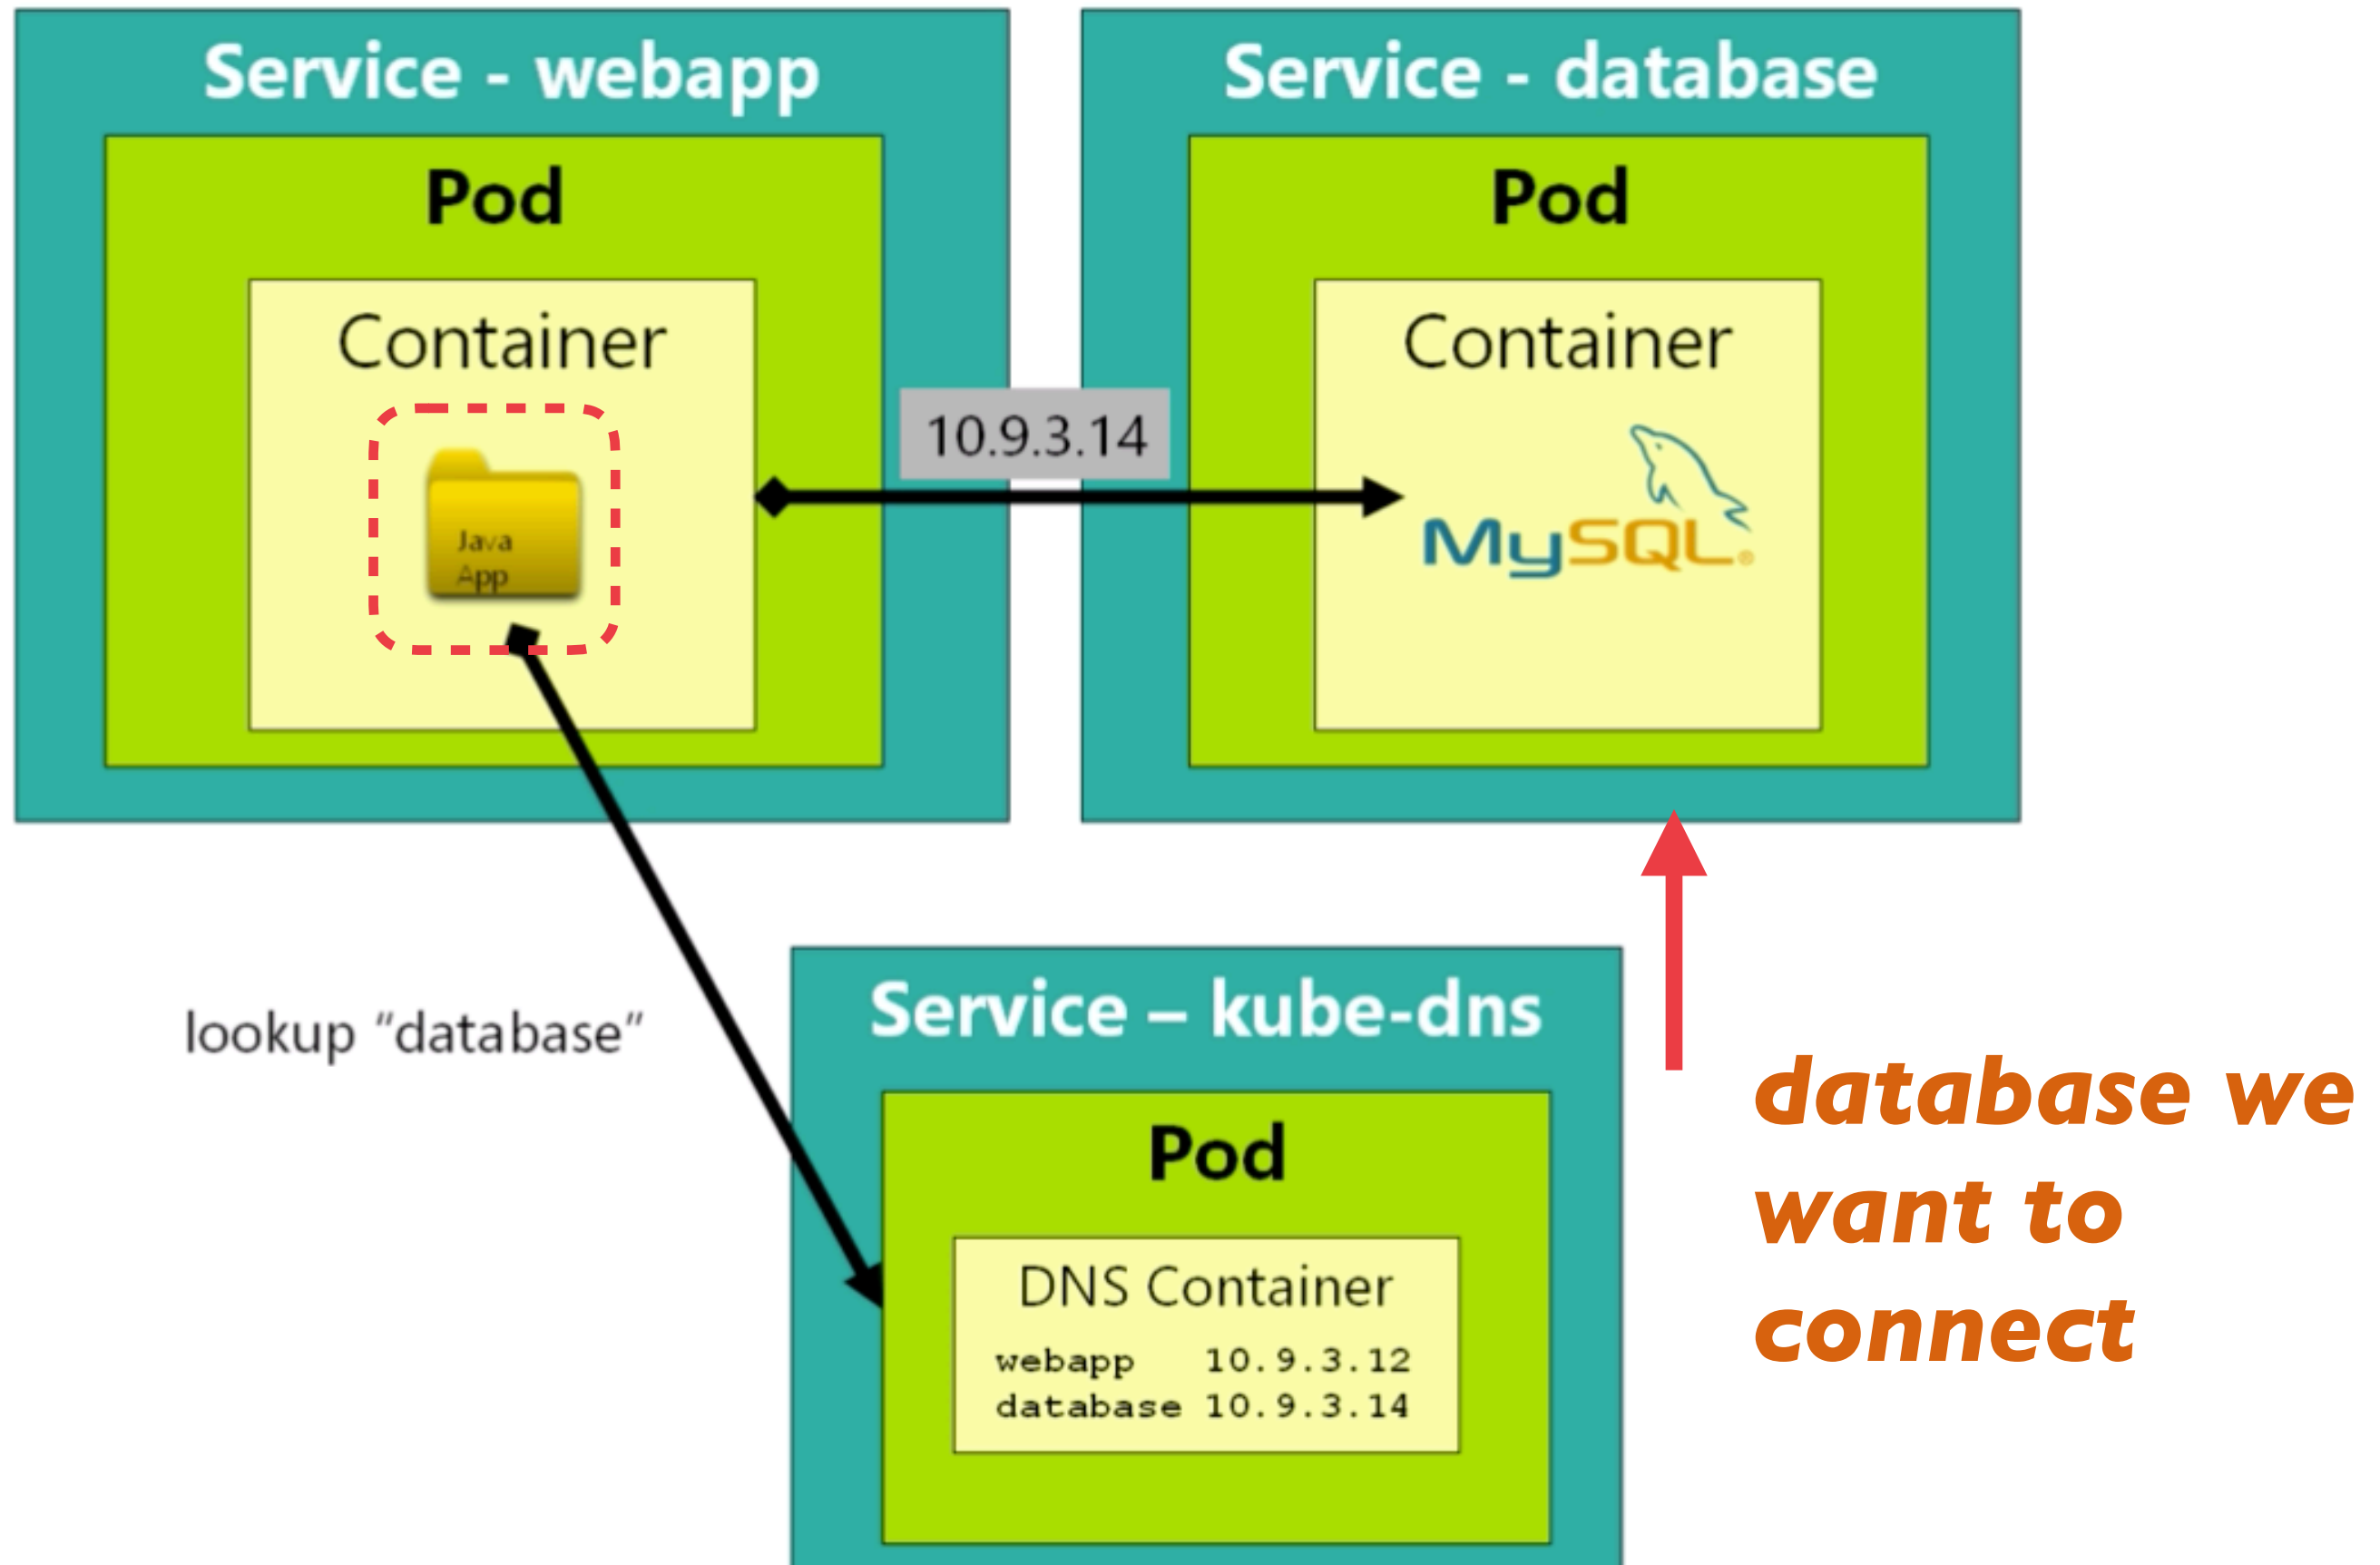
```
/ # apk add mysql-client
(1/6) Installing mariadb-common (10.1.32-r0)
(2/6) Installing ncurses-terminfo-base (6.0_p20171125-r0)
(3/6) Installing ncurses-terminfo (6.0_p20171125-r0)
(4/6) Installing ncurses-libs (6.0_p20171125-r0)
(5/6) Installing mariadb-client (10.1.32-r0)
(6/6) Installing mysql-client (10.1.32-r0)
Executing busybox-1.27.2-r7.trigger
OK: 53 MiB in 34 packages
```

- ▶ We are not installing a database, this is just a command line for my mysql

- ▶ If we try to use `mysql - mysql`

```
/ # mysql  
ERROR 2002 (HY000): Can't connect to local MySQL server through socket  
/ #
```

- ▶ We got this error because we don't have any database server running on this container
- ▶ We want to connect to a database running on a different container on a different pod - represented by the service database



- ▶ `mysql -h database -uroot -ppassword fleetman`

```
/ # mysql -h database -uroot -ppassword fleetman
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MySQL connection id is 2
Server version: 5.7.22 MySQL Community Server (GPL)

Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MySQL [fleetman]> |
```

- ▶ `create table testtable (test varchar(255));`

```
MySQL [fleetman]> create table testtable (test varchar(255));
Query OK, 0 rows affected (0.09 sec)

MySQL [fleetman]> |
```

► show tables

```
MySQL [fleetman]> show tables;
+-----+
| Tables_in_fleetman |
+-----+
| testtable           |
+-----+
1 row in set (0.01 sec)

MySQL [fleetman]> |
```

*Networking in kubernetes
is very easy when we
understand there is a DNS
service*

***this technique - service
discovery - is very
important to connect all
of our micro services
together***

- ▶ We discovered how to find any service by its name
- ▶ Using nslookup we could locate the database service through its name (nslookup database)
- ▶ But the service is not registered with the name database

```
/ #  
/ # nslookup database  
nslookup: can't resolve '(null)': Name does not resolve  
  
Name:      database  
Address 1: 10.104.224.29 database.default.svc.cluster.local  
/ #
```

- ▶ Its registered under its Fully Qualified Domain Name - FQDN:
`database.default.svc.cluster.local`
- ▶ What is happening when we do the lookup, is that it can't find "database" because it didn't matches the FQDN string

- ▶ `cat /etc/resolv.conf`

```
/ # cat /etc/resolv.conf
nameserver 10.96.0.10
search default.svc.cluster.local svc.cluster.local cluster.local
options ndots:5
/ #
```

- ▶ It says that in the event of the name not being found then it should try to append the string `default.svc.cluster.local` then if it's still not found append `svc.cluster.local` and then `cluster.local`

- ▶ The most important point here is the name after **database.**, in this case "**default**", because it relates to the namespace that that service is in

```
/ #  
/ # nslookup database  
nslookup: can't resolve '(null)': Name does not resolve  
  
Name:      database  
Address 1: 10.104.224.29 database.default.svc.cluster.local  
/ #
```

Polytechnic of
Porto

ESTG - School of
Management and Technology